



Monitoring Access to Shared Memory-Mapped Files

By

Christian Sarmoria, Steve Chapin

Presented At

The Digital Forensic Research Conference

DFRWS 2005 USA New Orleans, LA (Aug 17th - 19th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment. As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

<http://dfrws.org>

Monitoring Access to Shared Memory-Mapped Files

Christian G. Sarmoria, Steve J. Chapin

{cgsarmor , chapin} @ecs.syr.edu

Electrical Engineering and Computer Science Dept.

Syracuse University

NY, 13210

August 18, 2005



Overview

- Introduction
- Reconstruction of sequences of events
- Shared memory
- Granularity
- Page-level monitor
- Conclusions
- Q & A

Evidence Gathering

■ ***Post-mortem***

- Filesystem
- Memory dump
- Available logs

■ **Realtime**

- System, application, and network logs
- ***Record system objects and their interactions***
- Evidence is *potential*

Framework for Reconstructing Sequences of Events

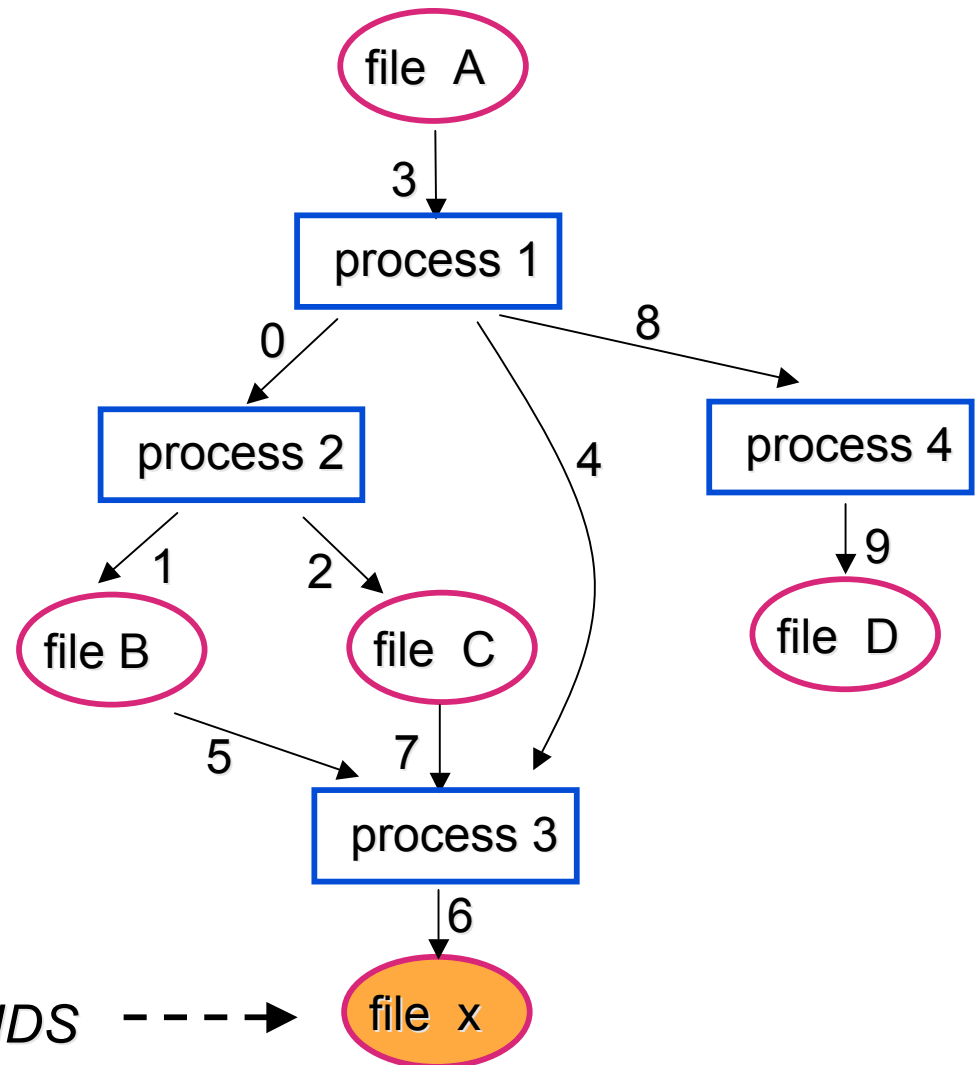
- Monitor system **objects** and **events**
- Gather potential evidence at runtime
- Detection point signaled by an IDS
- Build dependency graph
- Backtrack from detection point back to entry point of the attack

Framework for Reconstructing Sequences of Events

- **System objects:** Process, file, and filename.
- **Event:** Read, write, rename, create process, share memory, etc.
- **Role** of an object in an event: ***cause*** or ***effect***
- **Objects' dependencies:**
 - *Process-Process:* create, signal, share memory
 - *Process-file:* read, write, change attributes
 - *Process-filename:* rename, link, remove, etc.

Framework for Reconstructing Sequences of Events

- T₀: process 1 creates process 2
- T₁: process 2 writes file B
- T₂: process 2 writes file C
- T₃: process 1 reads file A
- T₄: process 1 creates process 3
- T₅: process 3 reads file B
- T₆: process 3 writes file X
- T₇: process 3 reads file C
- T₈: process 1 creates process 4
- T₉: process 4 writes file D



Reconstruction of Sequences of Events

☐ **Current Systems**

- ☐ SNARE (2003)
- ☐ Backtracker (2003)
- ☐ Forensix (2003)
- ☐ CIDS (2004)
- ☐ Improved Backtracker (2005)

☐ **System calls monitoring**

☐ **Shared memory?**

- ☐ System call event only (mmap/munmap)
- ☐ Aggregation and coarse granularity
- ☐ False dependencies
- ☐ Overlooked events

File access





- Current systems: `read` and `write` system calls
- Works for access to the file in the file system
- Once a file is memory-mapped we lose system calls support
- Process accesses a memory-mapped file through direct address space manipulation (pointers)
- We need to monitor memory *read* and *write* instructions

Shared-Memory Object (SMO)


- An *SMO* is a portion of memory address space allocated to a process and shared with other processes
- Roles:
 - *Cause*
 - *Effect*
- Process-SMO dependencies:
 - Process \Rightarrow SMO
 - SMO \Rightarrow Process

Granularity

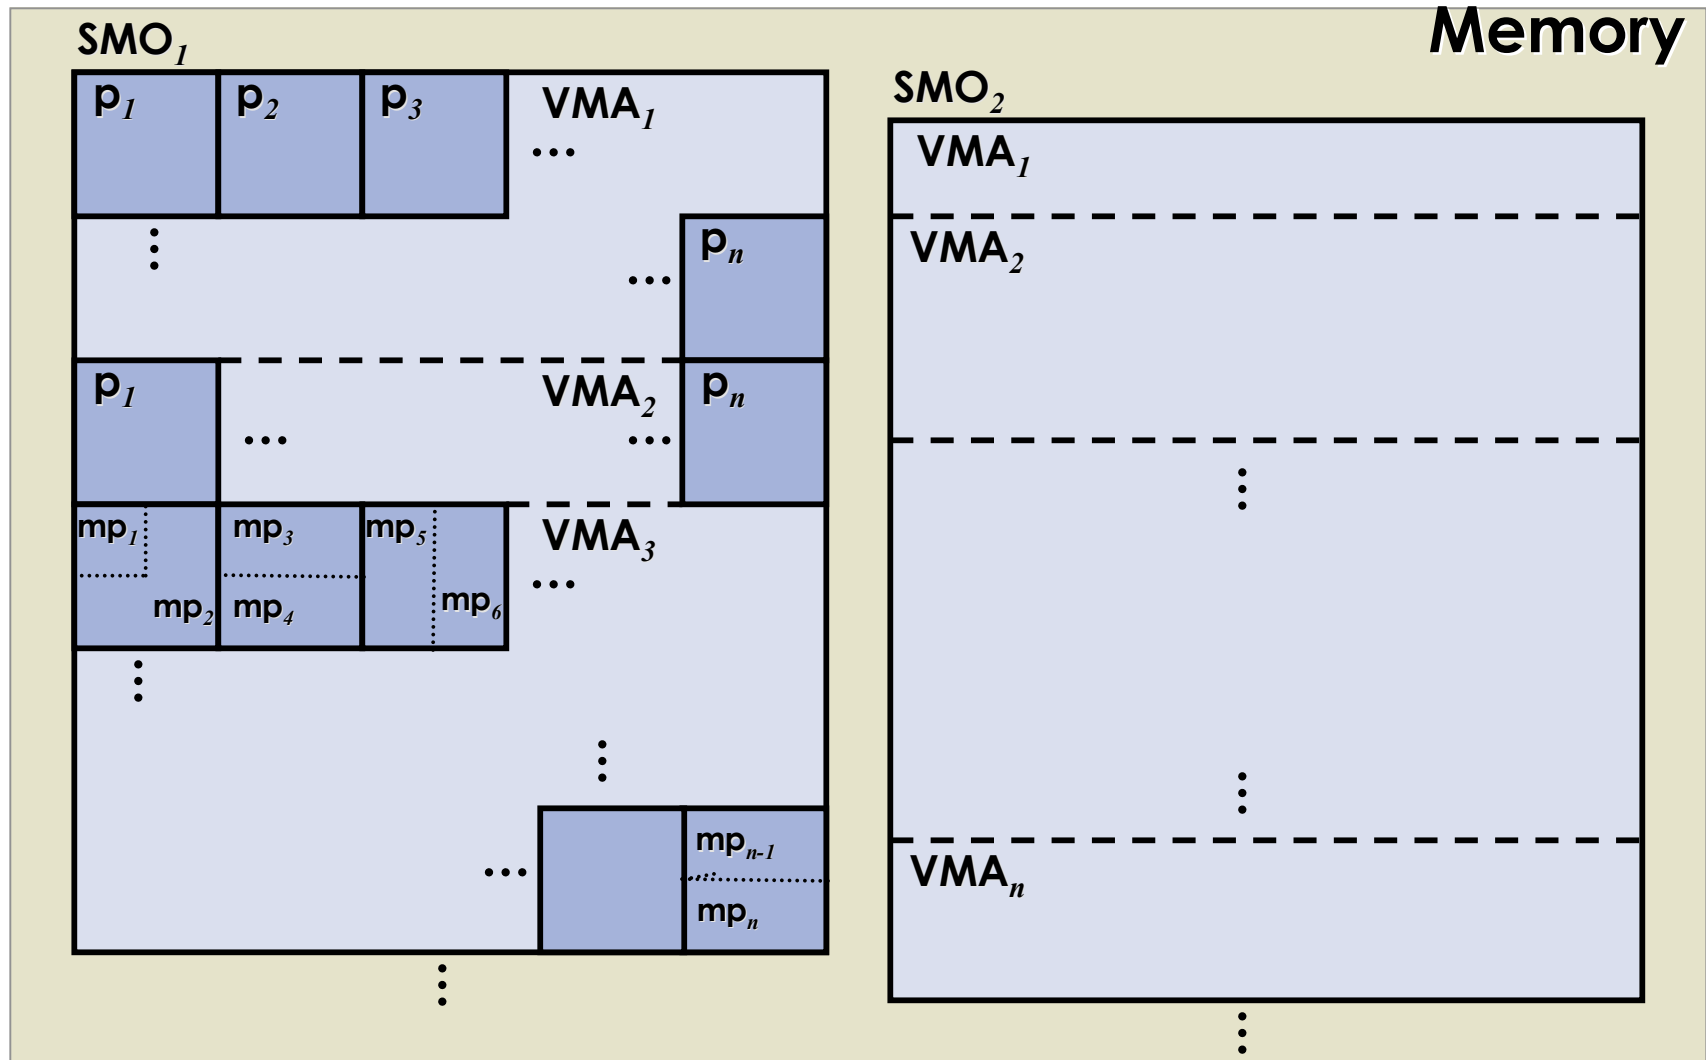
- SMO composed by constituent parts
- A *part* is a section of the address space occupied by the SMO

Constituent Parts	Size
 <i>Object</i>	The whole SMO
 <i>Memory Region (VMA)</i>	Runtime-OS-Defined set of contiguous memory frames
 <i>Page</i>	As defined in the OS, usually 4KB
 <i>Minipage</i>	$1 \text{ byte} \leq \textit{minipage} < \textit{page}$

Granularity



Granularity



Granularity: Properties

- Monitoring constituent parts eliminates a number of Process-Process false dependencies
- The larger the SMO the greater the benefit of applying granularity
- Monitoring at the *minipage* level has to deal with loss of paging hardware support
- Overhead

Technique

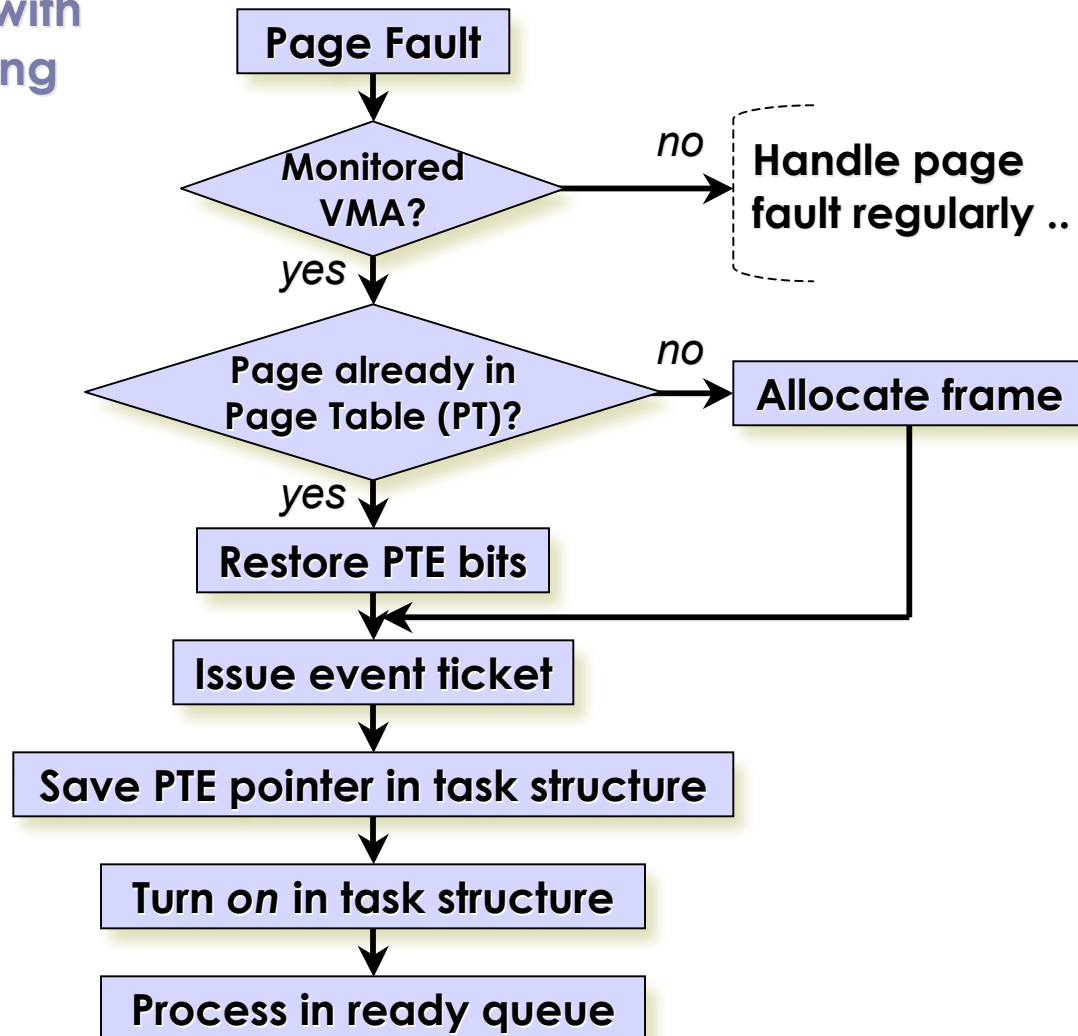
- Linux Kernel 2.4 on x86 architecture
- Does not require processes' source code
- Accesses to shared memory-mapped files at the page level of granularity
- Runtime SMO registration by intercepting `mmap` system call
- Monitor implemented in page fault handler and CPU scheduler
- Event logs saved in circular buffer in kernel space

Technique

- Memory pages where the SMO resides are set protected and read-only through Page Table Entry (PTE)
- Page fault alerts the monitor of SMO access event
- Race condition:
 - Time of page serviced and ready to use is no guarantee of read/write event at present time
 - CPU scheduler confirms event is effective

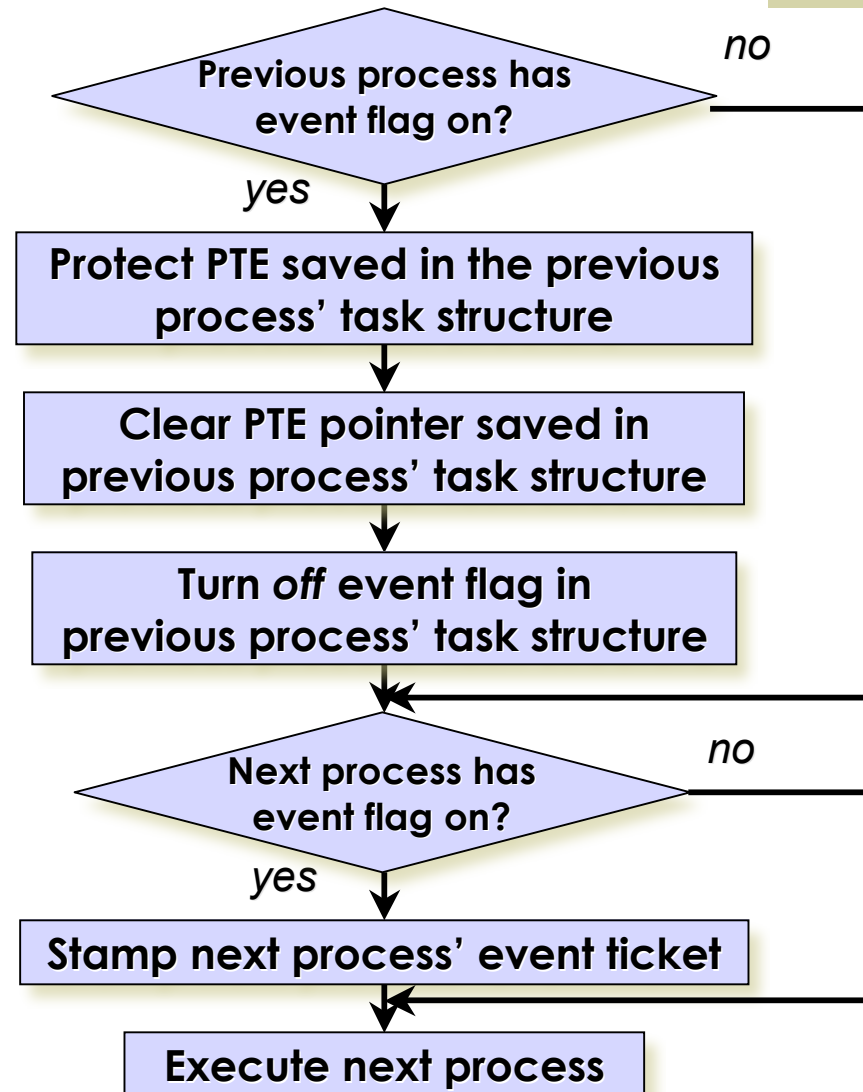
Page Fault Handler

Page Fault Handler with Page-Level Monitoring



CPU Scheduler

CPU Scheduler with Page-Level Monitoring



Synchronization

■ Race condition:

- P_1 and P_2 share a page
- P_1 page faults on it
- Page fault handler serves the page for P_1
- Monitor logs the event
- P_1 ready for execution, waiting for CPU
- P_2 page faults on same page
- Page fault handler serves the page for P_2
- Monitor logs event
- P_2 ready for execution
- P_2 gets CPU before P_1

Results

- Established Process-SMO dependencies for shared memory-mapped files
- x86 limitation: a read after a write in the same page during the same CPU time slice

Results

- Overhead:
 - **One page fault** for any number of **consecutive reads** in the same page during the same CPU time slice
 - **One page fault** for any number of **consecutive writes** in the same page during the same CPU time slice
 - **Two page faults** for any number of **consecutive reads followed** by any number **of consecutive writes** in the same page during the same CPU time slice

Conclusions

- Shared memory is a source of objects of different types which participate in events the same way as other objects do
- OS system call level limitation to monitor shared memory produces false dependencies and overlooked events in current reconstruction systems
- Monitoring SMOs can eliminate a number of false dependencies and reveal others
- Granularity provides a framework to trade off between accuracy and overhead

Questions...

