



Massive Threading: Using GPUs to Increase the Performance of Digital Forensics Tools

By

Lodovico Marziale, Golden Richard and Vassil Roussev

Presented At

The Digital Forensic Research Conference

DFRWS 2007 USA Pittsburgh, PA (Aug 13th - 15th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment. As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

<http://dfrws.org>

Massive Threading: Using GPUs to Increase the Performance of Digital Forensics Tools

Lodovico Marziale, **Golden G. Richard III**, Vassil Roussev



Me:

Professor of Computer Science

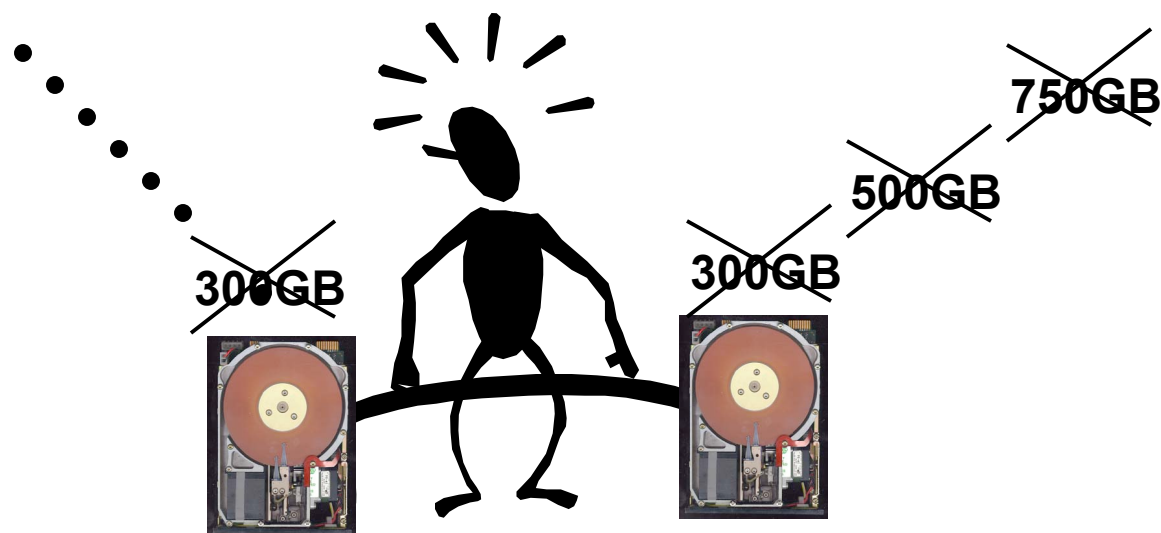
Co-founder, Digital Forensics Solutions

golden@cs.uno.edu

golden@digitalforensicssolutions.com

golden@digdeeply.com

Problem: (Very) Large Targets



- Slow Case Turnaround
- Need:
 - Better software designs
 - More processing power
 - Better forensic techniques

Finding More Processing Power

Filling this gap?

Graphics Processing Units (GPUs)?



Single CPU



Multicore CPUs



Clusters

Quick Scalpel Overview

- Fast, open source file carver
- Simple, two-pass design
- Supports “in-place” file carving
- “Next-generation” file carving will use a different model
 - Headers/footers/other static milestones are “guards”
 - Per-file type code performs deep(-er-er?) analysis to find file / fragment boundaries and do reassembly
- But that’s not the point of the current work
- Use Scalpel as a laboratory for investigating the use of GPUs in digital forensics
- First, multicore discussion

Multicore Support for Scalpel


- Parallelize first pass over image file
 - Thread pool: Spawn one thread for each carving rule
 - Loop
 - Threads sleep
 - Read 10MB block of disk image
 - Threads wake
 - Search for headers in parallel
 - Boyer-Moore binary string search (efficient, fast)
 - Threads synchronize then sleep
 - Selectively search for footers (based on discovered headers)
 - Threads wake
 - End Loop
 - Simple multithreading model yields ~1.4 – 1.7 X speedup for large, in-place carving jobs on multicore boxes
- first pass over image file
- Hard to find forensics software that **doesn't** need to do binary string searches

Multicore (2)

TABLE II
RESULTS FOR CARVING 100GB DISK IMAGE ON DUAL
PROCESSOR, DUAL CORE SUN ULTRA 40 (2.6GHz AMD
OPTERON 2218 PROCESSORS, 16GB RAM). 30 FILE TYPES,
~15M FILES CARVED. EACH RESULT IS THE AVERAGE OF
MULTIPLE, SEQUENTIAL RUNS.

Scalpel 1.60 “vanilla”	13067 secs
Scalpel 1.60 “new q”	8725 secs
Scalpel 1.70MT-multicore	4958 secs

blinds




GPUs?

- Multithreading mandatory for applications to take advantage of multicore CPUs
- Tendency to increase the number of processor cores rather than shoot for huge increases in clock rate
- So you're going to have to do multithreading anyway
- New GPUs are massively parallel, use SIMD, thread-based programming model
- Extend threading models to include GPUs as well?
- Yes. Why?

GPU Horsepower

GFLOPS $1.35\text{GHz} \times 128 \times 2 \text{ instructions per cycle} = \sim 345\text{GFLOPS}$



Filling the Gap: GPUs?

- Previous Generation
 - Specialized processors
 - Vertex shaders
 - Fragment shaders
 - Difficult to program
 - Must cast programs in graphical terms
 - Example: PixelSnort (ACSAC 2006)
- Current Generation
 - Uniform architecture
 - Specialized hardware for performing texture operations, etc. but processors are essentially general purpose

NVIDIA G80: Massively Parallel Architecture



~350 GFLOPS per card

**Can populate a single box with
Multiple G80-based cards**

**Constraints: multiple PCI-E 16
slots, heat, power supply**

8800GTX / G80 GPU

768MB Device Memory

**16 “multiprocessors” X 8 stream
processors**

**Total 128 processors, 1.35GHz
each**

**Hardware thread management,
can schedule millions of threads**

Separate device memory

DMA access to host memory

“Deskside” Supercomputing



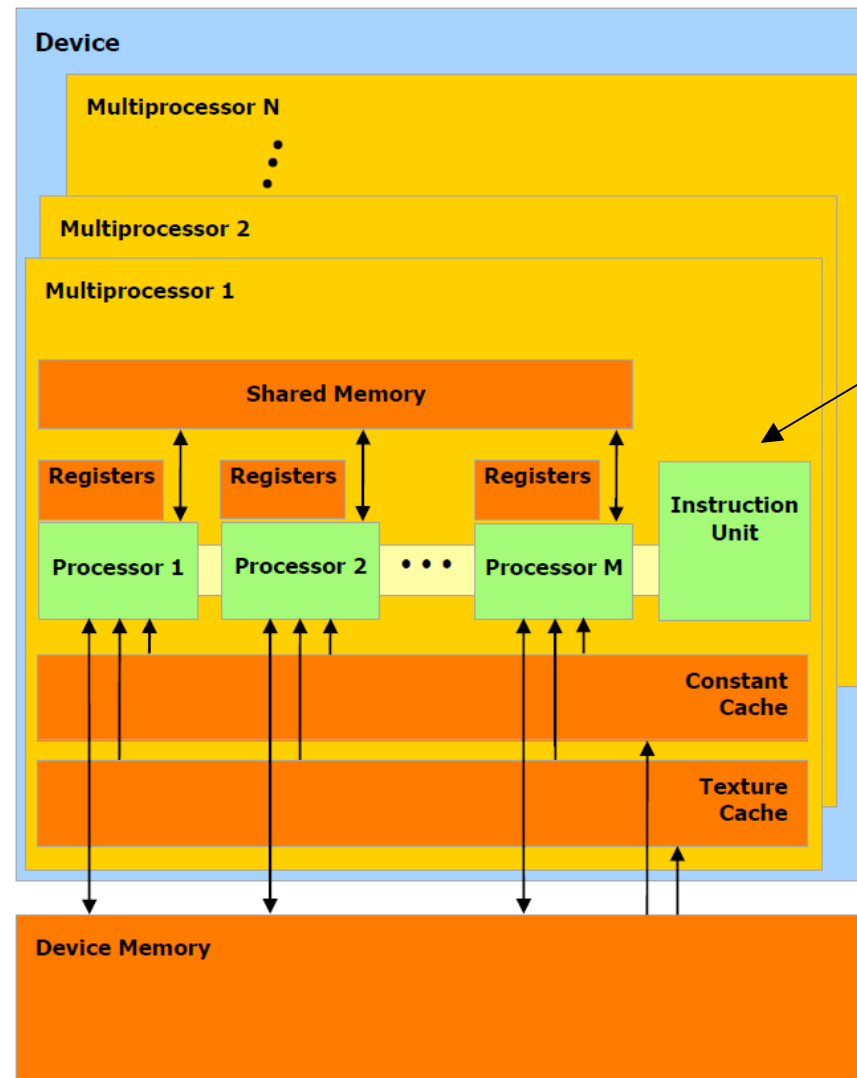
Dual GPUs

3 GB RAM

1 TFLOP

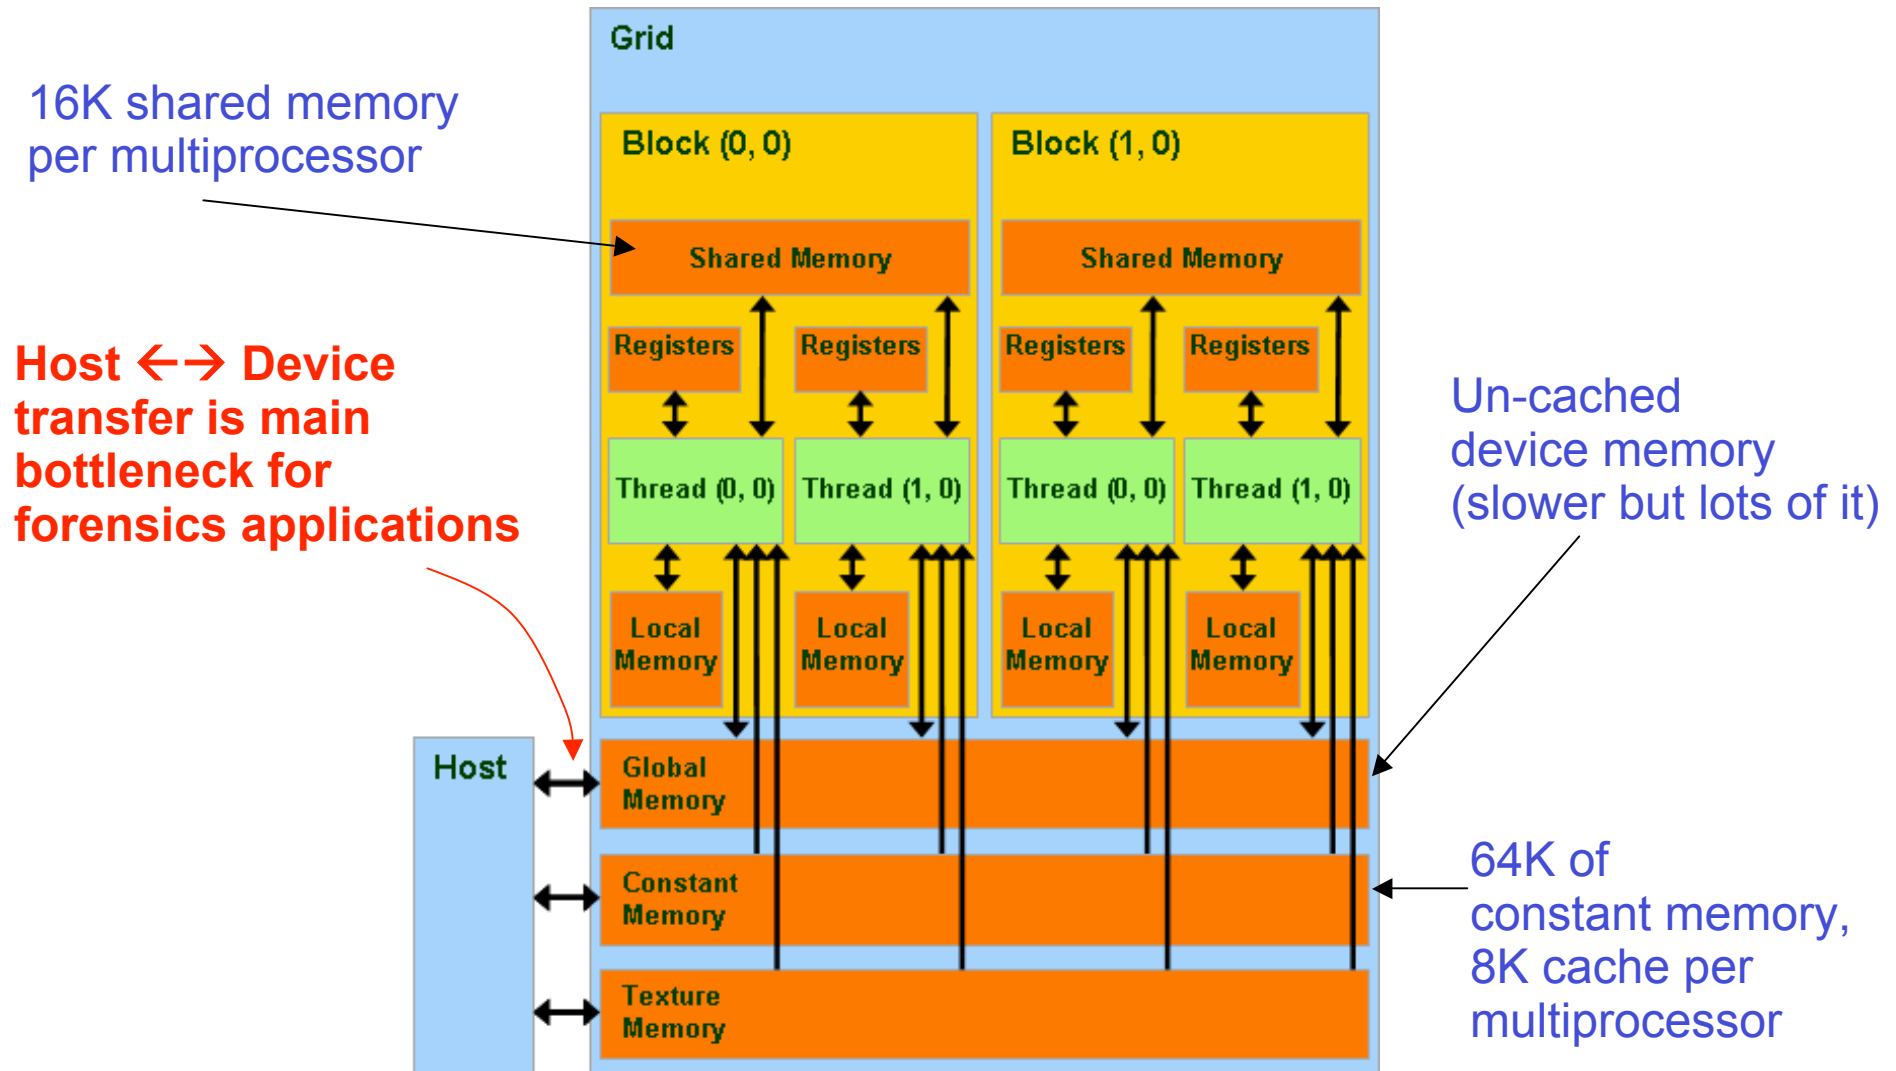
**Connects
via PCI-E**

G80 High-level Architecture



Shared instruction unit is reason that SIMD programs are needed for max speedup

G80 Thread Block Execution



NVIDIA CUDA

- **Common Unified Device Architecture**
- See the SDK documentation for details
- Basic idea:
 - **Code running on host has few limitations**
 - Standard C plus functions for copying data to and from the GPU, starting kernels, ...
 - **Code running on GPU is more limited**
 - Standard C w/o the standard C library
 - Libraries for linear algebra / FFT / etc.
 - No recursion, a few other rules
 - For performance, need to care about thread divergence (SIMD!), staging data in appropriate types of memory

Overview of G80 Experiments

- Develop GPU-enhanced version of Scalpel
- Target binary string search for parallelization
 - Used in virtually all forensics applications
- Compare GPU-enhanced version to:
 - Sequential version
 - Multicore version
- Primary question: Is using the GPU worth the extra programming effort?
- Short answer: Yes.

GPU Carving 0.2

- Store Scalpel headers/footer DB in constant memory (initialized by host), once
 - Loop
 - Read 10MB block of disk image
 - Transfer 10MB block to GPU
 - Spawn $512 * 128$ threads
 - Each thread responsible for searching 160 bytes (+ overlap) for headers/footers
 - **Simple binary string search**
 - Matches encoded in 10MB buffer
 - Headers: index of carving rule stored at match point
 - Footers: negative index of carving rule stored at match point
 - Results returned to Host
 - End Loop
- first pass over image file

GPU Carving 0.2: 20GB/Opteron

TABLE I

RESULTS FOR CARVING 20GB DISK IMAGE ON DUAL PROCESSOR,
DUAL CORE SUN ULTRA 40 (2.6GHz AMD OPTERON 2218
PROCESSORS, 16GB RAM). 30 FILE TYPES, ~3M FILES CARVED.
EACH RESULT IS THE AVERAGE OF MULTIPLE, SEQUENTIAL RUNS.

Scalpel 1.60 “vanilla”	2672 secs
Scalpel 1.60 “new q”	1784 secs
Scalpel 1.70MT-multicore	1054 secs
Scalpel 1.70MT-gpu-0.20	860 secs

GPU Carving 0.2: 100GB/Opteron


TABLE II

RESULTS FOR CARVING 100GB DISK IMAGE ON DUAL PROCESSOR, DUAL CORE SUN ULTRA 40 (2.6GHz AMD OPTERON 2218 PROCESSORS, 16GB RAM). 30 FILE TYPES, ~15M FILES CARVED. EACH RESULT IS THE AVERAGE OF MULTIPLE, SEQUENTIAL RUNS.

Scalpel 1.60 “vanilla”	13067 secs
Scalpel 1.60 “new q”	8725 secs
Scalpel 1.70MT-multicore	4958 secs
Scalpel 1.70MT-gpu-0.20	5185 secs

Cage Match!

(Or: The Chair Wants His Machine Back...)



Dual 2.6GHz Opteron (4 cores)
16GB RAM, SATA
Single 8800GTX

Vs.



Single 2.4GHz Core2Duo (2 cores)
4GB RAM, SATA
Single 8800GTX

GPU Carving 0.3

- Store Scalpel headers/footers in constant memory (initialized by host)
 - Loop
 - Read 10MB block of disk image
 - Transfer 10MB block to GPU
 - **Spawn 10M threads (!)**
 - Device memory staged in 1K of shared memory per multiprocessor
 - Each thread responsible for searching for headers/footers **in place** (no iteration)
 - **Simple binary string search**
 - Matches encoded in 10MB buffer
 - Headers: index of carving rule stored at match point
 - Footers: negative index of carving rule stored at match point
 - Results returned to Host
 - End Loop
- first pass over image file

GPU Carving: 20GB/Dell XPS

TABLE III

RESULTS FOR CARVING 20GB DISK IMAGE ON SINGLE PROCESSOR, DUAL CORE DELL XPS 710 (2.4GHz CORE2DUO PROCESSOR, 4GB RAM). 30 FILE TYPES, ~3M FILES CARVED. EACH RESULT IS THE AVERAGE OF MULTIPLE, SEQUENTIAL RUNS.

Scalpel 1.60 “new q”	1260 secs
Scalpel 1.70MT-multicore	861 secs
Scalpel 1.70MT-gpu-0.20	686 secs
Scalpel 1.70MT-gpu-0.30	446 secs

GPU Carving: 100GB/Dell XPS

TABLE IV

RESULTS FOR CARVING 100GB DISK IMAGE ON SINGLE PROCESSOR, DUAL CORE DELL XPS 710 (2.4GHZ CORE2DUO PROCESSOR, 4GB RAM). 30 FILE TYPES, ~15M FILES CARVED. EACH RESULT IS THE AVERAGE OF MULTIPLE, SEQUENTIAL RUNS.

Scalpel 1.60 “new q”	7105 secs
Scalpel 1.70MT-multicore	5096 secs
Scalpel 1.70MT-gpu-0.20	4192 secs
Scalpel 1.70MT-gpu-0.30	3198 secs

Bored GPU == Poor Performance

TABLE V

RESULTS FOR CARVING 500GB DISK IMAGE ON SINGLE PROCESSOR, DUAL CORE DELL XPS 710 (2.4GHz CORE2DUO PROCESSOR, 4GB RAM). 2 FILE TYPES, ~73,000 FILES CARVED.

Scalpel 1.60 "new q"	9946 secs
Scalpel 1.70MT-multicore	9922 secs
Scalpel 1.70MT-gpu-0.30	12168 secs



But this is NOT an appropriate model for using GPUs, anyway...



Discussion

- Host \leftrightarrow GPU transfers have significant bandwidth limitations
 - ~1.3GB/sec transfer rate (observed)
 - 2GB/sec (theoretical)
 - 3GB/sec (theoretical) with page “pinning” (**not** observed by us!)
- Current: Host threads blocked when GPU is executing
 - Host thread(s) should be working...
 - We didn't overlap host / GPU computation because we wanted to measure GPU performance in isolation
- Current: No overlap of disk I/O and compute
 - For neither GPU nor multicore version
- Current: No compression for host \leftrightarrow GPU transfers
- But...

Discussion (2)

- BUT:
 - GPU is currently using simple binary string search
 - Sequential/multicore code using optimized Boyer-Moore string search
- Despite this, GPU much faster than multicore when there's enough searching to do...
- Considering only search time, GPU > 2X faster than multicore even with these limitations

Discussion: 20GB

- Sequential:
 - Header/footer searches: 73%
 - Image file disk reads: 19%
 - Other: 8%
- Multicore:
 - Header/footer searches: 48%
 - Image file disk reads: 44%
 - Other: 8%
- GPU:
 - Total time spent in device <--> host transfers: 7%
 - Total time spent in header/footer searches: 24%
 - Total time spent in image file disk reads: 43%
 - Other: 26%

Conclusions / Future Work

- New GPUs are fast and worthy of our attention
- Not that difficult to program, but requires a different threading model
- Host \leftrightarrow GPU bandwidth is an issue
- Overcome this by:
 - Overlapping host and GPU computation
 - Overlapping disk I/O and GPU computation
 - Disk, multicore, GPU(s) should all be busy
 - Overlapping transfers to one GPU while another computes!
 - Compression for host \leftrightarrow GPU transfers
- Interesting issues in simultaneous use
 - Simple example: Binary string search: GPU better at **NOT finding things!**
 - **Reduces thread control flow divergence**

Je suis fini, Happy GPU Hacking...



Scalpel v1.7x (alpha) is available for testing

Must have NVIDIA G80-based graphics card

Currently runs only under Linux (waiting for CUDA gcc support under Win32)

Feel free to use this as a basis for development of other GPU-enhanced tools...

golden@cs.uno.edu