



Automated Mapping of Large Binary Objects Using Primitive Fragment Type Classification

By

**Gregory Conti, Sergey Bratus, Benjamin Sangster,
Roy Ragsdale, Matthew Supan, Andrew Lichtenberg,
Robert Perez-Aleman and Anna Shubina**

From the proceedings of

The Digital Forensic Research Conference

DFRWS 2010 USA

Portland, OR (Aug 2nd - 4th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment.

As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

<http://dfrws.org>

available at www.sciencedirect.comjournal homepage: www.elsevier.com/locate/diinDigital
Investigation

Automated mapping of large binary objects using primitive fragment type classification

Gregory Conti^{a,*}, Sergey Bratus^b, Anna Shubina^b, Benjamin Sangster^a, Roy Ragsdale^a,
Matthew Supan^a, Andrew Lichtenberg^c, Robert Perez-Aleman^a

^a United States Military Academy at West Point, West Point, NY, United States

^b Dartmouth College, Hanover, NH, United States

^c Skidmore College, Saratoga Springs, NY, United States

ABSTRACT

Keywords:

Binary mapping
Binary analysis
File carving
Hex editors
Classification
Reverse engineering

Security analysts, reverse engineers, and forensic analysts are regularly faced with large binary objects, such as executable and data files, process memory dumps, disk images and hibernation files, often Gigabytes or larger in size and frequently of unknown, suspect, or poorly documented structure. Binary objects of this magnitude far exceed the capabilities of traditional hex editors and textual command line tools, frustrating analysis. This paper studies automated means to map these large binary objects by classifying regions using a multi-dimensional, information-theoretic approach. We make several contributions including the introduction of the binary mapping metaphor and its associated applications, as well as techniques for type classification of low-level binary fragments. We validate the efficacy of our approach through a series of classification experiments and an analytic case study. Our results indicate that automated mapping can help speed manual and automated analysis activities and can be generalized to incorporate many low-level fragment classification techniques.

© 2010 Digital Forensic Research Workshop. Published by Elsevier Ltd. All rights reserved.

1. Introduction

Binary data is ubiquitous. Security analysts are constantly challenged to gain insight into large binary objects in the form of data and executable files, file systems, hibernation files, process memory, or network traffic. Although specialized tools exist for the study of common file and file system formats, the general case of unknown or malformed binary objects lacks all but the most basic tools. Hex editors are the most common tool employed to perform low-level analysis of binary objects;

however their utility drops off rapidly as file size increases. Analysts supplement hex editors with command line utilities such as *grep* to search for known or guessed patterns and strings and variants thereof to locate printable ASCII and Unicode strings, but these approaches are rudimentary and do not scale. Large binary objects, sometimes 2 GB or larger, are becoming more and more common, and an automated approach is necessary to facilitate effective exploratory analysis.

To address the problem of analysis of large binary objects, we present binary mapping techniques that compute start

The views expressed in this article are those of the authors and do not reflect the official policy or position of the United States Military Academy, the Department of the Army, the Department of Defense or the U.S. Government.

* Corresponding author.

E-mail addresses: gregory-conti@us.army.mil (G. Conti), sergey@cs.dartmouth.edu (S. Bratus), ashubina@cs.dartmouth.edu (A. Shubina), benjamin.sangster@us.army.mil (B. Sangster), roy.ragsdale@us.army.mil (R. Ragsdale), matthew.supan@us.army.mil (M. Supan), alichten@skidmore.edu (A. Lichtenberg), robert.perezalemany@us.army.mil (R. Perez-Aleman).

1742-2876/\$ – see front matter © 2010 Digital Forensic Research Workshop. Published by Elsevier Ltd. All rights reserved.

doi:10.1016/j.diin.2010.05.002

and stop offsets for each distinct region within a binary object as well as the region's probable primitive data type. We define primitive types as families of homogeneous data with closely related binary structure. Examples of primitive types include: US-ASCII encoded English text, 32-bit x86 protected mode machine code, random number sequences, image bitmaps; essentially any group of similar data, possibly encoded, compressed, or encrypted in the same way. However, one should not confuse an object of a primitive type with composite objects. For example, even simple file formats include a metadata header alongside the actual data payload. In this paper we treat distinctly structured regions as discrete types. Our primary use case for binary mapping is the exploration of large files and memory dumps. Binary maps, once generated, can then be used to guide manual analysis or to inform further automated analysis. For example, maps of binary objects can be used to assist humans in navigating to regions of interest. They can also be used to highlight certain desired regions and filter those that are less important to the analyst.

We make several contributions, including a generalized framework for mapping large binary objects, which will work with many low-level binary fragment classifiers. We also present a multi-dimensional information-theoretic fragment classifier tested using 14,000 primitive fragments of 14 distinct types. We validate our work through a series of classification experiments and an analytic case study employing an implementation of our mapping techniques. Our classifier does not assume a priori knowledge of file formats, such as magic numbers or header/footer signatures at known offsets, but instead relies upon the raw structure of the data to identify regions of interest. This property is important because many primitive types lack the consistent structure to support the use of naive signature matching and location semantics. Our framework is designed to support such tasks as rapidly analyzing undocumented file formats, identifying internal data structures that may be targets for fuzzing, exploring the behavior of applications generating binary data files, malware analysis, locating encrypted regions, and studying process memory dumps.

This paper is organized as follows. Section 2 places our research in the field of related work. Section 3 studies the internal structure of binary objects and discusses generalized binary mapping. Section 4 presents our analytic techniques. Sections 5 and 6 provide an evaluation of our approach through experimentation and an analytic case study. Finally, Section 7 presents our conclusions and promising directions for future work.

2. Related work

Work related to our study of binary mapping consists primarily of file classification and file carving research. In the file type classification field, Li conducted 1-gram analysis of byte values to create a “fileprint” and used fileprints to classify both complete and “truncated” files (Li et al., 2005). Stolfo used n -grams to classify files, including document files, containing embedded malicious software (Stolfo et al., 2005). Hickok used a combination of file extensions and magic numbers to rapidly classify files (Hickok et al., 2005). McDaniel used three

different algorithms based on byte frequency analysis, byte frequency correlation analysis, and file header/trailer analysis to perform file type identification (McDaniel and Heydari, 2003). Hall attempted file type identification using entropy and compressibility by employing a 90 byte sliding window (Hall and Davis). Karresand used an algorithm based on the rate of change of byte values and claimed a 99.2% detection rate of jpeg files, when taking into account signatures commonly found within such files (Karresand and Shahmehri, 2006).

File carving and the related area of file fragment reassembly are also related to our work. The work of Shanmugasundaram demonstrated that it is possible to reassemble document fragments by using context models to sequence fragments. Veenman explored the use of statistical techniques to classify file clusters by file type in order to support file recovery from cluster-based media, but examined files as a whole rather than considering their constituent parts (Veenman, 2007). Richard developed Scalpel, a high performance file carver, which used a database of header and footer signatures, at predictable offsets, to extract files. Scalpel is an evolution of the Foremost file carver, and Richard's research included an in depth performance comparison (Richard and Roussev, 2005).

Some researchers have explored the idea of classifying fragments. Calhoun attempted to classify fragments of .jpg, .gif, .pdf and .bmp files using algorithms based on Fisher's linear discriminant and longest common subsequences (Calhoun and Coles, 2008). Erbacher used a sliding window algorithm and graphical depictions of statistical values to visually identify characteristic signatures found in seven file types: .doc, .exe, .jpg, .pdf, .ppt, .xls, and .zip, but performed only a small study consisting of five files per type (Erbacher and Mulholland, 2007). Moody used a 256-byte sliding window to identify fragments extracted from 25 each of .bmp, .csv, .dll, .exe, .html, .jpg, .txt, and .xls files, again a very small sample (Moody and Erbacher, 2008). However, we believe that inclusion of container files, i.e., file formats commonly used to hold objects of many varied primitive types such as .pdf, .xls, .doc, and .ppt, in each of these analyses limits the effective application of their results to binary mapping because container file structure can vary widely based on the distribution of primitive types found within. Roussev made a similar point by suggesting that researchers must understand the “primitive” format of a fragment and whether the fragment is part of a compound file structure (Roussev and Garfinkel, 2009). Shamir studied the ability of an attacker to extract cryptographic keys from large hard drives using a 64 byte sliding window and counting the number of unique bytes in the window (Shamir and van Someren, 1999). Importantly, Garfinkel stressed the significance of locating and validating objects contained within container documents and complex file formats (Garfinkel, 2007).

The common trend in much of the preceding work is that researchers have incorporated only small numbers of files and fragments in their analyses, frequently relied on header and footer signatures at known offsets, or focused on extracting files from disk images. In many cases, file fragments, when used, were arbitrarily derived from complex container objects. The novelty of our work springs from countering these trends.

We do not seek to classify files themselves; instead we seek to map and classify regions *within* files, or other binary objects, with the goal of conserving analyst time and attention and informing further machine processing.

3. Mapping the internal structure of binary objects

To better understand the value of automated mapping, it is useful to visually examine the internal structure of an object, see Fig. 1. The image is a *byteplot*, a graphical depiction of a binary object, showing a Microsoft Word 2003 document (Conti and Dean, 2008; Conti et al., 2008). In a *byteplot*, each byte in the binary object is sequentially mapped to a pixel. The plotting of byte values in the object starts at the top left of the image. Subsequent byte values in the object are plotted from left to right, wrapping at the end of each horizontal row. The coloring scheme maps byte values from 0 (black) to 255 (bright white). As you examine the image, note the regions of similar structure. Our goal with automated mapping is (1) to locate the start and stop offsets of these distinct regions and (2) to




Fig. 1 – Internal structure of a 405 KB Microsoft Word 2003 document depicted using the byteplot visualization. Note the distinct regions of differing structure, including US-ASCII text (top) and compressed image (middle).

identify the region's primitive type. Manual inspection of the document using a hex editor allowed us to identify some regions as US-ASCII text, Basic Latin Unicode hyperlinks, a compressed PNG image, and numerous other data structures. However, our motivation for binary mapping is to avoid tedious manual inspection and create an effective automated approach. Because the number of primitive types is large and growing, a binary mapping framework must be flexible enough to incorporate new primitive data types as they are encountered or as an analyst desires their inclusion.

Accurate classification is a fundamental requirement of binary mapping. However, to understand classification requirements, we must first understand the types of content we desire to identify. The range of possible fragment types is large, but a useful way to think about the possibilities is in terms of entropy, the uncertainty (sometimes colloquially described as randomness), of fragments, see Fig. 2. High entropy types are indicative of compressed, encrypted, or random data; medium entropy types, such as machine code and written human languages, exhibit noticeable structure, and low entropy types may include uncompressed media and repeating values, such as is found in padding. We do not claim that the classification in Fig. 2 is complete, but instead include it here to demonstrate the hierarchy that is formed when identifying primitive types. As an example, a binary fragment may be identified as being of medium entropy and then further identified as a written language. The written language may then be categorized as to the encoding scheme and specific language that was used. As another example, a high entropy fragment may be encrypted, compressed, or random. If compressed, the region may be classified by the algorithm used to create it, such as Run Length Encoding (RLE) or Lempel Ziv Welch (LZW). A classifier used in binary mapping should attempt to accurately identify a fragment, but when this is not possible, the classifier should state that the region is unknown. We will discuss the practical construction of a fragment corpus and determining unknowns later in the paper. The following is a notional example of a binary map.

```
0000-03FF Data Structure (File Header?)
0400-07FF US-ASCII Text (English)
0800-9FFF Bitmap Image
8000-9FFF Variable Length Array
A000-BFFF Compressed Data (LZW?)
A000-BFFF Basic Latin Unicode (English)
C000-CFFF Unknown Region
D000-D02E Repeating Value (0xFF)
D400-D41C Encrypted Region (AES?)
E000-FFFF Basic Latin Unicode (Hyperlinks?)
...
```

The creation of such a map would allow analysts to quickly identify regions of interest for more detailed examination. In addition, a map to a binary object can serve as a navigational aid if integrated directly into interactive analysis tools, e.g., double clicking a region immediately jumps the tool's focus directly to the desired region. More importantly, a map can create an efficient interaction metaphor, allowing a user to filter or highlight various regions or classes of primitive types. Finally, the map can be modified by the analyst throughout the analytic process and perhaps even marked up with

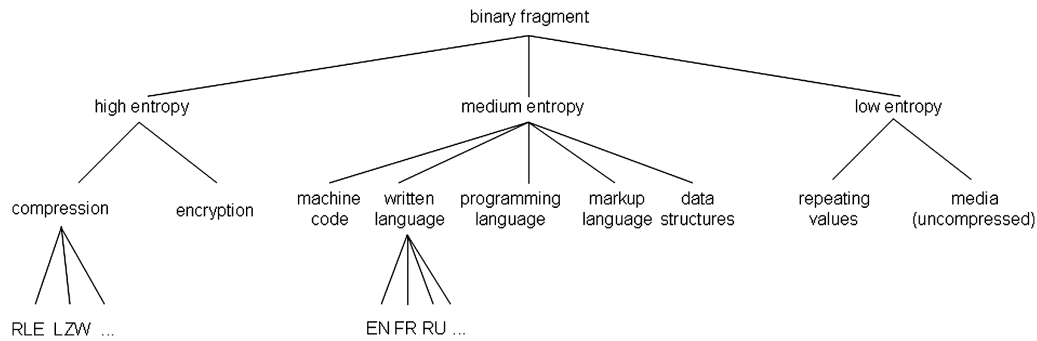


Fig. 2 – Example hierarchy of primitive data types commonly found within binary objects.

analyst notes. For example, the analyst may discover a new primitive type and feed its statistical signature back into the classification engine for future automated identification.

In some cases, the map will be small. For example, many image file formats contain a small number of differing regions, such as a header, footer, and compressed image payload. Even a small map can facilitate rapid analysis despite large file sizes. The analyst need only examine a few regions to gain rapid understanding of a file. In other cases, objects may be more complex, such as Microsoft Word binary data files or memory dumps. Automated mapping is useful in these cases as well, especially when compared to traditional manual analysis. Such objects can be mapped by the machine, which then could be used to guide the analyst to regions of interest.

It is important to note that similar regions within binary objects may differ significantly in size. They may be small, perhaps only tens of bytes, like the header of a PNG file, or very large, such as a 1 MB block of encrypted data. The size of the fragment is important because many classifiers, particularly those based on statistical techniques, require a minimum sample size to provide accurate results. We will discuss the implications of sample size and its impact on our use of a sliding window algorithm later in the paper.

4. Analytic techniques

To classify fragments we created statistical signatures of 14,000 fragments (1000 fragments of 14 commonly encountered primitive types). The size of each fragment was 1024 bytes, and they were collected in two ways. Some were collected directly from files known to consist of a single type, such as a file containing solely random numbers. In the case of files with headers and/or footers and a core payload of a desired primitive type, we extracted fragments from the middle of the file or, if possible, using knowledge of a region's exact location. To understand the statistical characteristics of each type and to facilitate classification, we carefully selected four statistical tests and used these tests to develop statistical signatures for each fragment.

4.1. Statistical tests

With the aim of characterizing the structure found in fragments, we considered many statistical tests, but chose four:

Shannon Entropy, Chi Square, Hamming Weight, and Arithmetic Mean. We chose these tests because we believed they would highlight statistical differences between primitive types in ways that would assist classification.

4.1.1. Shannon entropy

Shannon Entropy (H) is an established technique for measuring uncertainty, sometimes colloquially described as randomness, developed by Claude Shannon (Shannon, 1948).

$$H(X) = - \sum_{i=0}^{n-1} p(X_i) \log_b p(X_i)$$

In the case of byte-level entropy analysis, X is a random variable with 256 potential outcomes: $\{x_i: i = 0, \dots, 255\}$ and base $b = 10$. The probability mass function, $p(x_i)$ is the probability of byte value i within a given fragment. We explored several variations including the use of byte bigrams (256^2 possible outcomes) and byte trigrams (256^3 possible outcomes), ultimately choosing to use byte bigrams such that $\{x_i: i = 0, \dots, 65535\}$, i.e., treating each two byte sequence as a 16 bit value. We made this choice because we found greater accuracy using digrams than single byte values (unigrams) in initial experimentation.

4.1.2. Arithmetic mean

The arithmetic mean is simply the sum of the byte values in a given fragment divided by the fragment size.

4.1.3. Chi square

The Chi Square (X^2) Goodness of Fit Test is an effective means of measuring randomness and is sensitive to differences in random, pseudo random, and compressed data. We used the test to compare the observed distribution of byte values to a uniform random distribution.

$$X^2 = \sum_{i=0}^{n-1} \frac{(\text{observed} - \text{expected})^2}{\text{expected}}$$

For each byte value 0, ..., 255 we calculate the sum of the squared differences between the observed frequency of the byte in the sample and the expected frequency of the byte in a uniform random distribution divided by the expected frequency. We then calculate the upper probability of the Chi Square distribution using this Chi Square value and 255 degrees of freedom, which yields a value between 0 and 1.

As can be seen from Table 1, in cases where data was plain text (that is, 7-bit as opposed to the 16-bit uniform distribution it was compared to), the Chi Square test correctly indicated that the data was distributed very differently from uniform random data by returning very low values that got rounded to 0.

4.1.4. Hamming weight

Hamming Weight is determined by counting the number of non-zero symbols in a given alphabet. For our experiments we considered the alphabet to contain just binary data, i.e., an alphabet of zeroes and ones. We present Hamming Weight as the fraction of the total number of ones divided by the total number of bits.

While our primary experiments employed these four statistical measures, we do not claim they are necessarily the best possible set. We also explored the use of Monte Carlo Pi, Serial Correlation, and the Index of Coincidence (Friedman, 1922). Our initial results indicate that Monte Carlo Pi and Serial Correlation merit further exploration. However, we found that the Index of Coincidence is closely correlated with Shannon Entropy, but yields slightly less accurate classification results, and may not merit further exploration.

In order to support our classification algorithm, we normalized each range, when appropriate. For Shannon entropy and mean, we normalized each value to fall between 0.0 and 1.0 based on the estimated range of possible values yielded by each calculation. The mean has a minimum value of 0 and a maximum value of 255, so we divided each value by 255. Shannon entropy for a 1 KB sample has a minimum value of 0 and a maximum value of 10, so we divided each value by 10. Both Hamming weight and Chi Square values already fell between 0 and 1 so we did not normalize them.

4.2. Determining statistical signatures of fragments

For this phase of our research, we carefully collected the samples for our fragment corpus and calculated the statistical signature of each, using the previously discussed statistical tests. Summary statistics can be seen in Table 1 and individual fragment statistics are shown in Fig. 3.

4.2.1. Test corpus generation

In order to create the statistical signatures required to accurately map an arbitrary binary object, we needed a corpus of “clean” primitive fragments without any header and footer values. To this end, we chose source files from which we could extract a fragment of a primitive type without extraneous content such as metadata. To reduce the likelihood of accidentally including header or footer information in the primitive fragment samples, we extracted 1 KB samples from the middle of each file and ensured that each file was at least double this size, although most source files were significantly larger. We discuss additional details on test corpus generation in the following paragraphs.

4.2.1.1. Random. There are many potential sources of random (or pseudo random) data. We created our fragments using data made available at Random.org. We chose Random.org because its data is based on atmospheric noise and is understood to be a source of high quality random data.

4.2.1.2. Text. For textual data we selected text-only books from the Project Gutenberg 2006 DVD image (Gutenberg). Because each document included an identical header, we removed these headers in order to generate an independent set of documents. Note that some of these documents included non-English content, but all documents were encoded in US-ASCII.

4.2.1.3. Machine code. Machine code fragments were extracted from both Linux ELF and Windows PE files. ELF files were randomly chosen from a clean install of Ubuntu 9.10 (Karmic Koala), and Windows PE files were gathered from a Windows XP (Service Pack 3) workstation. From each of these files we extracted the .text section, i.e., the section that contains machine code. To confirm our belief that .text sections contained solely machine code, we manually inspected a number of samples. In each instance, ELF files appeared to contain just machine code, however PE files included a table of strings at the end of each .text section, perhaps due to the idiosyncrasies of a given compiler. To

Table 1 – Average (\times) and standard deviation (σ) of Shannon Entropy, Index of Coincidence, Hamming Weight and Arithmetic Mean for 1000 fragments in each category. (Fragment size = 1 KB).

	Mean		Shannon entropy		CHI SQUARE		Hamming weight	
	\times	σ	\times	σ	\times	σ	\times	σ
Random	127.4039	2.3436	9.9826	0.0055	0.4873	0.2968	0.5627	0.0050
Encrypt (AES256/text)	127.4778	2.3122	9.9830	0.0055	0.5008	0.2925	0.5627	0.0052
Compress (bzip2/text)	126.6846	4.2372	9.9802	0.0069	0.2118	0.2480	0.5597	0.0134
Compress (compress/text)	113.7279	8.8724	9.9662	0.0475	0.0681	0.1594	0.5316	0.0149
Compress (deflate/png)	121.7824	12.9482	9.7103	0.7053	0.0460	0.1294	0.5430	0.0444
Compress (LZW/gif/image)	113.7543	8.2331	9.9455	0.0551	0.0203	0.0932	0.5153	0.0265
Compress (mpeg/music)	126.2643	7.2295	9.8747	0.4421	0.0463	0.1260	0.5560	0.0245
compress (jpeg/image)	130.7620	12.7763	9.7314	0.8792	0.0647	0.1555	0.5744	0.0412
Encoded (base64/zip)	84.4643	0.7402	9.7672	0.0192	0.0000	0.0000	0.5306	0.0037
encoded (uuencoded/zip)	63.7171	0.6968	9.7026	0.0209	0.0000	0.0000	0.4991	0.0053
Machine code (linux elf)	116.4212	14.9786	7.6141	0.4381	0.0000	0.0000	0.4940	0.0429
Machine code (windows PE)	107.3952	18.4625	8.0671	0.7279	0.0022	0.0385	0.4819	0.0497
Bitmap	156.4776	69.1200	6.2298	3.6235	0.0000	0.0000	0.6635	0.1905
Text (mixed)	88.5252	7.4828	7.4389	0.2427	0.0000	0.0000	0.5140	0.0146

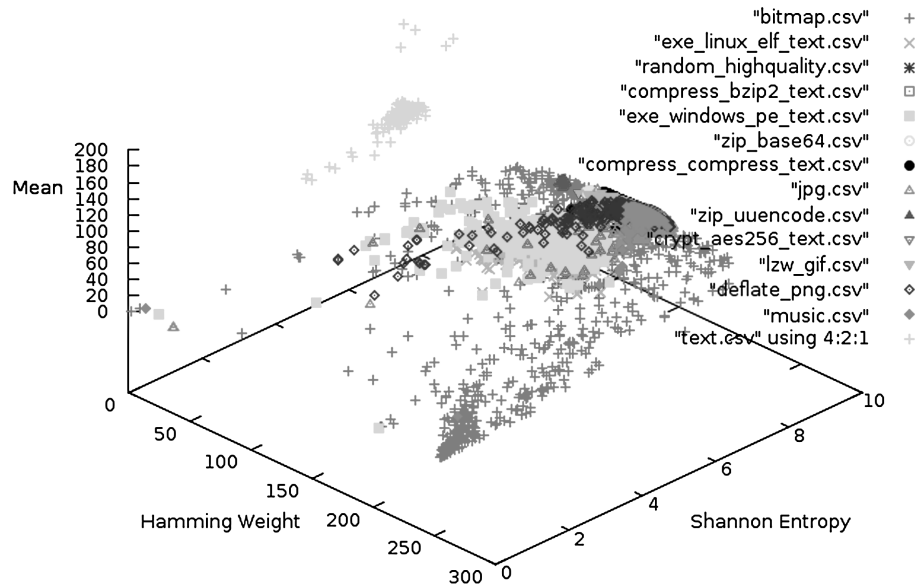


Fig. 3 – Distribution of data fragments by primitive fragment class. Note that the bitmap samples show little clustering, but the high entropy, text, encoded, and machine code primitive types are more densely clustered.

address this issue, we ensured that each extracted text section was at least 2048 bytes, from which we extracted the middle 1024 bytes.

4.2.1.4. Bitmap. Our bitmap fragments were extracted from image files which were created by converting JPEG compressed .jpg files to .bmp format using the batch convert capability of Adobe Photoshop. The original source files were collected from a variety of sources, including a college's internal use photo archive and personal photo archives generated by the researchers. We acknowledge that using these relatively limited image sources may tend to generate statistically similar fragments and suggest seeking to increase diversity for future work. However, as you examine the statistical summary provided in Table 1, note that we were still able to create a set of bitmap fragments with large standard deviations across the statistical tests, implying substantial diversity.

4.2.1.5. Compressed images, audio, and text. For compressed image examples we chose to extract fragments created by differing algorithms including Deflate (.png source files), LZW (.gif source files), and JPEG (.jpg source files). For compressed audio data we extracted fragments compressed using MPEG-1 Audio Layer 3 (.mp3 source files). In addition we created two additional sets of fragments by compressing Project Gutenberg text files. The first fragment class was created using the Linux compress command line utility (Lempel-Ziv coding) and the other using bzip2 (Burrows-Wheeler block sorting and Huffman coding).

4.2.1.6. Encoded. We generated Base64 and Uuencoded samples using the GNU uuencode command line utility and .zip files from the Project Gutenberg 2006 DVD. We chose to encode .zip files instead of .txt files because Base64 and

Uuencoding are commonly used to convert binary data to text for transmission using text based protocols.

4.2.1.7. Encrypted. There are myriad ways to encrypt data. For our experiments we chose to use a popular algorithm (AES) and key length (256 bits) to encrypt Project Gutenberg text files using the Linux mcrypt utility. We made the decision to use AES to provide a reasonable representation of encrypted data for our experiments. However, exploration of differing algorithms, key lengths, and source data (such as compressed data) merits additional investigation.

When dealing with fragments we believe it is important to consider encryption, compression, and encoding *algorithms* to be an appropriate way to categorize such fragments, as it is the combination of the underlying algorithm and the source data it acts upon that generates the unique binary primitive type irrespective of a given file format. In addition, you may note that some primitive types are based on the same source data, e.g., our corpus includes compressed and encrypted fragments based on the same source text documents. Because of this, one might assume a statistical bias toward similarity between primitive types classes. We believe potential similarity is advantageous as we are attempting to differentiate *between* types. A bias toward similarity increases the difficulty of classification, making good classification results harder to obtain. However, we believe this issue is substantially mitigated by the significant transformations that occur when compressing, encrypting, and encoding the data.

5. Fragment classification

To create a classifier we combined the statistical signatures developed in the previous section with the k -nearest neighbor (k -NN) algorithm. k -nearest neighbor is a well understood

classification algorithm which uses distance between a given item and its k -nearest neighbors to perform classification. For example, if $k = 3$ and an unknown item has type1, type2, and type1 as its three nearest neighbors, the unknown item will be classified as type1.

Distance is measured using a distance metric. We evaluated two, Manhattan distance and Euclidean distance. Both metrics scale to N dimensions and perform best when values in each dimension are normalized, which we have done. In our experiments $N = 4$ because we are measuring the distance between fragments based on four statistical dimensions: mean, Shannon Entropy, Chi Square, and Hamming Weight. Ultimately we chose Euclidean distance because our initial analysis found slightly increased accuracy over Manhattan distance. When comparing two fragments, F1 and F2, of four dimensions, Euclidean distance works as follows. For both F1 and F2 we create a feature vector populated by the fragment's normalized mean, Shannon Entropy, Chi Square, and Hamming Weight values as shown in the following example:

F1 = (0.1, 0.7, 0.2, 1.0)
F2 = (0.2, 0.8, 0.6, 0.4)

We then calculate the distance (D) between the fragments by taking the square root of the sum of the squared differences between each pair of values.

$$D(F1, F2) = \sqrt{(0.1-0.2)^2 + (0.7-0.8)^2 + (0.2-0.6)^2 + (1.0-0.4)^2}$$

$$D(F1, F2) = 0.54$$

In our experiments we calculated the distance between each fragment and the other 13,999 fragments in the corpus. We classified a fragment based on the most frequently occurring type of its k -nearest neighbors. Ties were broken by choosing the type with the minimum total distance from the unknown fragment. During experimentation we tested values of $k = 1$ to $k = 25$ and found that values of $k \geq 3$ provided approximately 1% greater overall accuracy, but that values of $k = 4$ to $k = 25$ did not noticeably increase accuracy beyond this point. In order to minimize processing requirements we chose to use $k = 3$. Our classification results are shown in Table 2. Rows indicate the actual type of the fragment and columns indicate the guessed fragment type. An optimal solution would include 1.0 (e.g., 100% accuracy) along the diagonal and 0.0 in all other cells. When examining the matrix, note that confusion is localized within clusters of

Table 2 – Confusion matrix of classification results on a scale of 0 (no matches) to 1 (all matches). Exact classification outcomes are shown on the diagonal. Note the increased accuracy within the outlined regions: high entropy, Uuencoded, Base64 encoded, machine code, text and bitmap. Rows indicate actual fragment type and columns are the guessed fragment type.

	random	encrypt(AES256/text)	compress(bzip2/text)	compress(compress/text)	compress(LZW(gif)/image)	compress(mpeg/audio)	compress(deflate(png)/image)	compress(jpeg/image)	encode(base64/zip)	encode(uuencode/zip)	machine code(linux elf)	machine code(windows PE)	text	bitmap
random	.375	.370	.141	.018	.004	.022	.029	.041	0	0	0	0	0	0
encrypt(AES256/text)	.363	.386	.133	.019	.003	.024	.026	.044	0	0	0	.002	0	0
compress(bzip2/text)	.160	.163	.306	.078	.049	.073	.072	.097	0	0	0	.002	0	0
compress(compress/text)	.022	.030	.072	.588	.176	.040	.035	.031	0	0	0	.002	0	.004
compress(LZW(gif)/image)	.009	.007	.054	.148	.661	.041	.056	.024	0	0	0	0	0	0
compress(mpeg/audio)	.033	.036	.093	.031	.048	.455	.160	.130	0	0	0	0	0	.014
compress(deflate(png)/image)	.030	.037	.081	.027	.061	.177	.424	.101	0	0	.007	.043	0	.012
compress(jpeg/image)	.055	.054	.119	.031	.039	.116	.115	.441	0	0	.006	.009	0	.015
encode(base64/zip)	0	0	0	0	0	0	0	0	1	0	0	0	0	0
encode(uuencode/zip)	0	0	0	0	0	0	0	0	0	1	0	0	0	0
machine code(linux elf)	0	0	0	0	0	0	0	0	0	0	.823	.166	0	.011
machine code(windows PE)	0	.003	.001	.002	.001	0	.020	.002	0	0	.224	.721	.012	.014
text	0	0	0	0	0	0	0	0	0	0	0	.007	.987	.006
bitmap	0	0	0	.008	.002	.020	.034	.032	.006	.007	.024	.030	.012	.825

statistically similar fragment types (outlined in the table) especially in the high entropy (random, encrypted and compressed) and machine code (ELF and PE) categories. The greatest classification error is found in the high entropy group with individual accuracy rates ranging from 0.306 (bzip2/text) to 0.661 (LZW/gif). However, very little confusion occurs outside this group (0.017 on average). Using this insight we can determine accuracy rates for each cluster of similarity:

• Random/Compressed/Encrypted	98.55%
• Base64 Encoded	100%
• Uuencoded	100%
• Machine Code (ELF and PE)	96.7%
• Text	98.7%
• Bitmap	82.5%

When considering how to employ our classifier in binary mapping applications, we could choose to map regions based on the above categories, rather than attempt more precise classification. Without sufficiently high classifier accuracy, a mapping program must cope with the likelihood of frequent misclassifications even when mapping a single homogenous region. However, when a classifier has reasonable accuracy inside a cluster, as is the case for ELF and PE (82.3% and 72.1%, respectively) this information could optionally be passed to the user.

6. System implementation

To test the efficacy of binary mapping we implemented a binary mapping tool, *binmap*, using our classifier and a 1 KB sliding window. We chose a 1 KB window because it exactly matches the fragment size used to generate the statistical signatures previously discussed. A differing window size would increase the likelihood of less accurate classification because our calculations, especially Shannon Entropy and Chi Square, vary significantly based on sample size. *Binmap* takes as input any arbitrary binary object, including, but not limited to, complex data files, executable files, process memory dumps, and hibernation files.

Written in Perl, *binmap* performs mapping by sliding the window incrementally through the file, using a step size of 512 bytes. It starts at offset zero and continues until reaching the end of the file. At each step, we used our classifier to identify the contents of the window. If at the end of the file, there were not a full 1024 bytes remaining, we did not attempt classification. To perform classification, *binmap* compared the statistical signatures of the contents of the window against the previously derived 14,000 binary fragment signatures using the same *k*-NN algorithm (*k* = 3) and Euclidean distance metric described earlier. As the window slid through the file, *binmap* calculated the start and stop offset of each differing type of region in the following categories: random/compressed/encrypted, Base64 encoded, Uuencoded, machine code, text, bitmap, and unknown. Upon encountering a different region, *binmap* would record the current offset and consider the region contiguous until it detected a transition to a different type. At this point, *binmap* would output the start and end offset and the data type. This process would begin again by recording the start offset of the new

region and continue until the next transition. Because an arbitrary binary object would likely contain types that were not previously analyzed, *binmap* classifies a region as unknown if the contents of the window were sufficiently distant from an existing fragment signature. We chose the following distance thresholds as the maximum acceptable distance for each given type: Random/Compressed/Encrypted (0.362), Base64 Encoded(0.005), Uuencoded(0.125), Machine Code(0.740), Text(0.432) and Bitmap(0.190). We derived these values by determining the most distant successful match in each category from our previous experiments and set this value as the maximum acceptable distance.

7. Case study

To explore the effectiveness of *binmap* we used it to analyze two complex binary objects, a 10.3 MB Microsoft Word 2003 document and a 41.4 MB Firefox browser process memory dump. We chose these two objects because they contained a diverse range of primitive types, but were larger than would be practical to analyze using manual inspection alone. We used *binmap* to create a map of each file and then attempted to determine the maps' accuracy by manual inspection. This analysis highlighted important advantages and challenges of binary mapping which we report below.

Because the files were complex, each map contained numerous entries, 353 in the Word document and 20,175 in the process memory dump. Depending on the analyst's goal, very high fidelity maps may be desired. In other cases however, the analyst may be interested in seeking out only major regions. Toward this end, we modified *binmap* to filter out regions that fell below a user selectable minimum size threshold. Upon rerunning *binmap* with a minimum size threshold of 16 KB, the Word document contained 78 map entries and the process memory dump contained 153 entries, much more manageable figures, particularly for the process memory dump. Depending on the type of binary object being examined and analyst goals, we found that thresholds of 4 KB to 32 KB, when combined with the grouping of primitive types we discussed earlier, worked well in highlighting an object's high-level structure. Increasing the minimum size also served to increase accuracy, as this approach required longer sequences of identically classified regions and minimized the impact of misclassification at transitions between disparate primitive types. Another related approach, which we did not explore, is to sort map entries based on size, allowing the analyst to examine the largest regions first. We believe an interactive binary map which is closely integrated with graphical byteplot and hexadecimal views to be another promising approach. We also suggest that researchers explore alternative ways of visualizing mapping data, such as tree-maps or color coding of byteplots, to assist in analysis (Treemap Project Home Page).

Binmap performed strongly when detecting transitions between regions, but classification results were mixed. One of the challenges of creating a classifier in the laboratory and then applying it in real-world use is coping with the wide range of primitive types the classifier was not trained to handle. When *binmap* encountered regions it had been trained to

detect it performed well. Unsurprisingly, it performed less well when encountering a previously unstudied primitive type, even one with a closely related structure. We believe our thresholds for the unknown category to be too strict. Instead of categorizing new regions as unknown, binmap often confused these regions with bitmaps. We believe this is because our set of bitmap fragment signatures appeared quite similar to data structures found in both the word document and the process memory dump. Misclassification also occurred with padding, i.e., regions of identical byte values, such as xFF and x00. Upon reflection, this confusion makes sense because regions of padding would appear identical to all black and all white regions found in our bitmap fragments. One future optimization is to create regular expression signatures that would detect padding, before attempting statistical classification. In general however, we believe that our approach of including a diverse range of examples in each primitive type category was beneficial to classification accuracy. For example, if a few of our machine code examples inadvertently contained some textual strings, these statistical signatures might assist matching less clean examples found in real-world data, particularly at the transition between two different primitive types. Thus the use of many statistical signatures makes binmap resistant to errors in corpus construction. Binmap did perform well when classifying bitmaps of a 24-bit RGB structure, as it was trained to do, but frequently misclassified bitmaps constructed using other encoding schemes. We believe this issue can be overcome by including statistical signatures for other appropriate formats such as those found in programming literature (PixelFormat Enumeration). Bitmap did perform well when identifying US-ASCII encoded text, but ran into difficulties when classifying 16-bit Basic Latin Unicode and HTML encoded text. Again, we believe this can be addressed by explicitly creating representative statistical signatures for these new primitive types. Binmap also performed well when classifying random/compressed/encrypted regions and was able to detect transitions between sequences of embedded compressed images.

The common trend in binmap's classification accuracy is that it performed well when classifying primitive types it was trained for, but often confused similar but distinct primitive types, such as US-ASCII, Basic Latin Unicode and HTML. We believe a reasonable solution is to increase the diversity of the primitive type signatures employed in classification to incorporate a wider range of possibilities and we then believe our approach will work well with other types of fragments. Techniques such as weighting, clustering, Markov Chains and simulated annealing should also be explored as ways to improve classification. As we mentioned earlier in the paper binary mapping is not wedded to any one classification technique. We believe a more desirable classification solution is one that includes regular expression pattern matching and statistical signature classification combined with knowledge of file format semantics, if applicable in the given context. Researchers should also consider techniques that operate at the individual bit level and do not assume byte-level encoding schemes. Two desirable goals of classification in support of binary mapping is to precisely identify transitions between differing regions and to classify primitive types with the smallest sample size possible.

8. Conclusions and future work

In this paper we made several contributions: a statistical analysis of 14,000 low-level binary fragments, the presentation and evaluation of a classifier for identifying 14 primitive binary fragment types, and the implementation of a binary mapping application. The statistical techniques we employed included Shannon Entropy, Hamming Weight, mean, and Chi Squared, which proved useful for analyzing low-level binary data without relying on possibly non-existent or untrustworthy metadata. The classifier used the *k*-nearest neighbor algorithm and Euclidean distance to identify unknown fragments and accurately placed them into the following groups: Random/Compressed/Encrypted, Base64 Encoded, Uuencoded, Machine Code, Text, and Bitmap. We then developed a binary mapping implementation and used it to study two complex binary objects.

Binary mapping is valuable to analysts because it provides rapid insight into the structure of large, potentially massive, binary objects by identifying the location of the varying types of data they contain. Mapping can be employed in such tasks as memory forensics and the analysis of undocumented file types, but bears promise for most tasks requiring low-level analysis of binary data. Binary mapping may also be employed in interactive applications, such as hex editors, visual analytic tools and GUI-based forensic toolkits. The resultant mapping data can be used to highlight areas of interest for the analyst, allow rapid navigation to these regions, and facilitate filtering, ultimately allowing the analyst to cope with very large binary objects at a level exceeding the capabilities of today's low-level analysis tools. In addition, binary mapping generalizes well and can be used with a different classifier or combination of classifiers to increase accuracy.

For future work we recommend studying the combination of statistical classification techniques and traditional signature matching, perhaps using a plug-in architecture, to increase the accuracy of classification. Also, we believe that extending the breadth of the classification capability to include a wider array of primitive data types will increase the utility of binary mapping. Finally, we recommend seeking to counter ways in which an adversary might obfuscate data to force incorrect classification.

REFERENCES

- Calhoun W, Coles D. Predicting the types of file fragments. In: Digital forensics research conference; 2008.
- Conti G, Dean E. Visual forensic analysis and reverse engineering of binary data. Black Hat USA; 2008.
- Conti G, Dean E, Sinda M, Sangster B. Visual reverse engineering of binary and data files. In: Workshop on visualization for computer security; 2008.
- Erbacher R, Mulholland J. Identification and localization of data types within large-scale file systems. In: Systematic approaches to digital forensic engineering; 2007.
- Friedman William. The index of coincidence and its applications in cryptology. Department of Ciphers. Publication 22. Geneva, Illinois, USA: Riverbank Laboratories; 1922.
- Garfinkel S. Carving contiguous and fragmented files with fast object validation. In: Digital forensics research conference; 2007.

- Gutenberg. The CD and DVD project. Project Gutenberg, http://www.gutenberg.org/wiki/Gutenberg:The_CD_and_DVD_Project.
- Hall G, Davis W. Sliding window measurement for file type identification. In: White paper, ManTech cyber solutions international.
- Hickok D, Lesniak D, Rowe M. File type detection technology. In: Midwest instruction and computing symposium; 2005.
- Karresand M, Shahmehri N. File type identification of data fragments by their binary structure. In: IEEE information assurance workshop; 2006.
- Li W, Wang K, Stolfo S, Herzog B. Fileprints: identifying file types by *n*-gram analysis. In: IEEE information assurance workshop; 2005.
- McDaniel M, Heydari M. Content based file type detection algorithms. In: Hawaii international conference on system sciences; 2003.
- Moody S, Erbacher R. SÁDI – statistical analysis for data type identification. In: Systematic approaches to digital forensic engineering; 2008.
- PixelFormat Enumeration. Net framework developer center. Microsoft developer network, <http://msdn.microsoft.com/en-us/library/system.drawing.imaging.pixelformat.aspx>.
- Richard G, Roussev V. Scalpel: a frugal, high-performance file carver. In: Digital forensics research workshop; 2005.
- Roussev V, Garfinkel S. File fragment classification – the case for specialized approaches. In: Systematic approaches to digital forensics engineering; 2009.
- Shamir A, van Someren N. Playing hide and seek with stored keys. In: International conference on financial cryptography; 1999.
- Shannon C. A mathematical theory of communication. Bell System Technical Journal July, October, 1948;27:379–423. 623–656.
- Stolfo S, Wang K, Li W. Fileprint analysis for Malware detection. In: Workshop on rapid Malcode; 2005.
- Treemap Project Home Page. Human computer interaction lab, <http://www.cs.umd.edu/hcil/treemap/>.
- Veenman C. Statistical disk cluster classification for file carving. In: Symposium on information assurance and security; 2007.
- Gregory Conti** is an Academy Professor and Director of West Point's Cyber Security Research Center. His research includes online privacy, web-based information disclosure, security data visualization, and usable security. He is the author of Security Data Visualization (No Starch Press) and Googling Security (Addison–Wesley). His work can be found at www.gregconti.com and www.rumint.org.
- Sergey Bratus** is a Research Assistant Professor at Dartmouth College, affiliated with the Institute for Security, Technology, and Society (ISTS). Dr. Bratus is interested in a broad range of practical operating systems and network security topics, and frequents hacker conferences.
- Anna Shubina** is a Postdoctorate Fellow at the Institute of Security, Technology, and Society at Dartmouth College. Her primary interest and the subject of her doctoral thesis is Privacy in modern information systems.
- Ben Sangster** is an Assistant Professor of Computer Science at the United States Military Academy, West Point, NY. His research includes binary object identification in support of information assurance, behavior-based information security, and virtualization of the computer science curriculum.
- Roy Ragsdale** is a recent graduate of the United States Military Academy where he majored in computer science. He is currently serving as a commissioned officer in the United States Army Military Intelligence Corps.
- Matthew Supan** is a recent graduate of the United States Military Academy where he majored in computer science. He is currently serving as a commissioned officer in the United States Army Signal Corps.
- Andrew Lichtenberg** is currently a student at Skidmore College where he is a dual major in Mathematics and Computer Science.
- Rober Perez-Aleman** is currently a cadet at the United States Military Academy where he is majoring in Mathematics.