



Robust Linux Memory Acquisition with Minimal Target Impact

By

Johannes Stuettgen and Michael Cohen

Presented At

The Digital Forensic Research Conference

DFRWS 2014 EU Amsterdam, NL (May 7th - 9th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment. As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

<http://dfrws.org>



LMAP

Robust Linux Memory Acquisition
with Minimal Target Impact

Johannes Stuettgen and Michael Cohen



Outline

- Problem statement
- Anatomy of a linux kernel module
 - How is a LKM loaded and linked
 - Why does it need to be compiled for a specific kernel config and version
- Existing techniques to load incompatible modules
- Technical details of a new, more stable approach
- Demo
- Challenges and Future work



Problem statement

- Current linux memory acquisition tools need to be compiled to the exact kernel version they need to run on.
 - This is difficult for incident response since we don't always know which kernel we will be responding to.
 - Some solutions pre-build a large library of commonly used standard distribution kernels.
 - But this does not solve the problem of customized or hand built kernels.
- The last thing you want to worry about during an incident is whether you have the exact kernel version in your library!

Product features (Second Look Incident Response Edition)

- Memory acquisition and analysis for all 32- and 64-bit x86 (i386/i486/i586/i686, amd64/x86_64) Linux systems running 2.6- and 3-series kernels. This includes:
 - Debian 4, 5, 6, and higher;
 - Fedora 2 through 19, and higher;
 - Red Hat Enterprise Linux (RHEL) and
 - Ubuntu 4.10 through 13.10, and higher
 - and other distributions.

I wonder if this will fit my incident?



What is a Linux Kernel Module?

- Not an executable!
- Relocatable ELF object. Basically the same as an .o file
- Loading a kernel module is actually more similar to **static linking** of an object file - not dynamic linking.

/bin/ls

```
Class: ELF64  
Data: 2's complement, little endian  
Version: 1 (current)  
OS/ABI: UNIX - System V  
ABI Version: 0  
Type: EXEC (Executable file)  
Machine: Advanced Micro Devices X86-64  
Version:  
Entry point address: 0x1  
Start of program headers: 0x4045a4  
Start of section headers: 64 (bytes into file)  
Start of section headers: 104048 (bytes into file)
```

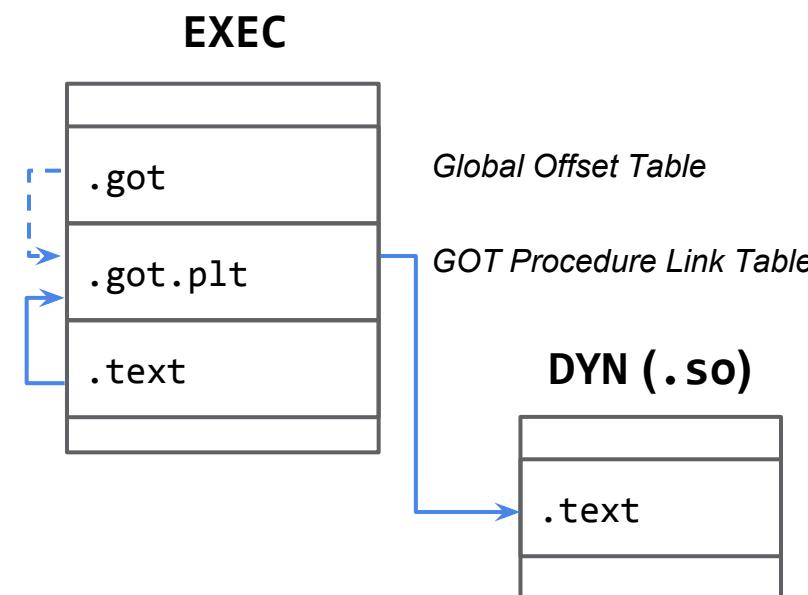
pmem.ko

```
Class: ELF64  
Data: 2's complement, little endian  
Version: 1 (current)  
OS/ABI: UNIX - System V  
ABI Version: 0  
Type: REL (Relocatable file)  
Machine: Advanced Micro Devices X86-64  
Version:  
Entry point address: 0x1  
Start of program headers: 0x0  
Start of section headers: 0 (bytes into file)  
Start of section headers: 11376 (bytes into file)
```

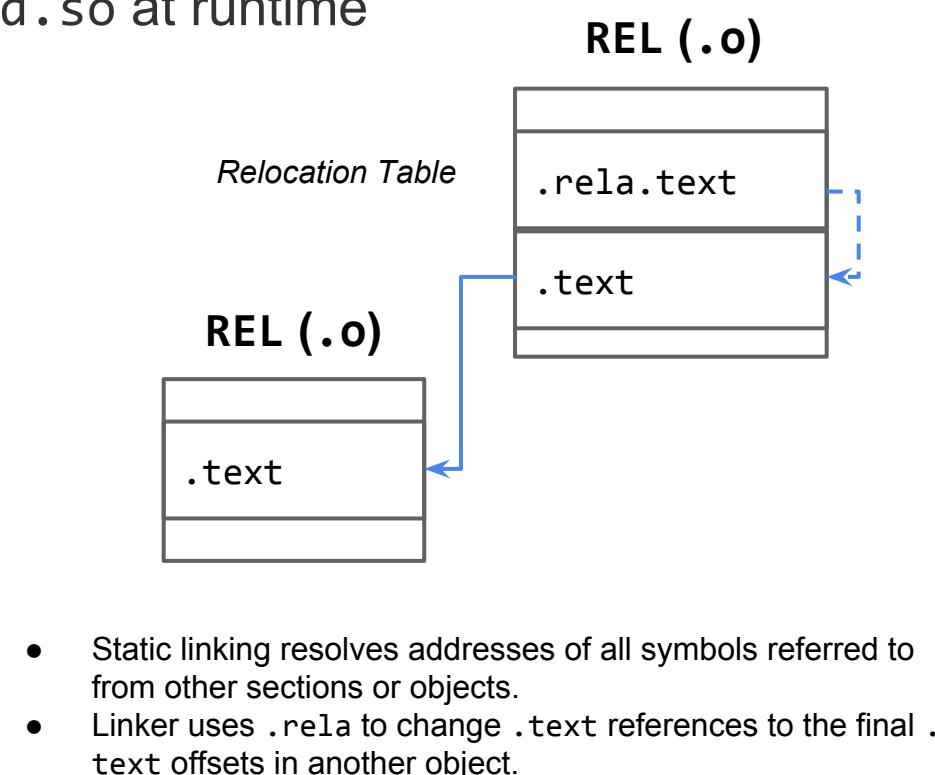


Dynamic vs. Static Linking

- Static linking usually performed when building an executable
- Dynamic linking performed by ld.so at runtime



- .got holds offsets of symbols resolved by linker at runtime.
- Jumps into the .got.plt use .got to jump to referenced symbol



- Static linking resolves addresses of all symbols referred to from other sections or objects.
- Linker uses .rela to change .text references to the final .text offsets in another object.



Static Linking

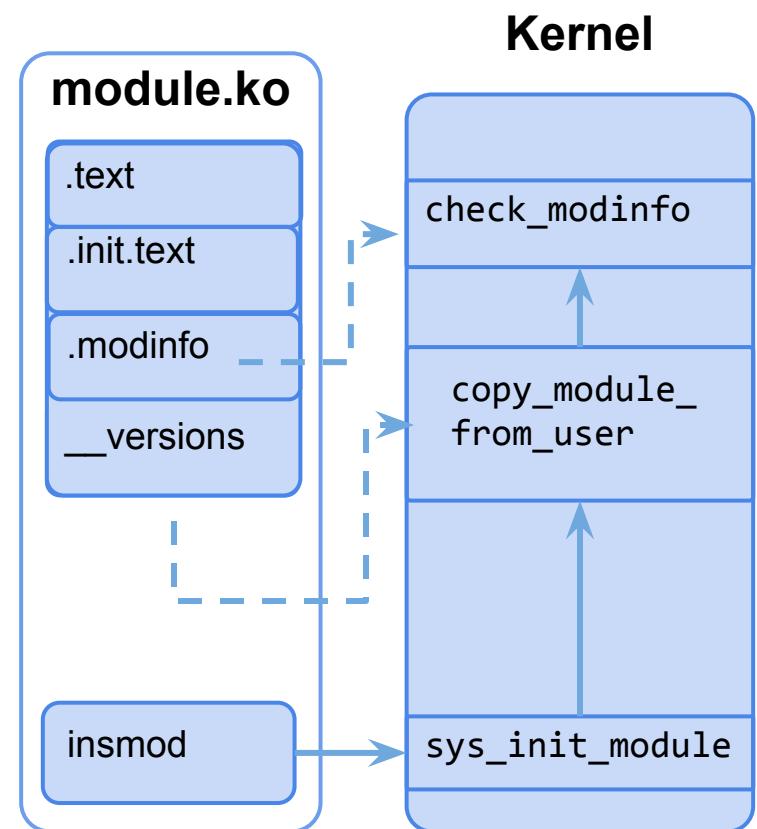
- Object contains list of relocation entries
- Linker iterates over relocation entries
 - Resolves relocated symbol
 - Patches object directly with address of symbol
- Relocations also used for data references

Relocation section '.rela.text' at offset 0x2c48 contains 134 entries:					
Offset	Info	Type	Sym.	Value	Sym. Name + Addend
000000000005	004100000002	R_X86_64_PC32		0000000000000000	_cond_resched - 4
000000000005e	004800000002	R_X86_64_PC32		0000000000000000	__register_chrdev - 4
000000000008a	004d00000002	R_X86_64_PC32		0000000000000000	__unregister_chrdev - 4
000000000015b	003800000002	R_X86_64_PC32		0000000000000000	pte_mmap - 4
0000000000166	003800000002	R_X86_64_PC32		0000000000000000	pte_mmap - 4
0000000000188	00050000000b	R_X86_64_32S		0000000000000000	.rodata + 18



Loading an LKM in Linux

1. insmod called from user-space
 - prepares module image
 - performs `init_module` system call
2. Kernel allocates memory and copies the module in
3. Kernel verifies module is compiled with identical config and version:
 - Checks module compatible kernel version in `.modinfo`
 - Verifies each imported symbol's version in `_versions` if `CONFIG_MODVERSIONS` is enabled





Loading an LKM in Linux continued

4. Kernel resolves relocations by directly patching references to symbols
5. Kernel calls `init_module` function in the module

module.ko

```
Relocation section '.rela.gnu.linkonce.this_module'
  Offset      Type      Sym. Name + Addend
000000000150 R_X86_64_64 init_module + 0
```

```
struct module __this_module
__attribute__((section(".gnu.linkonce.this_module")))
{
    .name = KBUILD_MODNAME,
    .init = init_module,
    ...
};
```

Kernel

```
int apply_relocate_add(...) {
    for (i = 0; ... ; i++) {
        ...
        val = sym->st_value + rel[i].r_addend;
        ...
        switch (ELF64_R_TYPE(rel[i].r_info)) {
            ...
            case R_X86_64_64:
                *(u64 *)loc = val;
            ...
        }
    }
}
```

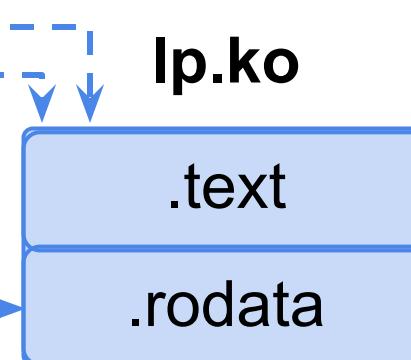
```
/* This is where the real work happens */
static int do_init_module(struct module *mod) {
    ...
    /* Start the module */
    if (mod->init != NULL)
        ret = do_one_initcall(mod->init);
    ...
    /* Now it's a first class citizen! */
    mod->state = MODULE_STATE_LIVE;
    ...
}
```



How do LKMs communicate with user-space?

- Typical LKMs use file operations to communicate with user-space:
 - Create a `file_operations` struct with function pointers to implemented functions registered with this device.
 - Register a device inode under `/dev/`

```
static const struct file_operations  
lp_fops = {  
    .owner      = THIS_MODULE,  
    .read       = lp_read,  -----  
    .llseek     = noop_llseek, -----  
};
```



Modules have a static fops struct in the `.rodata` segment containing function pointers to `.text`.



How to load LKM into different kernel?

1. The kernel config option "CONFIG_MODULE_FORCE_LOAD" allows modules without valid vermagic to be loaded.
 - a. In practice few kernel enable this because it is quite dangerous.
 - b. "*Forced module loading [...] is usually a really bad idea.*"
2. It is possible to trick the kernel into loading arbitrary modules by patching the .modinfo and __versions sections
 - a. "Steal" valid signatures from other existing modules.
 - b. Copy them into arbitrary module to make the version check pass

Contents of section .modinfo:

```
0000 6c696365 6e73653d 47504c00 7665726d  
0010 61676963 3d332e38 2e302d33 332d6765  
0020 6e657269 6320534d 50206d6f 645f756e  
0030 6c6f6164 206d6f64 76657273 696f6e73  
0040 20006465 70656e64 733d0073 72637665  
0050 7273696f 6e3d3631 32323832 38414436  
0060 32414537 43413544 36364241 3100
```

```
license=GPL.verb  
agic=3.8.0-33-ge  
neric SMP mod_un  
load modversions  
.depends=.srcve  
rsion=6122828AD6  
2AE7CA5D66BA1.
```



So we have execution, now what?

- By cheating the version check, we can get execution in the `init_module` code. However, this is not enough
- We still depend on the struct layout of the kernel which we used to compile the module.

```
/* This is where the real work happens */
static int do_init_module(struct module *mod)
{
    ...
    /* Start the module */
    if (mod->init != NULL)
        ret = do_one_initcall(mod->init);
    ...
    /* Now it's a first class citizen! */
    mod->state = MODULE_STATE_LIVE;
    ...
}
```

```
struct module {
    enum module_state state;
    struct list_head list;
    char name[MODULE_NAME_LEN];
    ...
#ifndef CONFIG_UNUSED_SYMBOLS
    ...
#endif
#ifndef CONFIG_MODULE_SIG
    bool sig_ok;
#endif
    ...
    /* Startup function. */
    int (*init)(void);
    ...
}
```



What actually needs to be done

- We need to be able to insert arbitrary code into the running kernel
 - Pass version check.
 - Get code execution through a valid `struct module->init`
- We need to predict the layout of the kernel's data structures
 - We definitely need `struct module`
 - We also need `struct file_operations` for communication with user-space.
 - Finally, we need other API's structs as needed by the acquisition tool
 - The more complex work we do in the driver, the more APIs we will use which makes this harder (e.g. copy to network, copy to disk etc).
- Ideally we want to use as few APIs as possible in the driver!
 - No disk I/O in the driver
 - No usage of memory mapping and enumeration API's
 - No memory allocation!



LKM infection to the rescue

- An LKM is a relocatable ELF and can be linked to other objects
- Easy to add code and data to it
 - ‘`ld -r parasite.ko host.ko`’
- Known and stable rootkit persistence technique
 - Phrack #61 (2003) File 0x0a by Truff: “*Infecting loadable kernel modules*”
 - Phrack #68 (2012) File 0x0b by Styx: “*Infecting loadable kernel modules kernel versions 2.6.x/3.0.x*”
 - Used by the Adore-ng rootkit (and probably many others)
- Several ways to divert execution flow to injected code
 - Most commonly done by renaming symbols
- We’re more interested in the structs a host module has to offer
 - Contains valid struct module for running kernel
 - Has relocations that “know” about struct layout



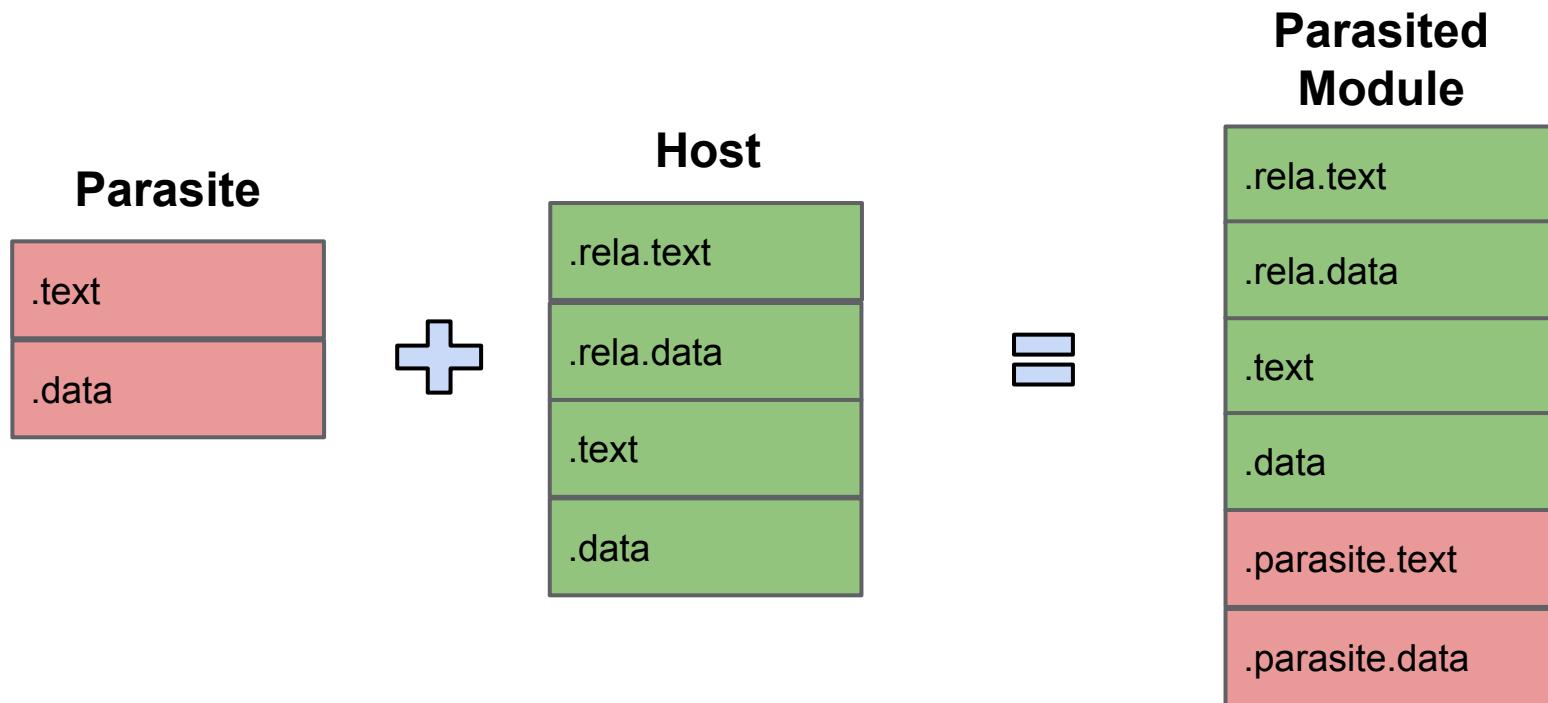
Selection of Host Module

- Scan `/lib/modules/`uname -r`/` for suitable modules.
- What makes a module suitable?
 - Needs to contain the following symbols:
 - `struct file operations`
 - `read` and `llseek` methods
 - Has to import `register_chrdev` to register it's fops struct
 - legacy method for registering character devices
 - More widely supported on older kernels than `cdev_init`
 - needs to have relocations from `read` and `seek` functions into
`static struct file_operations`
- In practice there are many suitable modules on most systems:
 - `lp.ko`, `irnet.ko` etc.



Inject our code into the module

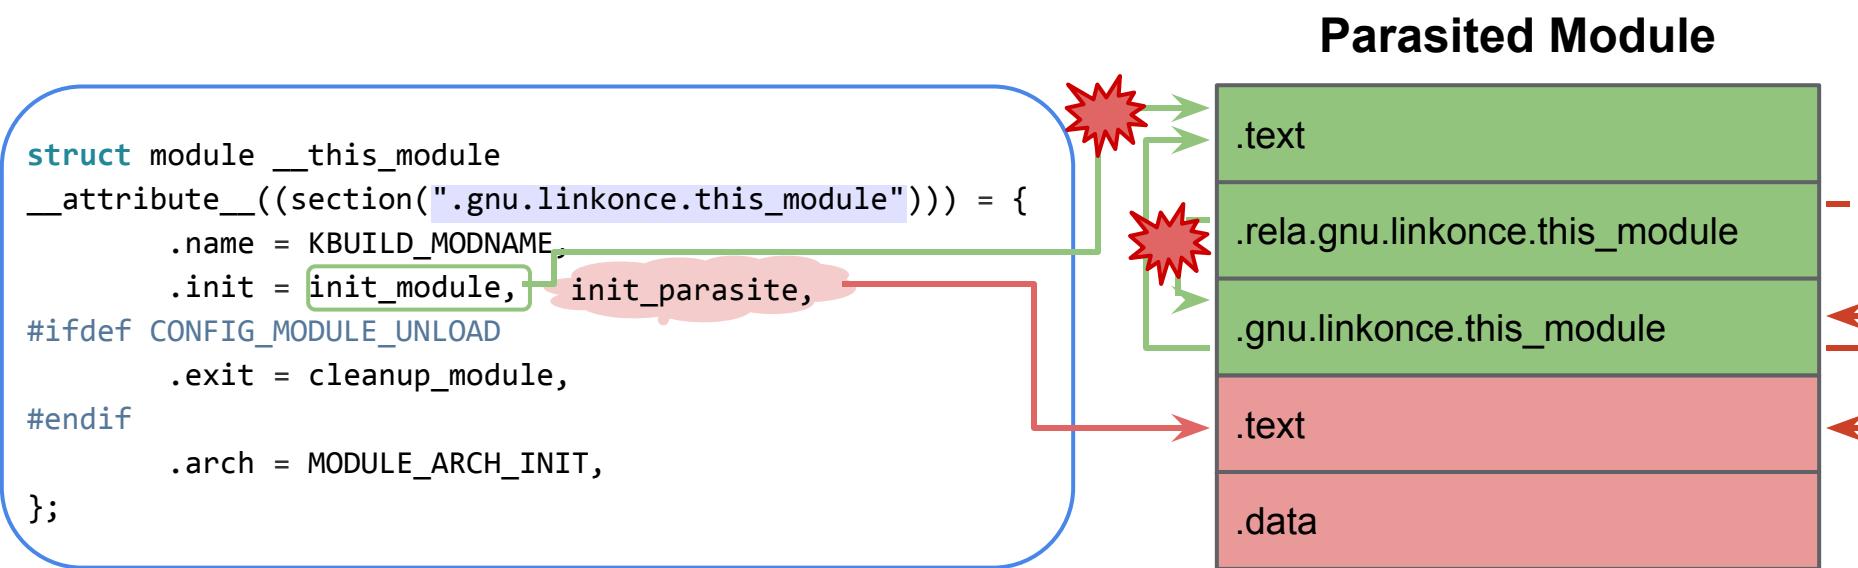
- Copy .text and .data sections from the parasite into the host





Redirect execution to parasite code

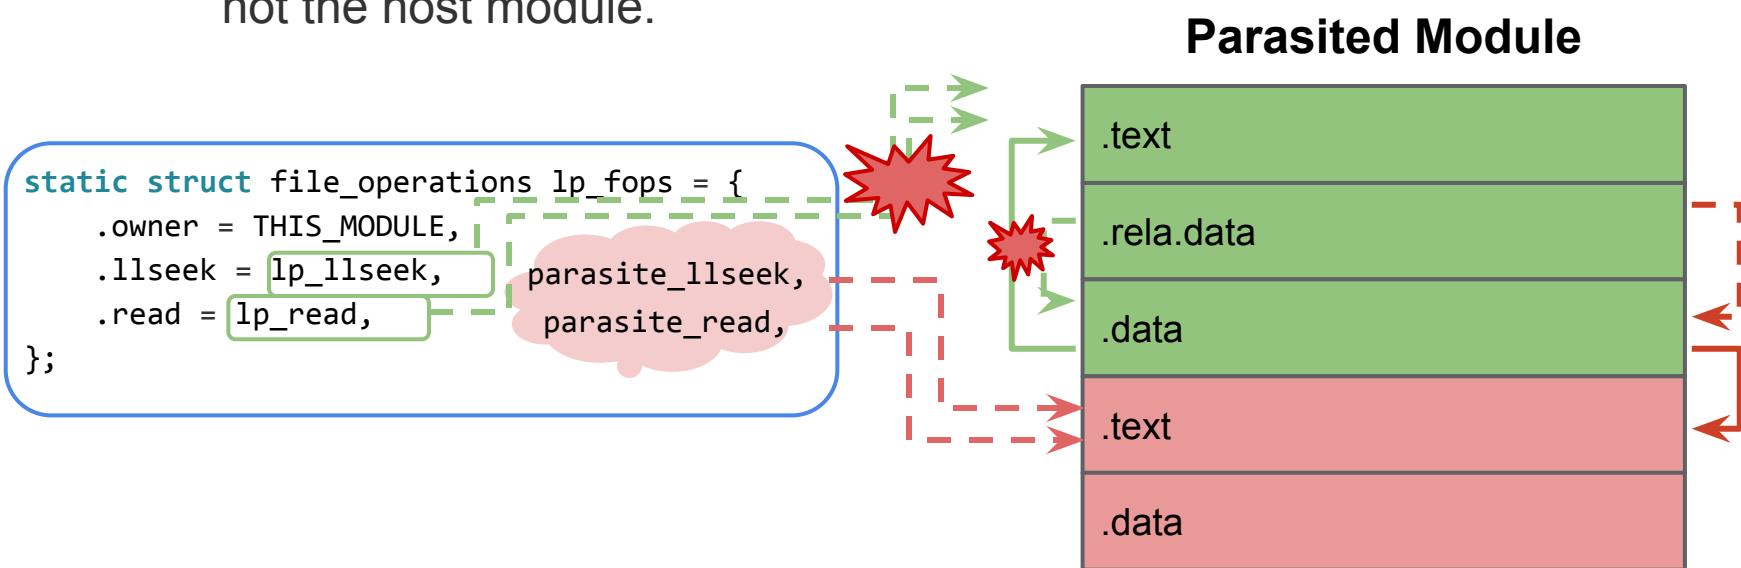
- Fix up relocations to put the location of our code into `_this_module->init`, instead of the location of `module_init`.
- We don't need to know anything about the layout of struct module, the kernel linker will handle it for us





Redirect File Operations

- Fix up the relocations to point at our code for struct file_operations.
 - All read and seek operations on the device file call the parasite now, not the host module.





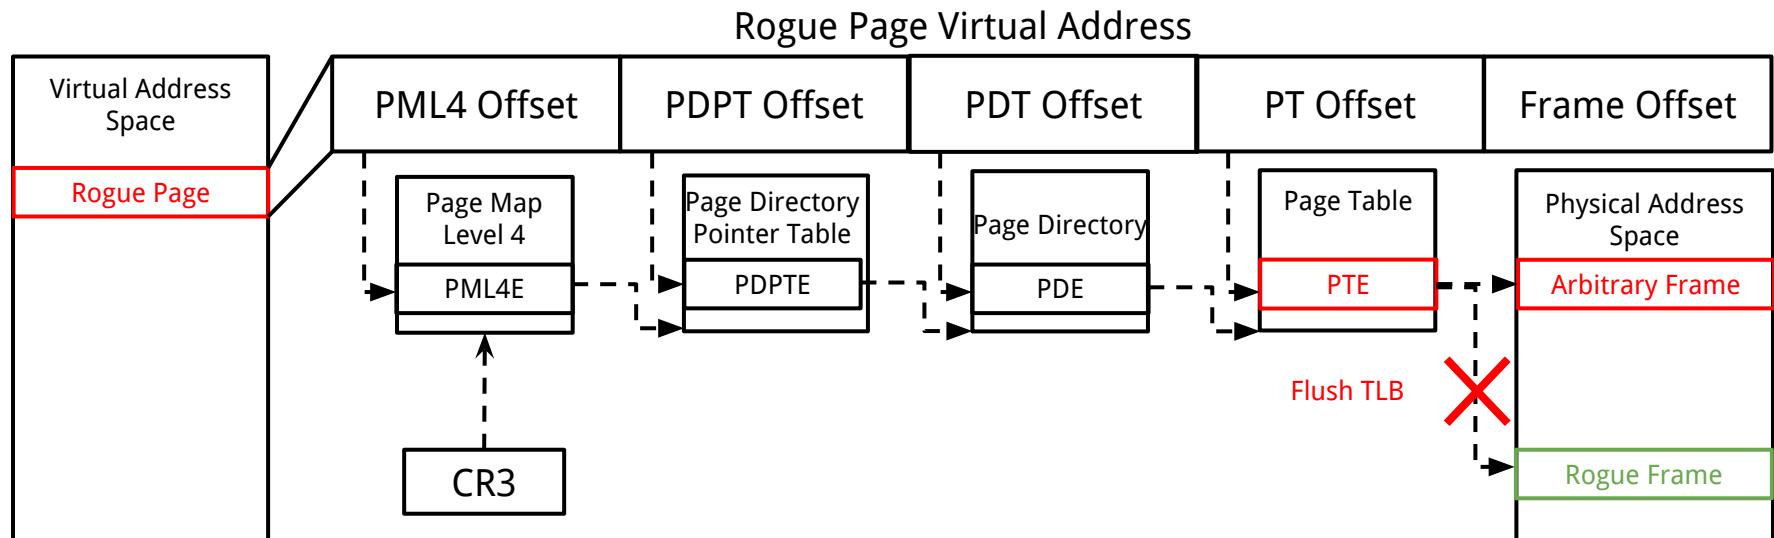
API Independent Acquisition Code.

- Even after getting code execution we need to ensure we are not using the wrong APIs:
 - Internal kernel struct layouts could change between versions.
 - Kernel APIs can change between versions.
- If we declare a dependency on an exported symbol we need to also declare a version magic for it.
 - This means we need to search existing modules for this symbol.
(To copy the magic hash from their `__versions` section.)
 - It may not exist on the running system.
- Ideally a memory acquisition module should not depend on any kernel APIs at all!



PTE based acquisition methods.

- Manually parse and manipulate the page tables.
 - No dependencies on mapping API's.
- Remap page in data segment to any physical address as needed.
 - No need for memory allocation.
- Get data to user-space via read calls on device inode.
 - User supplied buffer, no memory allocation in driver.





Demo

- Minimal version of pmem_pte
- Packaged with a statically linked loader (LMAP)
- Build on an Ubuntu 12.10 system with kernel 3.5.0-41-generic
- Dynamically finds, injects and patches a suitable host
- In this example a Debian 7.0 with Kernel 3.2.0-4-amd64

```
root@victim:~# ./lmap -a dump
[+] Scanning modules in /lib/modules/3.2.0-4-amd64 for suitable host
[+] Found suitable host /lib/modules/3.2.0-4-amd64/kernel/drivers/char/lp.ko
[+] Using underscore prefixes and noop_llseek (kernel 3.x) relocation hooking method (2)
[+] Successfully injected parasite into /lib/modules/3.2.0-4-amd64/kernel/drivers/char/lp.ko
[+] Injected pmem module has been loaded
[+] Starting to dump memory
[+] [0000000000010000 - 000000000009efff] [WRITTEN]
[+] [0000000000100000 - 000000007ffeffff] [WRITTEN]
[+] Acquired 524159 pages (2146955264 bytes)
[+] Size of accessible physical address space: 2147418112 bytes (2 segments)
[+] Successfully wrote elf image of memory to dump
```



Challenges

- Some APIs and interfaces change over time
 - Character devices don't use `llseek` in `struct file_operations` on kernels < 2.6.37.
 - Breaks random access on these systems.
 - Kernel symbol names change sometimes
 - `copy_to_user` vs `_copy_to_user`
- Code stability is a big issue
 - Crashes have fatal consequences
- 32-Bit support
 - Different relocations and data types, should be pretty straight forward.



Future Work

- Android Memory Acquisition.
 - Cellphone vendors often don't release their kernel sources.
 - Adapt existing implementation to ARM style relocations.
- Relocations contain implicit information about struct layout.
 - This is also very interesting for analysis.
 - Can derive config settings from relocation offsets.
 - Build a partial Volatility profile from this.

000000000080	004300000001 R_X86_64_64	0000000000000000	__this_module
000000000088	005700000001 R_X86_64_64	0000000000000000	noop_llseek
000000000090	000100000001 R_X86_64_64	0000000000000000	lp_read
000000000098	000100000001 R_X86_64_64	0000000000000000	lp_write
0000000000c0	000100000001 R_X86_64_64	0000000000000000	lp_ioctl
0000000000c8	000100000001 R_X86_64_64	0000000000000000	lp_compat_ioctl
0000000000d8	000100000001 R_X86_64_64	0000000000000000	lp_open
0000000000e8	000100000001 R_X86_64_64	0000000000000000	lp_release



The End

Slides: <http://goo.gl/4hbPCJ>

Code: <http://rekall-forensic.com>

