Image-Based Kernel Fingerprinting

By

Vassil Roussev, Irfan Ahmed and Thomas Sires

Presented At

The Digital Forensic Research Conference

DFRWS 2014 USA Denver, CO (Aug 3rd - 6th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment. As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

http:/dfrws.org

dfrws'14: aug 3-7, 2014

image-based kernel fingerprinting

Vassil Roussev, Irfan Ahmed, Thomas Sires

University of New Orleans

vassil@roussev.net

dfrws'14: aug 3-7, 2014

image-based kernel fingerprinting

a 'MALKOVICH, MALKOVICH' story

VASSIL ROUSSEV, IRFAN AHMED, THOMAS SIRES

UNIVERSITY OF NEW ORLEANS

VASSIL@ROUSSEV.NET

why are we here?

GIVEN A RAM DUMP, OR VM SNAPSHOT

- # find the exact kernel version
  - » so that we can do proper memory analysis

WINDOWS → NOT A PROBLEM

LINUX?

- # thousands of 'stock' kernels (ubuntu, red hat, …)
- # custom kernels
- # different architectures: x86, amd64, arm5/6/…

BOTTOM LINE:

- # we need something robust that requires no reversing

first thought:
if you have a hammer → find nails

quick solution sketch

SDHASH CAN CORRELATE ANY TWO BLOBS BASED ON COMMONALITY:

```
# sdhash vmlinux-* > kernels.sdbf
# sdhash server.vmem > ram.sdbf
# sdhash -c kernels.sdbf ram.sdbf
```

PROBLEM SOLVED?

how is identifying a ***linux*** kernel a MALKOVICH, MALKOVICH problem?

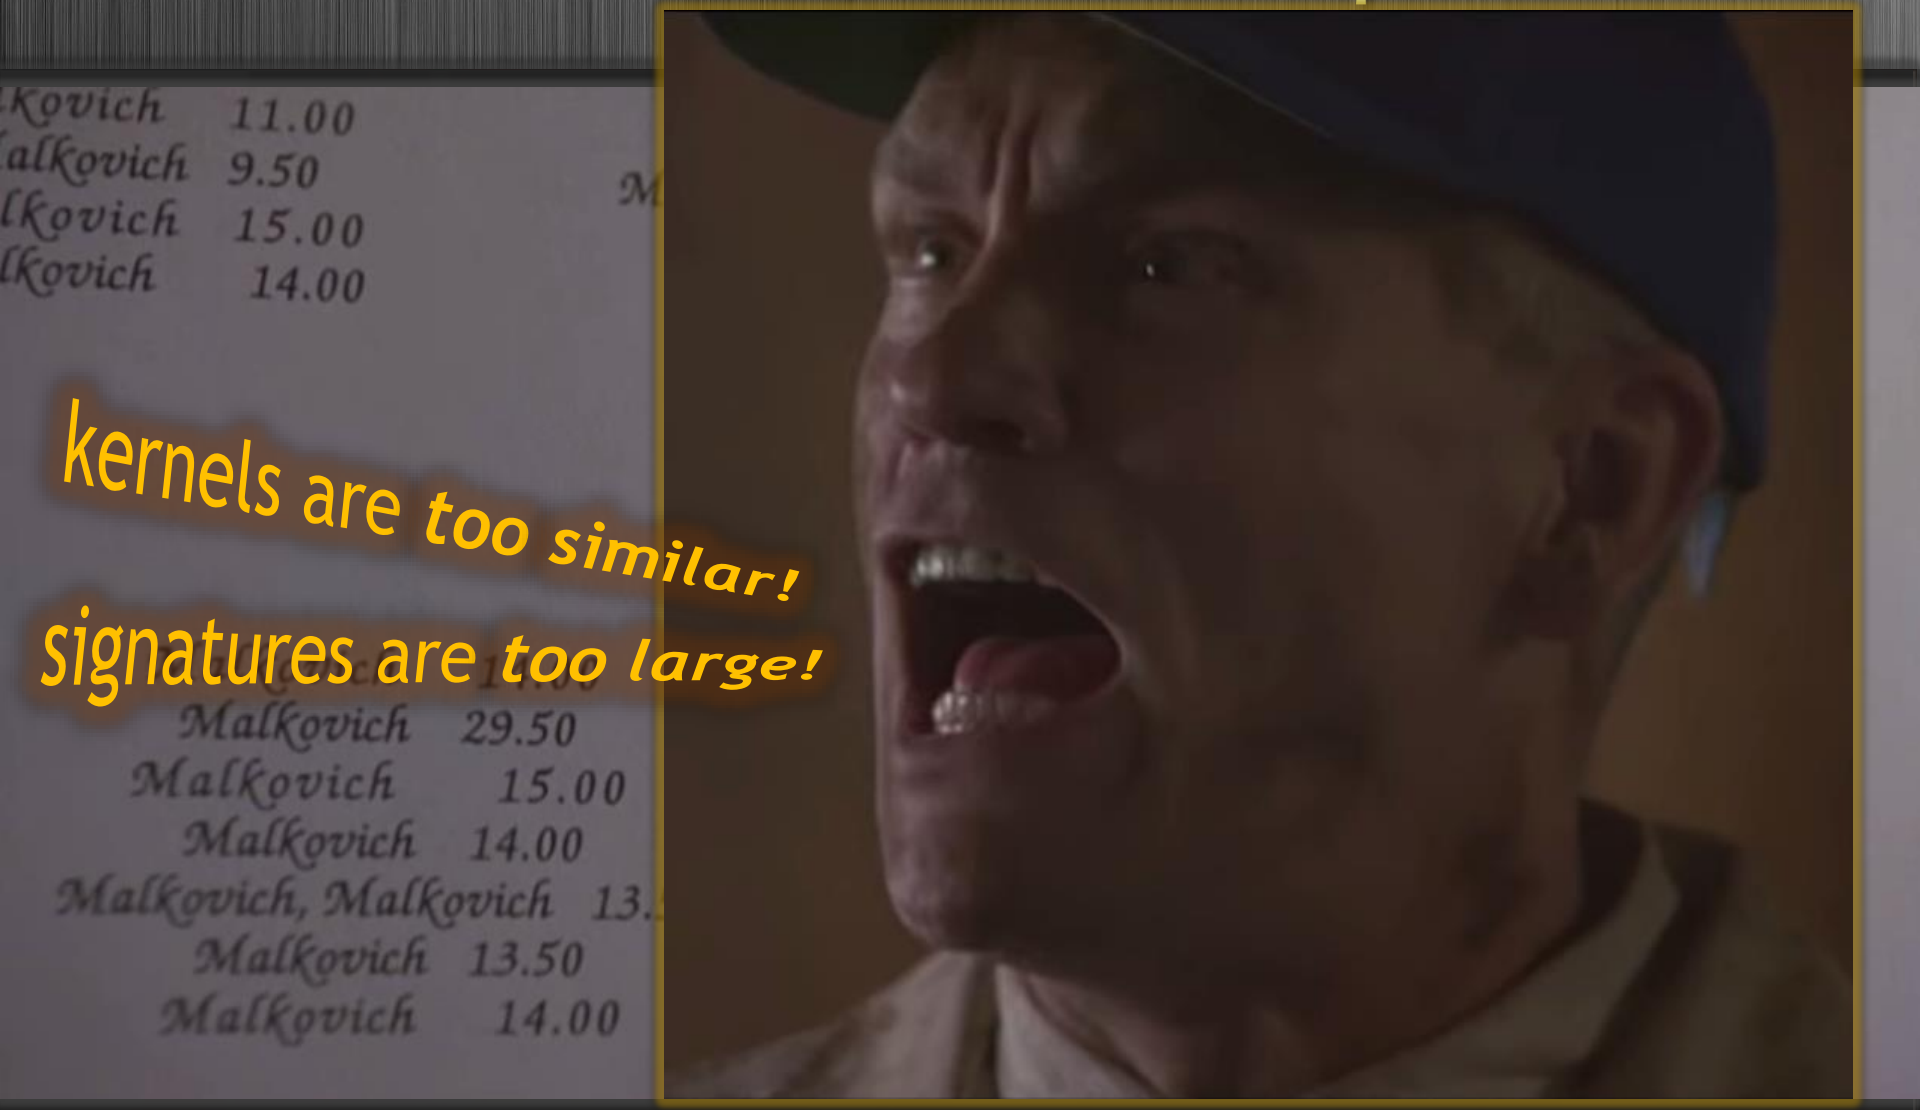

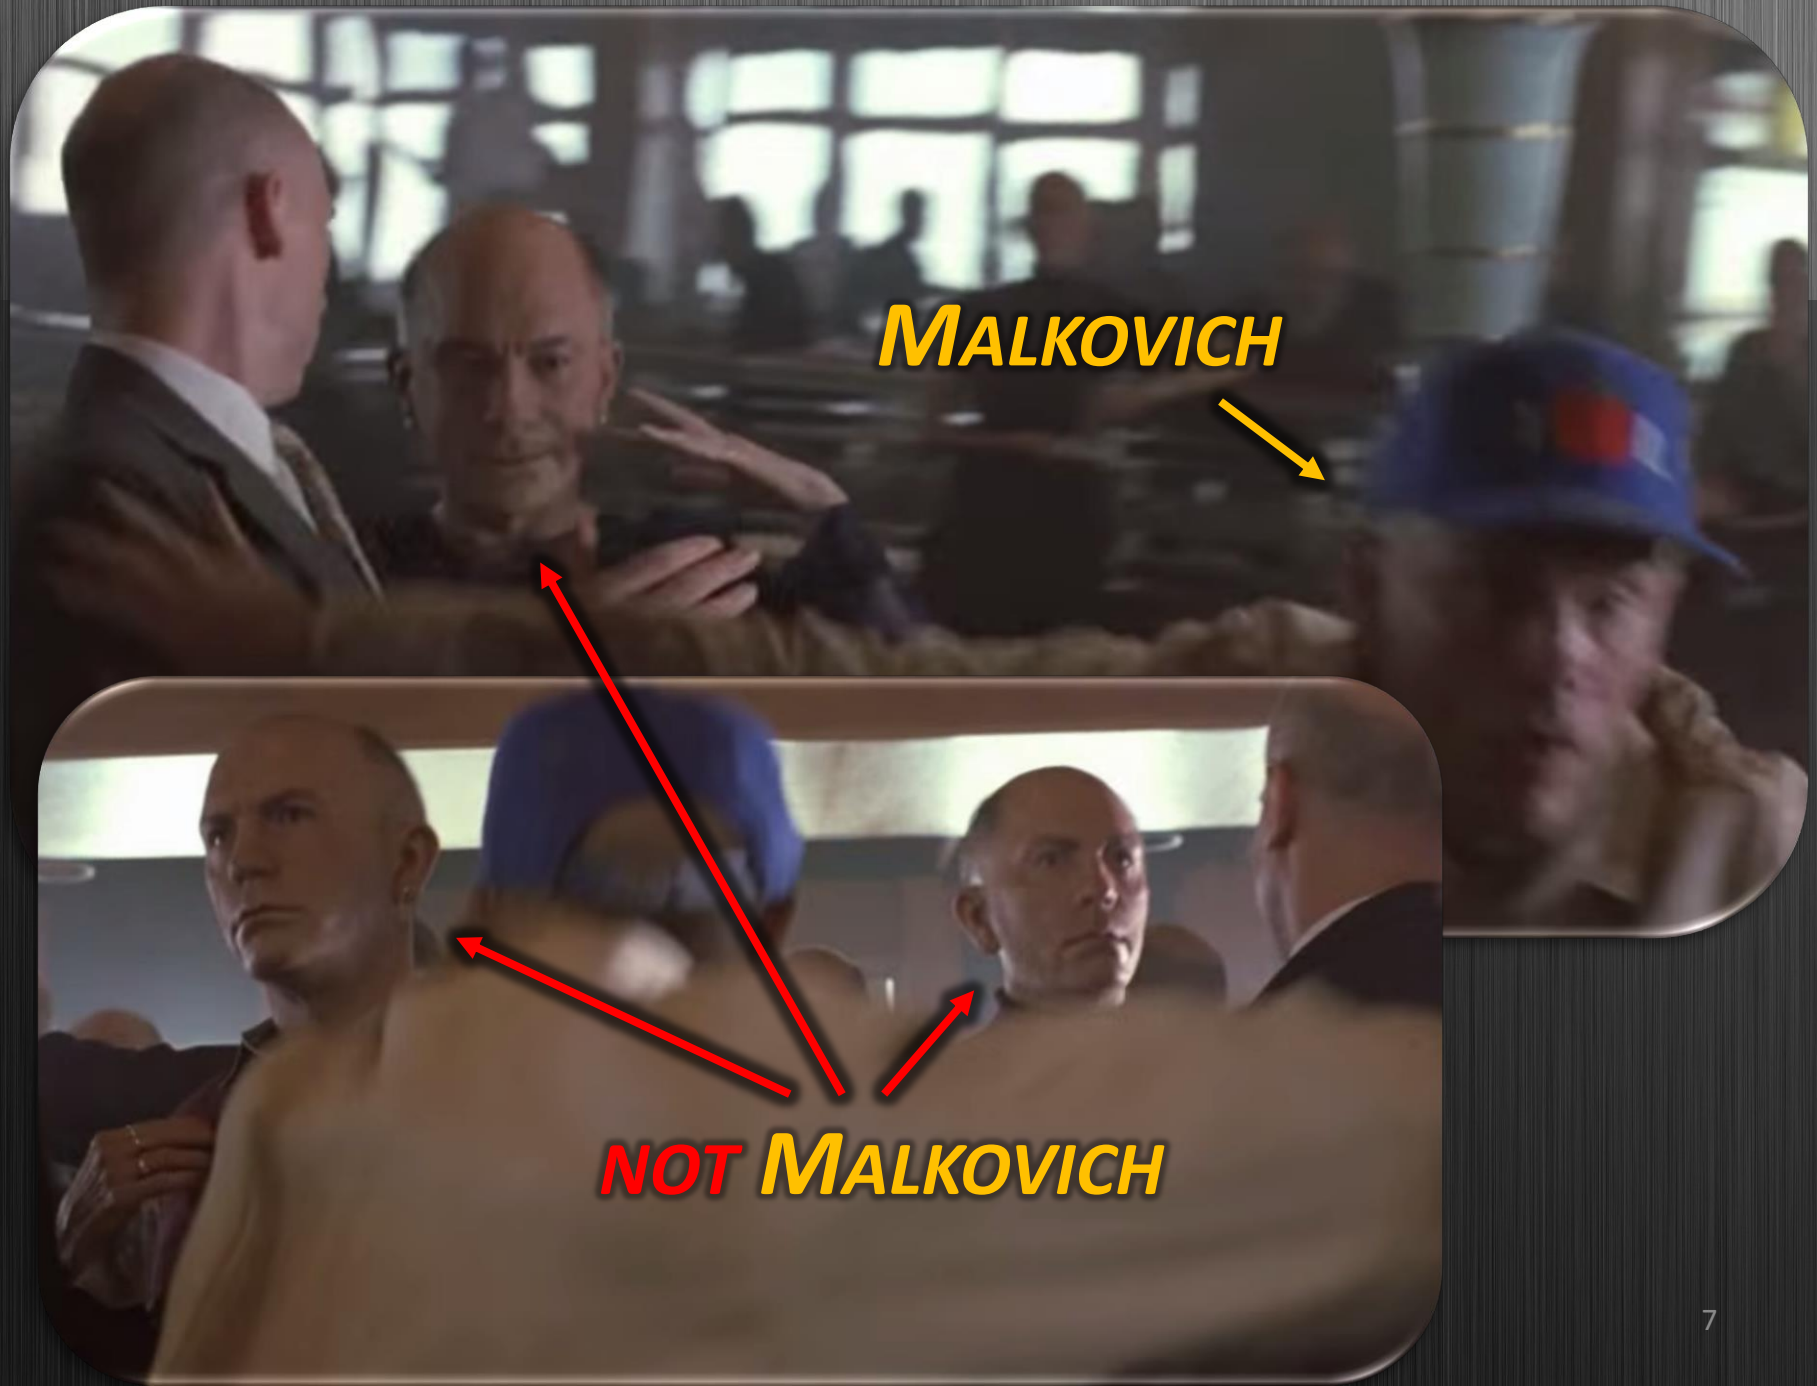
MALKOVICH
NOT MALKOVICH

second thought: remove rough commonality, ***then*** apply hammer

sdkernel

1. Obtain a set of kernel images that are of interest
2. Split the images into 4K/16-blocks &
   # eliminate all repetitive blocks
3. Create a sample of the blocks of the desired size from each image
   # 100x4K, or 25x16K
4. Create a block-based sdhash similarity digest

how unique are *linux* kernels?

http://security.ubuntu.com/ubuntu/pool/main/l/linux/:

- 943 packages
- 300 *generic* (150 32-/64-bit each)
- 288 *unique* kernels (144 each)

Kernel range	*Samples*	*Ubuntu version*
2.6.32-21 – 2.6.32-56	36	10.04
3.0.0-12 – 3.0.0-32	21	11.04
3.2.0-23 – 3.2.0-59	35	12.04
3.5.0-17 – 3.5.0-46	28	12.10
3.8.0-19 – 3.8.0-35	17	13.04
3.11.0-12 – 3.11.0-17	5	13.10
3.13.0-7 – 3.13.0-8	2	13.10

measuring uniqueness

- SLICE KERNEL IMAGE INTO 4K BLOCKS
  - # (rationale: it's a *ramfs* image)
- (CRYPTO-) HASH
- REMOVE REPETITIVE BLOCKS
  - # results in *signature base*
- COUNT UNIQUE BLOCKS AS A ROUGH MEASURE OF UNIQUENESS
  - # any shared block is useless as part of a signature

uniqueness: i386 vs amd64 kernels

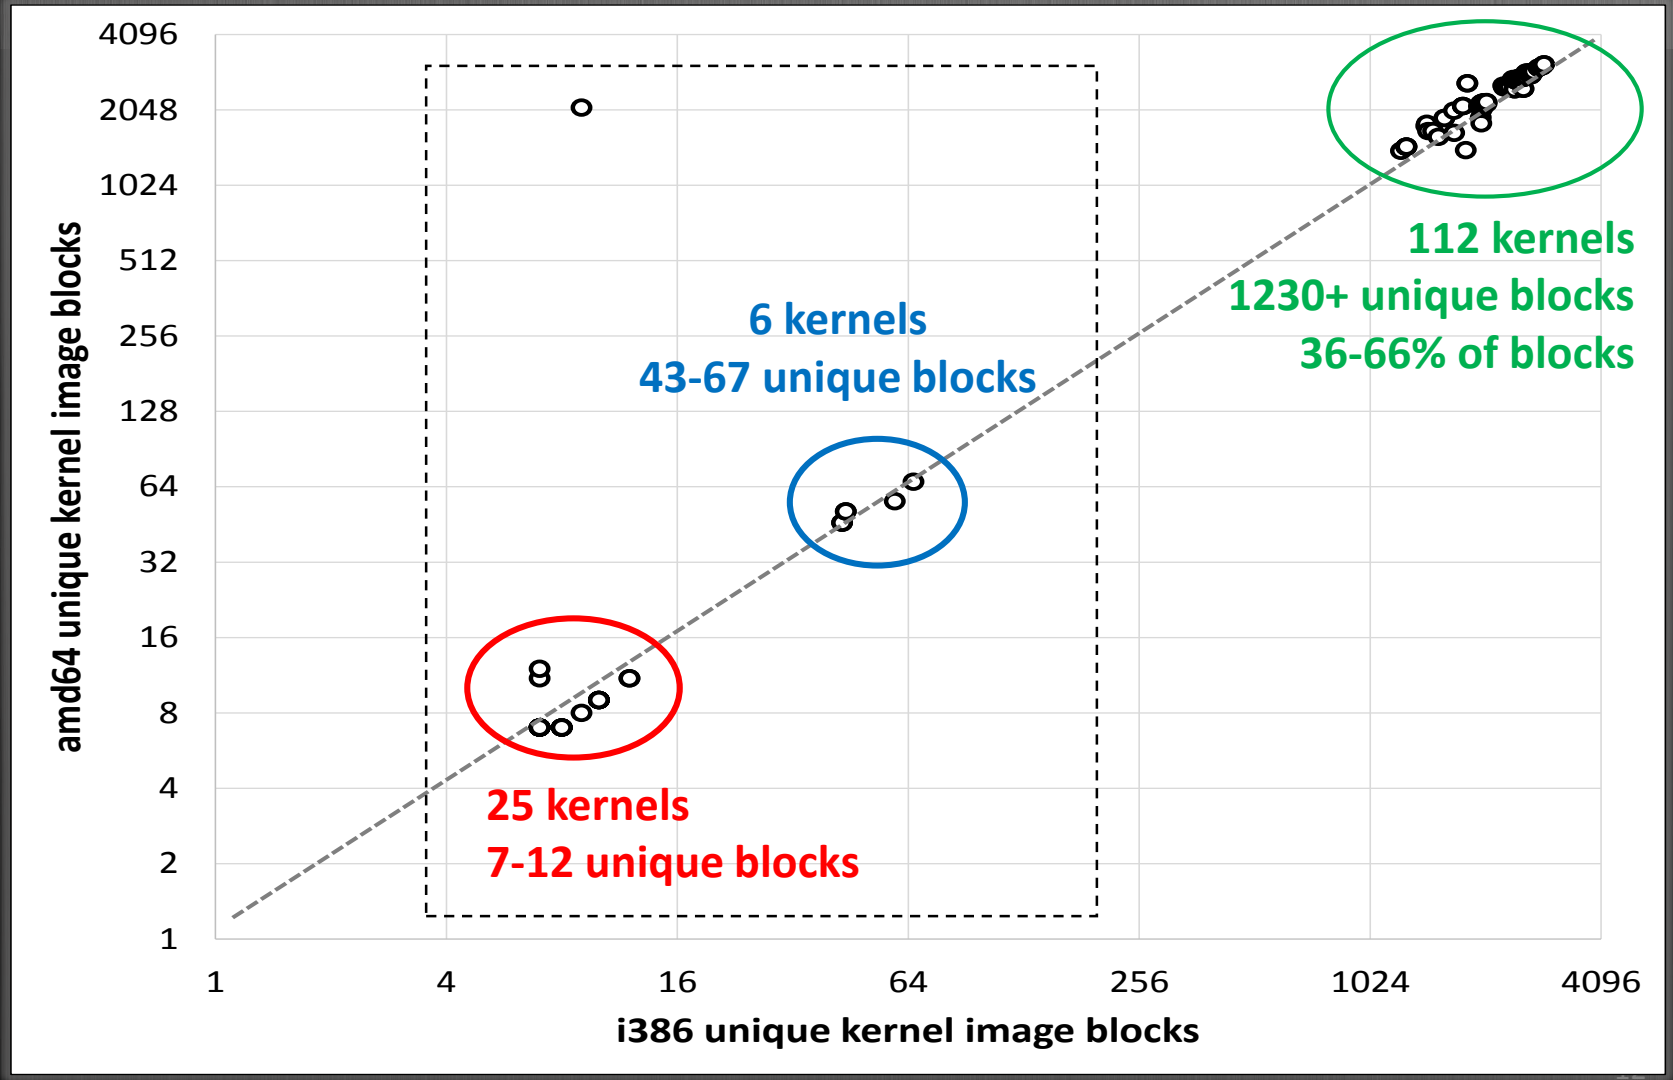

unique blocks of select kernels

Kernel	i386	amd64	Kernel	i386	amd64
2.6.32-47	7	11	3.5.0-32	7	7
2.6.32-48	7	12	3.5.0-33	7	7
2.6.32-53	59	56	3.5.0-34	7	7
2.6.32-54	66	67	3.5.0-38	43	46
3.2.0-44	8	7	3.5.0-39	43	46
3.2.0-45	8	7	3.5.0-41	7	7
3.2.0-47	8	7	3.5.0-42	7	7
3.2.0-48	8	7	3.8.0-19	12	11
3.2.0-49	8	7	3.8.0-20	10	9
3.2.0-50	7	7	3.8.0-21	12	11
3.2.0-51	7	7	3.8.0-22	10	9
3.2.0-52	9	2095	3.8.0-23	10	9
3.5.0-28	10	9	3.8.0-24	10	9
3.5.0-29	9	8	3.8.0-25	10	9
3.5.0-30	10	9	3.8.0-28	44	51
3.5.0-31	7	7	3.8.0-29	44	51

measuring selectivity

GIVEN SIGNATURE BASES $b_1, \ldots, b_n$ ***SELECTIVITY*** *IS:*

$$sel(b_i) = 100 - \max\left(sdhash_{4k}(b_i, b_j) : 1 \leq j \leq n, j \neq i\right)$$

RATIONALE

shrink a confusion matrix row to a single number:

	b_1	b_2	b_3	b_4	b_5	b_6	sel(b_i)
b_1		60	40	28	18	12	40
b_2	60		53	20	15	13	40
b_3	40	53		51	37	17	47
b_4	28	20	51		49	22	49
b_5	18	15	37	49		62	38
b_6	12	13	17	22	62		38

selectivity: amd64 signature bases

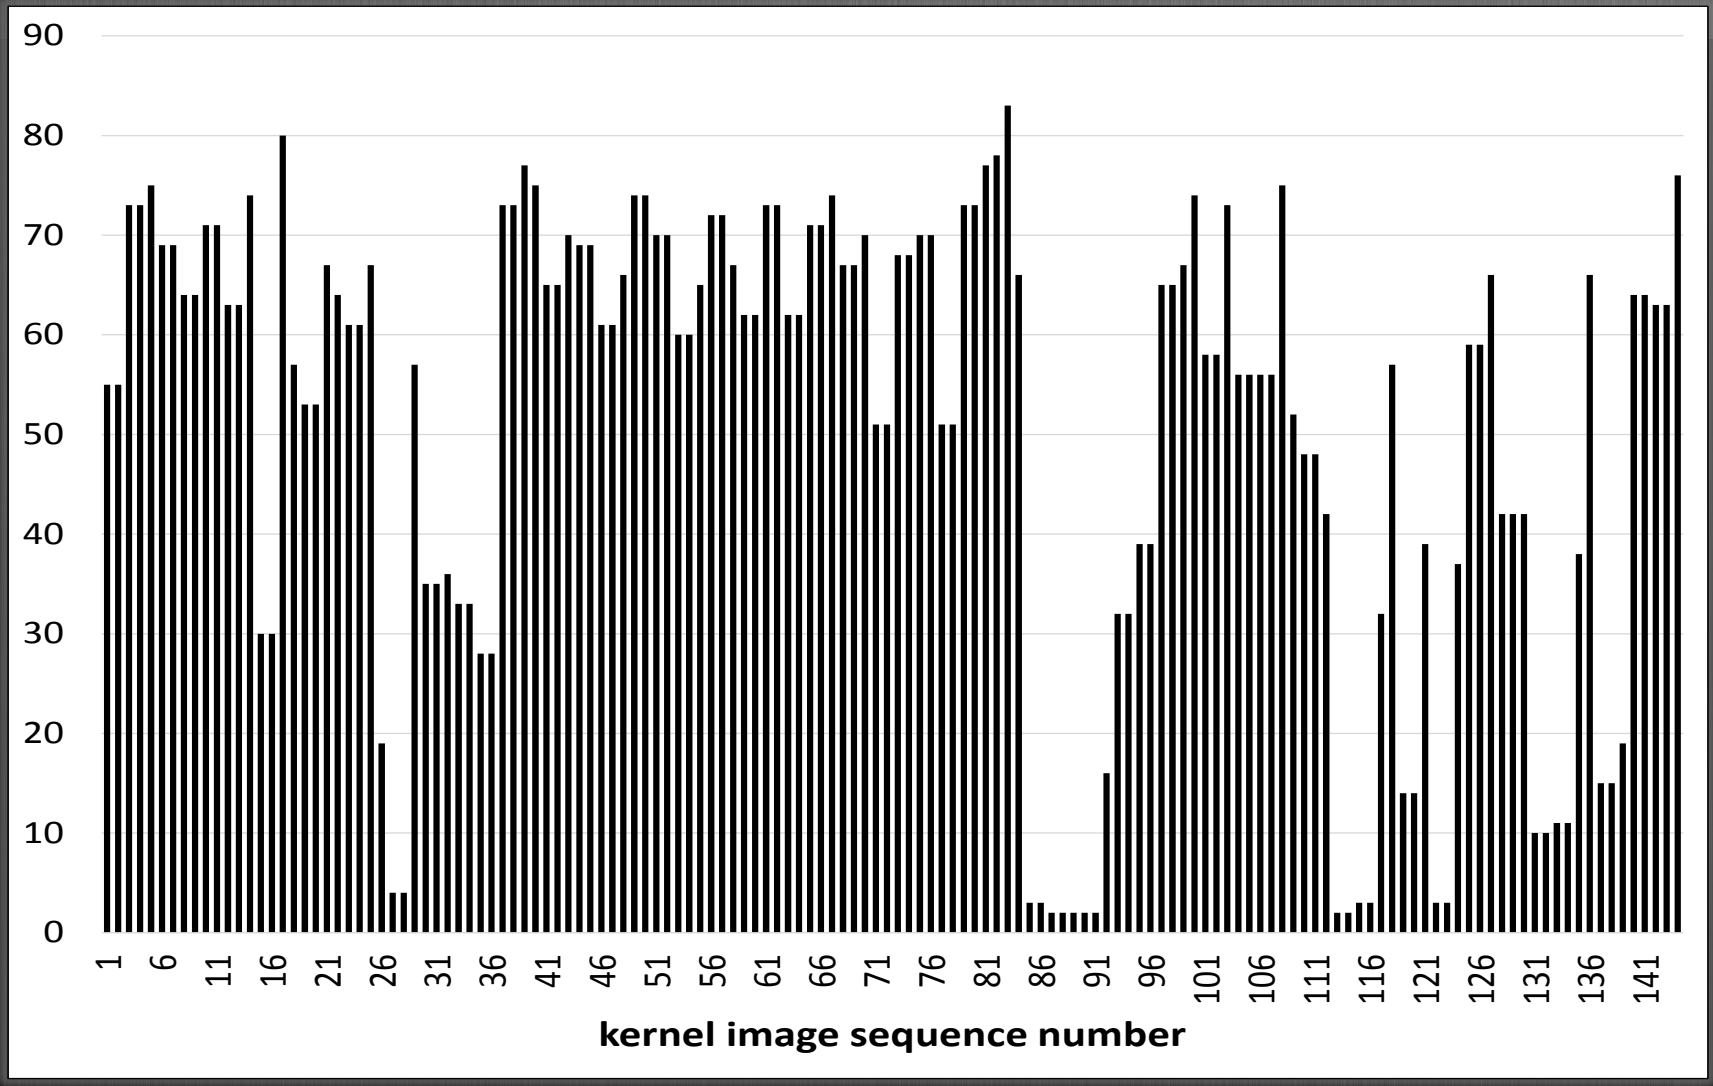

relative selectivity: original vs. sampled

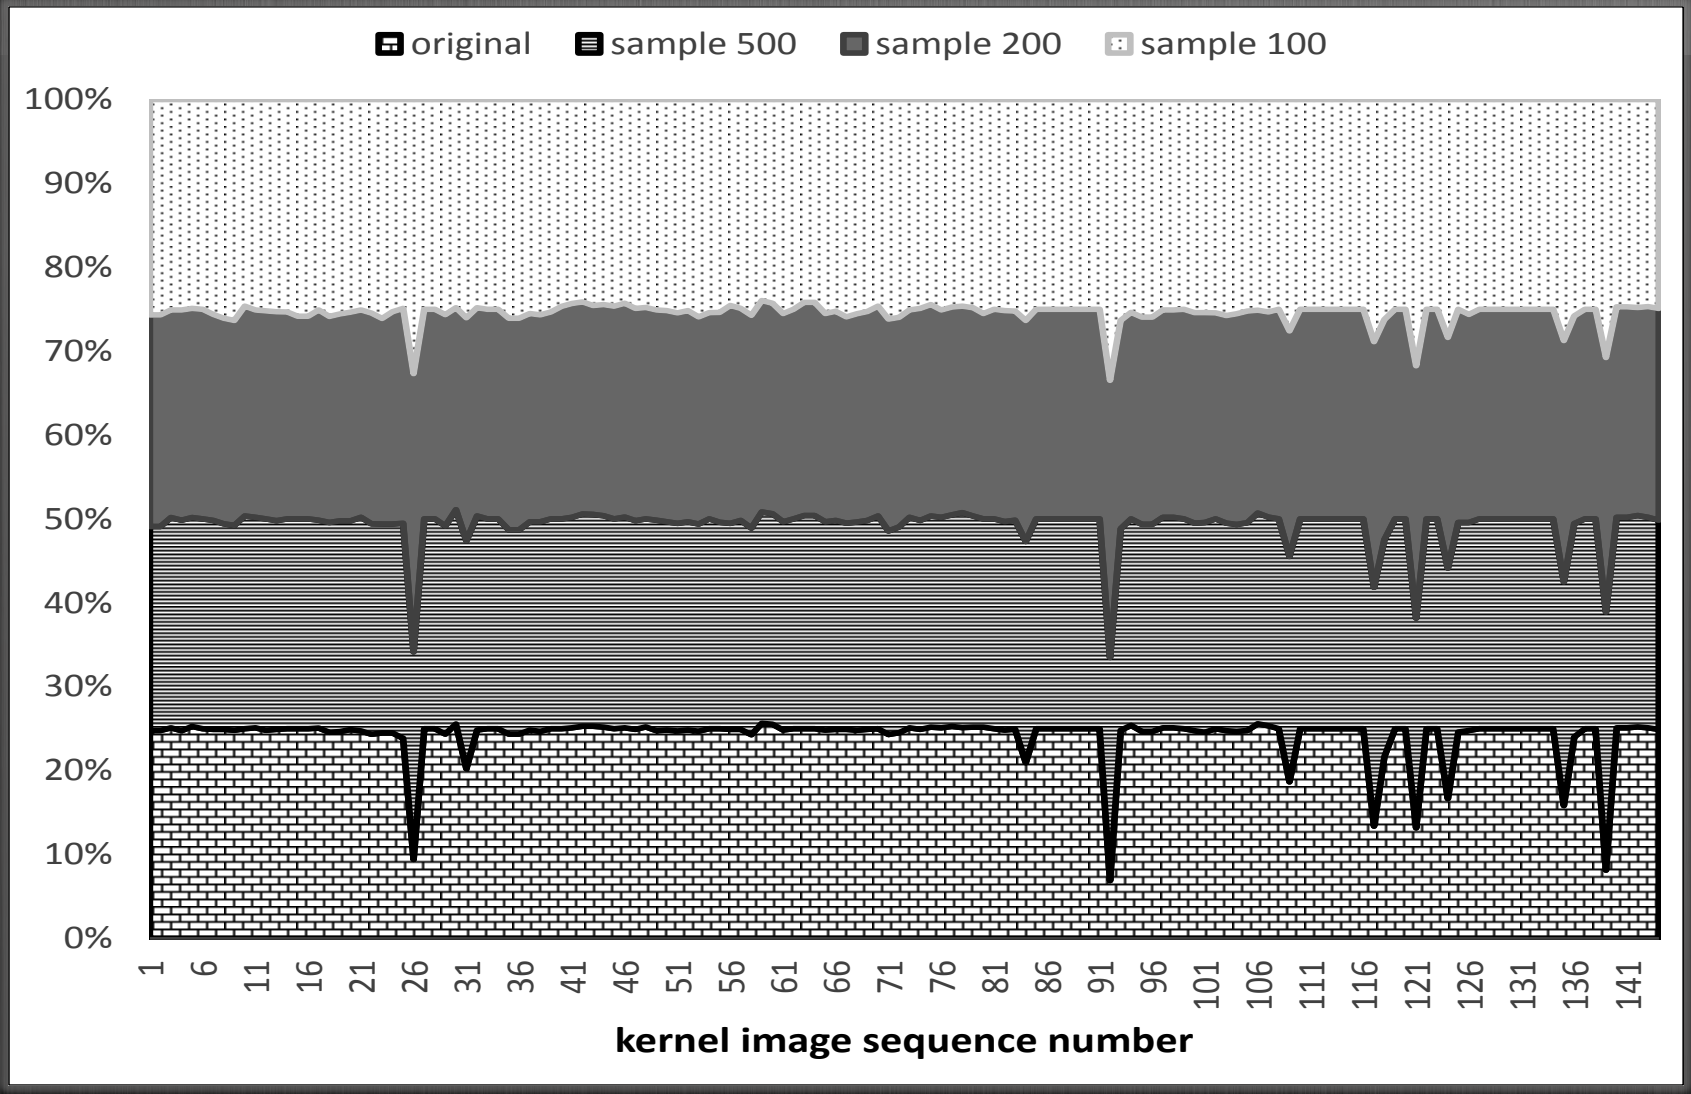

effects of sampling: amd64

OS version	$sdhash_{base}$	sel_{base}	$sdhash_{16k}$	sel_{16k}
2.6.32-35	67	33	61	32
2.6.32-42	67	26	61	26
2.6.32-48	67	18	61	9
2.6.32-54	67	0	61	0
3.2.0-31	68	23	62	24
3.2.0-38	68	41	62	40
3.2.0-45	68	0	62	0
3.2.0-53	68	17	62	18

effects of sampling: arm5

OS version	$sdhash_{base}$	sel_{base}	$sdhash_{16k}$	sel_{16k}
3.12.00	82	22	77	22
3.12.01	81	21	77	20
3.12.02	68	21	65	29
3.12.03	71	24	70	35
3.12.04	87	14	85	17
3.12.05	87	15	85	17
3.12.06	71	26	72	30
3.12.07	71	19	69	26
3.12.08	75	25	68	24
3.12.09	77	36	70	37
3.12.10	71	22	70	37

effects of sampling: windows

	xp.2	xp.3	vista.0	vista.1	vista.2	win7.0	win7.1	win8.0	win8.1
xp.2-os	13	14	13	14	14	13	13	13	13
xp.2-pa	**44**	21	14	14	14	14	13	13	13
xp.3-os	16	17	13	14	13	13	13	13	13
xp.3-pa	22	**46**	14	15	14	13	13	13	13
vista.0-os	11	12	17	16	15	13	13	13	12
vista.0-pa	12	12	**36**	19	19	14	14	13	13
vista.1-os	12	12	16	19	17	15	15	13	13
vista.1-pa	12	12	16	**64**	19	16	16	13	14
vista.2-os	12	12	17	20	19	16	16	14	14
vista.2-pa	11	12	16	19	**37**	13	13	13	13
win7.0-os	12	12	15	16	15	17	18	13	14
win7.0-pa	11	12	14	15	14	**40**	18	13	13
win7.1-os	11	12	14	15	14	17	20	13	13
win7.1-pa	11	11	13	14	14	18	**32**	13	13
win8.0-os	11	11	12	13	12	12	12	**34**	14
win8.1-os	10	11	12	13	13	13	13	16	**55**

sdkernel summary

New approach to kernel ID

- fully automated, zero reverse engineering
- uses only on-disk kernel image to build signatures
- accurate & robust
- works across different architectures (x86/arm)
- efficient → signature can produced from 25 blocks
  - |signature| = 25x256 = 6,400 bytes

Survey of Linux kernels

- varying degree of similarity
- build options **substantially** influence outcome
  - more than (neighboring) code versions
  - x86: *generic* vs. *lowlatency*; arm5: *default* vs. *qemu*
  - custom kernels should be quite unique

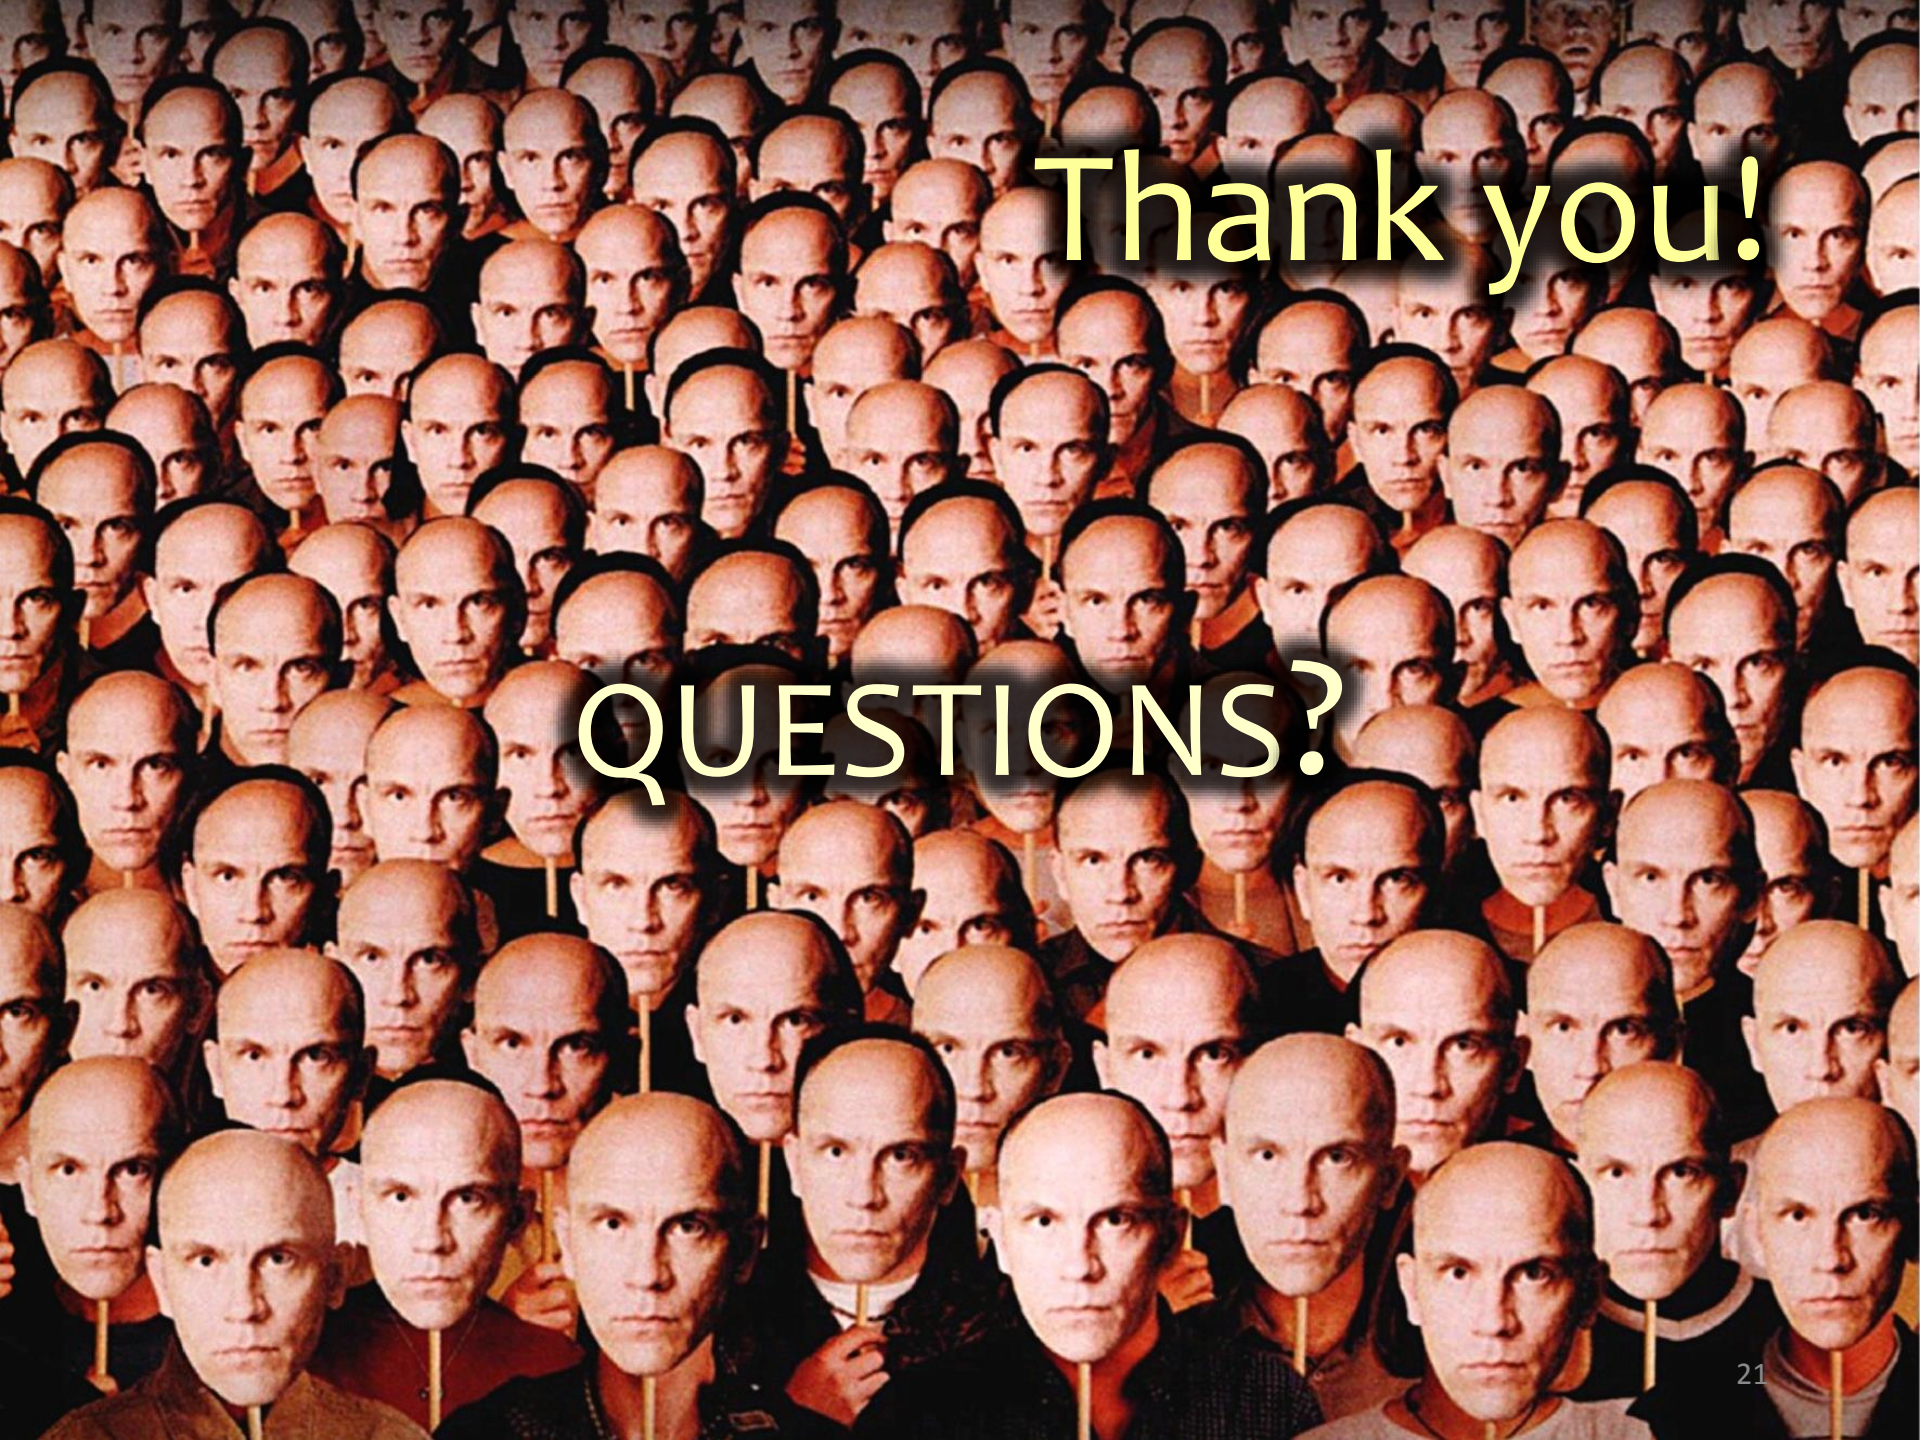
Thank you!
QUESTIONS?