



Paper-A Scalable File Based Data Store For Forensic Analysis

By

Flavio Cruz, Andreas Moser and Michael Cohen

Presented At

The Digital Forensic Research Conference

DFRWS 2015 EU Dublin, Ireland (Mar 23rd- 26th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment. As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

<http://dfrws.org>


A scalable file based data store for forensic analysis

Flavio Cruz, Andreas Moser, Michael Cohen

GRR

- Distributed Incident Response Framework
 - Agent running on client machines
 - Server collects information from agents running remotely and stores it in the data store
- A large number of machines implies huge amounts of forensic data

GRR Overview



AFF4 Object Space

- AFF4 is an object store abstraction
- Made of many objects:
 - Object: `aff4:/C.34a62f06/fs/os/boot.ini`
 - An object has a set of attributes
 - Attribute: set of pairs (value, timestamp)
- AFF4 namespace universe is assumed to be incomplete
 - AFF4 objects cannot be enumerated
 - We must store references to children of objects as an attribute of the object

Data Store Requirements

- Single object access
- Support for both synchronous and asynchronous operations
- Object locking
- Concurrent: multi-process and multi-threaded
- Time-stamped attributes

Data Store

- Due to the characteristics of AFF4, most key-value “NoSQL” tend to be a good fit
- Available options in GRR:
 - Mongo (objects as documents)
 - MySQL (single table where each row represents a timestamped value)
- Those data stores have shown some performance problems

SQLite Data Store


- Partition the AFF4 object space into several files
 - Each file may have multiple objects
 - Attributes of a single object are stored in a single file
- Files are independent databases since each operation only works with an object
- Granularity can be customized

Granularity: Path Configuration


- **aff4:/C.34a62f06/fs/os/boot.ini**
 - C.34a62f06.sqlite
- **aff4:/C.34a62f06/fs/os/cmd.exe**
 - C.34a62f06.sqlite
- **aff4:/blobs/2af3012a**
 - blobs/2af3012a.sqlite
- **Pathing configuration:**
 - (?P<path>C\\.\\{1,16\\}?)(\$|/.*) - One file per client
 - (?P<path>blobs/[^/]+).* - One file per blob
 - (?P<path>hunts/[^/]+).* - One file per hunt
 - (?P<path>[^/]+).* - One file per top-level directory

SQLite Data Store

SQLite Data Store.




Directory of files



Distributed data store

- While the SQLite data store has great performance, it will not scale well because all the server processes must run on a single machine
 - A single machine can only hold so much data
- Partition the SQLite files across N servers


Partitioning the directory



Architecture


GRR Processes

Data Server Group




Mapping AFF4 objects

- We need to map objects to data servers
 - Solved problem: objects to files
 - We need to map files to data servers



Consistent hashing

- Server hash space ranges from 0 to 2^{64}
- 'aff4:/C.34a62f06/fs/os/boot.ini' -> 'C.34a62f06.sqlite'
- hash('C.34a62f06.sqlite') -> 0x10000000




How is the mapping configuration used?


- Data Master:
 - Startup: check if mapping configuration exists
 - Create it if that is not the case
- Data Server:
 - Read master location
 - Register with master and receive mapping configuration
- Client:
 - Read list of servers
 - Ask a random server for the mapping
 - Master has a limited role

Data Store Overview

Data Server Client




Data Server




Micro Benchmarks

- Measured the performance of the new data store using micro-benchmarks
 - Benchmark script performs a sequence of data store operations
- Used two scenarios

(1) Objects with few attributes





(2) Objects with many attributes



End to End Benchmarks


- Execute several “flows” on multiple clients running on a single machine
- We measured the scalability of the data stores by increasing the number of:
 - Workers
 - Frontend servers
 - Data servers

End to End Benchmarks




End to End Benchmarks

2 Data Store Servers



8 Data Store Servers



Using the Data Store with NSRL

- NSRL is a library of hashes of files from known software packages
- GRR can de-duplicate files from multiple clients
- Idea: use the NSRL dataset to avoid storing well known files
 - Need to populate the data store with the NSRL library

Storing the NSRL library

- New AFF4 object that represents a NSRL file: `aff4:/files/nsrl/sha1-hash`
- Use the 3 characters of the sha1-hash for the path configuration
 - `000.sqlite ... fff.sqlite`: 4096 x 33MB files
- Took 5 hours to import and files are easily re-sharded if needed

Using the NSRL library on Windows 7

Statistic	With NSRL	Without NSRL
Files found	1605	1605
Files skipped	1245	2
Files fetched	360	1603
Data store size	148MB	314MB
Client sent	117MB	243MB
Client received	5MB	8MB
Client time	293s	400s

Thanks for your time!


Flavio Cruz

<fmfernand@cs.cmu.edu>

Authentication

- Data servers use an username/password when registering with the master
- After data server registration, the master also sends the encrypted client credentials to the data server
- Client credentials are a set of (user, password, permissions)

Scalability Results



Server processes

- Server processes share the data store and are allowed to manipulate it.
 - Frontend server
 - Worker: used by the server to perform tasks retrieved from a data store queue
 - Enroller: special worker responsible for the initial enrollment of clients
 - Admin UI: user interface that allows users to interact with the system

Data Store Operations

- MultiSet(subject, values, sync, replace)
- MultiResolveRegex(subject, predicate_regex, timestamp)
- MultiResolve(subject, predicates, timestamp)
- DeleteSubject(subject)
- DeleteAttributes(subject, predicates, start, end, sync)
- DeleteAttributesRegex(subject, regexes)
- Transaction(subject, lease_time)
- Flush()
- Size()

Sqlite Data Store

tbl	
* <u>subject</u>	varchar
*predicate	varchar
*timestamp	big integer
°value	blob

lock	
* <u>subject</u>	varchar
*expires	big integer
*token	big integer


Asynchronous Requests

- Each client may create up to N connections to each data server
- When a data store request is made, we pick the connection with the least number of pending requests
- If request is synchronous, we write it to the socket and wait until we receive a reply
- If request is asynchronous, we do not wait for a response


Threading

- Data servers create one thread per client connection
- Improved CPU utilization since network I/O and data store I/O takes time
- The GIL may still become a problem since each data server can only use one core at a time
 - Solution: use more data servers


1 worker




2 workers



4 workers



Push Benchmarks (Time)





Re-balancing

- Two-phase commit:
 - Data servers send files to target data servers
 - Sent files are kept in temporary storage
 - And that also includes list of files to remove
 - All data servers apply the transaction
 - Move files to real places
 - Delete unneeded files
- Even if the second phase fails, we can always resume the process
 - recover <transaction-id>


Communication

- Mostly using HTTP, except for the the data store requests
- After the initial handshake, client and server “upgrade” to a simple socket protocol, which allows multiple requests at once and faster serialization
 - protocol buffers


Blobs




Blobs




End to End Benchmarks (Size)




Manager




Old distribution



New distribution




End to End Benchmarks



Re-sharding

- What if we want to add more data servers?
 - Need new mapping configuration
 - Need to move files because some of them will be in the wrong server
- Manager solves this problem
 - Creates new mapping configuration
 - Files are moved automatically from server to server
 - Only works in “offline” mode

Adding new server



Removing data servers

- To delete a server, we first re-balance and set the target server hash range to $[x, x[$
 - Afterwards it is just a matter of stopping the server and removing the server from the configuration

(2) Few objects and many attributes

