



Database Forensic Analysis Through Internal Structure Carving

By

James Wagner, Alexander Rasin and Jonathan Grier

Presented At

The Digital Forensic Research Conference

DFRWS 2015 USA Philadelphia, PA (Aug 9th - 13th)

DFRWS is dedicated to the sharing of knowledge and ideas about digital forensics research. Ever since it organized the first open workshop devoted to digital forensics in 2001, DFRWS continues to bring academics and practitioners together in an informal environment. As a non-profit, volunteer organization, DFRWS sponsors technical working groups, annual conferences and challenges to help drive the direction of research and development.

<http://dfrws.org>

Database Forensic Analysis through Internal Structure Carving

James Wagner

Dr. Alexander Rasin

Jonathan Grier





Roadmap

- ❖ Motivation
 - ❖ File carving (for databases)
 - ❖ Stochastic analysis
- ❖ Results
- ❖ Future work

Motivation

- A lot of data lives in databases
- Recovery
 - Rely on “safe” storage (or backups)
 - Inconsistent / commercial
- Monitoring
 - Logs and profiling
 - Inspect DB connections
 - e.g., IBM Guardium

Database Storage

Table Rows

RID	Name	Position	FavoriteTown
1	Stan	Professor	Westwood
2	Ugur	Professor	Providence
3	Tom	Professor	Providence
4	Alex	Student	Providence
5	Stan	Professor	Cambridge
6	Andy	Student	Providence

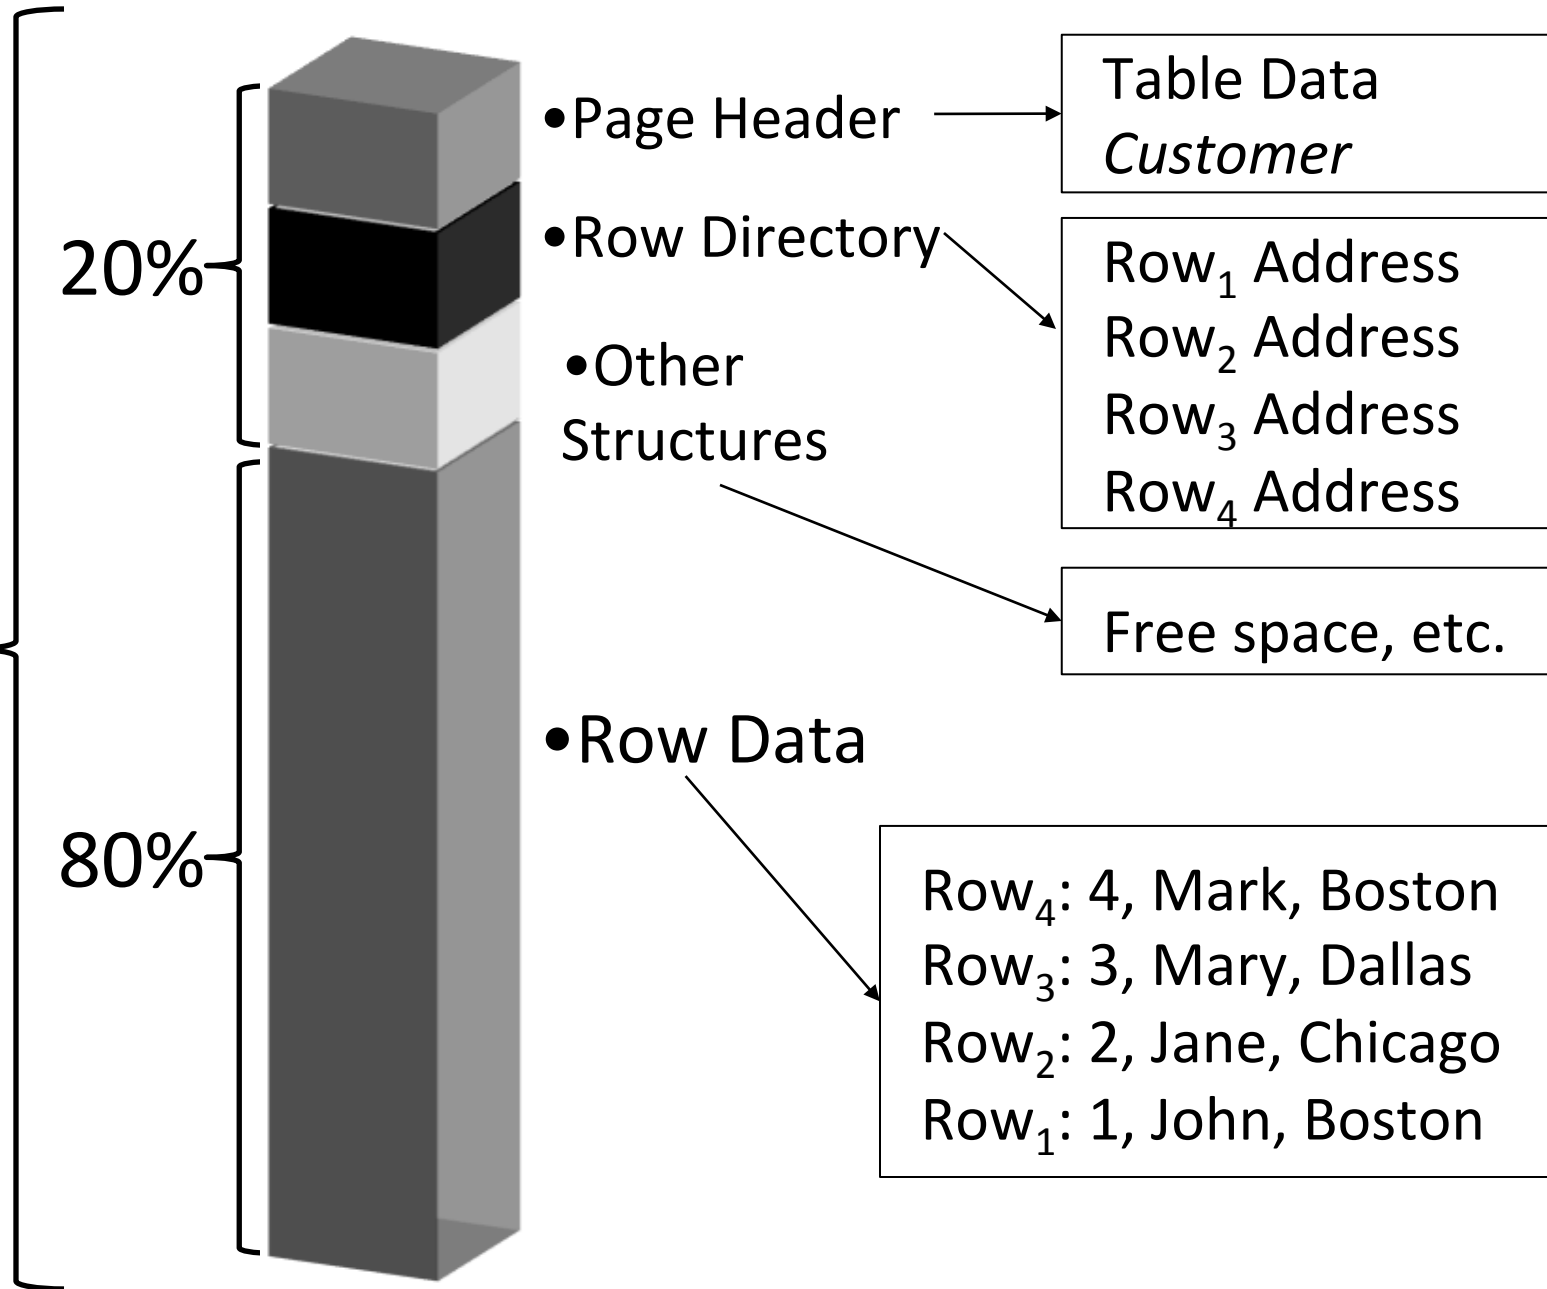
Logical

Physical

Disk Pages

#1	1	Stan	Professor	Westwood
	2	Ugur	Professor	Providence
#2	3	Tom	Professor	Providence
	4	Alex	Student	Providence
#3	5	Stan	Professor	Cambridge
	6	Andy	Student	Providence

**Database
Page**



Indexes and other Structures

Index on
Favorite
Town

Cambridge	5
Providence	2
Providence	3
Providence	4
Providence	6
Westwood	1

RID	Name	Position	FavoriteTown
1	Stan	Professor	Westwood
2	Ugur	Professor	Providence
3	Tom	Professor	Providence
4	Alex	Student	Providence
5	Stan	Professor	Cambridge
6	Andy	Student	Providence

MV

Logical

Physical

#1	Cambridge	5	Providence
	2	Providence	3

#2	Providence	4	Providence
	6	Westwood	1

#1	1	Stan	Professor	Westwood
	2	Ugur	Professor	Providence

#2	3	Tom	Professor	Providence
	4	Alex	Student	Providence

#3	5	Stan	Professor	Cambridge
	6	Andy	Student	Providence

#1

#2

#3



Roadmap

❖ Motivation

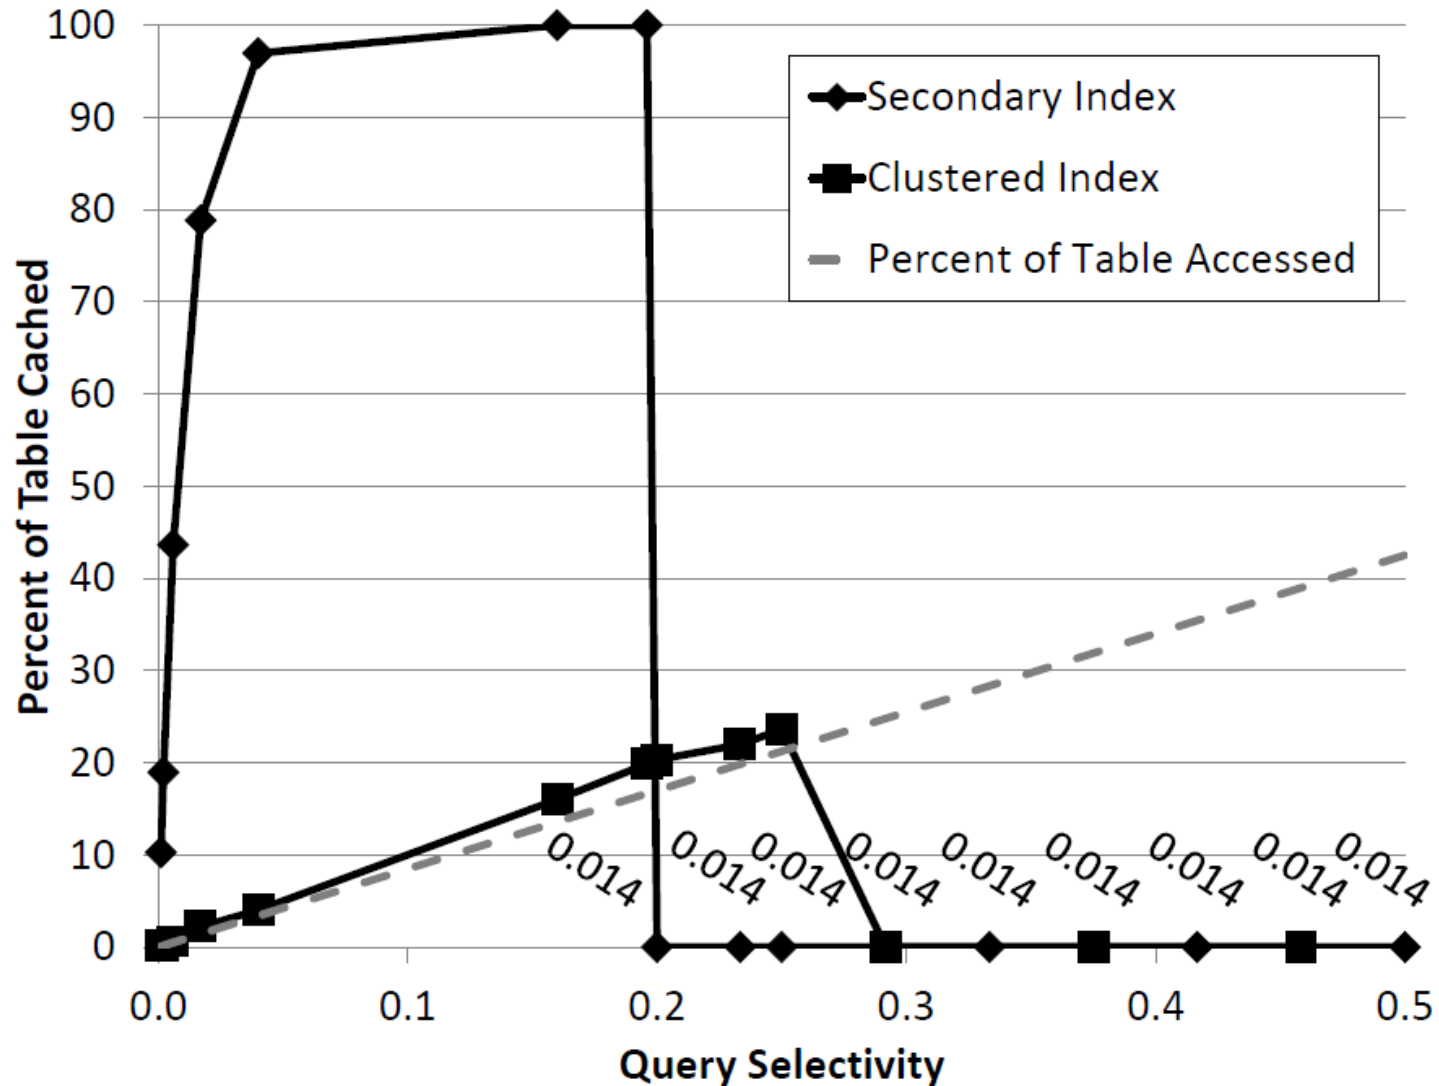
- ❖ File carving (for databases)

❖ Stochastic analysis

- ❖ Results

- ❖ Future work

Database Caching Policies

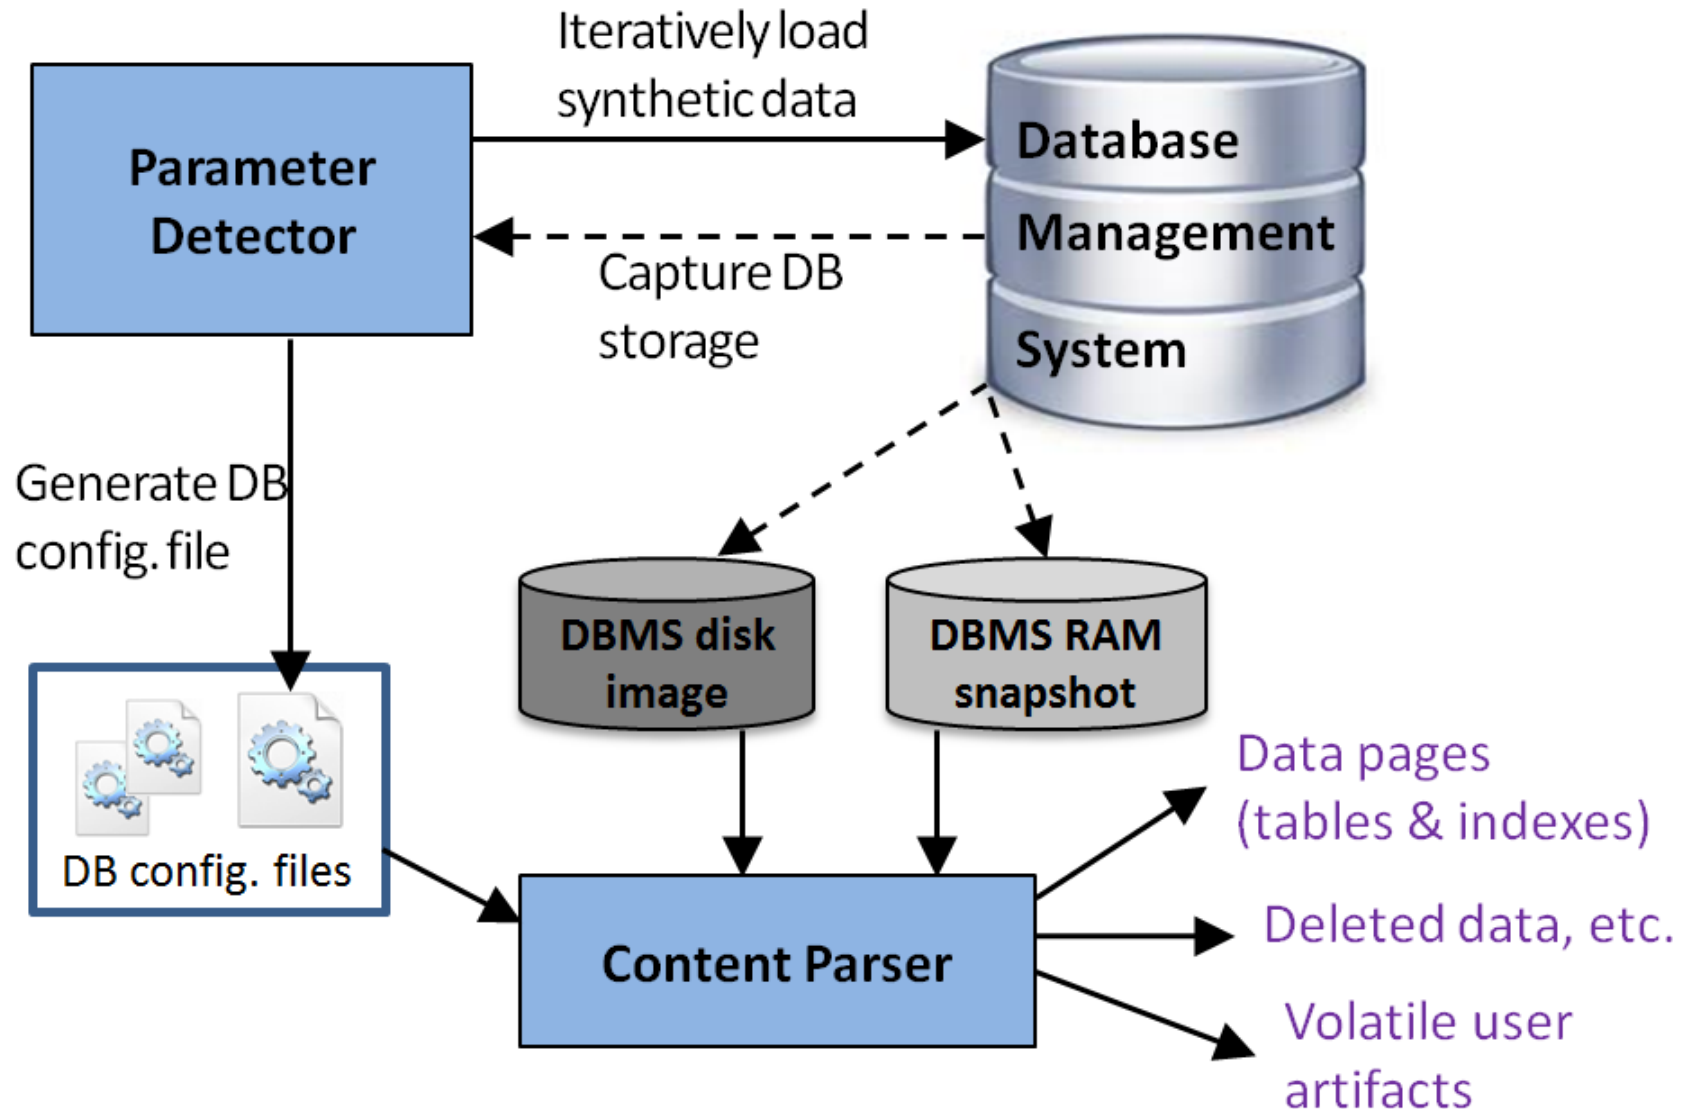




Roadmap

- ❖ Motivation
 - ❖ File carving (for databases)
 - ❖ Stochastic analysis
- ❖ Results
- ❖ Future work

Architecture



	Oracle	PostgreSQL	SQLite	Firebird	DB2	SQLServer	MySQL	Apache Derby
Structure Identifier	Yes	No	Yes					No
Unique Page ID	Yes							No
Row Dir. Sequence	Top-to-bottom insertion					Bottom-to-top insertion		
Row Identifier	No	Yes	No				Yes	
Column Count	Yes			No		Yes	No	Yes
Column Sizes	Yes				No		Yes	
Column Directory	No				Yes		No	
Numbers w/Strings	Yes				No		Yes	

DBMS Versions

Hard to get
some older
DB versions



Different
parameters


DBMS Version	Testing OS	Buffer Size(MB)	Page Size(KB)
Apache Derby 10.10	Linux	400	4
Apache Derby 10.5	Linux	400	4
DB2 Express-C 10.5	Linux	400	4
Firebird 2.5.1	Linux	400	8
Firebird 2.1.7	Windows	400	8
MySQL Server 5.1.73	Linux	800	16
MySQL Server 5.6.1	Windows	800	16
Oracle 11g R2	Windows	800	8
Oracle 12c R1	Windows	1200	8
PostgreSQL 7.3	Linux	400	8
PostgreSQL 8.4	Linux	400	8
PostgreSQL 9.3	Windows	800	8
SQLite 3.8.6	Linux	2	1
SQLite 3.8.7	Windows	2	1
SQLServer 2008 Enterprise	Windows (Linux)	800	8


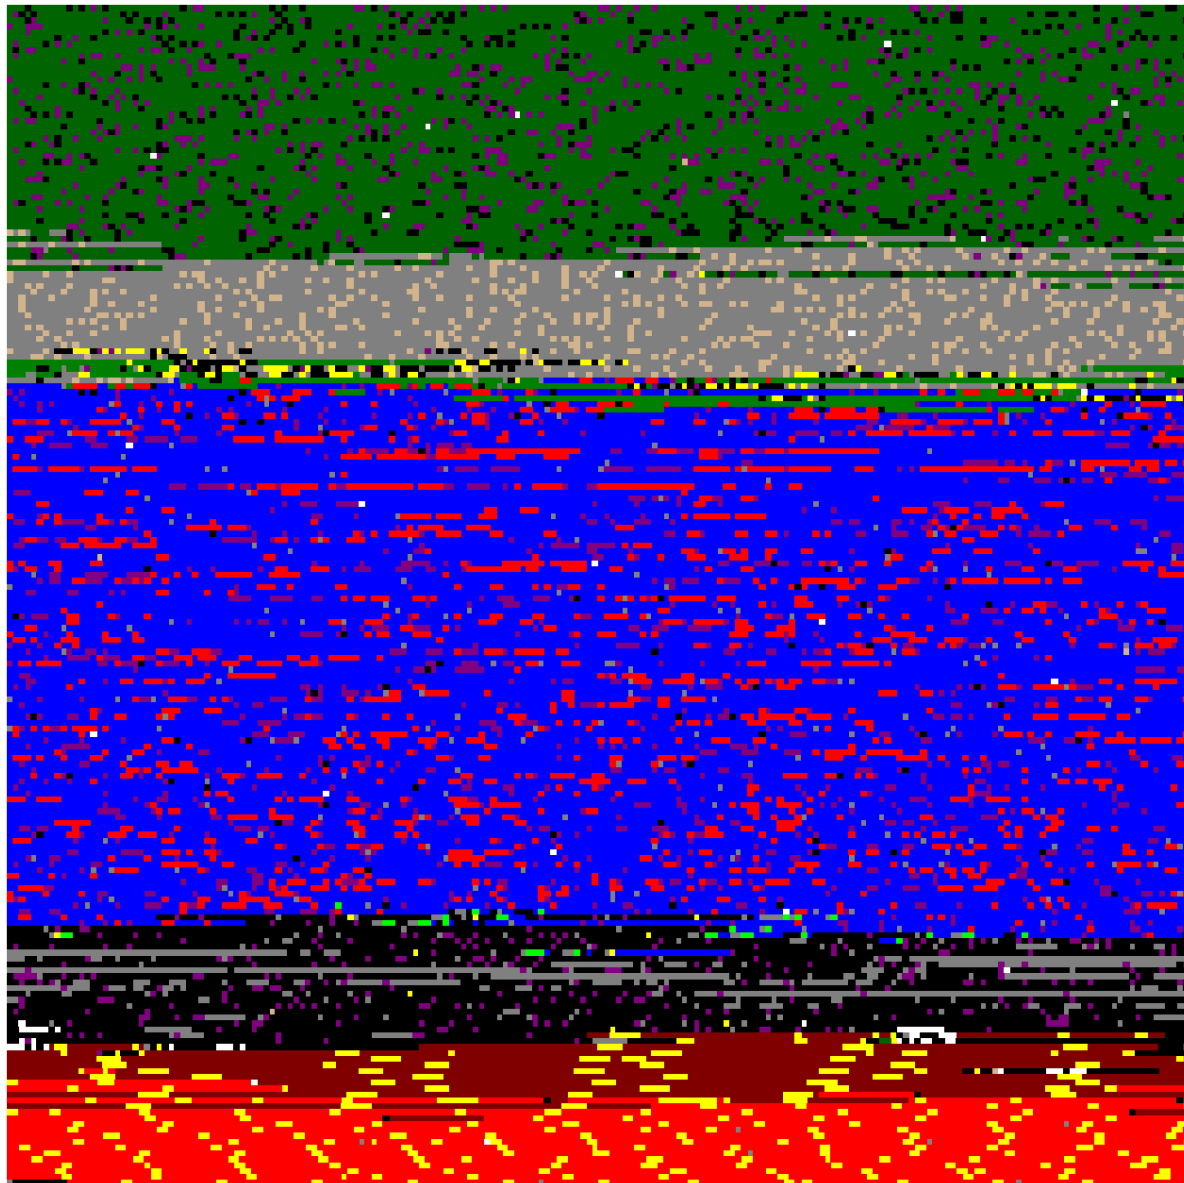
Recovering Corrupted Data



- An AWS instance + PostgreSQL
 1. Load SSBM benchmark data
 2. Delete the rows, then **delete all Postgres files**

	Damage = 0%	Damage = 10%	Damage = 25%
<i>Dwdate</i>	35 (100%)	31 (88.6%)	20 (57.1%)
<i>Supplier</i>	565 (100%)	455 (80.5%)	326 (57.7%)
<i>Customer</i>	1915 (100%)	1559 (81.4%)	1075 (56.1%)
<i>Part</i>	8659 (100%)	6969 (80.5%)	4864 (56.2%)
<i>Lineorder</i>	115 K (100%)	104 K (89.9%)	87 K (75.2%)
TOTAL	416 K (100%)	375 K (89.9%)	312 K (74.9%)










Data Modification

- DELETE/UPDATE/INSERT
- When is a value **really** deleted?
- Example
 - Customer table (Phone# column)
 - Indexed Phone#

Row Deletion

Event	Table		Index	
	HDD	RAM	HDD	RAM
Time ₀			✓	
Delete a Phone #	✓	X	✓	✓
Run a few of queries	✓	X	✓	
Many more queries	X		✓	
Rebuild the index	X			
Rebuild the table				



Roadmap

- ❖ Motivation
 - ❖ File carving (for databases)
 - ❖ Stochastic analysis
- ❖ Results
- ❖ Future work

Future Work

- Column-stores and Key-value stores
- Automated database deconstruction
- Database performance tuning
- Monitoring user behavior (DB cache)
- Independent database audit/verification

