



Advancing the AFF4 to the challenges of Volatile Memory & Single Hashes

Dr. Bradley Schatz
Director, Schatz Forensic

v1.0 – DFRWS - 2017

© Schatz Forensic 2017

Agenda

- AFF4 background
- What can be improved with hashing in AFF4?
- Our refined single hashing proposal: the *Linear Chunk Hash*

Traditional forensic imaging introduces unnecessary delays

- Physical acquisition delays analysis by hours
 - Doesn't scale to multi-core/RAID/SSD/NVMe
 - Doesn't allow non-linear writes (concurrent imaging & analysis)
- Possible solution: logical imaging/triage
 - loses context and unallocated

Traditional forensic imaging doesn't scale

- ZIP/Deflate compression is really CPU expensive
 - (~ 40MB/s/core for high entropy)
- Hashing is CPU expensive
 - (~ 600MB/s/core for 2016 i7)
- ***Linear bitstream hashing isn't parallelizable***
 - ***Ceiling of around 600MB/s on current generation i7 CPU's***

The AFF4 removes the bottleneck from forensic imaging

- Snappy compression is CPU inexpensive
 - (> 1000 MB/s/core)
- Hashing is still CPU expensive
 - (~ 600 MB/s/core for 2016 i7)
- AFF4 block hashing is parallelizable
 - Seeing > 1000 MB/s on current generation CPU's

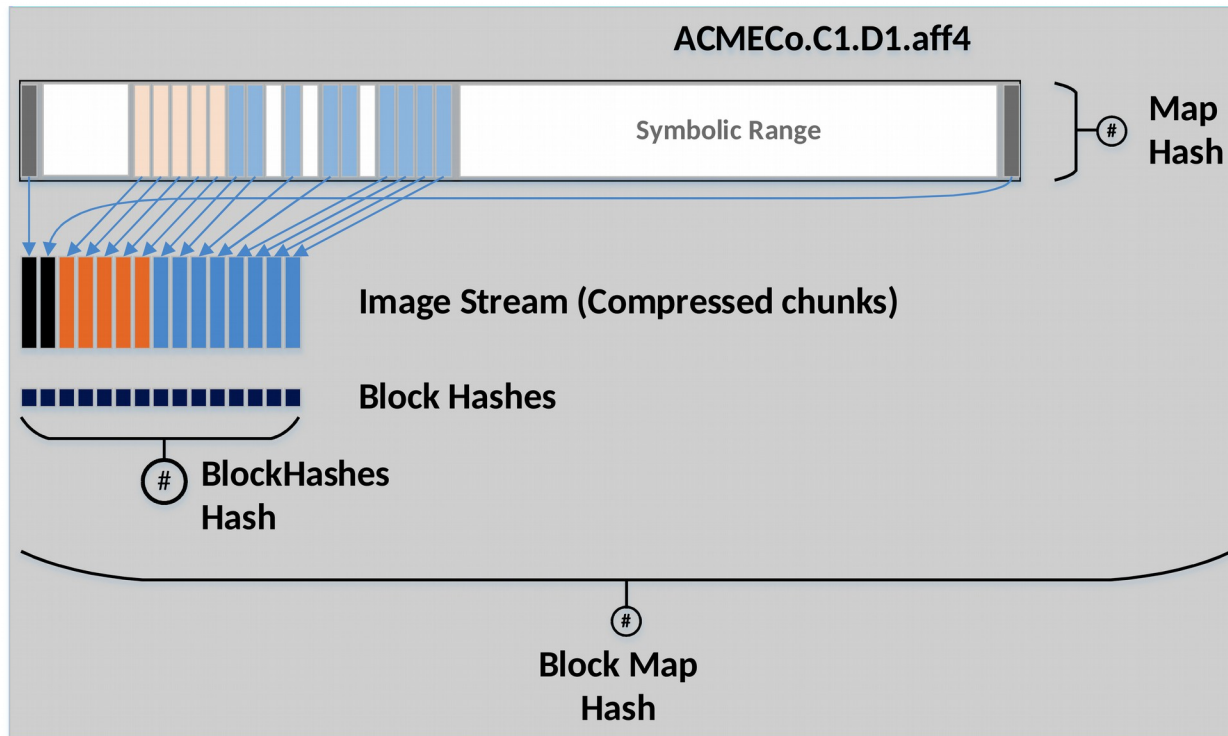
The AFF4 removes enables acquisition to occur at the same time as analysis

- Non-linear, potentially partial compressed block stream
- Hashing mechanism that enables incremental hash calculation

AFF4 Standard v1.0 Hashing

- Defines a hash construction called the **Block Map Hash**
 - Single hash for image
- Hybrid of:
 - Block (segment) hashing for concrete data chunks
 - Non concrete data chunks (eg. Sparse) protected by hash on map
 - Merkel tree like hash hierarchy


AFF4 Standard v1.0 Hashing



Desirable properties of hashing

| | Linear Bitstream Hash | AFF4 Block Map Hash |
|--|-----------------------------|------------------------|
| Hash is concise | Yes | Yes |
| Hash of image matches hash calculated when imaging | Yes | Yes |
| Changed image results in failed hash match | Yes | Yes |
| Successive hashes of same source result in same hash | Yes | <i>Maybe</i> |

Parallelism may lead to differing hashes



Can we generate a block based hash that is not subject to reordering ?

- Hypothesis 1: Reorder the lower Block Hashes based on their upper address
 - Works well for images of disks with no read errors.
 - Doesn't work so well for:
 - Disk images with read errors
 - *Memory images with chunk sizes > memory page size*

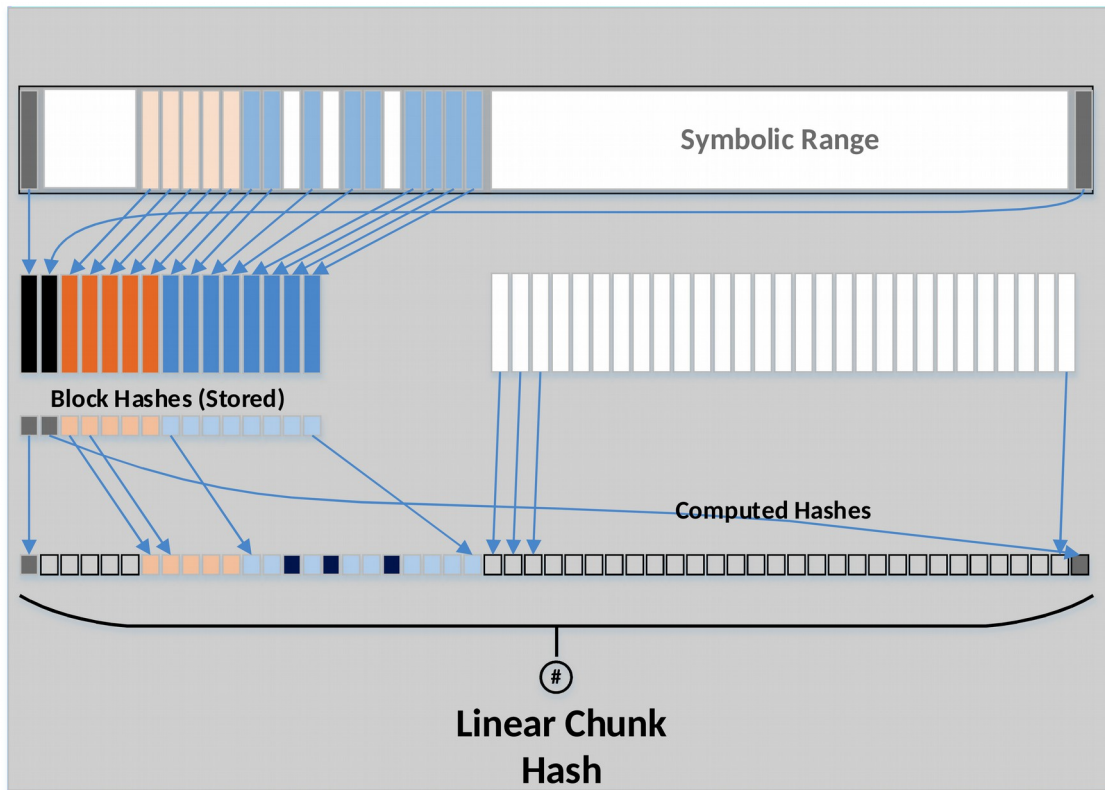
Sub chunk size data introduces further complication

- High performance requires (chunk size > disk sector size)
 - IO Throughput - 1MB for disk IO
 - Memory consumption- Maps
- ddrescue bad sector imaging algorithm requires chunks of 512B or 4KB when the chunk size might be many multiples of this
- RAM physical address space holes lead to sub chunk sized data runs

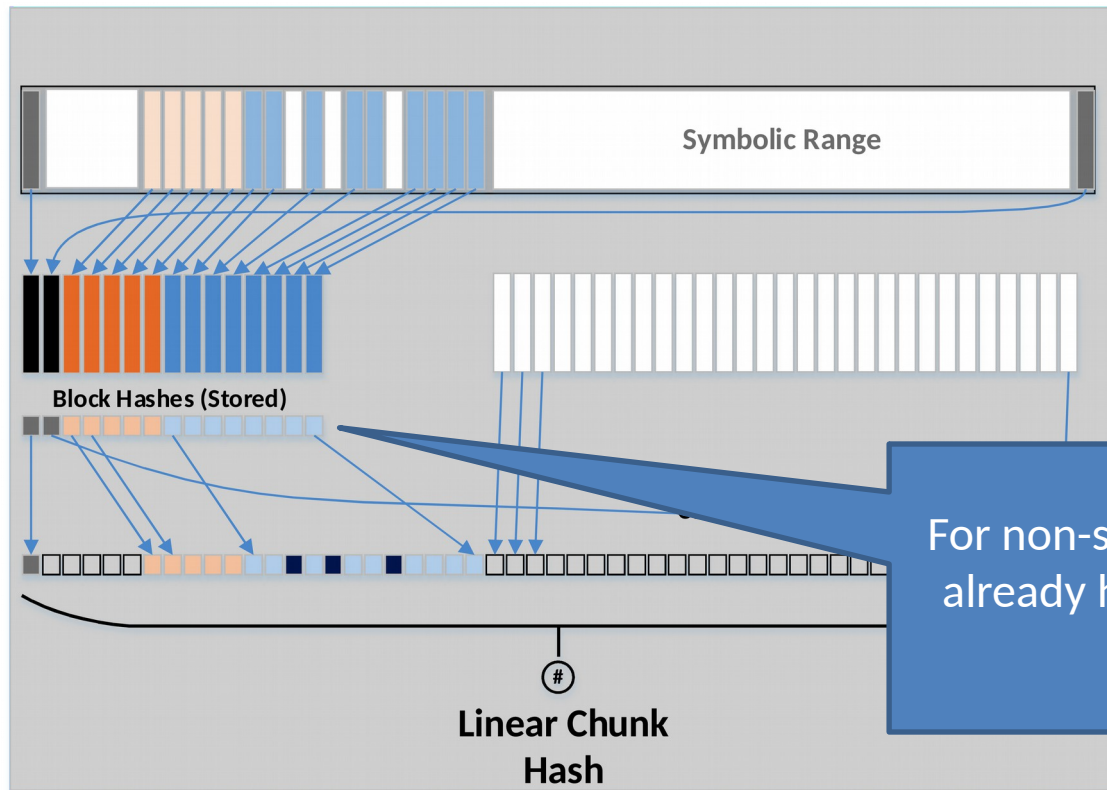
Hypothesis 1a: introduce sub chunk encoding rules

- 4k sector gets written to first 4k of chunk, rest zero filled.
- -or-
- 4k sectors coalesced

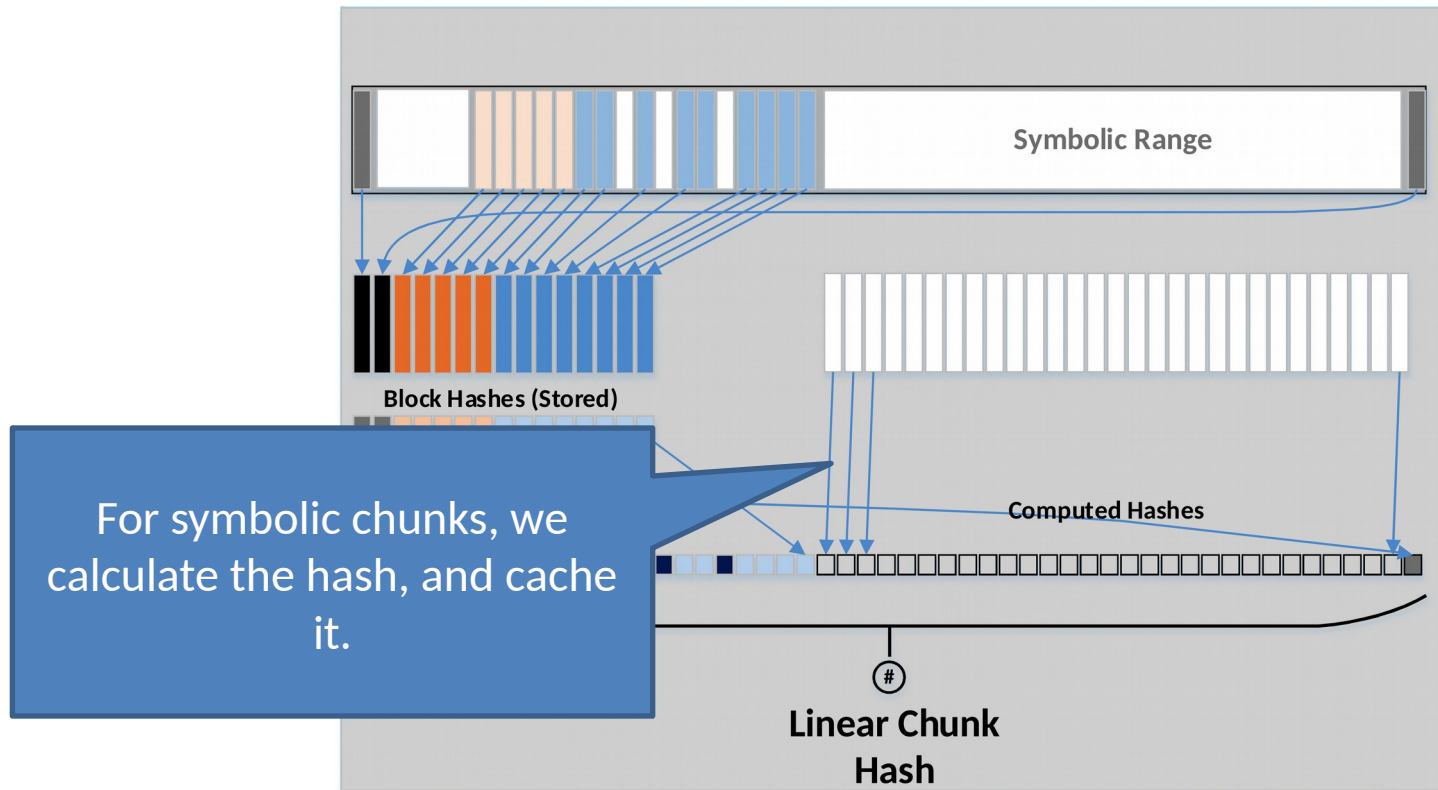
Hypothesis 2: hash all chunks of the source



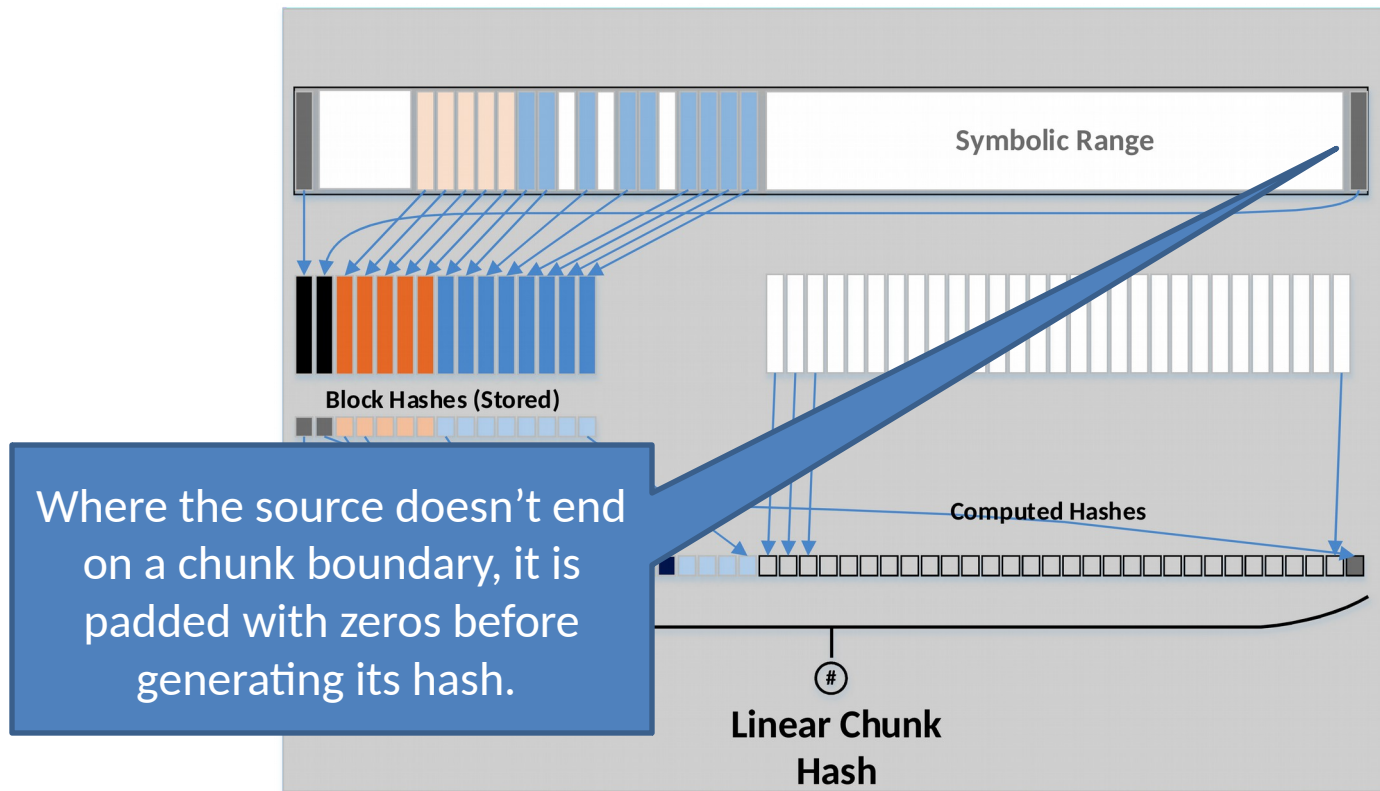
Hypothesis 2: hash all chunks of the source



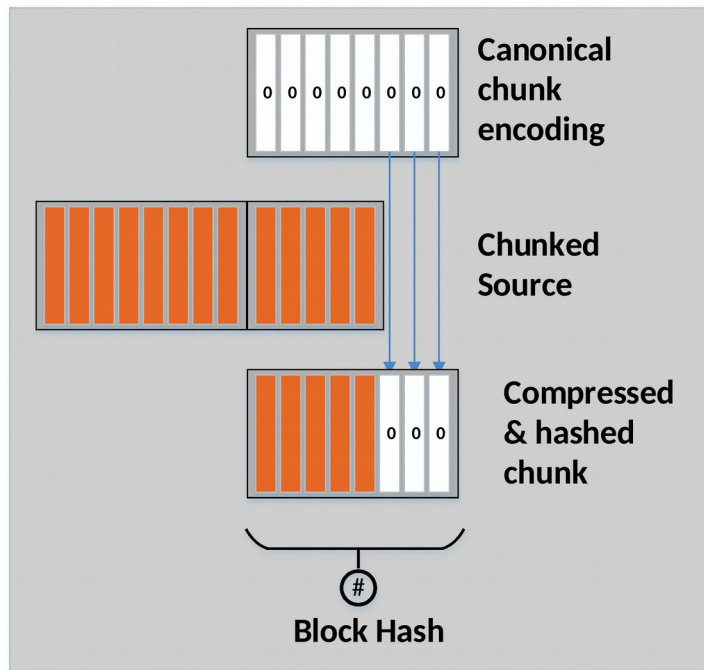
Hypothesis 2: hash all chunks of the source




Hypothesis 2: hash all chunks of the source

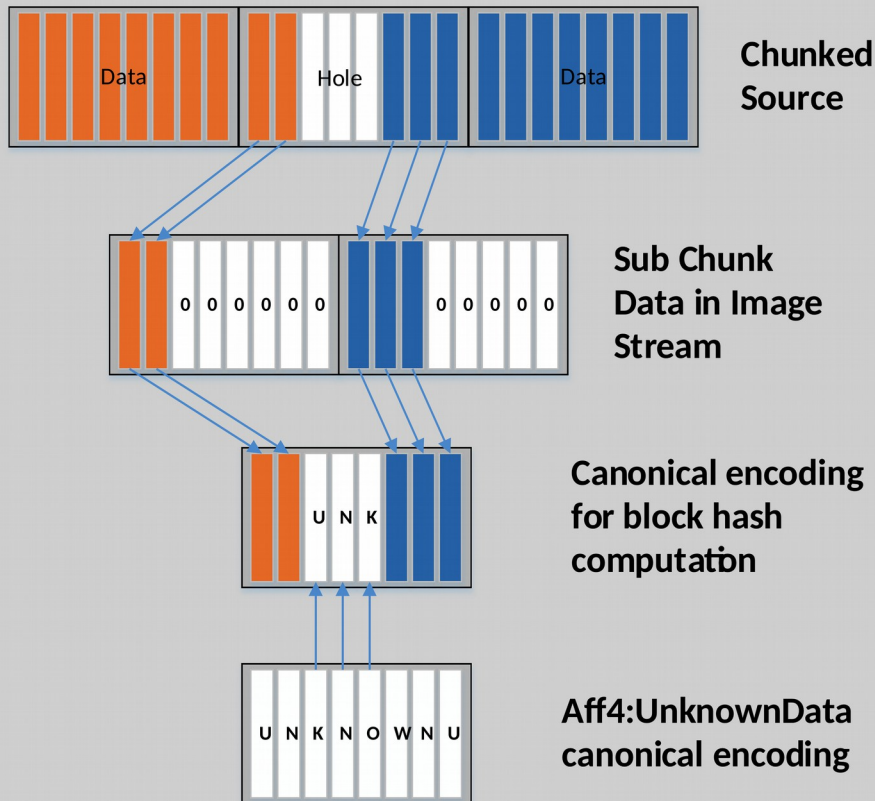


Sub chunk sized end of stream is zero filled



Hypothesis 2: hash all chunks of the source

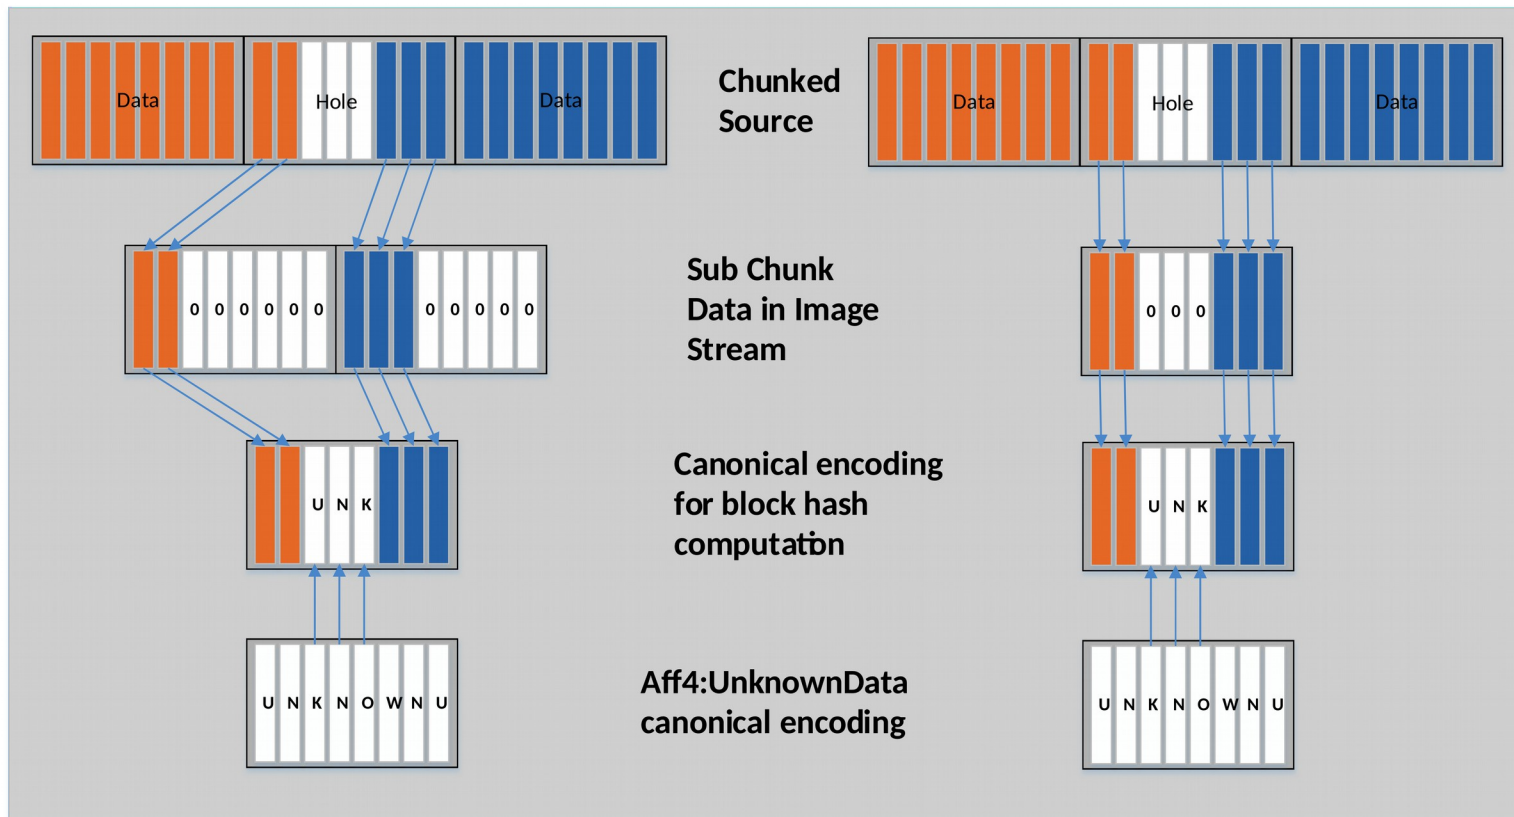




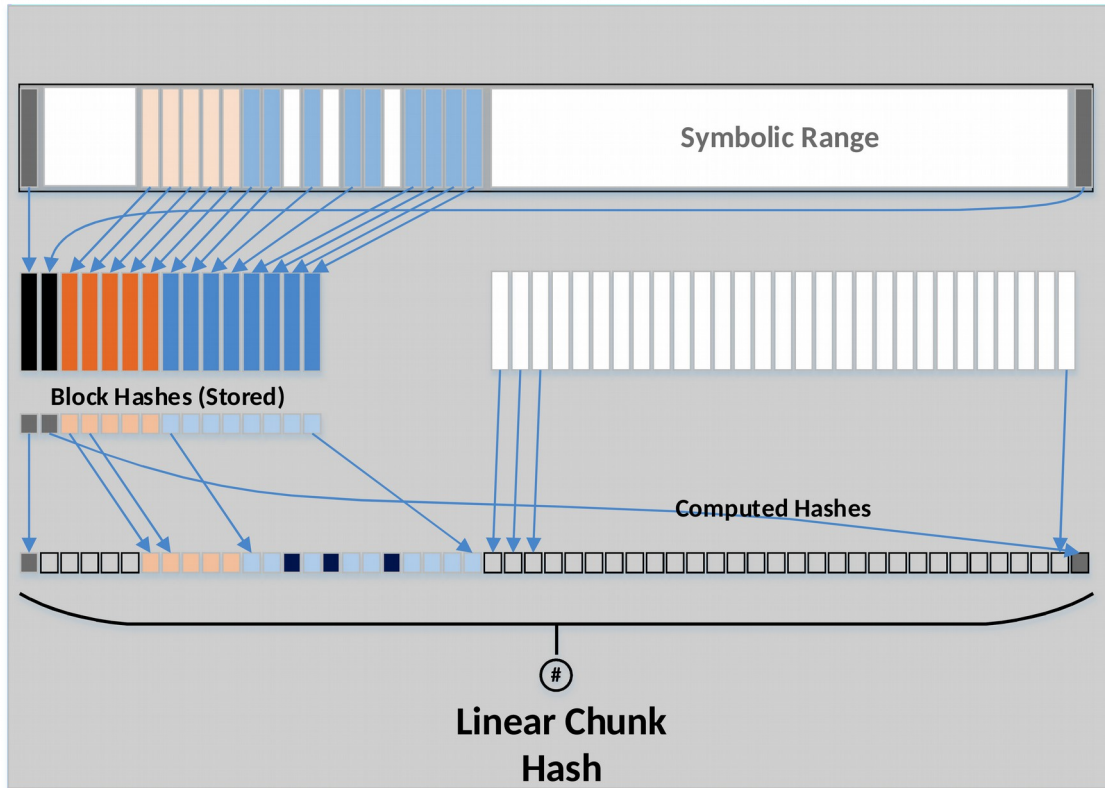
Sub chunk symbolic encoding rules

- Partial chunks stored padded
- Block hash computed from canonical block encoding

Alternate part chunk storage schemes result in the same hash




Hypothesis 2: hash all chunks of the source



Desirable properties of hashing

| | Linear Bitstream Hash | AFF4 Block Map Hash | AFF4 Linear Chunk Hash |
|--|-----------------------|---------------------|------------------------|
| Hash is concise | Yes | Yes | Yes |
| Hash of image matches hash calculated when imaging | Yes | Yes | Yes |
| Changed image results in failed hash match | Yes | Yes | Yes |
| Successive hashes of same source result in same hash | Yes | Maybe | Yes |



Contact

Dr Bradley Schatz

[https://evimetry.com/
bradley@evimetry.com](https://evimetry.com/bradley@evimetry.com)

@blschatz

@wirespeed4n6