
Extending The Sleuth Kit and its Underlying Model for Pooled Storage File System Forensic Analysis

**Fraunhofer Institute for Communication,
Information Processing and Ergonomics**

Jan-Niclas Hilgert*

jan-niclas.hilgert@fkie.fraunhofer.de

Martin Lambertz

martin.lambertz@fkie.fraunhofer.de

Daniel Plohmann

daniel.plohmann@fkie.fraunhofer.de

Digital Forensic Analysis



File Systems

- Define how data is **read from** and **written to** a storage device
- Utilize **metadata** to keep order of the data stored
- File systems **differ** in many aspects

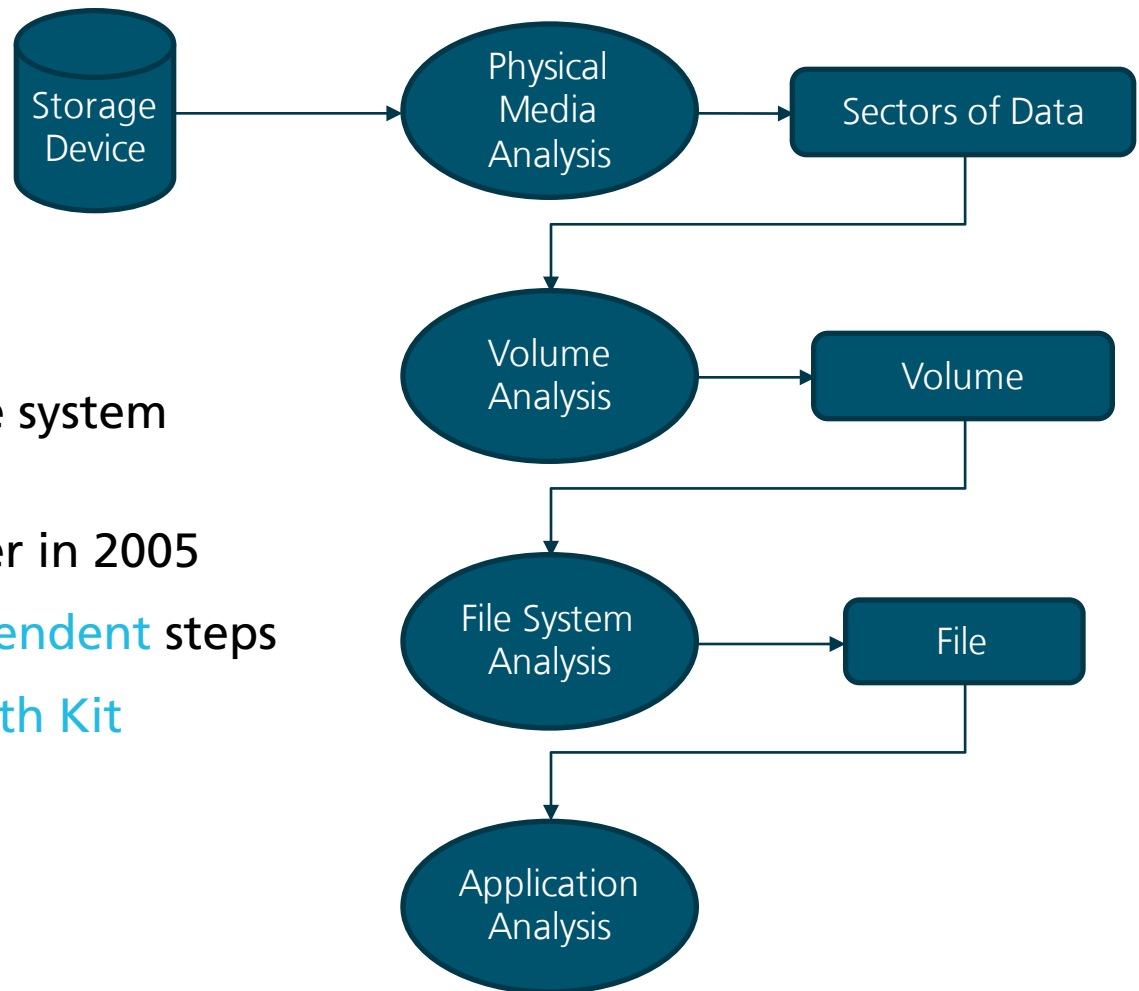


- Extensive background knowledge is required for a forensic analysis
 - But not always existent

File System Forensic Analysis

De-Facto Standard Model

- Universal model for a file system forensic analysis
- Presented by Brian Carrier in 2005
- Consists of four interdependent steps
- Implemented in The Sleuth Kit



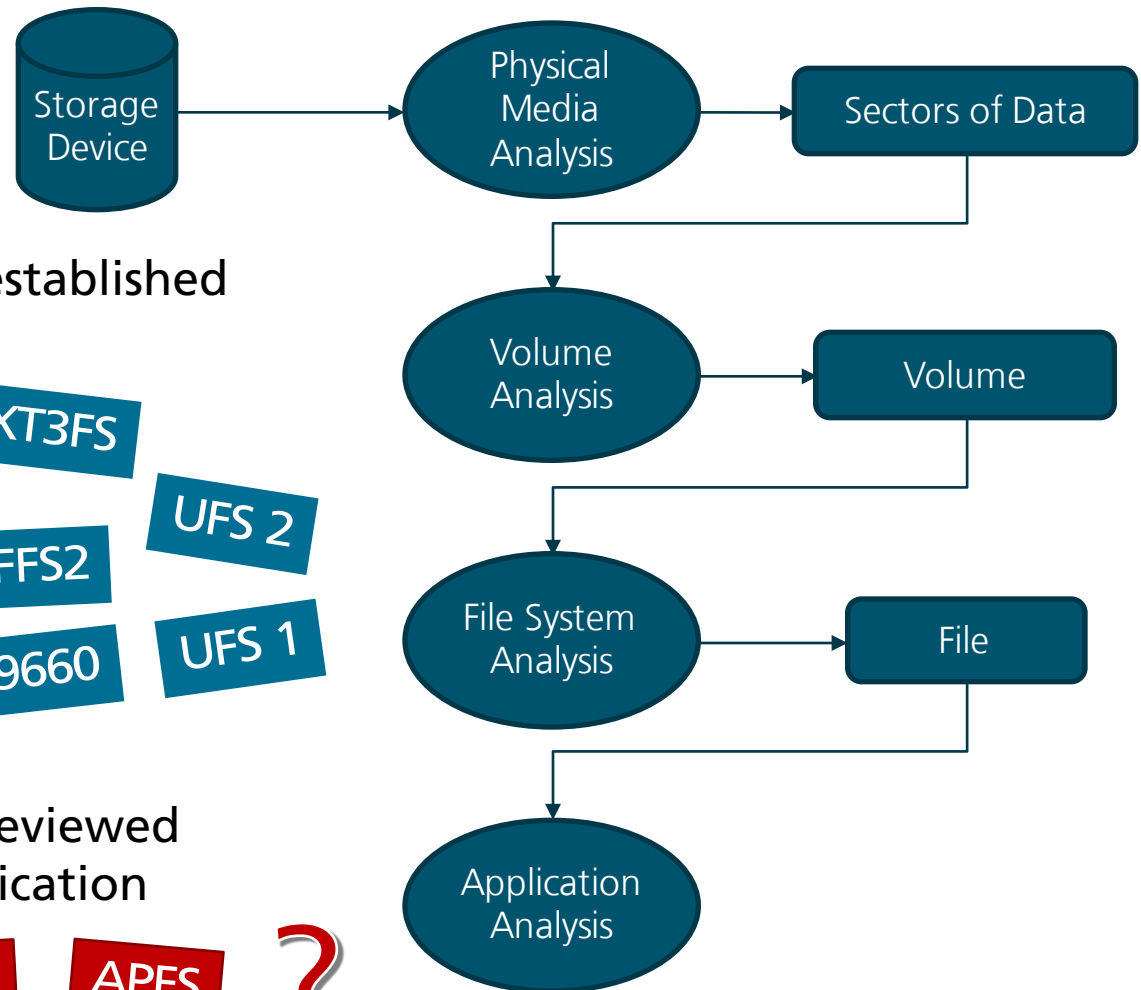
File System Forensic Analysis

De-Facto Standard Model

- Works great on a lot of established file systems



- Problem: It hasn't been reviewed or changed since its publication



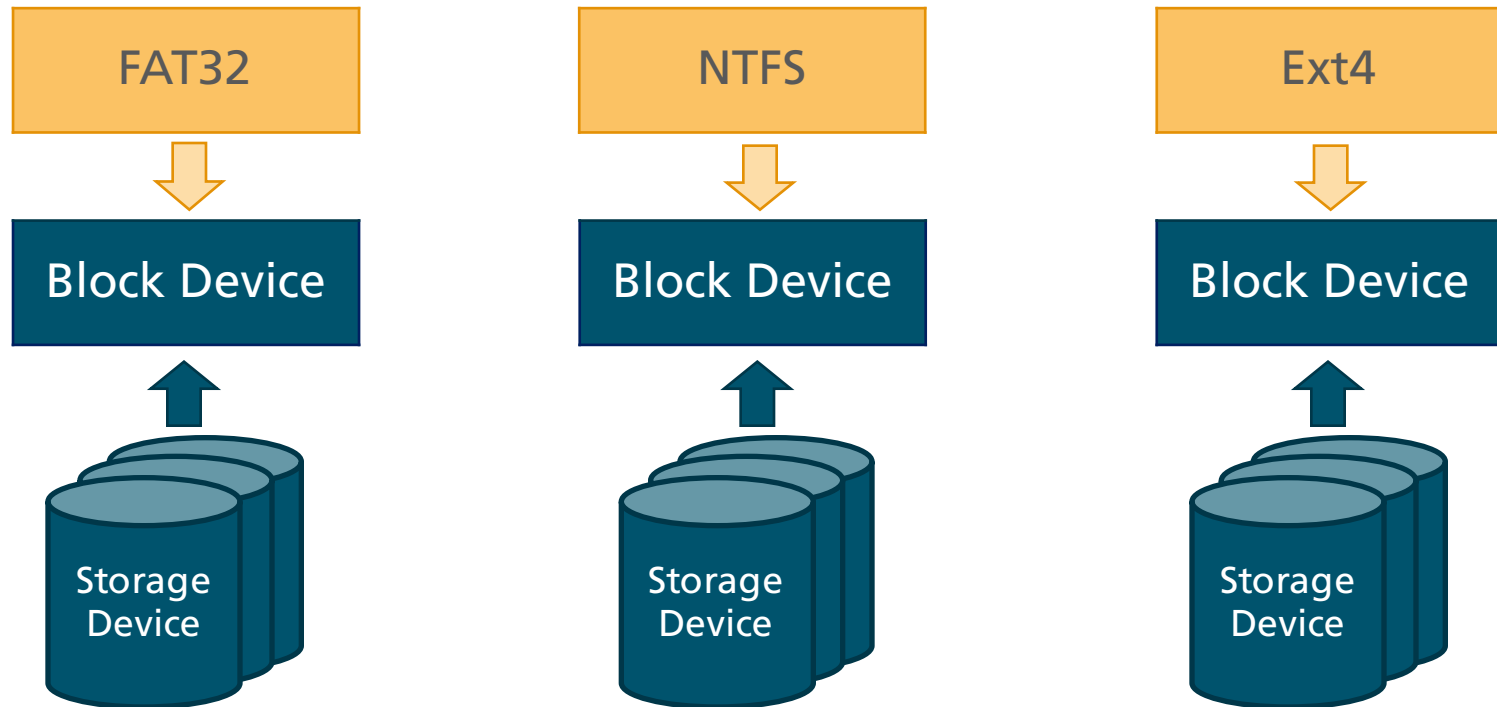
File System Forensic Analysis

Trying TSK on ZFS

```
jhilgert:ZPool$ fsstat disk1  
Cannot determine file system type
```

Pooled Storage File Systems


Old File System Mapping



- Storage devices (hard drives, solid state drives ...) are somehow used to create **new block devices**
- One file system is assigned to one of these block devices

Pooled Storage File Systems


New Approach



- Storage Devices are combined to form a storage pool
- Its configuration depends completely on the implementation
- File systems share the available space from the storage pool


Extension of the Standard Model

Requirements – Detect Pool Members




Extension of the Standard Model

Requirements – Detect Pool Members




Extension of the Standard Model

Requirements – Detect Pool Configurations




Extension of the Standard Model

Requirements – Analyze a Complete Pool




Extension of the Standard Model

Requirements – Access Correct Offsets on Physical Disks




Extension of the Standard Model

Requirements – Access All Filesystem Data



Extension of the Standard Model

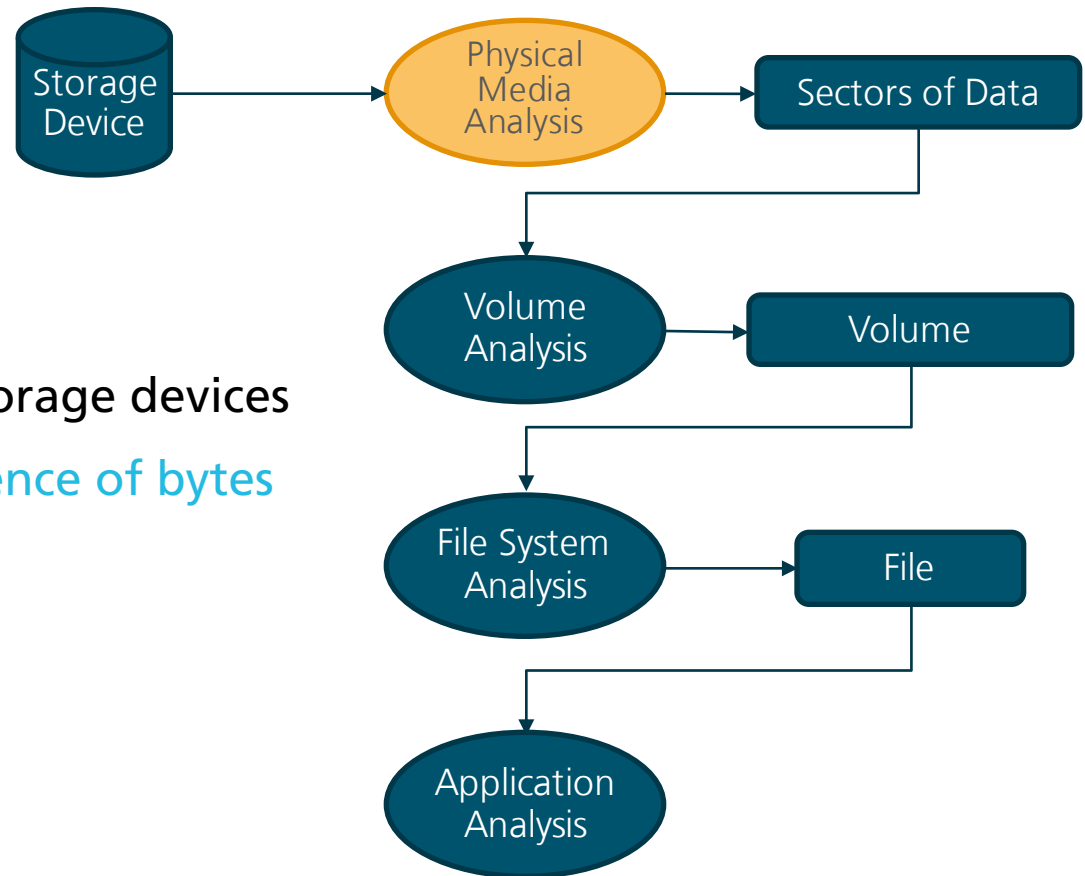
Requirements – Deal with Incomplete Pools



File System Forensic Analysis

Physical Media Analysis

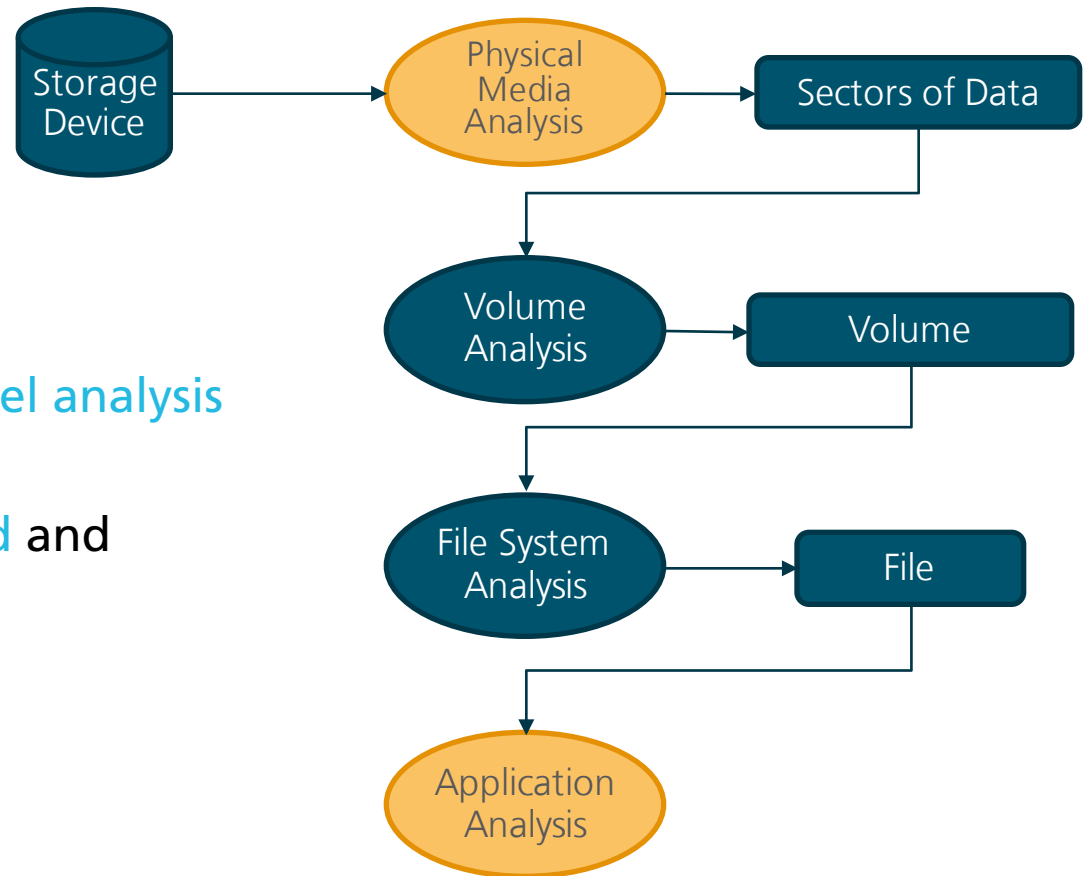
- Acquisition of data from storage devices
- Data is only seen as a sequence of bytes
 - No change necessary



File System Forensic Analysis

Application Analysis

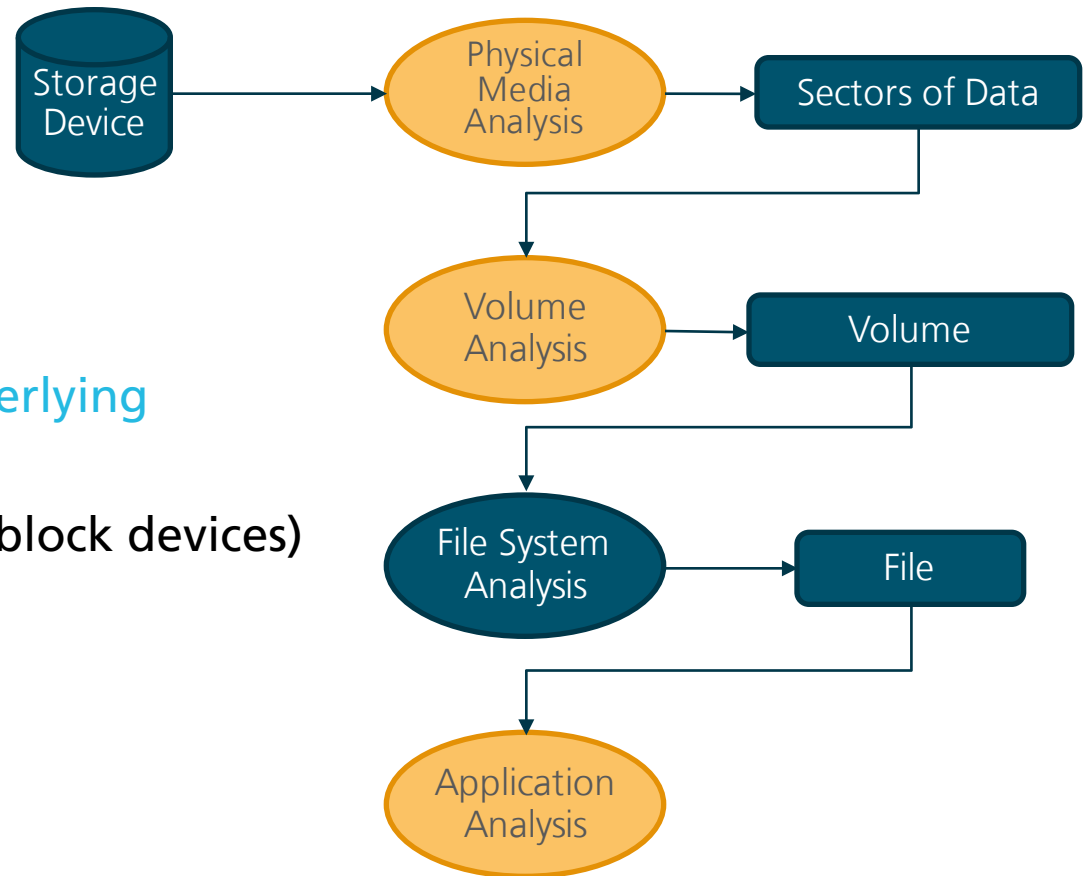
- Performs an **application-level analysis** of the acquired files
- Works on already **recovered** and **collected files**
 - No change necessary



Extension of the Standard Model

Volume Analysis

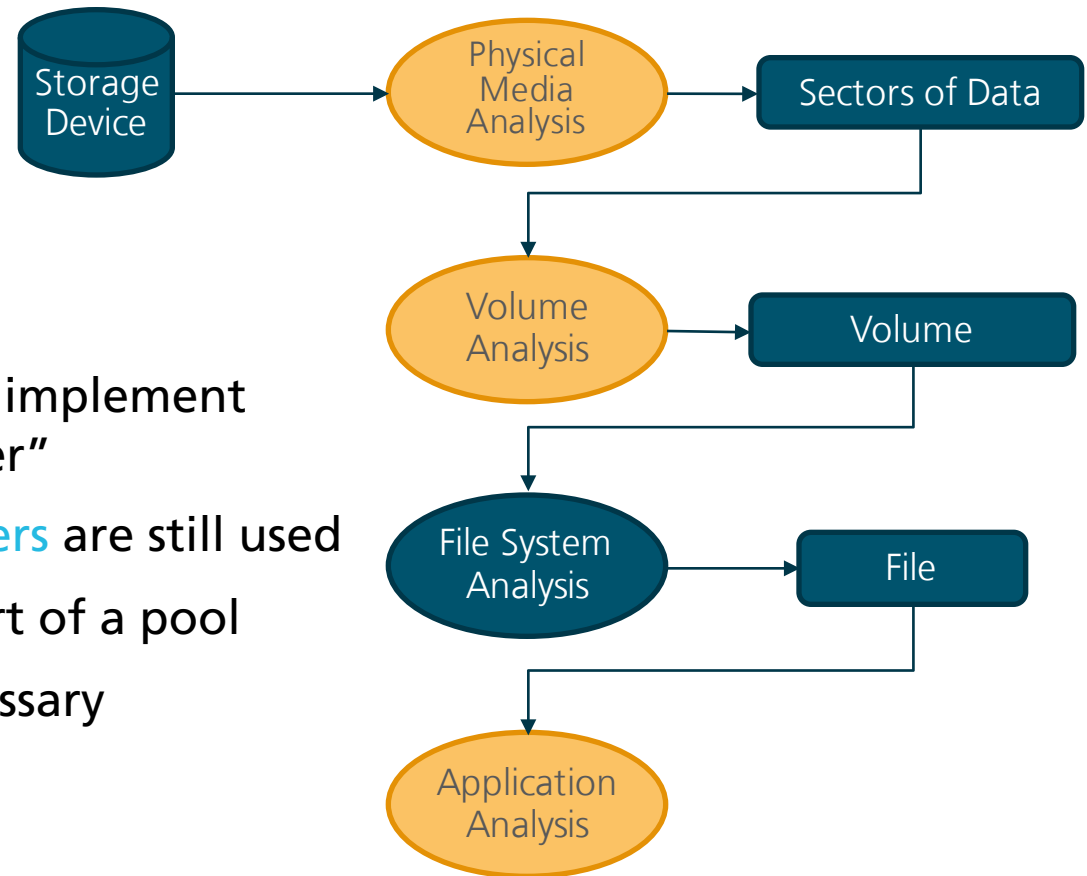
- Data is searched for its **underlying volume structure**
- Returns detected volumes (block devices)



Extension of the Standard Model

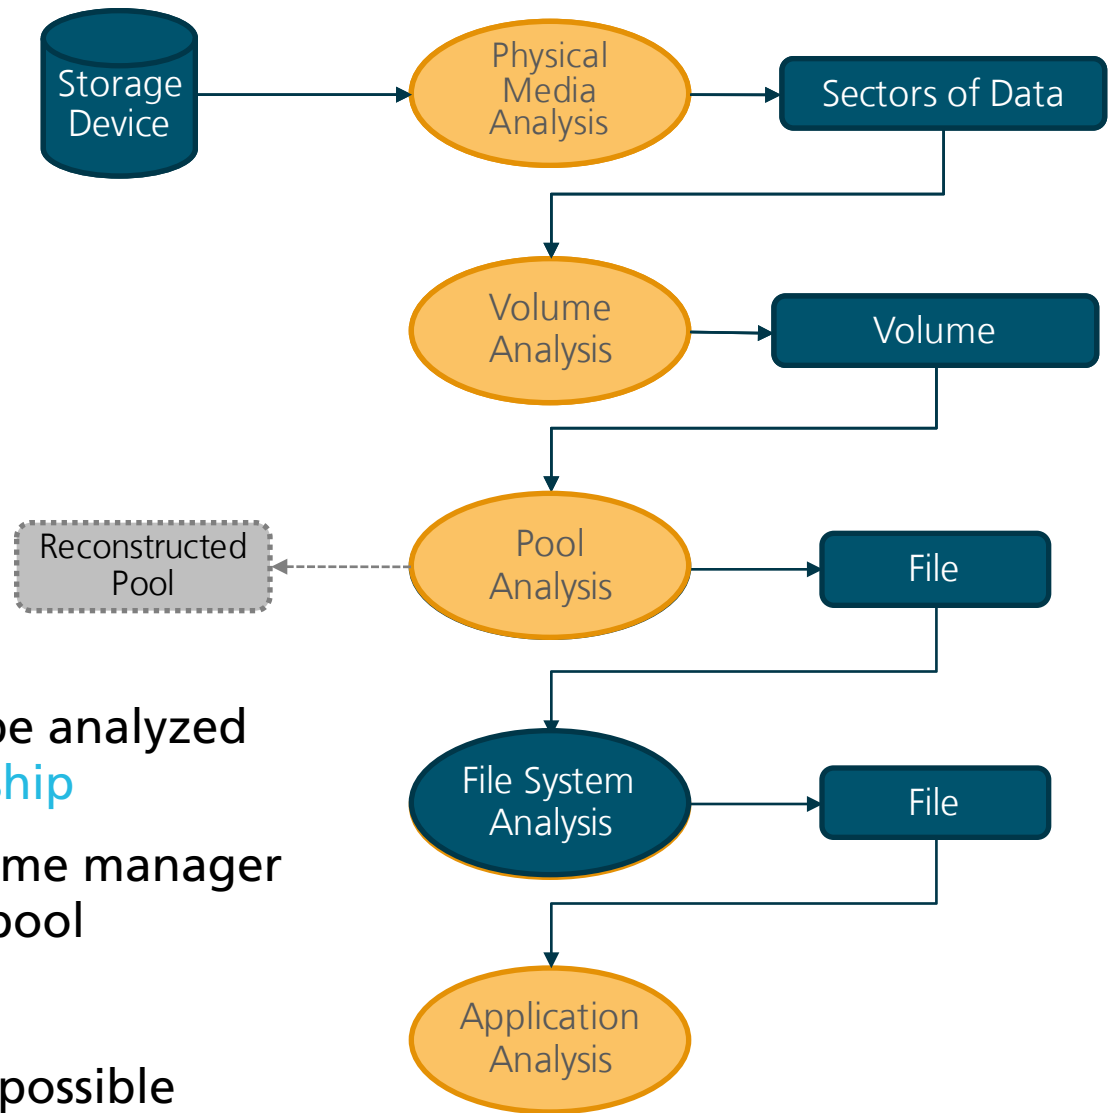
Volume Analysis

- Pooled storage file systems implement their own “volume manager”
- Established volume managers are still used
 - Volumes can also be part of a pool
- Volume analysis is still necessary



Extension of the Standard Model

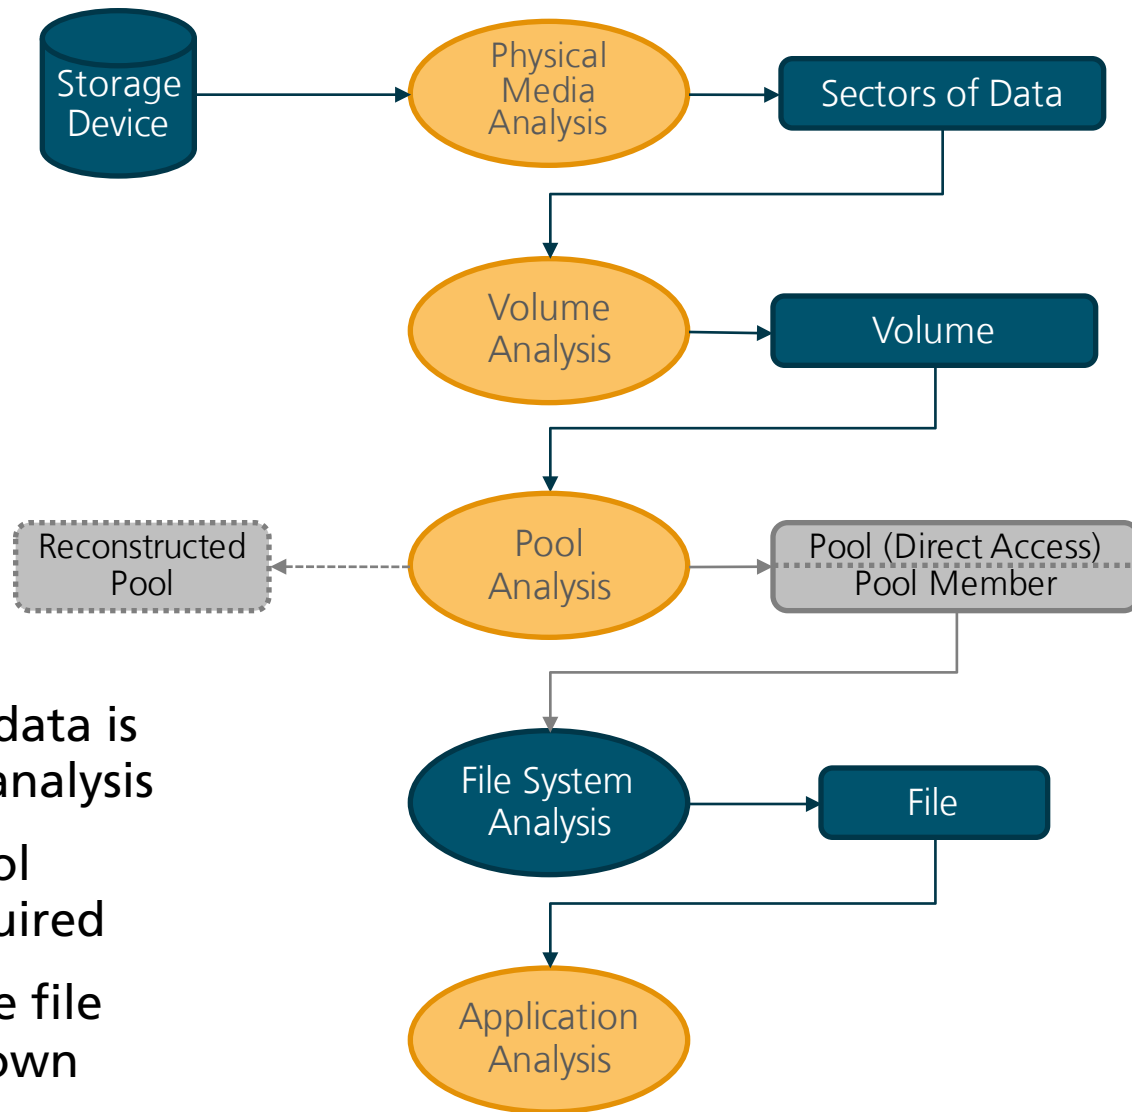
Pool Analysis



- Volumes or data need to be analyzed for possible **pool membership**
- Using the file systems volume manager results in a reconstructed pool
 - Only high-level access
 - No file system analysis possible

Extension of the Standard Model

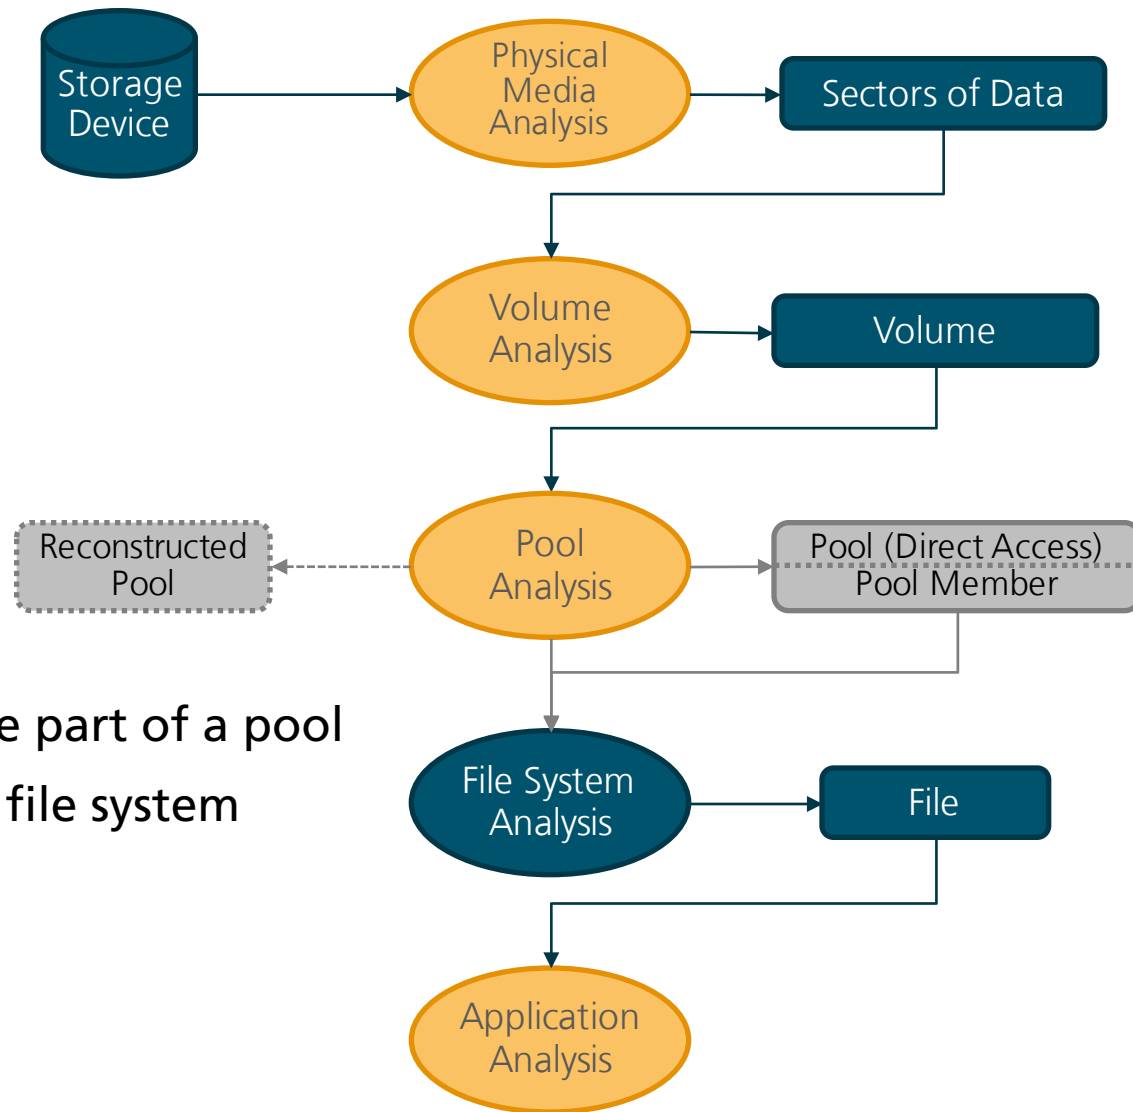
Pool Analysis



- File system data and metadata is essential for a file system analysis
- Direct access to the pool and its members is required
- Mapping schema of the file system needs to be known

Extension of the Standard Model

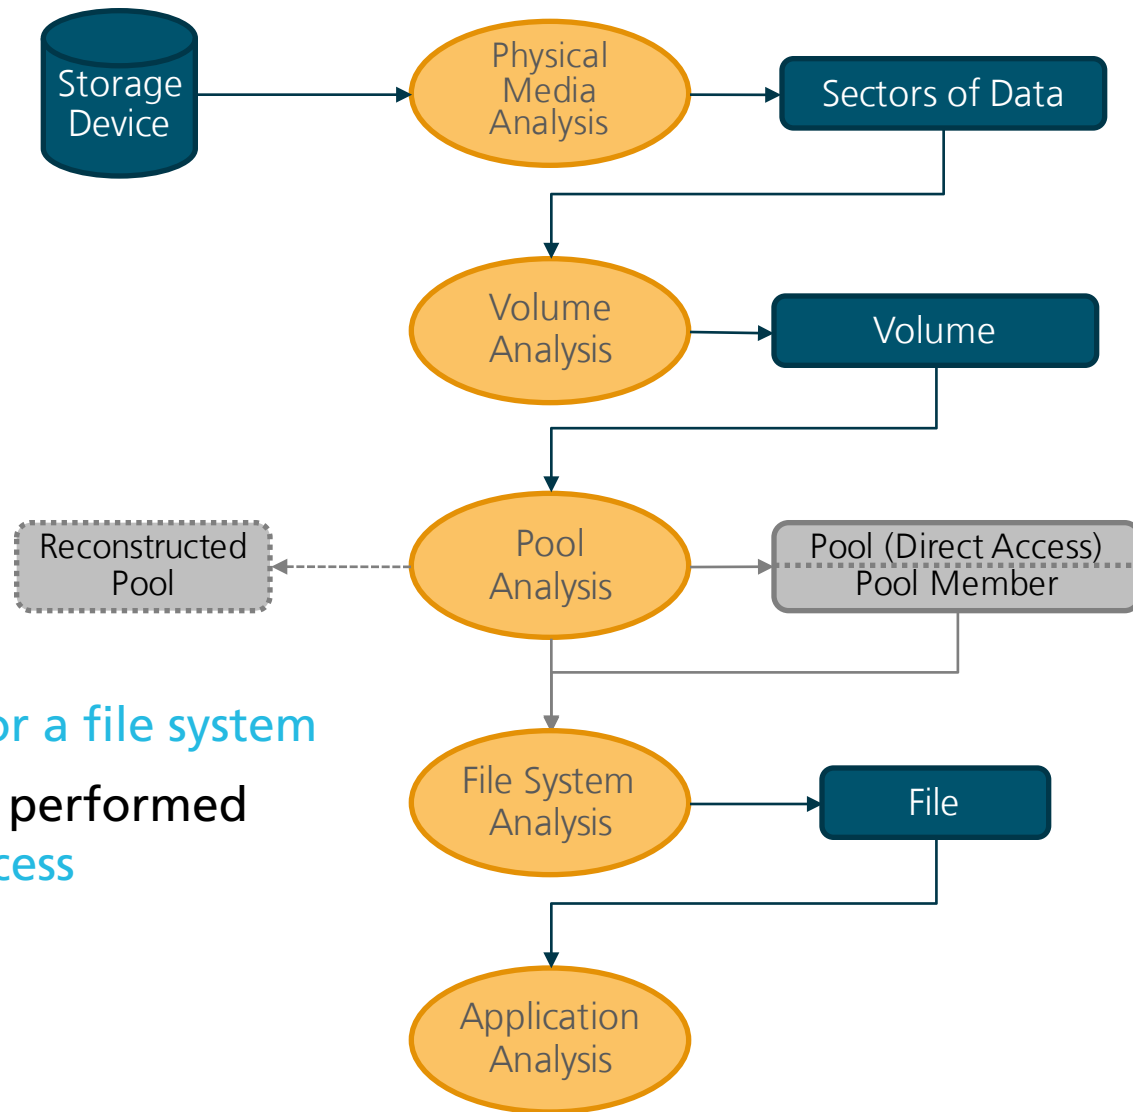
Pool Analysis



- Volumes do not need to be part of a pool
- **Important:** Pool analysis is file system dependent

Extension of the Standard Model

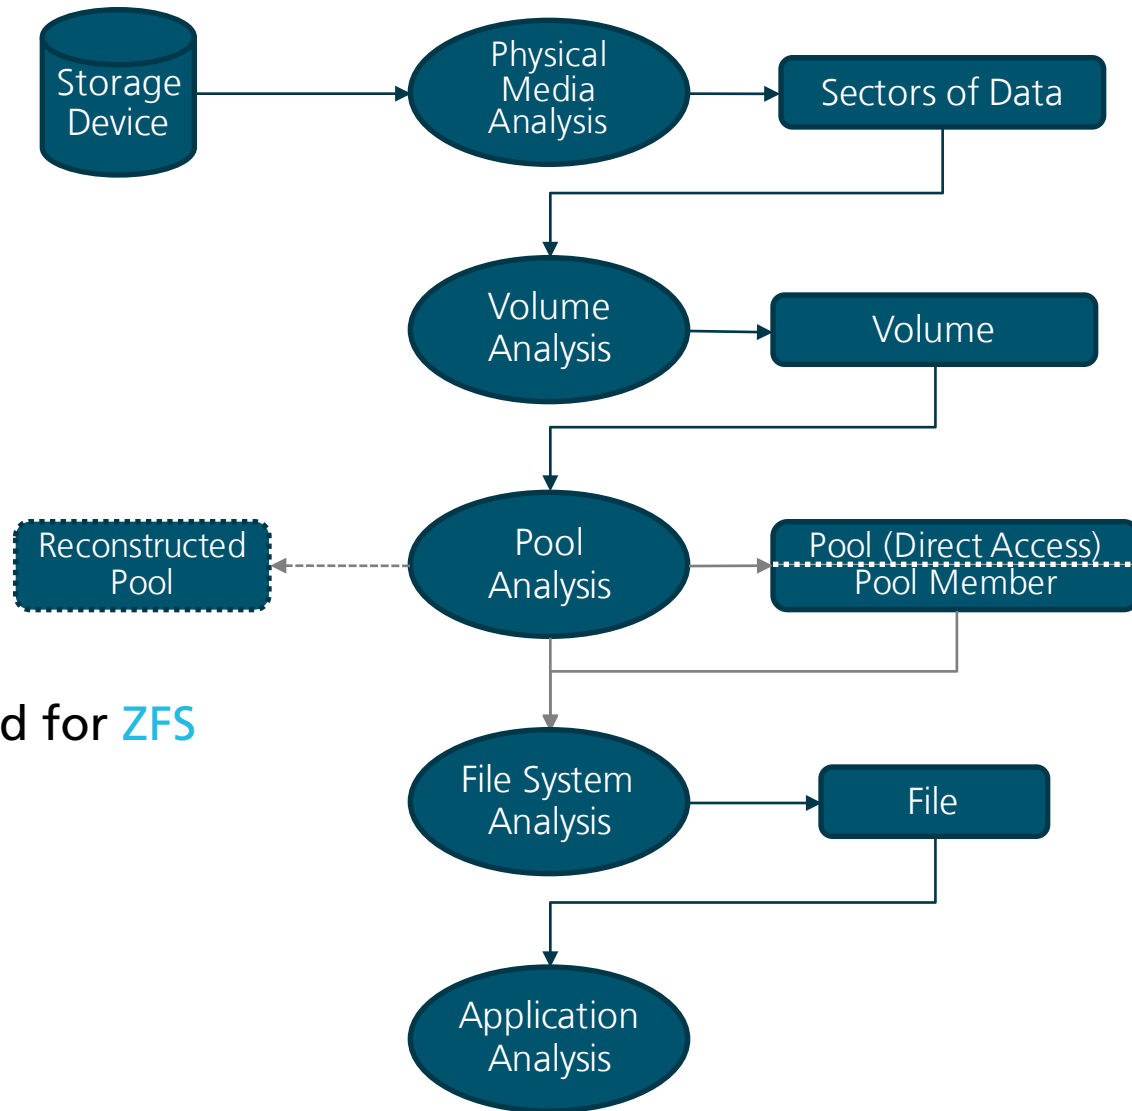
File System Analysis



- Each volume is searched for a file system
- File system analysis can be performed on the pool with direct access

Extension of the Standard Model

Extended Model




- Prototypically implemented for ZFS

Pooled Storage File Systems

ZFS – Volume Management


- First presented in 2003 (more than a decade ago)
- Available for major platforms like Solaris, FreeBSD, Linux, MacOS
- ZFS combines multiple volumes into a storage pool
- Pool members are identified by four vdev-labels



Pooled Storage File Systems

ZFS – Volume Management

- Pools in ZFS consist of one or more **top-level virtual devices (vdevs)**
- Data is striped across all of these top-level vdevs

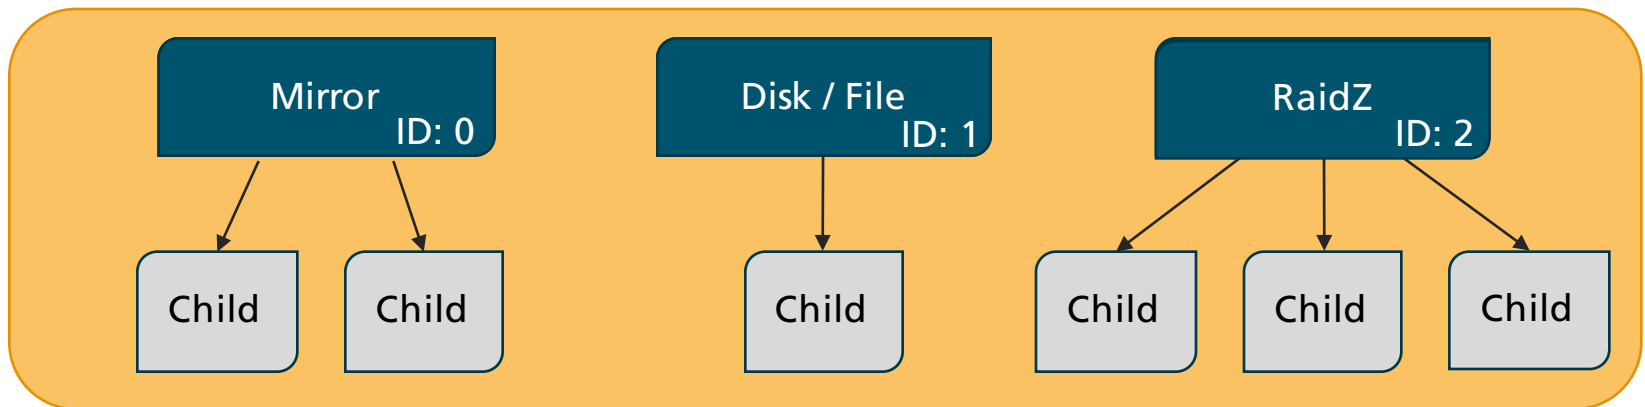


Pooled Storage File Systems

ZFS – Volume Management

- Top-level vdevs are made up of children
- Different types of top-level vdevs are supported by ZFS
- Each top-level vdev is addressed using its ID


ZFS Pool




Pooled Storage File Systems

ZFS – General Structure

- Filesystem is represented in a **tree-structure**
- **Überblock** is the head of this tree
- Utilizes the **copy-on-write** principle




- Block pointers refer to variable-sized blocks of data
- Data is addressed by using **Data Virtual Addresses (DVAs)**



Evaluation

Evaluation Pool

- Use multiple numbers of disks in multiple configurations
 - Two simple disks
 - One mirror with two children
 - One raidz1 with three children



Contribution & Future Research

- Extension of The Sleuth Kit for pooled storage file systems
- Digital forensic analysis of ZFS
- Accessing older versions of the ZFS copy-on-write tree
- Dealing with incomplete pools

- Further investigations on other ZFS structures:
 - ZFS Intent Log (ZIL)
 - L2ARC
- Implementations for other pooled storage file systems
 - BTRFS, ReFS, APFS

Thanks for your attention!

<https://github.com/fkie-cad/sleuthkit>