

# Linux Memory Forensics: Dissecting the User Space Process Heap

Frank Block (a,c) , Andreas Dewald (b,c)

a: ERNW GmbH, Heidelberg, Germany

b: ERNW Research GmbH, Heidelberg, Germany

c: Friedrich-Alexander University Erlangen-Nuremberg (FAU), Germany

# Agenda

## **Introduction**

Motivation

Heap Fundamentals

Evaluation

Demo

Conclusion

## Introduction

- Most of the previous work in the memory forensics area focused on information residing in the kernel space.
- The user space is however also a rich source of artifacts/data, valuable for the investigator, especially the heap.
- Might contain for example passwords or (en/de)rypted data from
  - legitimate software (Mail Client, Password Manager, ...)
  - malware (passwords used for encrypted communication, decrypted binary data, ...)

## Introduction

- For Windows, there is a plugin to analyze the heap in a more detailed way, thanks to the research by Michael Cohen.
  - <http://www.rekall-forensic.com/docs/References/Papers/p1138-cohen.pdf>
- But there was no public available plugin for any heap implementation for Linux (at least as far as we know).

## Introduction

- Existing plugins that investigate the Linux process memory space for artifacts/data, normally search the whole memory dump or the heap/stack as a big blob for a given pattern.
- Depending on the data of interest, there are however not always specific patterns marking data of interest, or these patterns are yet not known to the investigator.

# Agenda

Introduction

**Motivation**

Heap Fundamentals

Evaluation

Demo

Conclusion

## Challenge: Pattern-less data

- Data of interest, which is not identifiable by a special pattern and also not easily distinguishable from data surrounding it.
  - An IP address stored as an integer
  - A password, not consisting of printable characters
  - ...

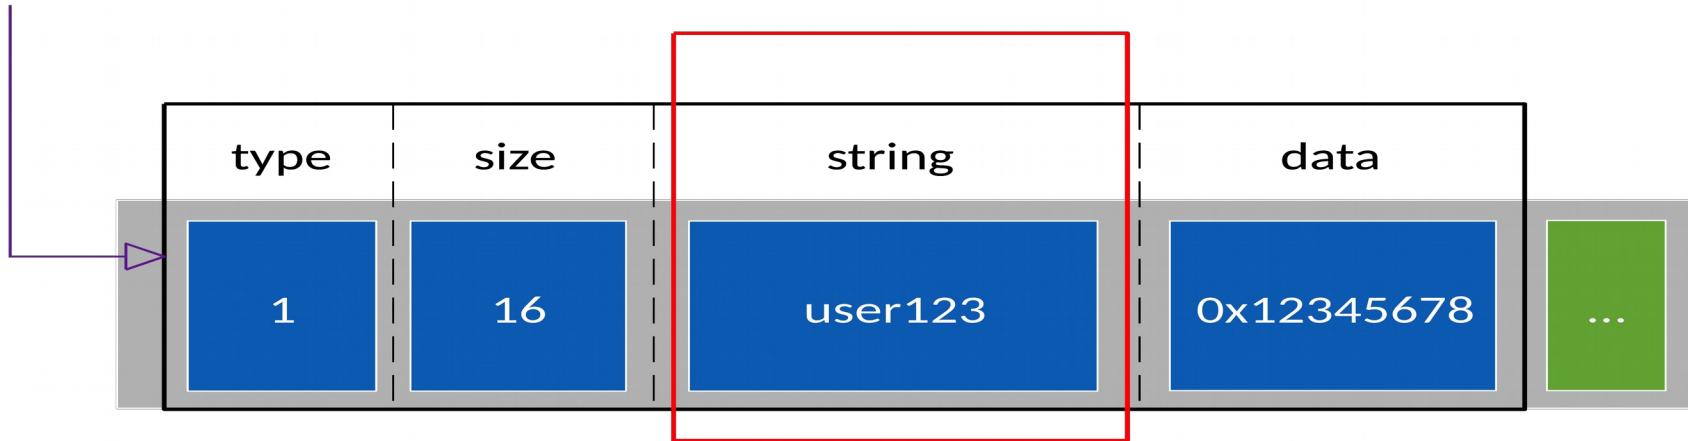
9dcfa4b5f5a9de67**2db4728b05de2197**05deca9575e4f9471a51e5a8627

## Challenge: A pointer to a Struct





## Challenge: A pointer to a Struct



## Goals

- Understand and document **Glibc's** Heap Implementation for forensic purposes.
- Implementation of Rekall plugins that assist an investigator in the analysis process.

# Agenda

Introduction

Motivation

**Heap Fundamentals**

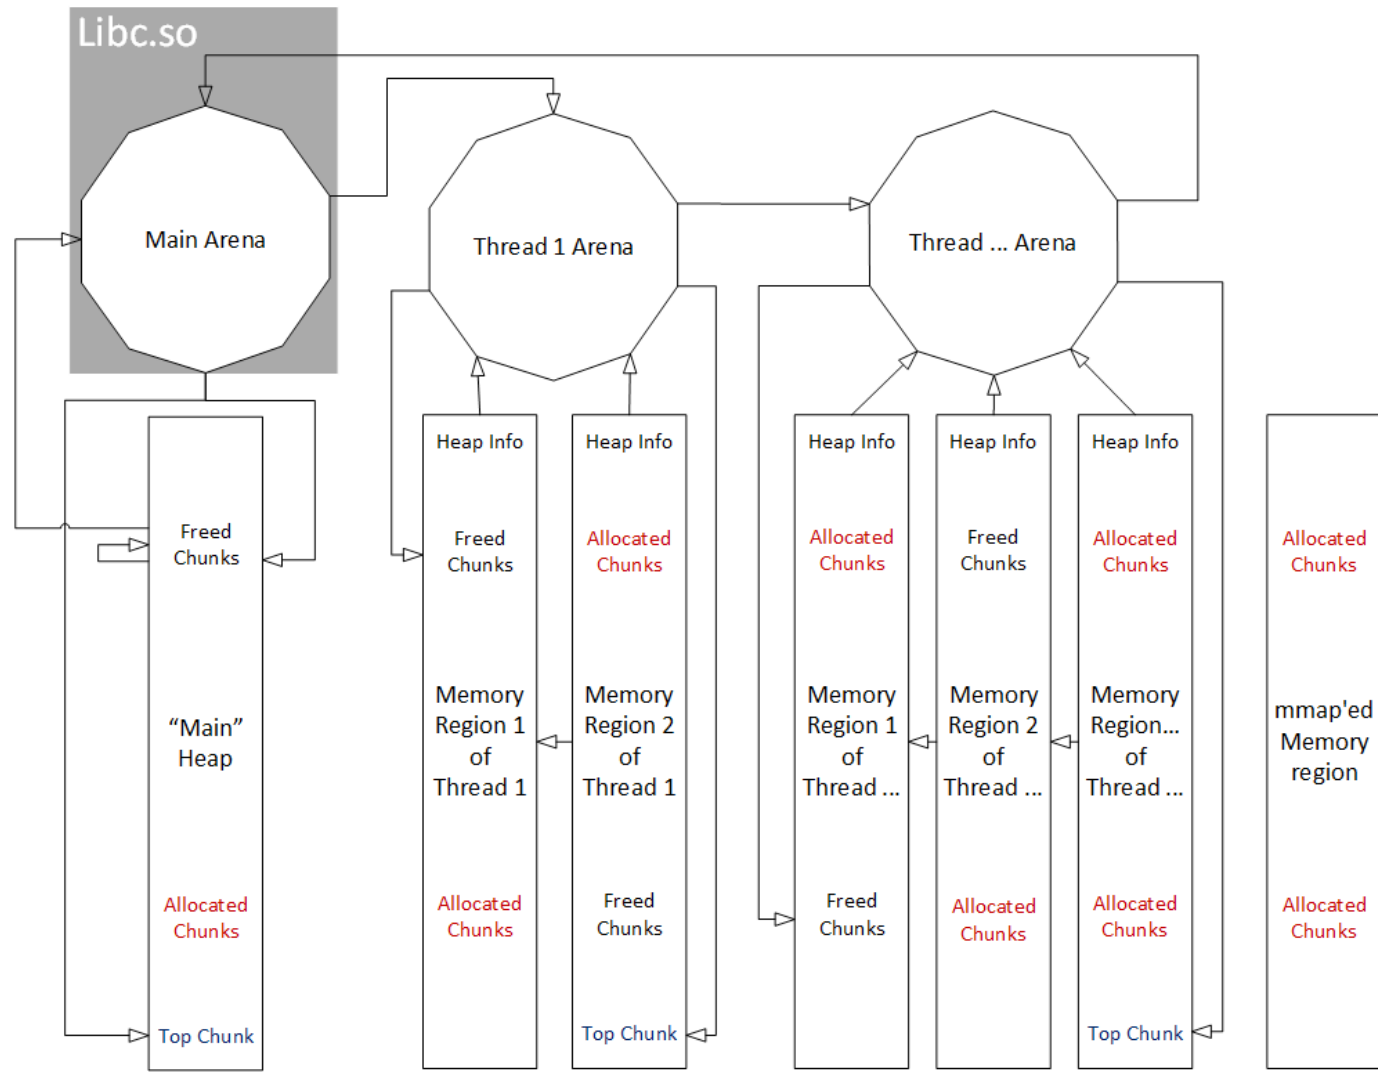
Evaluation

Demo

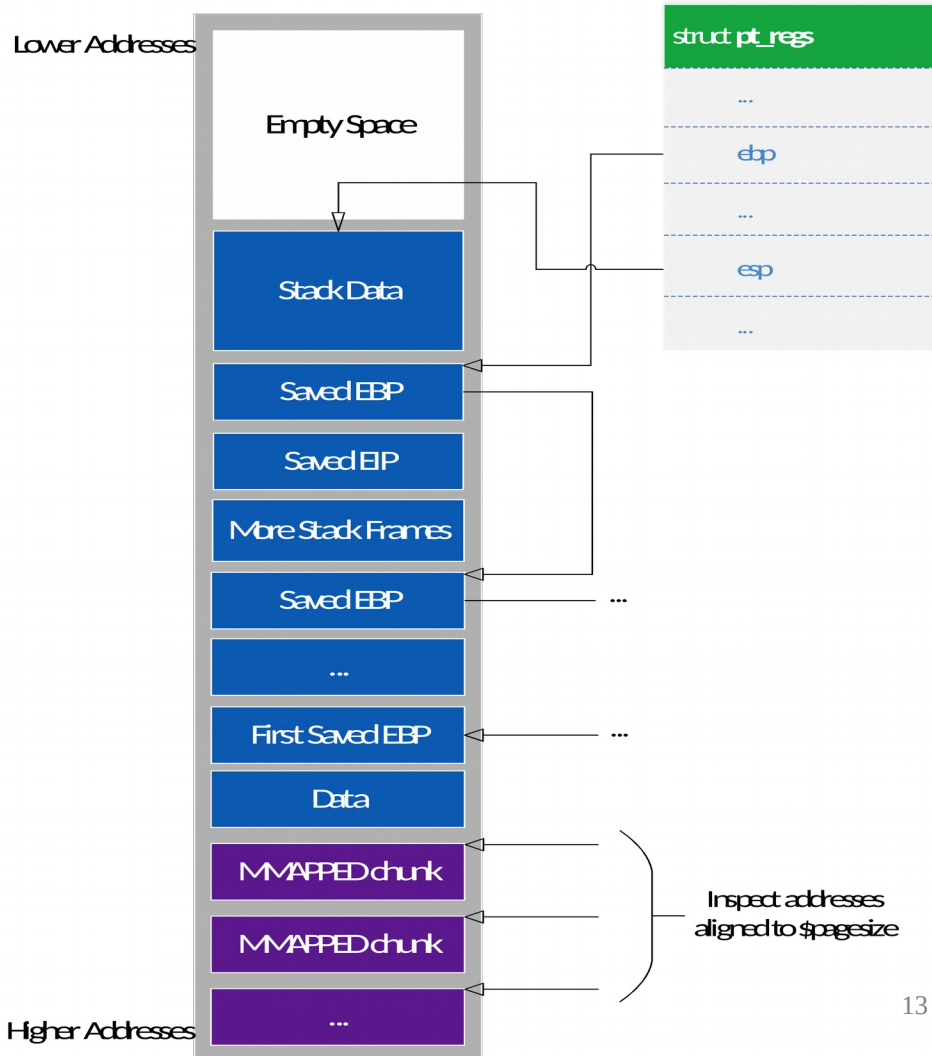
Conclusion

# Heap Fundamentals

## Heap Overview



## EBP unrolling

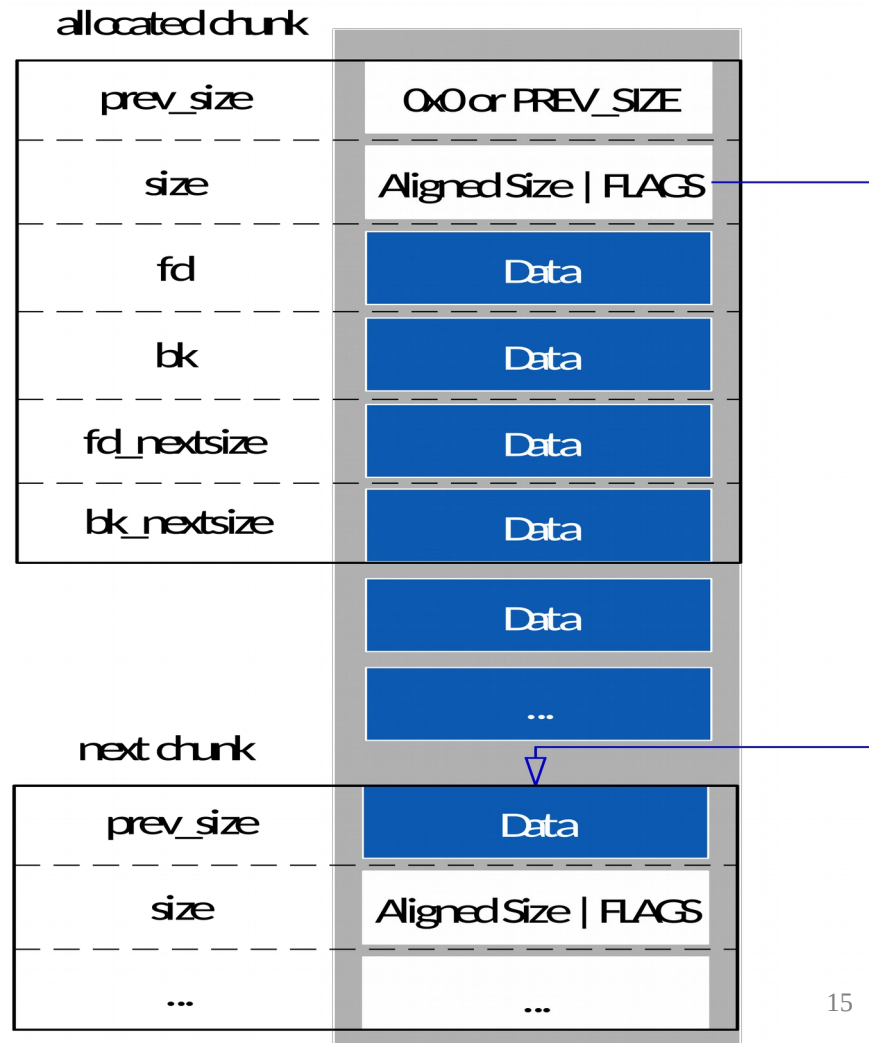


## MMAPPED Chunks

- The identified MMAPPED chunks can be verified, when the offset to the *malloc\_par* struct is given (must be gathered from the Glibc library of the target system).
- This struct contains the number and size of all MMAPPED chunks.

Allocated chunk

In memory



Small bin chunk

prev_size	Data
size	Aligned Size   FLAGS
fd	Next bin chunk
bk	Previous bin chunk
fd_nextsize	Data
bk_nextsize	Data

next chunk

prev_size	PREV SIZE / Overwritten Data
size	Aligned Size   FLAGS
...	...

Large bin chunk

prev_size	Data
size	Aligned Size   FLAGS
fd	Next bin chunk
bk	Previous bin chunk
fd_nextsize	0x0 or next chunk with differing size
bk_nextsize	0x0 or prev. chunk with differing size

next chunk

prev_size	PREV SIZE / Overwritten Data
size	Aligned Size   FLAGS
...	...



## Another Challenge: Extracting Strings from the Heap

**“He said:** I will kill him”

- might become

**non\_printable\_chars** “ I will kill him”

**“Please, do not** kill him”

- might become

**non\_printable\_chars** “I’ll kill him”

# Agenda

Introduction

Motivation

Heap Fundamentals

**Evaluation**

Demo

Conclusion

## Evaluation

- Test Environment: Arch Linux x86/x64, Kernel Version 4.4.5-ARCH with Glibc Versions 2.20 – 2.24
- Self written programs, which use Glibc's *mallinfo* function.
- Analyzing all processes running in the test environments.
- Reading a lot of source code. Mainly from:
  - malloc/malloc.c
  - malloc/malloc.h
  - malloc/arena.c

## Evaluation

- Background checks are performed during the execution of each plugin.
  - For each chunk: Size/address alignment, allocation status, flags, ...
  - Some special tests for MMAPPED chunks, and a comparison to the information from the *malloc\_par* struct.
  - Does following all chunks in a memory region lead to the expected end?
  - Comparison of the size information from all arenas with all objects from the heap.

Demo Time?

# Agenda

Introduction

Motivation

Heap Fundamentals

Evaluation

Demo

**Conclusion**

## Conclusion

- The plugins ease the analysis process and serve more reliable information.
- They enable to gather all parts belonging to the heap (e.g. also hidden MMAPPED chunks).
- Limitations
  - Only Glibc support (officially, versions 2.20 – 2.25).
  - MMAPPED chunks without debug symbols can't reliably be identified.
  - Swapped pages
- Future Work
  - Add linked list detector to the heaprefs plugin.
  - Keep up with new Glibc versions.

## Links

- The paper:
  - <https://authors.elsevier.com/sd/article/S1742287617301895>
- The technical report (contains more heap details):
  - <https://opus4.kobv.de/opus4-fau/frontdoor/index/index/docId/8340>
- The plugins are merged in the Rekall project:
  - <https://github.com/google/rekall>
- Blogpost, containing a brief manual:
  - <https://insinuator.net/2017/07/release-of-glibc-heap-analysis-plugins-for-rekall/>



Thank you for your Attention!

Questions? Feedback? Suggestions? Criticism?