

CGC Monitor

A vetting system for the DARPA Cyber Grand
Challenge

Michael F. Thompson & Timothy Vidas (@tvidas)

CGC Monitor: Presentation outline

DARPA Cyber Grand Challenge overview

Motivation for infrastructure integrity assurances (proactive forensics)

Software vetting on a full system emulator

Running a computer backwards to analyze vulnerabilities

CTF?

What is CTF in this context?

- A cyber security based Capture-the-Flag contest (aka exercise, event, game)
- Typically these contests involve demonstrating proficiency or excellence in one or more areas of computer and network security
- There are different models for architecting these contests, which can stress different skills, lend to particular objectives
- Increasingly popular, common

It is not:

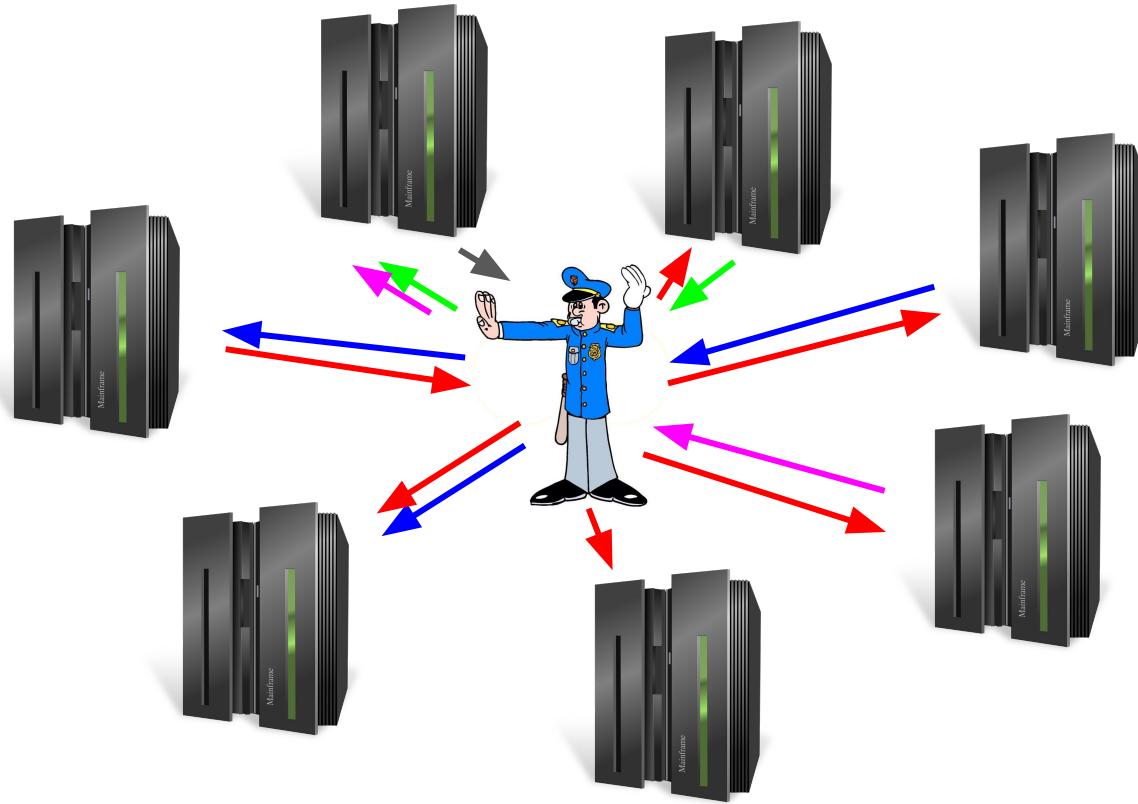
- A game kids play with physical flags on hills
- A first-person shooter video game CTF (usually)
- Focused in the field of Social Engineering
- A hackathon

Though there are certainly similarities to these other games.

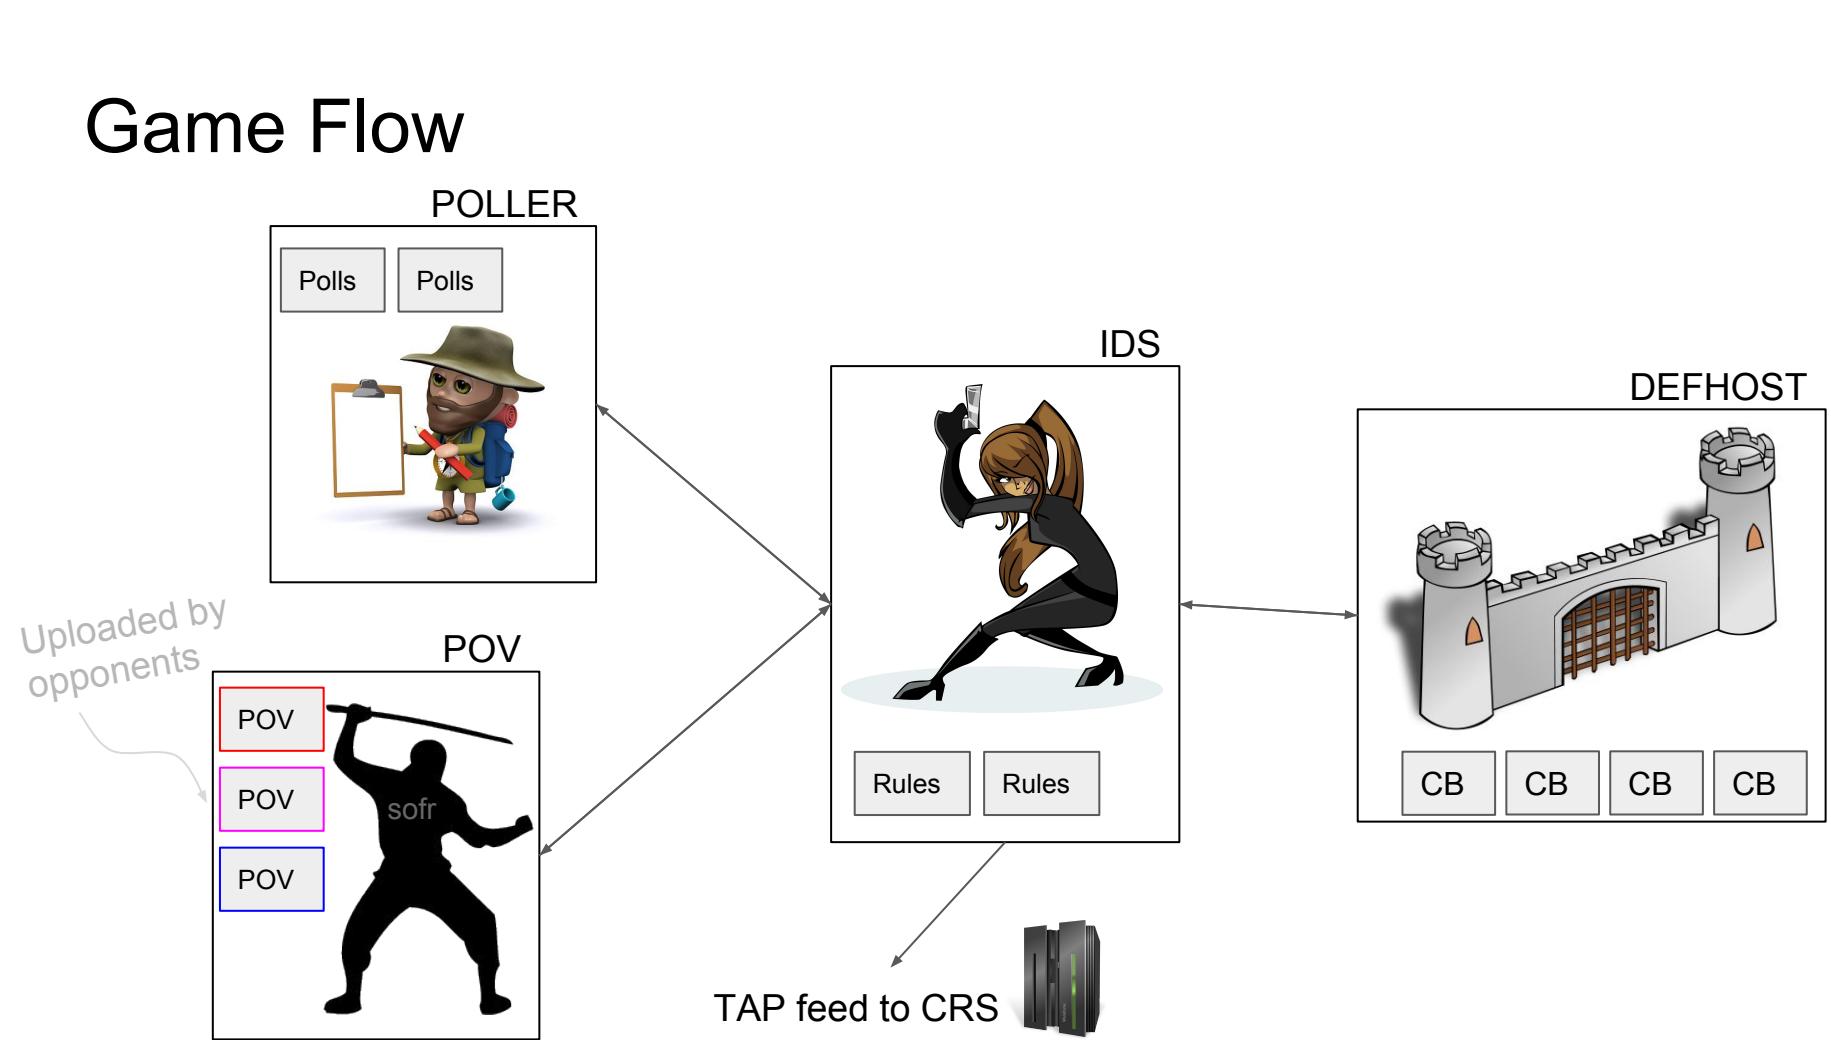
Today, the characters “CTF” are appended to many contests, in most cases this simply means “contest,” sometimes there are flags involved

Game Flow

- **Brokered game**
 - Infrastructure mediates everything
 - API designed for autonomous systems
- **Download** binary software
- **Upload** binary software (replacements)
- Register “moves” against targets



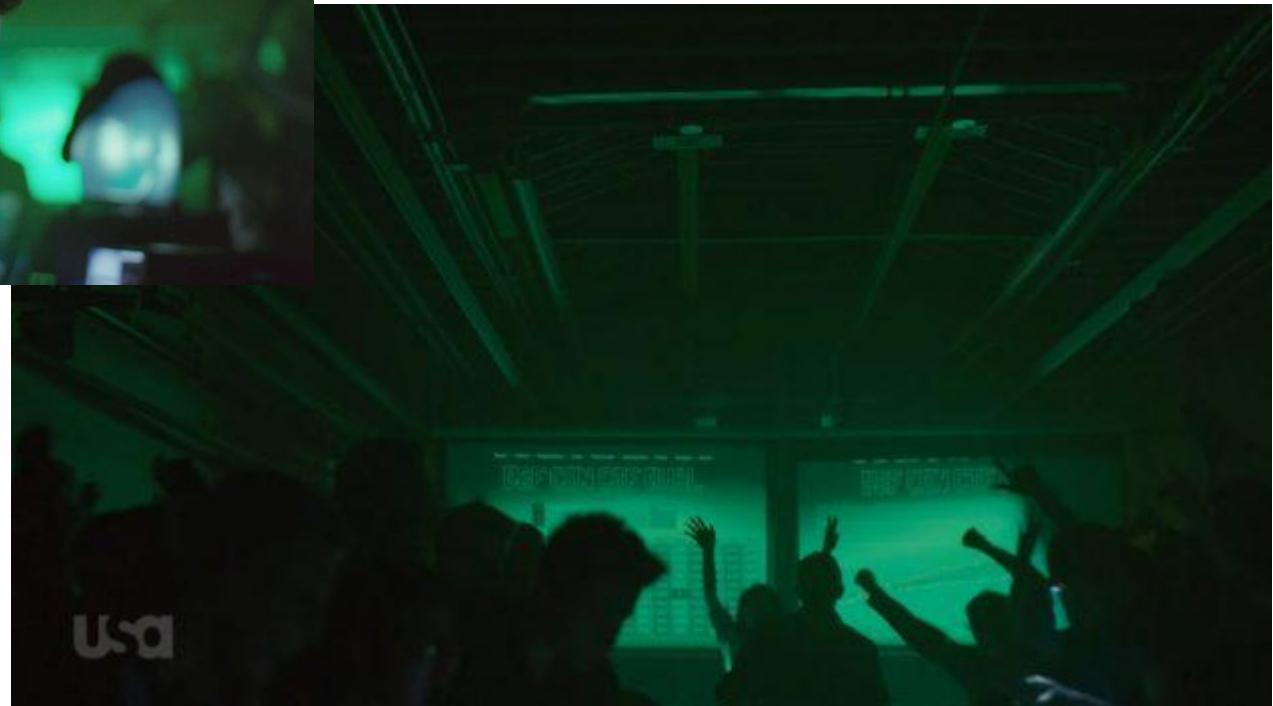
Game Flow

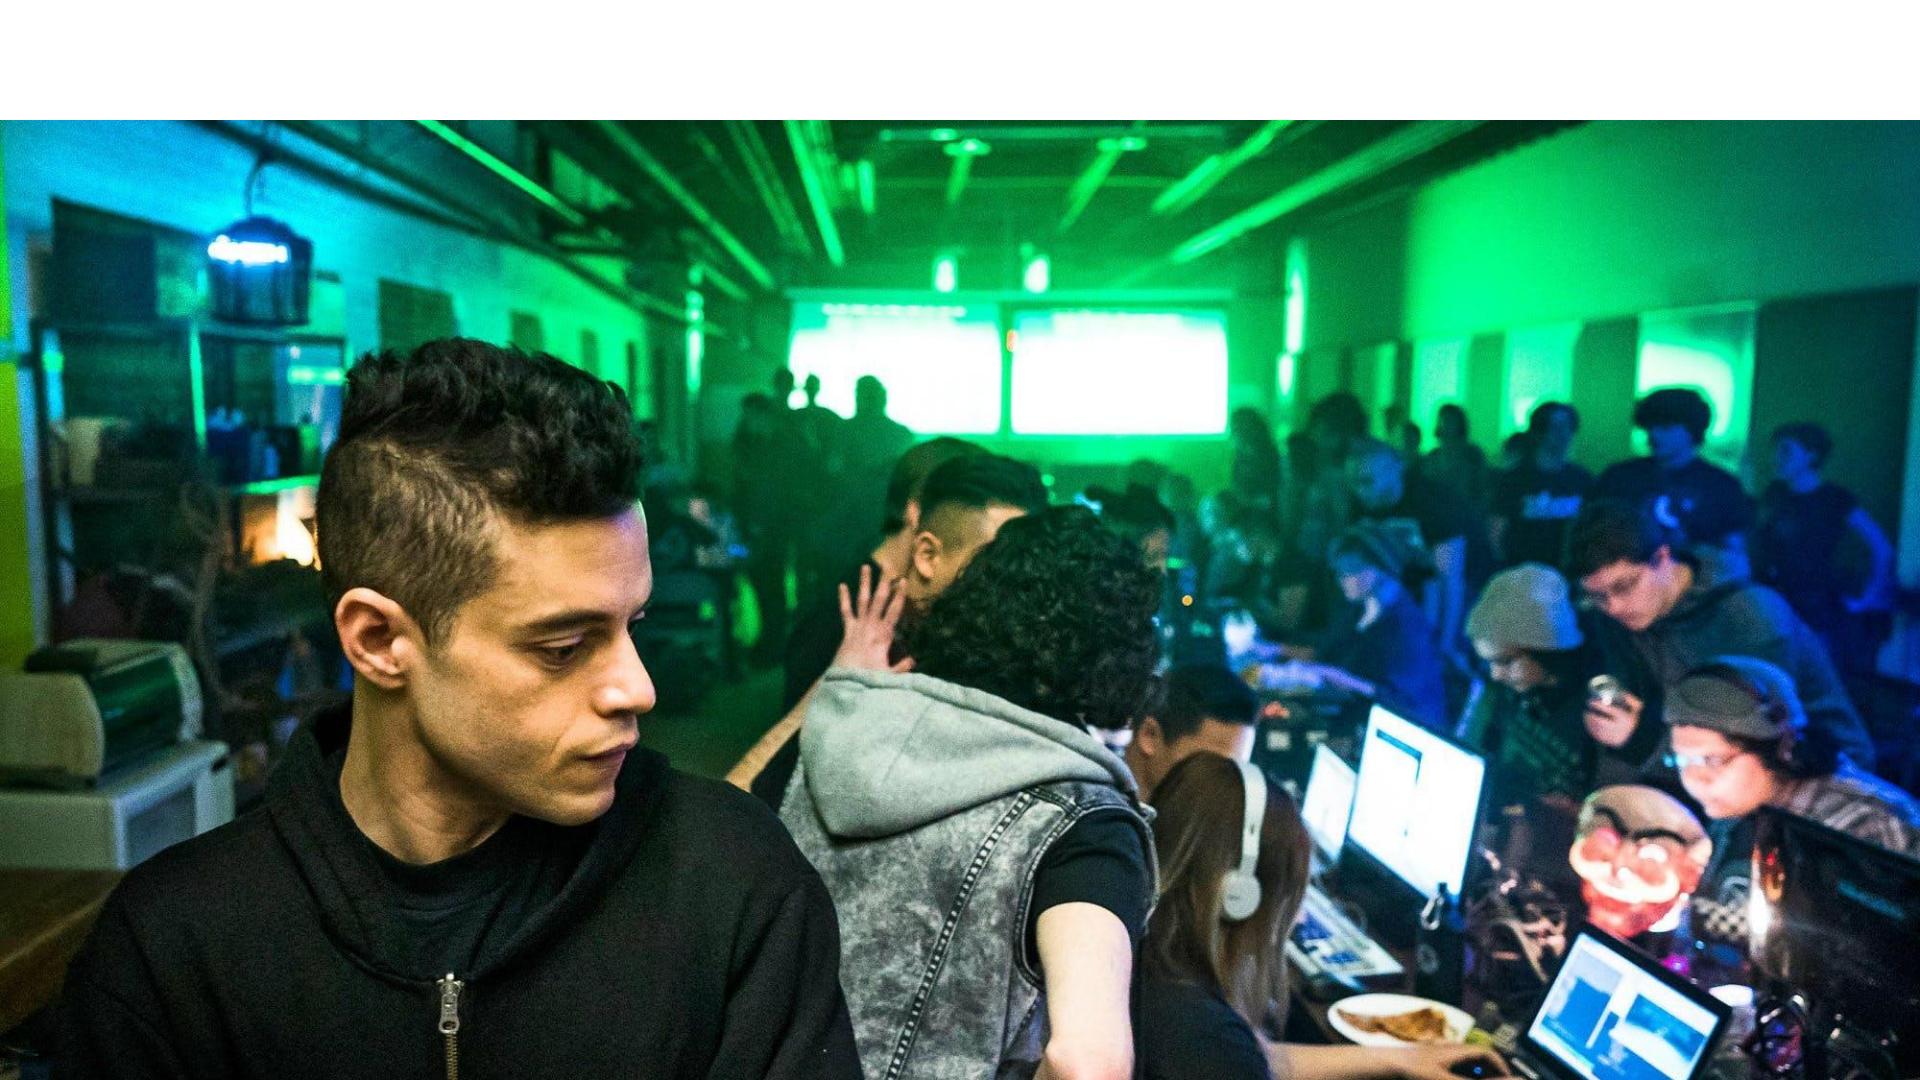


CTF: Hollywood style (well, USA Network)



USA Network 2017





CTF: real life

DEF CON 2002



DEF CON 2016

DEF CON's CTF is often cited as the
“world series” or “superbowl” of CTFs



CTF: real life

A screenshot of a CTF scoreboard interface. At the top right, it says "Score: 4000" and "Logout: TheUnc". Below this is a navigation bar with five categories: "Binary Leetness", "Forensics", "Real World", "Potent Pwnables", and "Trivia". The main area contains a grid of 25 boxes representing challenges. The grid is organized into five rows and five columns. The first four columns have a green background, while the fifth column has a dark grey background. The values in the boxes are as follows:

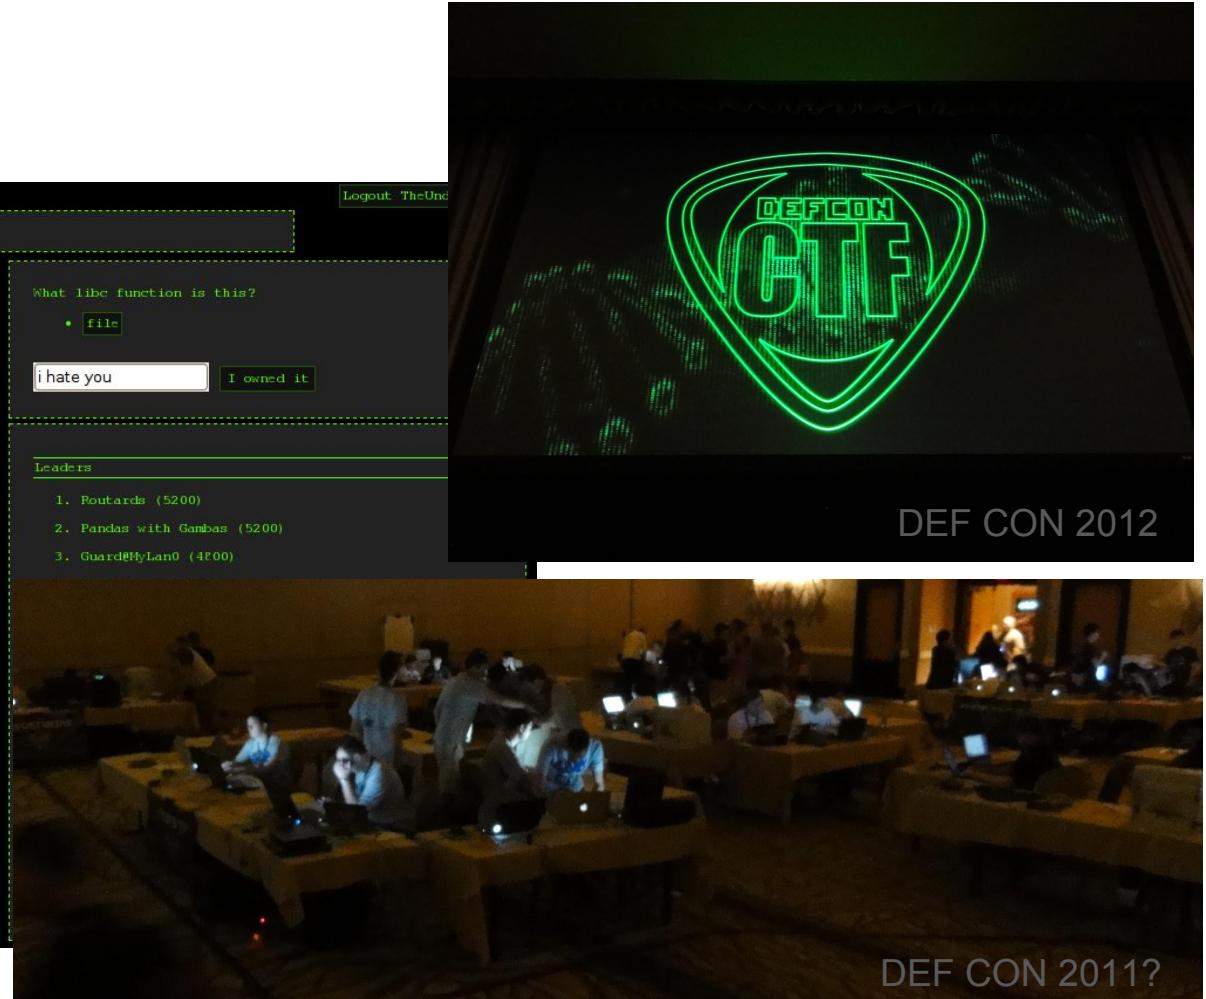
Row\Col	1	2	3	4	5
1	100	100	100	100	100
2	200	200	200	200	200
3	300	300	300	300	300
4	400	400	400	400	400
5	500	500	500	500	500

Below the grid, there is a question: "What libc function is this?" with a radio button next to "file". A text input field contains "i hate you" and a button labeled "I owned it". At the bottom, there is a "Leaders" section with the following list:

1. Routsards (5200)
2. Pandas with Gambas (5200)
3. Guard@MyLan0 (4800)

DEF CON 2008

DFRWS 2018



DEF CON 2011?

CGC?

Could a purpose-built super computer play in DEF CON's Capture-the-flag (CTF)?

Autonomous...

- Binary analysis
- Binary patching
- Vulnerability discovery
- Service Resiliency (availability)
- Network Defense (IDS)

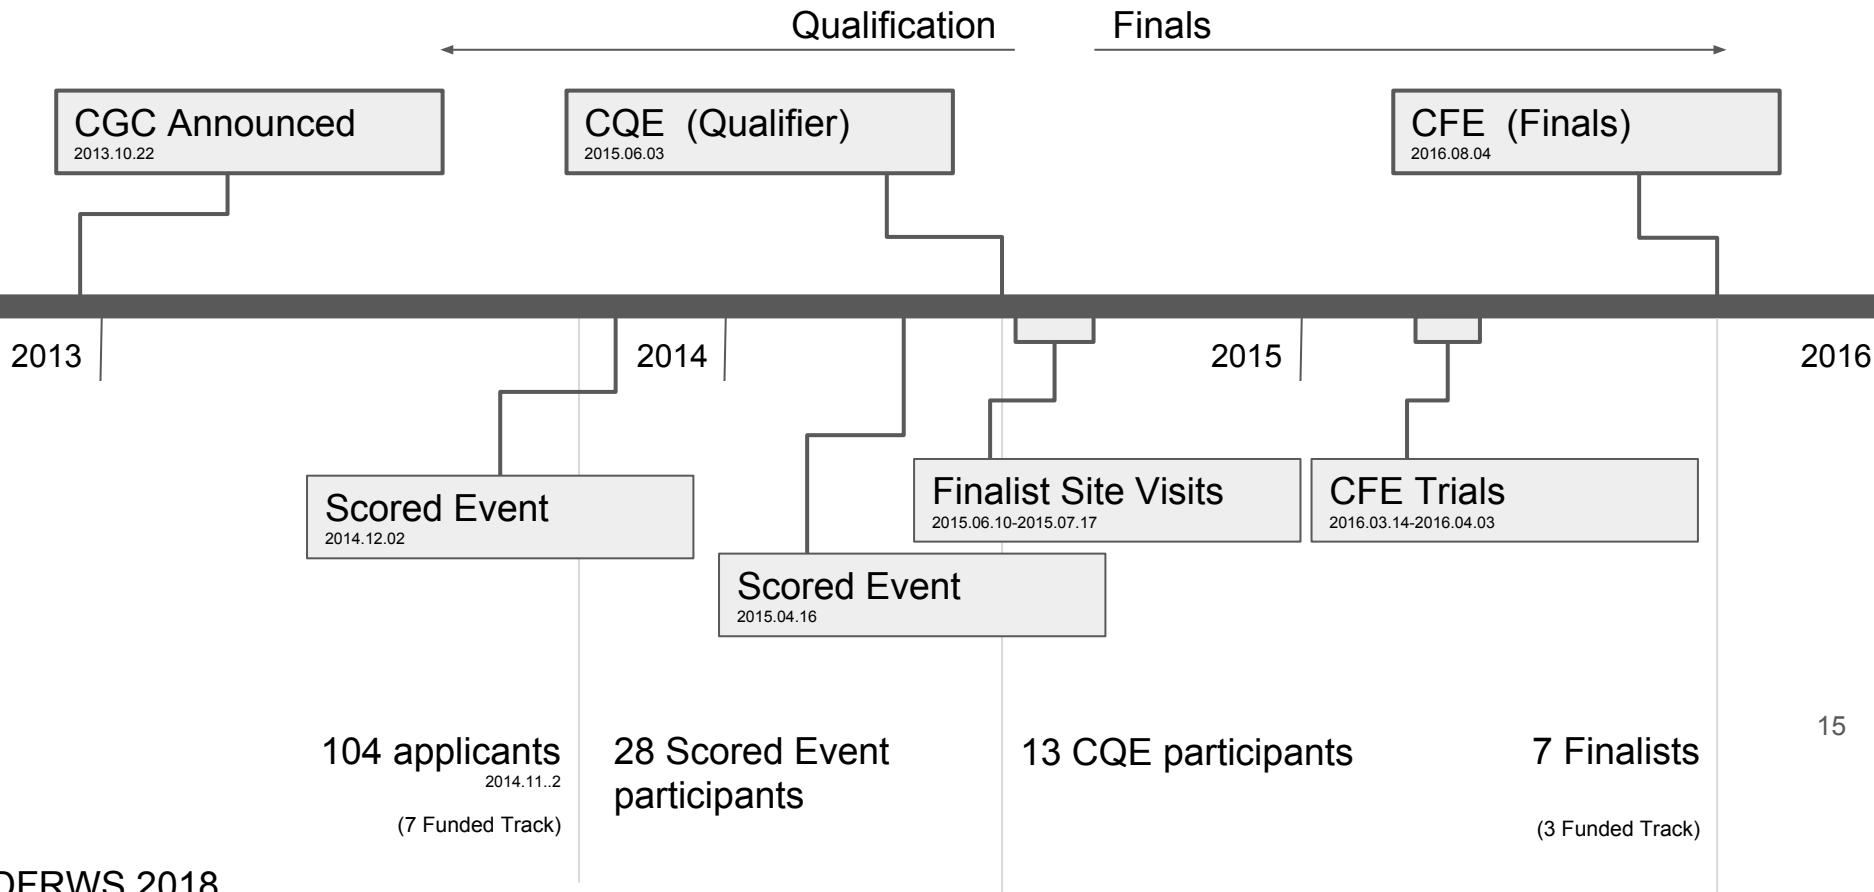
CGC: Real life



CGC: Real life



Competition Overview



Building the Competition

- Design concerns from the outset
 - Repeatability
 - Anyone should be able to verify CFF results
 - Competition integrity
 - Concerns with running competitor-provided code (POV/RCB)
 - Concerns with parsing competitor-provided data (IDS filters)
 - Data collection
 - Desire to publish corpus to serve as a reference for program analysis going forward

Competition Integrity

- Given the amount of prize money at stake, integrity of the competition was a grave concern and drove many design decisions
- Randomness was limited and/or made to be deterministically pseudorandom
- However, **nobody** should be able to predict aspects of CFE
 - The entire event was seeded with input from DARPA and all competitors (XORed) (Collected between June 10-17, 2016)
 - To ensure that DARPA did not select a particular input after knowing all competitor inputs DARPA's input was cryptographically committed to early (June 10, 2016)
- Similarly, the CFE event plan (including challenge set schedule) was committed to on Aug 2, 2016)
 - Organizers could not change the schedule in order to influence the event outcome

Q185: What were the competitor team TeamPhrases used to contribute to the calculation of the master seed?

A185: The TeamPhrases solicited from finalists and used according to A176 of the FAQ are published in the below JSON:

https://github.com/CyberGrandChallenge/Event-FAQ/blob/master/event_faq.md

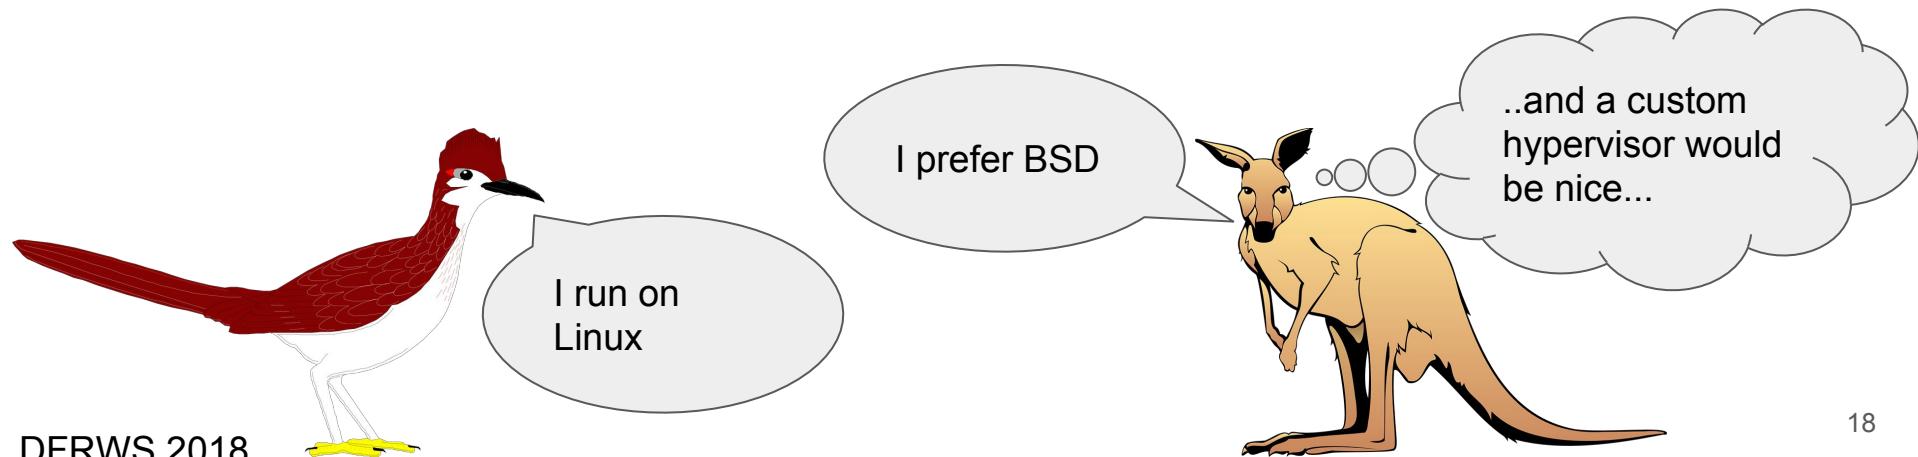
http://archive.darpa.mil/cybergrandchallenge_competitorsite/Files/CGC_FAQ.pdf

Competition Integrity

7 system calls

_terminate, transmit, receive, fdwait,
allocate, deallocate, random

- Committed to kernels versions released prior to announcement of CGC
- Designed DECREE syscall environment / file format to reduce attack surface
- All game infrastructure components released to the public had private internal implementations
 - Notably, CFE ran on 64-bit FreeBSD 10 with a custom hypervisor module



Competition Integrity

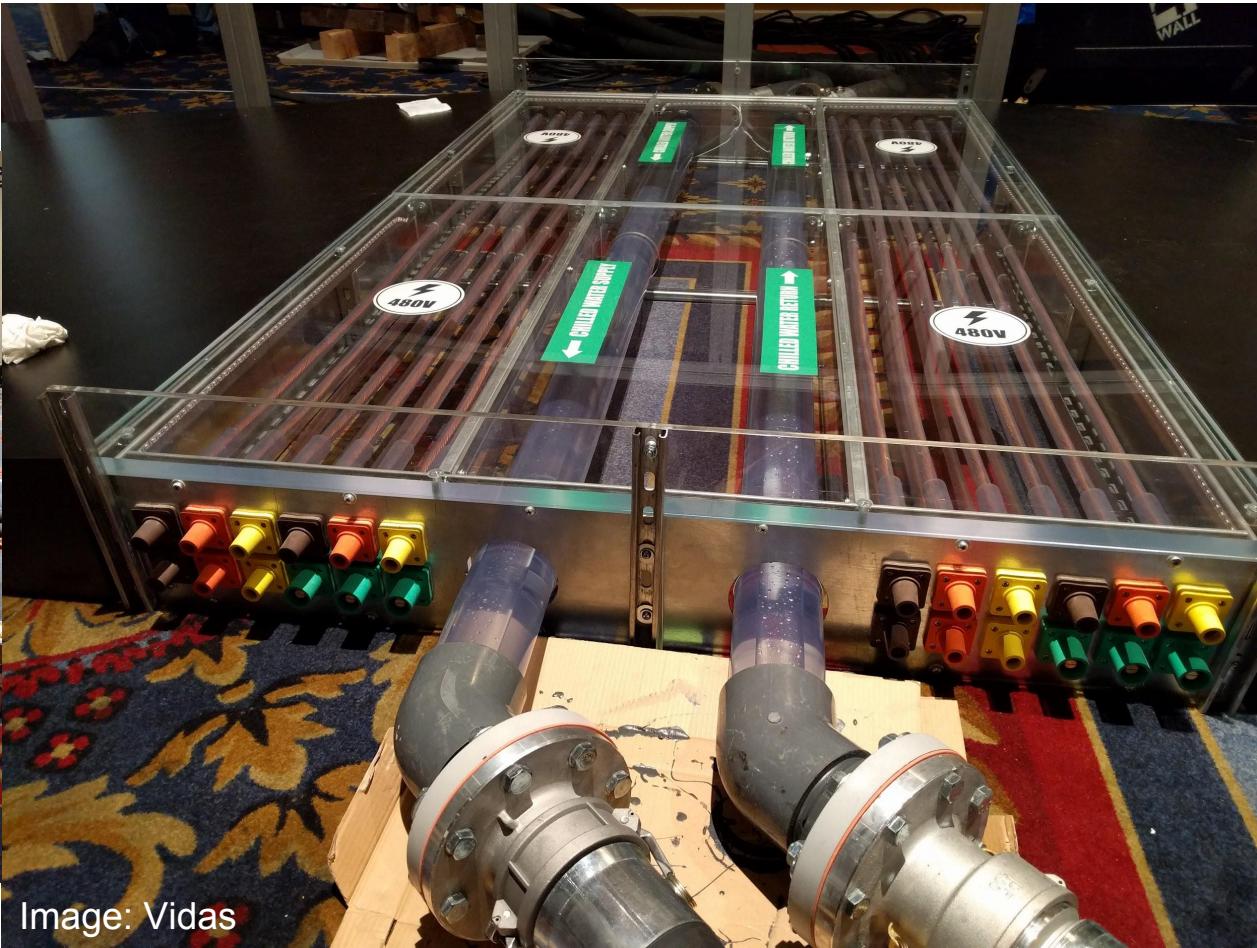
- Air Gap
 -



Image: Vidas

Competition Integrity

- Air Gap
 - Power, cooling



Competition Integrity

- Air Gap
 - One-way data



Competition Integrity

- Competitors were required to be autonomous, organizers weren't
 - Referees
 - However, air gap
-
- Redundant HW
 - Power/cooling
 - Monitoring



Image: DARPA

Competition Integrity: Forensics

- Real-time forensics harness to vet software
 - Monitor OS for execution & data integrity
 - Built upon a full system emulator (Simics)
 - High fidelity x86 model from Intel
- Evaluated non-trusted code (POV/RCB) for attempts to breakout of DECREE environment
- Analyst replay tool
 - Replay any CFE session via IDA Pro gdb client
 - Reverse execution & scoring event detection

CGC Monitor vetted all competitor submissions

CGC infrastructure duplicated on the Simics full system simulator

- Multiple components; all game services

- Monitor OS for execution & data integrity

- High fidelity x86 model from Intel

GCG Monitor built upon Simics primarily from breakpoints and callbacks

- Implementation similar to dynamic VM introspection

- No monitoring functions execute on monitored systems

- Built custom “OS awareness” subsystem based on OS internals

- Variations for 32/64 bit Linux and FreeBSD (and combinations thereof)

Implemented on 32 blade servers with multiple instances of CGC systems

What was monitored?

Competitor-supplied software

- Proofs of vulnerabilities executing on PoV throwers

- Replacement challenge binaries on the defended host

- IDS subsystem while consuming competitor's filters

While scheduled for execution:

- Kernel ROP -- execution of a "ret" not following a "call"

- Page tables allocated to the kernel

- Process credentials -- e.g., effective user ID

- Unexpected code sections -- e.g., process create while an RCB runs

Artifacts generated by monitoring

Anomalous events from kernel monitoring

Full execution traces, including data references

System call logs, including all parameters

Successful Proofs of Vulnerabilities (PoVs) against services

ROP or stack area execution in services

Faults in services, e.g., segmentation violations leading to crashes

CGC Analysis Tool: Running a computer backwards

Real world analogy: Your hybrid fuzzer found a vulnerability: But what is the bug?

Competitors found 20 vulnerabilities in 82 challenge sets. But what flaws?

Analysis of effective patches would not help: they were all generic

Instrumented the full system simulator for analysis of application exploitation

Automatically detect a successful exploit and pause the session

Analyst can then use reverse execution to track the bug

IDA Pro debugger client as a front end to the CGC Monitor

Library function Data Regular function Unexplored Instruction External symbol

Debug View Structures Enums General registers Call Stack

IDB View-EIP

```

08048F71 mov [esp], eax
08048F74 call free
08048F79 mov eax, [ebp+var_C]
08048F7C mov [esp], eax
08048F7F call free
08048F84 jmp loc_8048F93

```

Graph overview

The screenshot shows the IDA Pro debugger interface. The main window displays assembly code in the 'IDB View-EIP' tab. A specific instruction at address 08048FAB is highlighted with a blue selection bar. Below the assembly view, the status bar indicates: '100.00% (4329,1346) (436,326) 00000FOB 08048FAB: sub_8048DC0+1EB (Synchronized with EIP)'. The bottom left pane shows a command-line history with ROP gadgets and their cycles. The bottom right pane shows the 'Call Stack' and 'General registers' tabs.

Registers

EAX 08060E3C	.bss:dword_8060E3C	OF 0
ECX B7FFF00C	MEMORY:B7FFF00C	DF 0
EDX 74ED1016	MEMORY:74ED1016	IF 1
EBX 00000000	MEMORY:saved_fp	TF 0
ESP BAAAABFC	MEMORY:BAAAABFC	SF 0
EBP BAAAAFF4	MEMORY:BAAAAFF4	ZF 1
ESI 08048950	do_stats	AF 0
EDI 00000000	MEMORY:saved_fp	PF 1
EIP 08048FAB	sub_8048DC0+1EB	CF 0
EFL 00000246		

Call Stack

Hex View-1

Output window

```

_start+1
rop:0x8048309 cycle:befaa0a
rop:0x4e7c7450 cycle:befaa0c
Type 1 POV eip:0x4e7c7450 ebp:0x584d455f cycle:befaa0c

```

signalclient back from cont
use monitor uncall function
signalclient call setAndDis for 0x8048fab
signalclient return to main

DRKWS 2010

28

IDA Pro Extensions for Reverse Execution

Reverse (e.g., until a breakpoint is encountered)

Step backwards over or into a function

Reverse to cursor

Reverse to previous write of selected register or address

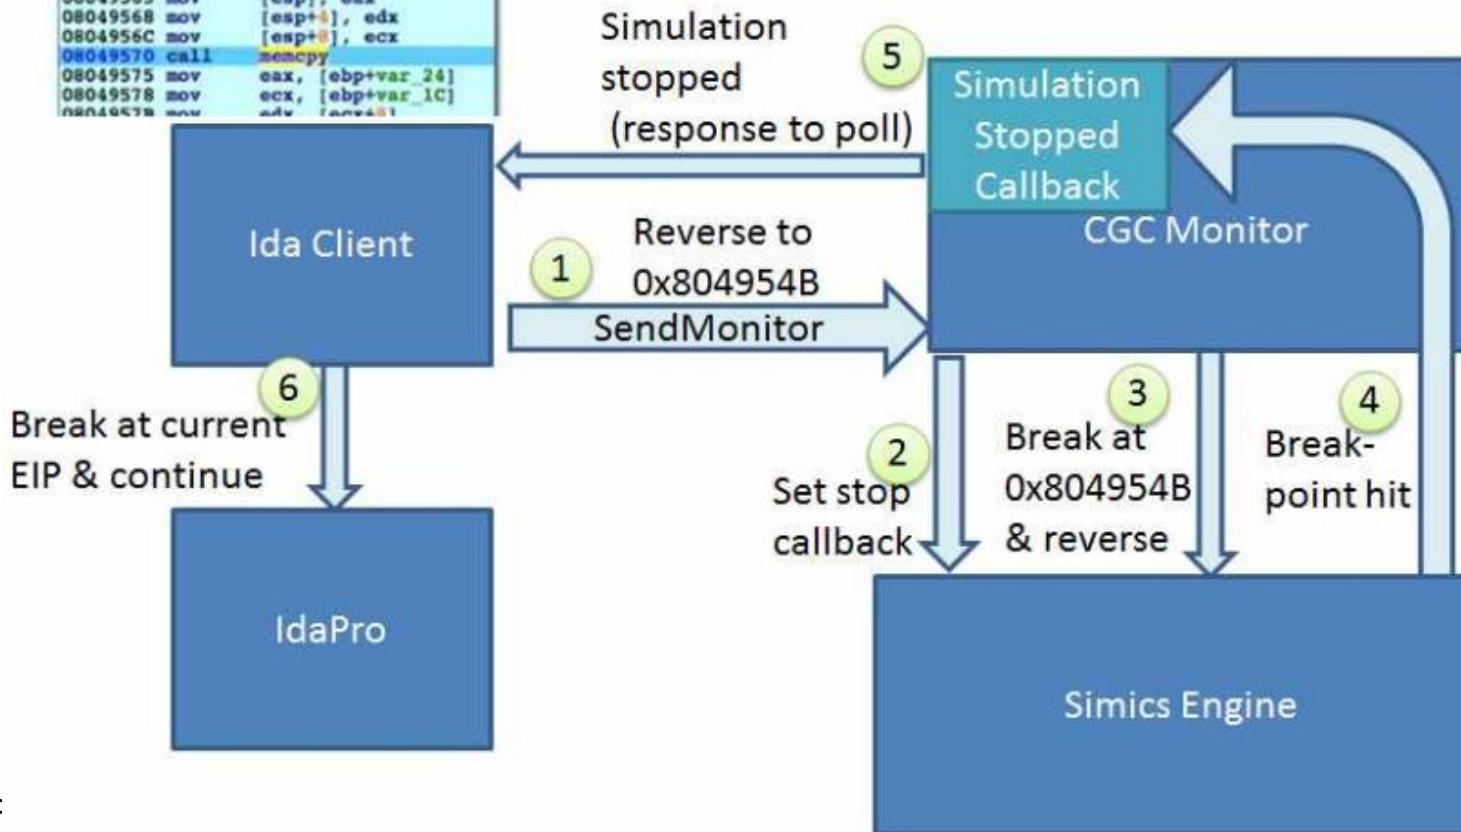
Set or jump to an execution bookmark managed by the user

Back trace the source of data in address or register

Often leads all the way back to the syscall that received data

Halts on computed assignments (e.g., addition -- but not increment)

```
0804954B l 804954B:  
0804954B    mov    eax, [ebp+var_34]  
0804954E    mov    [ebp+var_24], eax  
08049551    mov    eax, [ebp+var_30]  
08049554    mov    ecx, [ebp+var_1C]  
08049557    mov    ecx, [ecx+0Ch]  
0804955A    mov    edx, [ebp+var_1C]  
0804955D    mov    edx, [edx+4]  
08049560    add    edx, ecx  
08049562    mov    ecx, [ebp+var_24]  
08049565    mov    [esp], eax  
08049568    mov    [esp+4], edx  
0804956C    mov    [esp+8], ecx  
08049570    call   _memcpy  
08049575    mov    eax, [ebp+var_24]  
08049578    mov    ecx, [ebp+var_1C]  
0804957B    mov    edx, [ebp+var_1C]
```



Simics illusion of reverse execution

Resource intensive, enable only for analyst sessions

Records “micro-checkpoints” referenced during reverse execution

Iterates from checkpoints, running forward until “most recent” breakpoint

Warning: backwards progression is not serial

Callback for one breakpoint may be invoked many times

Breakpoint callbacks are therefore not useful when reversing

Associate callback with simulation “stop” event

Then figure out where you are and why you stopped

Analysis of CGC Final Event

82 Challenge sets, having 109 intended vulnerabilities

20 challenge sets had working POVs in CFE

Half of these working POVs were not what the author intended

Six were different vulnerabilities (2 services exploited via same bug)

Four were the intended bug, exploited via an simpler alternate path

All exploits of each challenge set used the same vulnerability and path

Fully automated back trace of data

Back trace sources of data, e.g., to a receive syscall (like backwards taint analysis)

Corrupted return addresses

Corrupt values of call registers

Executable payloads

General register values negotiated in Type 1 PoVs

The source of protected memory addresses

Traces available in the CGC Corpus at: <http://www.lungetech.com/cgc-corpus/>

Future Work & Availability

Extend for general application environments (currently DECREE)

Package Analysis Tool as a remotely accessible service

CGC Monitor at <https://github.com/mfthomps/cgc-monitor>

BYOS (bring your own Simics)

Analysis results at <https://github.com/mfthomps/CGC-Analysis>