

Leveraging Relocations in ELF-binaries for Linux Kernel Version Identification

Manish Bhatt, Irfan Ahmed

University of New Orleans

DFRWS 2018

Outline

- Introduction
- Design Rationale
- Implementation
- Evaluation
- Possible Optimization Measures
- Conclusion
- Acknowledgement

Introduction

- Forensics Challenges in analysing Linux memory dumps:
 - **Address Randomisation**
 - **Kernel Version Identification**
- Proposed Solution:
 - Using Relocation Entries
 - **codeid-elf**

Design Rationale (1)

- S be one of the obtained signatures. $S = \langle ptr, offset \rangle$
- Let $P(i)$ refer to the i^{th} page in memory.
- Base address in Memory be B_m
- Base address in disk be D_m , then the following is true:

$$B_m - D_m = RandomizedOffset = ptr - S.pointer$$

Design Rationale (2)

In Memory Kernel Code

1b001000: 8b0d **c01e b71b** f686 1102 0000 4075 160f@u..
1b001010: 0115 **c61e b71b** b818 0000 008e d88e c08e
1b001020: e08e e88e d08d a100 0000 40fc 31c0 bf00@.1...

In Disk Kernel Code

00000000: 8b0d **c01e b701** f686 1102 0000 4075 160f@u..
00000010: 0115 **c61e b701** b818 0000 008e d88e c08e
00000020: e08e e88e d08d a100 0000 40fc 31c0 bf00@.1...

Relocation Entries

No.	Offset	Pointer	Symbol
1	c1000002	c1000002	initial_stack
2	c1000012	c1000002	boot_gdt_descr

Design Rationale

No.	In Disk	In Memory	Difference
1.	0x01b71ec0	0x1bb71ec0	0x1a000000
2.	0x01b71ec6	0x1bb71ec6	0x1a000000

Fingerprinting and Derandomization

- Generate Signatures for all known versions.
- Match **all the signatures** to a memory dump for **version identification**.
- Match **single page** for address derandomization.

Data: output of readelf -r program

Result: List of $\langle \text{offset}, \text{pointer} \rangle$ signature tuples
initialization;

filehandler = Open Output File;

while *every line in the output* **do**

if (*line contains "R_386_32"*) **then**

 filehandler.write(line.offset+ " " + line.value) ;

else

 continue;

end

end

Algorithm 1: Algorithm to extract signatures.

Implementation (1)

- Leveraging known utilities for signature extraction: **readelf**
- Implemented in a mixture of C and Python
- Only relevant relocation entry is **R_386_32**
 - According to standards, direct replacement of relocated address
 - As opposed to the **PC** relative address

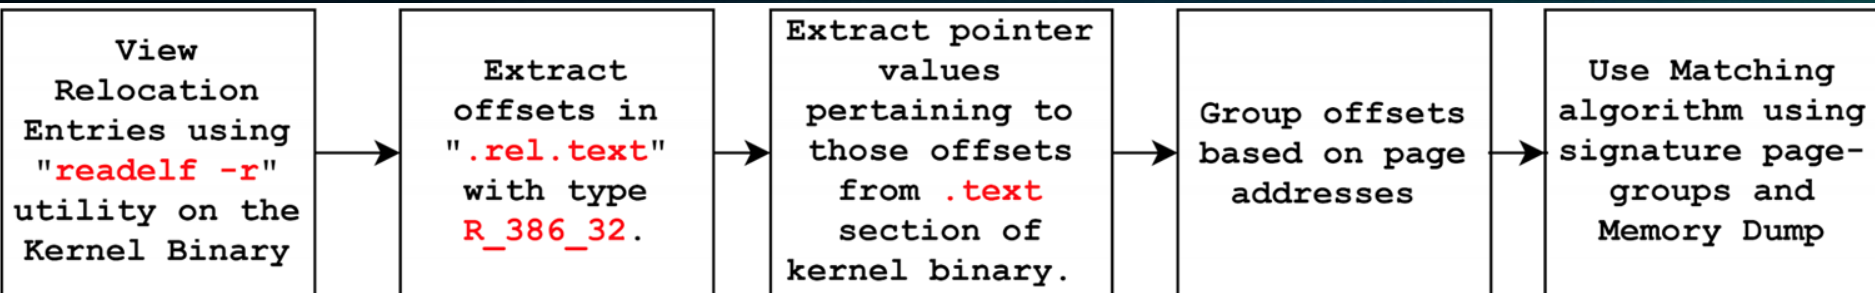
Implementation (2): PAGE DETECTION

- Could have just done the whole executable at once. Why?
 - All pages were not detected
 - Some pages no relocs (less than 2%)
 - Some relocs changed at runtime
 - Therefore, page-level analysis was performed

Implementation (3): Version Detection

- Run for all known versions
- Pick the version with the highest number of page matches
- Just want derandomization?
 - Just extract a single page or a couple of pages
 - Faster as opposed to trying the whole executable

Implementation (4): Workflow



Gathering Data

- Avoid compiling a lot of kernels
- Easier to create the **vmlinux** executable inside the kernel
- Memory snapshot via vmware **vmss2core** utility
- Created 22 different kernel versions this way
- Standard kernels used to evaluate

Custom Kernels?

- Are they different? As long as
 - They have elf binaries
 - They have relocs of type **R_386_32**
- Fingerprinting the executable using executable structure and properties
- Expectedly kernel independent

Evaluation (1): Metrics

- **Relocation Prevalence** and **Code Coverage?** Snip:

Vmlinux Kernel Version	Relocation Prevalence	Ratio of Pages with no Relocation Entries
4.4.0-22	55.4482	0.0050
4.4.0-24	55.4314	0.0056
4.10.0-14	54.0464	0.0047
4.10.0-14-lowlatency	53.2292	0.0051
4.10.0-19	54.0825	0.0042
4.10.0-19-lowlatency	53.2381	0.0061
4.10.0-20	54.0825	0.0042
4.10.0-21	53.9475	0.0070
4.10.0-22	53.9418	0.0051
4.10.0-24	53.9442	0.0042
4.10.0-26	53.9584	0.0037
4.10.0-27	53.9584	0.0037

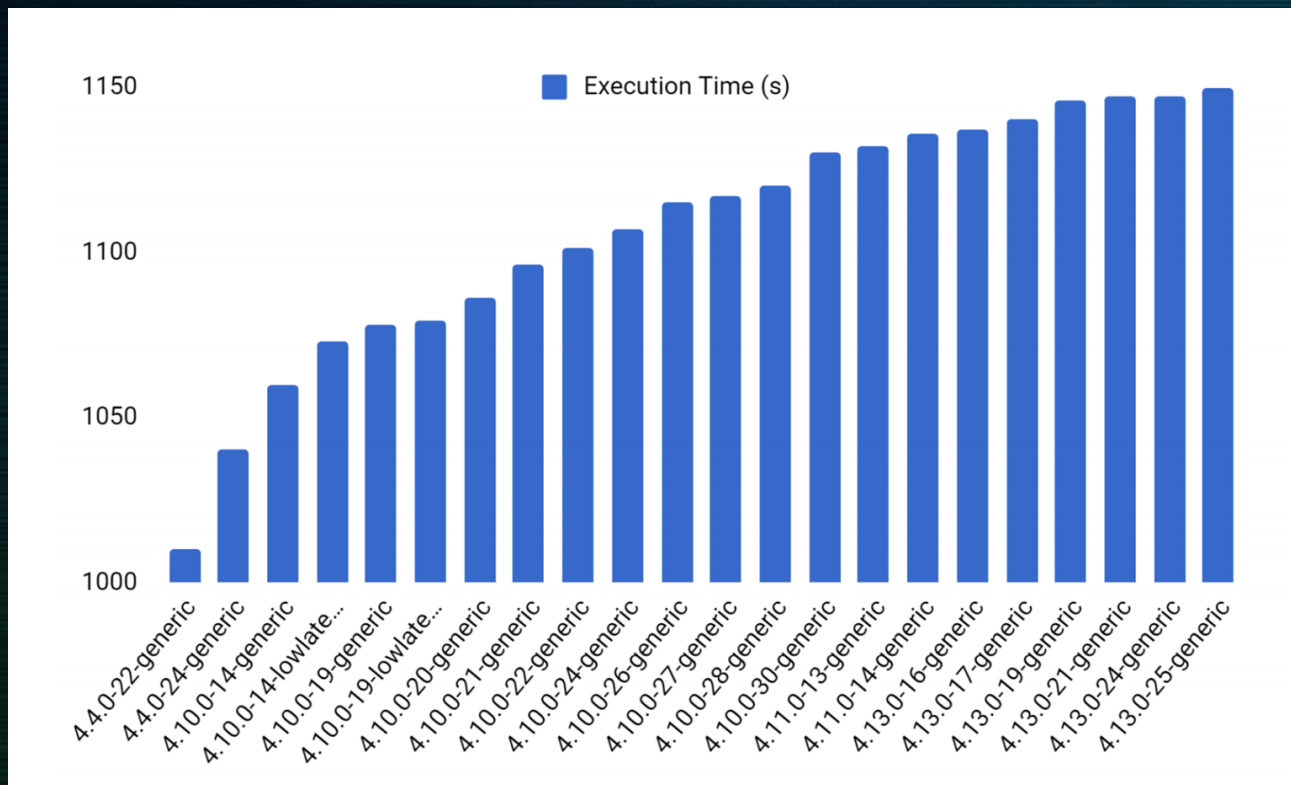
Evaluation (2): Page Hit Rate and Detection

Vmlinux Kernel Version	Number of Generated Signatures	Size of Ker- nel Code in MB	Total Num- ber of Pages in code	Page Rate	Hit	Base Address of Kernel	Derandomisation	Kernel Iden- tification
4.4.0-22	108734	7.6640	1961	99.47		0xc100000	√	√
4.4.0-24	108812	7.6708	1963	99.43		0xc100000	√	√
4.10.0-14	113984	8.2387	2109	99.52		0xdc00000	√	√
4.10.0-14- lowlatency	113059	8.2976	2124	99.48		0xc600000	√	√
4.10.0-19	114060	8.2419	2109	99.57		0xd300000	√	√
4.10.0-19- lowlatency	113131	8.3010	2125	99.38		0xc800000	√	√
4.10.0-20	114060	8.2419	2109	99.57		0xd100000	√	√
4.10.0-21	114207	8.2701	2117	99.29		0xd000000	√	√
4.10.0-22	114195	8.2702	2117	99.48		0xcc00000	√	√
4.10.0-24	114200	8.2704	2117	99.57		0xc900000	√	√
4.10.0-26	114230	8.2707	2117	99.62		0xc800000	√	√
4.10.0-27	114230	8.2707	2117	99.62		0xda00000	√	√
4.10.0-28	114230	8.2708	2117	99.62		0xdb00000	√	√
4.10.0-30	114232	8.2713	2117	99.43		0xca00000	√	√

Evaluation (3): Fingerprint Similarity

Kernel A	Kernel B	Number of Common Signatures	Sim _{AB}
4.10.0-21-generic	4.10.0-28-generic	23	0.0201
4.10.0-21-generic	4.10.0-22-generic	2107	1.8449
4.10.0-21-generic	4.10.0-14-generic	69	0.0604
4.10.0-21-generic	4.10.0-24-generic	2097	1.8361
4.13.0-25-generic	4.13.0-21-generic	99	0.0882
4.10.0-28-generic	4.10.0-30-generic	32607	28.545
4.10.0-28-generic	4.10.0-20-generic	20	0.0175
4.10.0-28-generic	4.10.0-27-generic	29530	25.8514
4.10.0-28-generic	4.10.0-26-generic	2273	1.9898
4.10.0-30-generic	4.10.0-27-generic	29285	25.6364
4.10.0-30-generic	4.10.0-24-generic	102	0.0893
4.10.0-30-generic	4.10.0-26-generic	2163	1.8935
4.13.0-16-generic	4.13.0-21-generic	60	0.0534
4.10.0-20-generic	4.10.0-22-generic	149	0.1306
4.10.0-27-generic	4.10.0-22-generic	105	0.0919
4.10.0-27-generic	4.10.0-14-lowlatency	43	0.0376
4.13.0-19-generic	4.13.0-17-generic	214	0.1904

Evaluation (4): Performance Measures



Optimisation Measures

- Resample Fingerprints
 - Intelligently create a subset of signatures
 - Decreasing the number of signatures
 - Without sacrificing code coverage
 - Decreasing relocation prevalence
- Implementing completely in C
- For kernels, scan just areas where kernels are likely to be found. (Kernel Ranges)

Takeaways

- Very simple to understand and implement
- Low number of pages with no relocation entries
- High code coverage and prevalence
- Fingerprints are mostly unique
 - In experiments, except for **4.11.0-14-generic** and **4.11.0-13-generic**
- Technique works for both derandomization and version identification

Comparison to Similar Works

Method	Similarity of Derived Signatures	Signatures Derived From Mem- ory Dumps	ASLR Sup- port	Version Detection Accuracy
OS-Sommelier	Low	Yes	✓	100%
k-BinID	High	No, derived at Run-Time using VMI	✓	100%
Image-Based Kernel-Fingerprinting	High/Low depending on Similarity in implementation	No	✗	100%
codeid-elf	Low	No	✓	100%

Conclusion

- High page hit rate
- Always derandomized addresses correctly
- Over 99% page hit rate
- Similarity of Relocations Low
- Versions detected correctly every time
- Extendable to other executables

Future Works

- Extend the relocs method to other user-space binaries
- Test it on other custom kernels
- Extend this to other architectures like ARM
- Use this in preliminary Malware Analysis

Acknowledgements

This work was supported by NSF grant #1623276

Questions?

