

Tomasz Tuzel

Who Watches the Watcher?

Detecting Hypervisor Introspection from
Unprivileged Guests



Overview

“The Cloud”

- Numerous organizations are moving en masse to **the cloud**
 - It's easier to manage
 - It's easier to scale
- This results in a loss of control of **physical hardware** and **privileges on the system**
 - Everything is tied to guest virtual machines

Overview

Hypervisors (Virtual Machine Monitors)

- The root-of-trust is the hypervisor (virtual machine monitor)
 - It can introspect on guests with **few restrictions**
 - It can introspect on guests with **little evidence of apparent action**

Overview

Hypervisors (Virtual Machine Monitors)

- This trade-off puts organizations and individuals in a difficult situation
 - They have **sensitive data and processes**
 - Compromises are **expensive and dangerous**

Overview

Related Work

- **Compromise of or introspection by** a hypervisor has been a known issue
- Previous work on hypervisor detection has focused on **hardware or software artifacts**
 - This is more oriented towards **detection** of past events rather active **introspection**

Motivation

- A hypervisor's activities are **not entirely invisible** to its guests
- **System performance is impacted** as the hypervisor seeks to provides an environment which is functionally equivalent to native hardware
- Increases in an instruction's **execution time and/or cache artifacting** can provide evidence of a hypervisor's intervention

Implementation

- A **test framework** emulates inappropriate introspection
- A **monitoring module** on the guest employs a set of sensors to detect malicious behavior

Test Framework

- The *Xen Project* was used as the hypervisor
- Modifications were made to support hypercalls which would toggle introspection capabilities
- Modifications were made to support instructions that *Xen Project* does not support VM-exiting for

Monitoring Module and Sensors

- **Instruction intercession sensors** detect intercession of instruction execution using intercession timing
- **Active memory intercession sensors** detect the hypervisor actively interceding in memory access operations
- **Passive memory monitoring sensors** detect when a hypervisor accesses memory externally to a guest
- **Non-temporal access sensors** are not completed, but could be used to detect hypervisor introspection via non-caching page mappings or non-temporal instructions

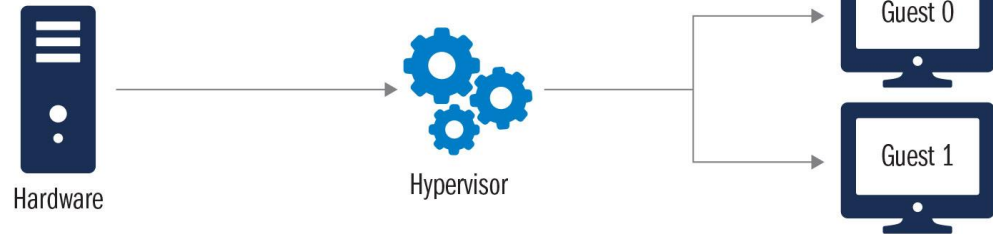
Background

Hypervisor

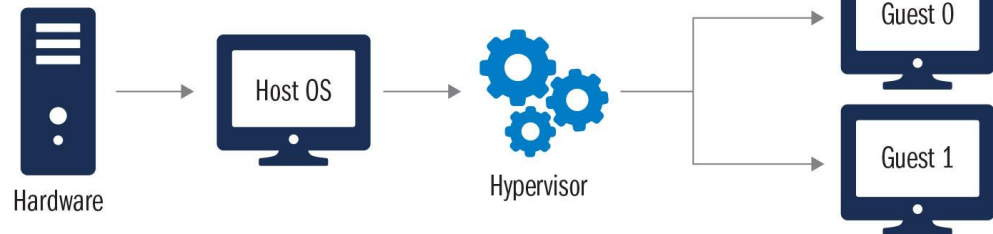
- Privileged software that handles **execution** and **isolation** of guest virtual machines

HYPERVISOR TYPES

TYPE 1 HYPERVISOR



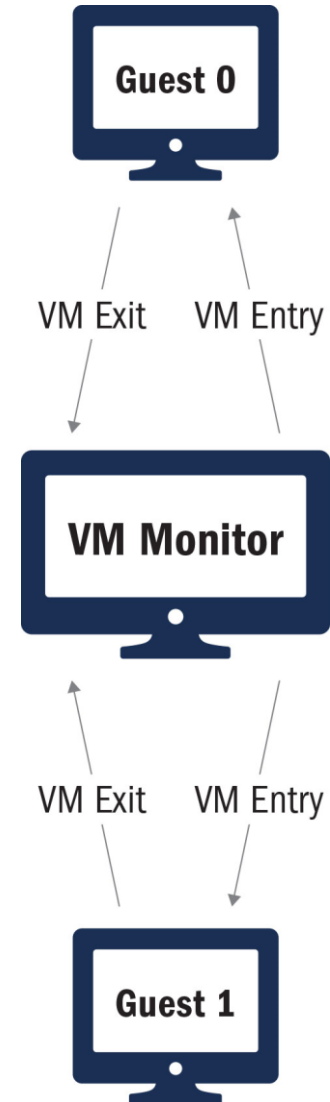
TYPE 2 HYPERVISOR



Background

Virtual Machine Exits (VM-Exits)

- A hypervisor must be able to intercede in guest execution as necessary
- Guest execution is **paused** and execution is **handed to the hypervisor**
- Guest state information is stored in the Virtual Machine Control Structure (VMCS)
 - Resides in memory
 - Exits require that **state information is saved off during transition**
 - As such, an exit incurs significant overhead



Background

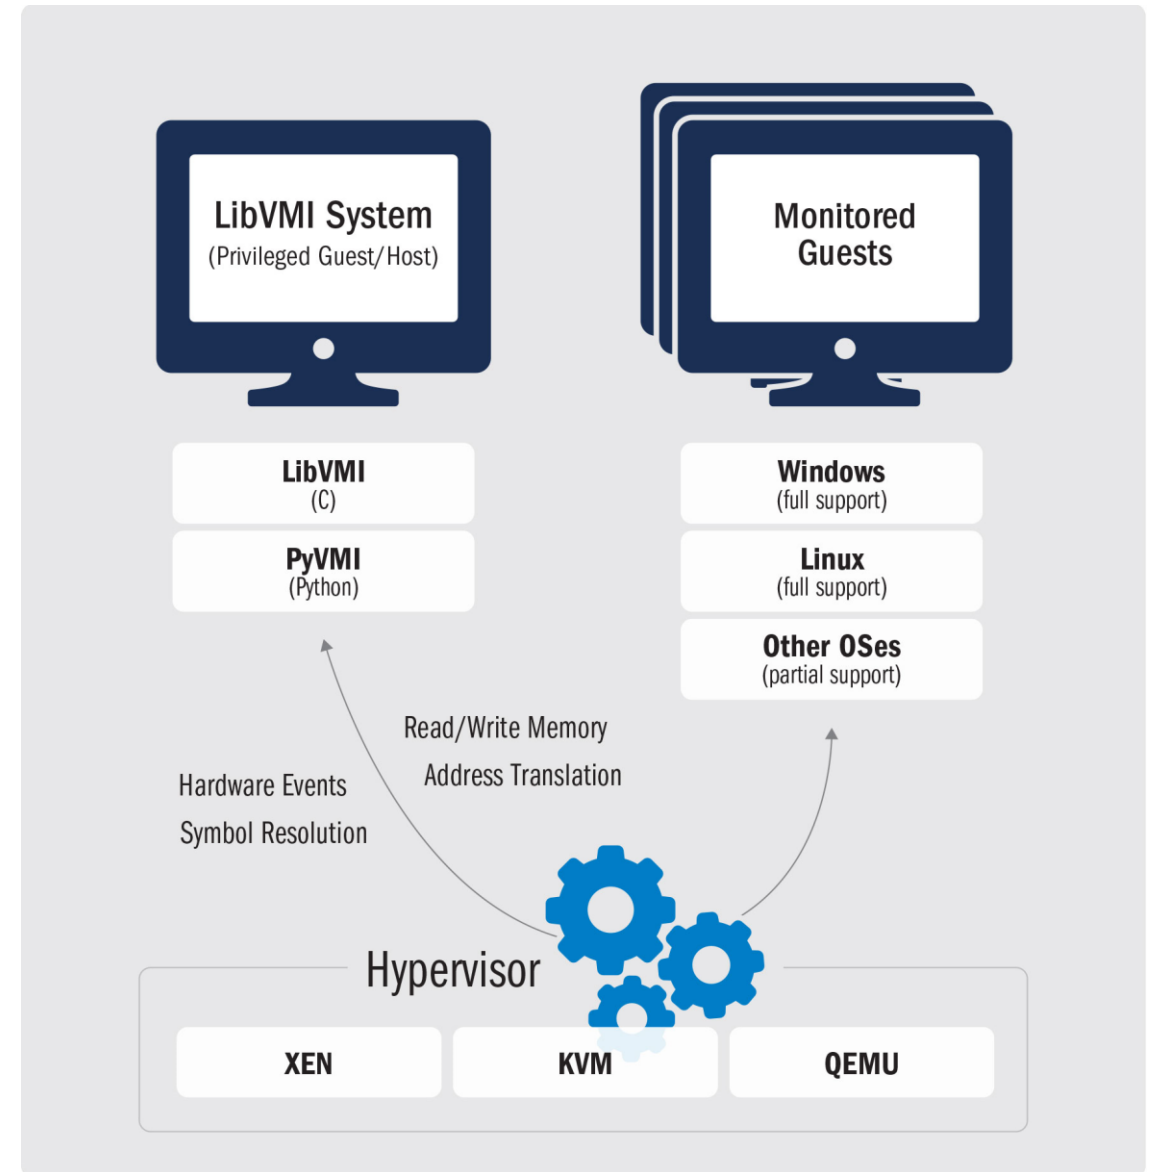
Timers & Timing Methods

- Time Stamp Counter (TSC)
- High Precision Event Timer (HPET)
- Thread racing

Background

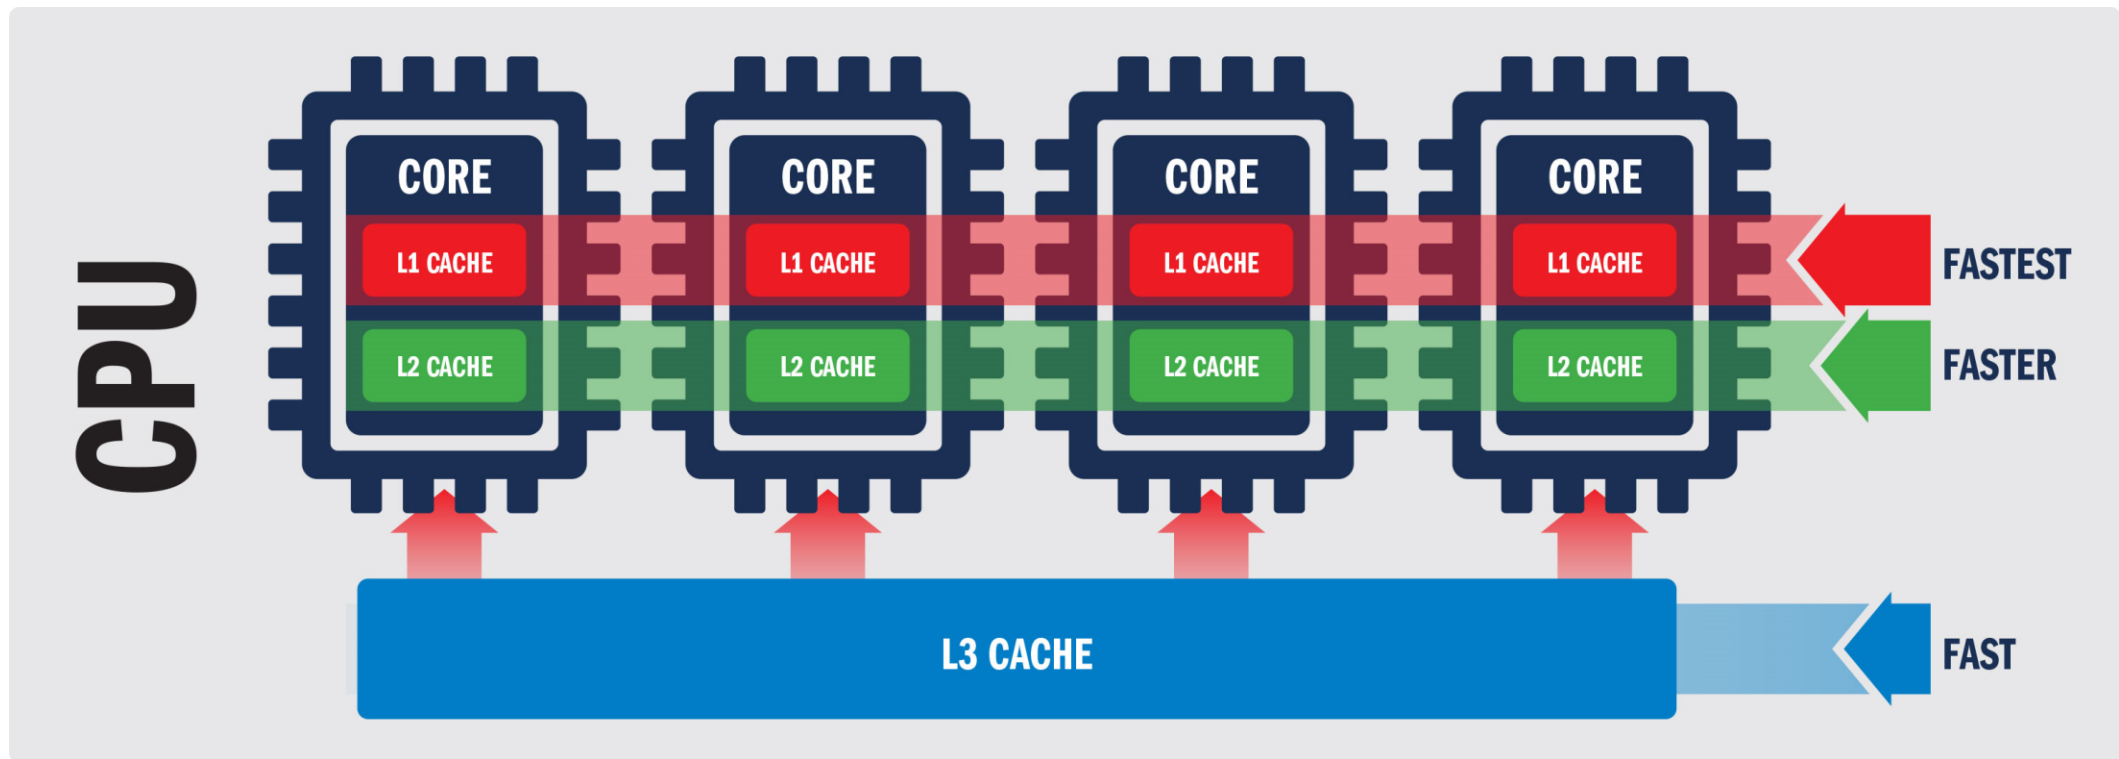
LibVMI

- C and Python library that enables introspection on virtual machines



Background

Cache

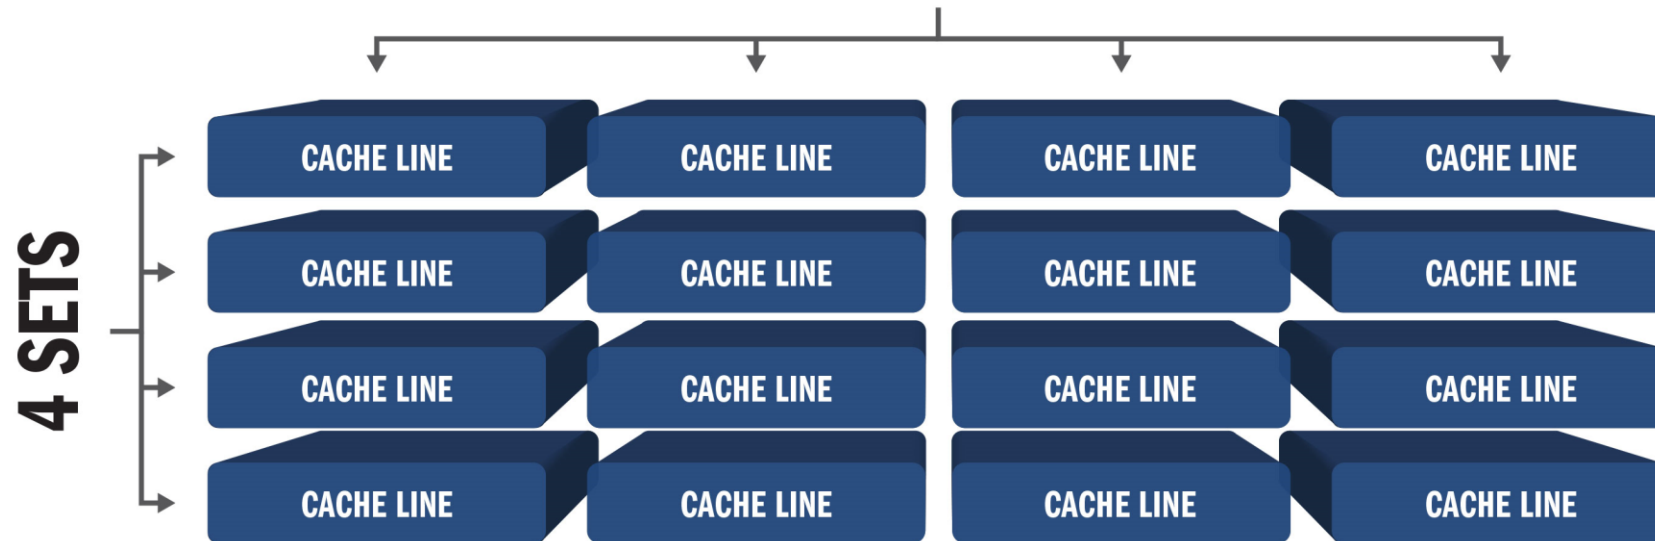


Background

Cache

SET ASSOCIATIVE CACHE

4-WAY



Instruction Intercession Sensors

- A hypervisor can trap on certain instructions to:
 - **Modify** the guest's execution behavior
 - Determine when a guest **performs** various operations
 - **Extraction** of information
- Timing can be used to determine that this is happening, and how much work the hypervisor performed while trapping

Instruction Intercession Sensors

- Wall timing via the HPET
 - Manipulation of the HPET would result in **noticeable interruptions** in operations to the observer
- Thread racing
 - Parallel threads run, one of which executes NOPs, with number of executions counted upon completion

Active Memory Intercession Sensors

- Virtual to physical memory mappings can be marked to **cause a VM-exit** (using LibVMI)
- VM-exit incurs a **sizable overhead**, and resultant **large timing increase**

Passive Memory Monitoring Sensors

- A more stealthy hypervisor can choose to **map the guest's physical pages** into other contexts
- Time required to access a memory line can be used
- Using Flush+Reload:
 - The memory line of interest is flushed from the cache
 - A period of time passes to allow for potential access to the memory region to occur
 - The memory line is reloaded, and the access is timed
- An **decrease in timing** is indicative of introspection

Non-Temporal Access Experimentation

- Non-temporal, streaming, and vector instructions have cache-coherence side-effects, despite bypassing the cache
 - A non-temporal instruction reading/writing a populated page **triggers a cache flush**
- Intel's Page Attribute Table (PAT) allows specifying per-page caching behavior
 - Passive mappings can become **non-cache-interacting**

Analysis & Classification of Results

- A variety of machines are used which had different processor generations
- For instruction intercession, a baseline is derived from non-exiting instructions
 - Adequately high values can be flagged as potentially exiting, or of interest
- For memory intercession, baselines are observable from adjacent non-introspected memory

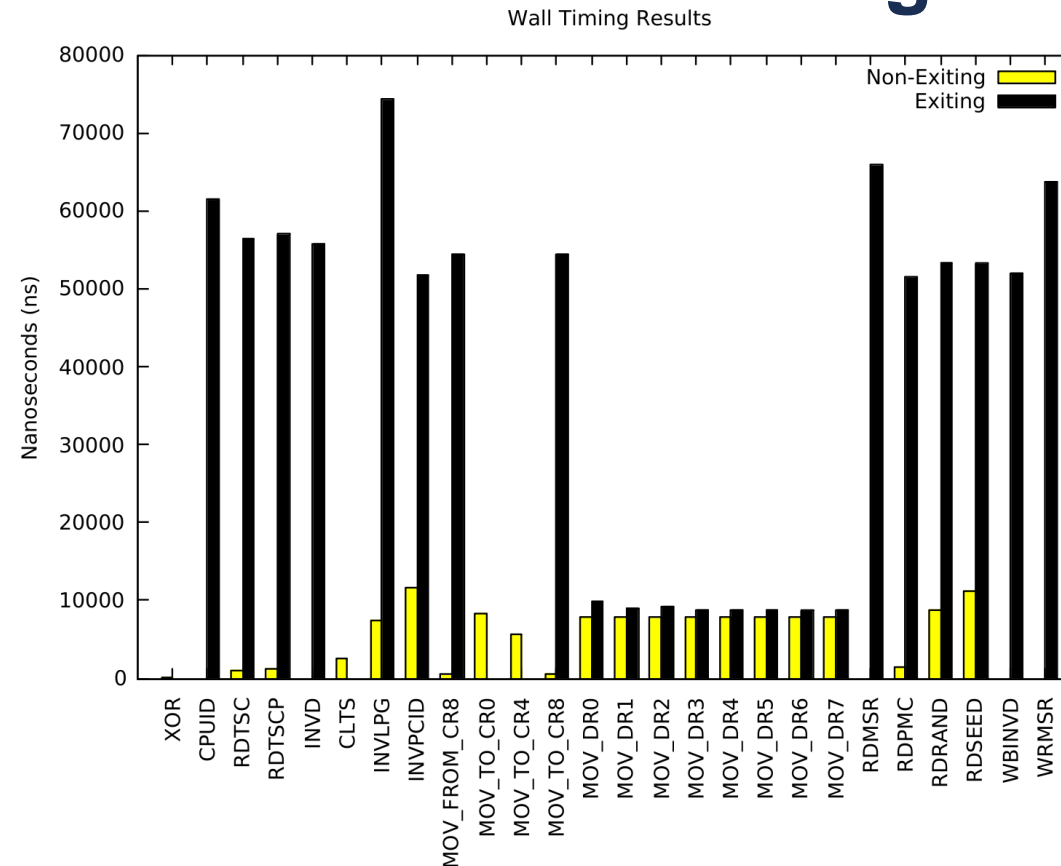
Results

Instruction Intercession

- Trapping by the hypervisor immediately returns, thus delivering the minimal possible impact on timing

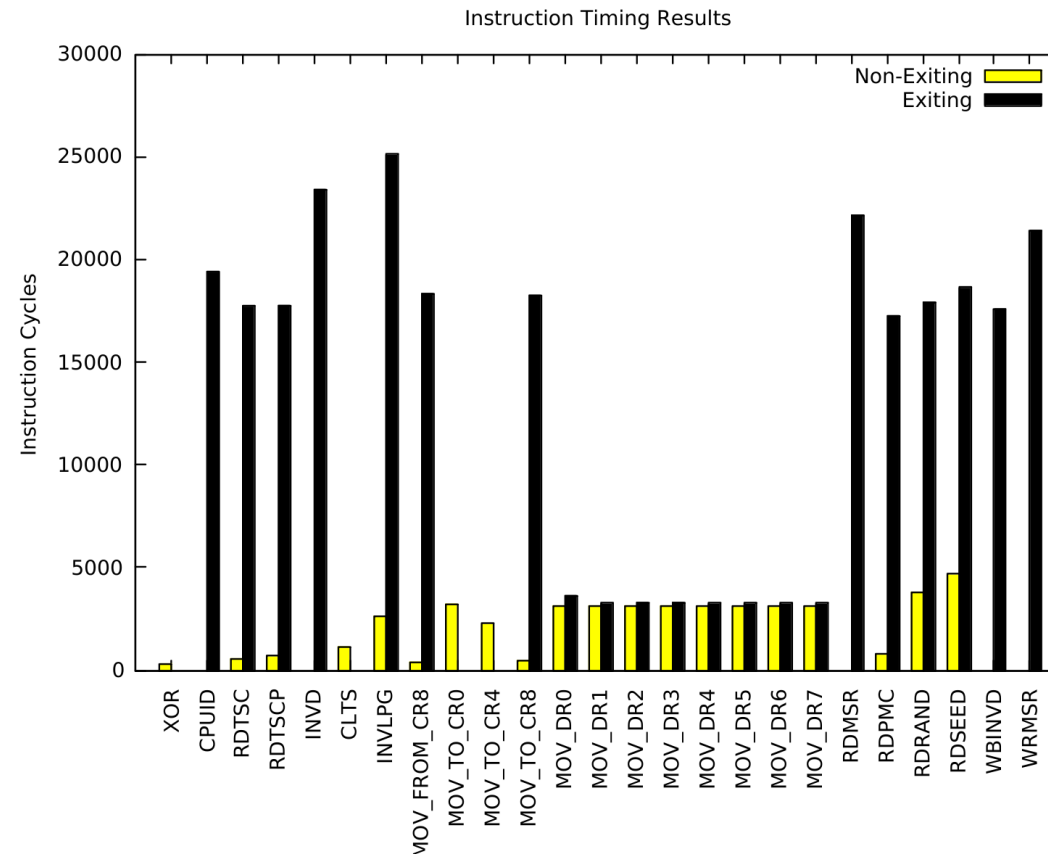
Results

Instruction Intercession – Wall Timing



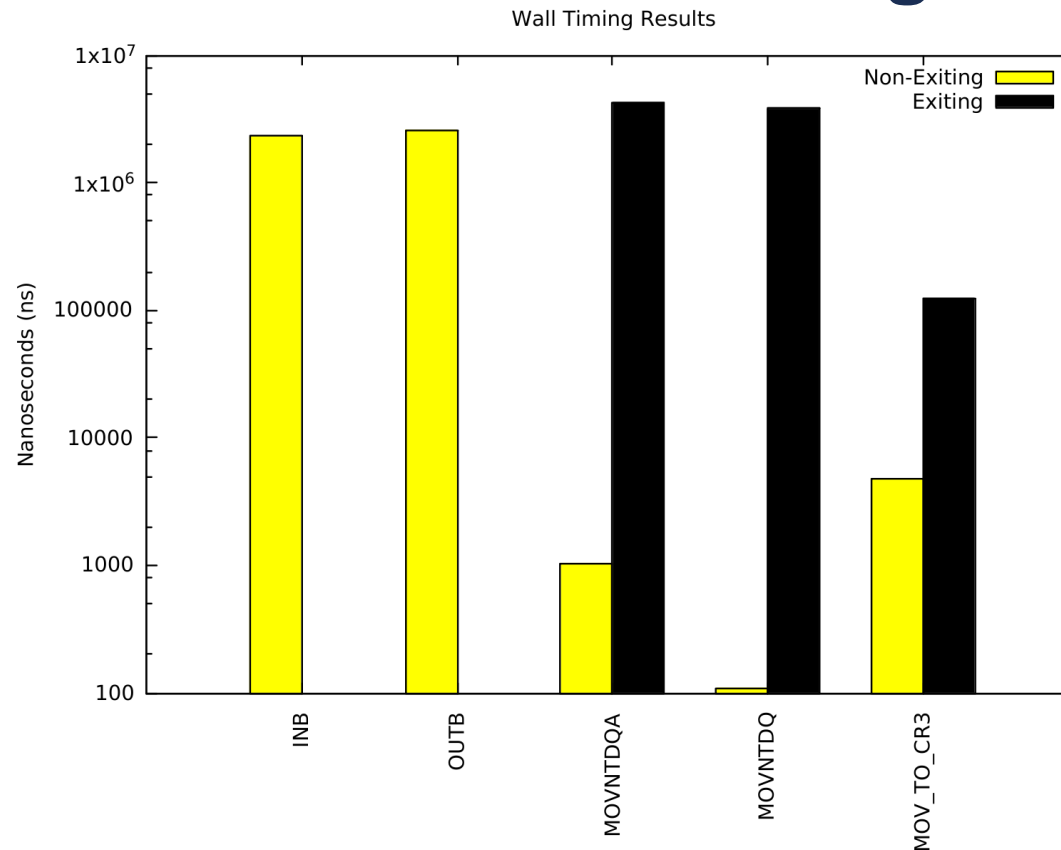
Results

Instruction Intercession – Instruction Timing



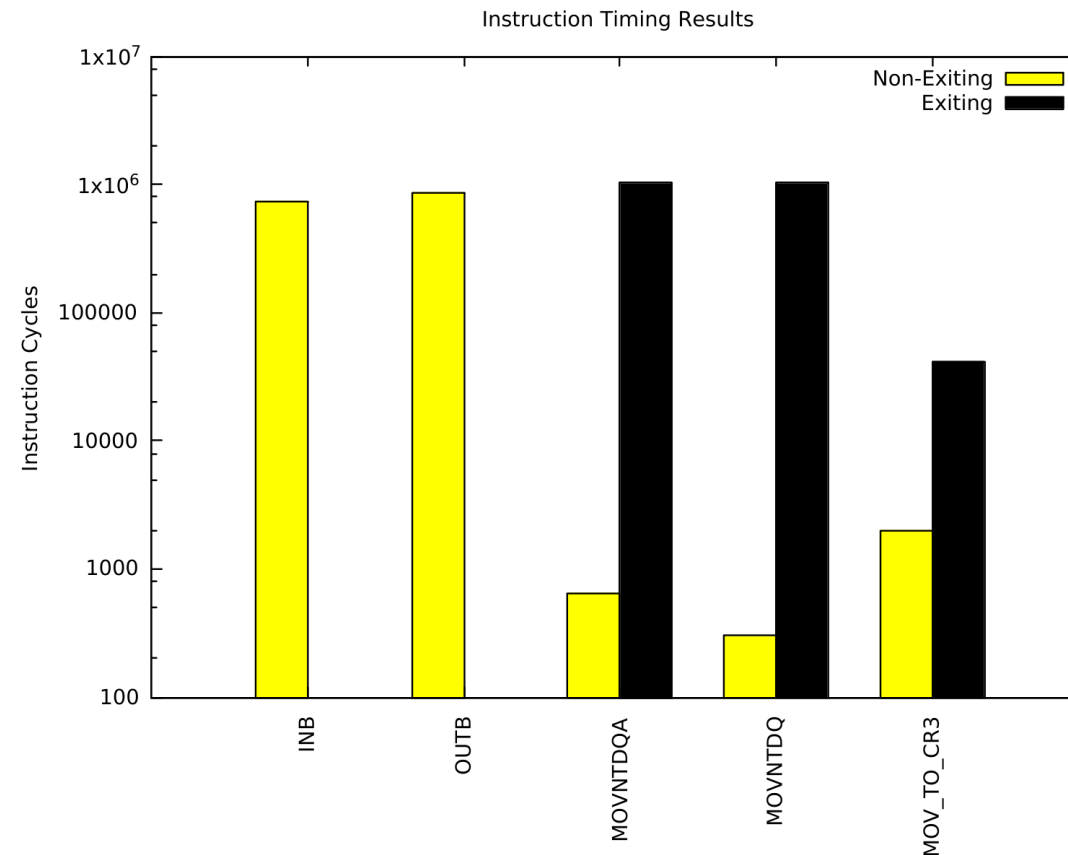
Results

Instruction Intercession – Wall Timing



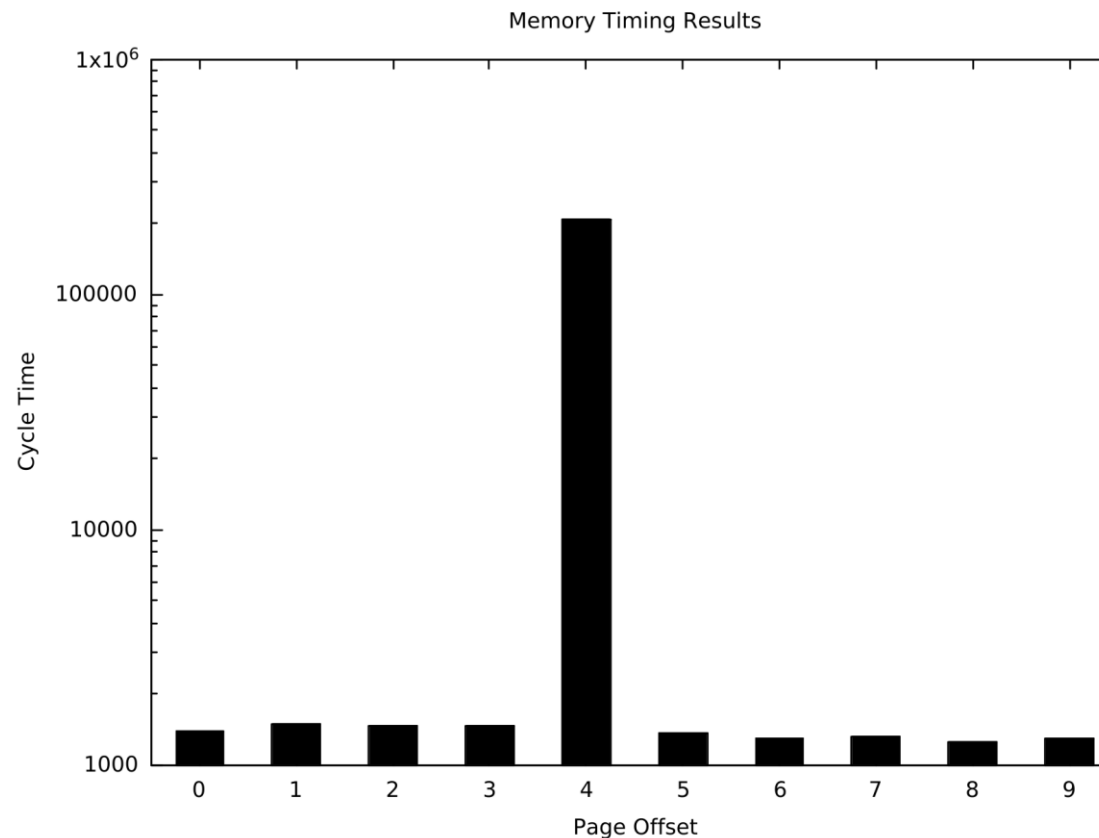
Results

Instruction Intercession – Instruction Timing



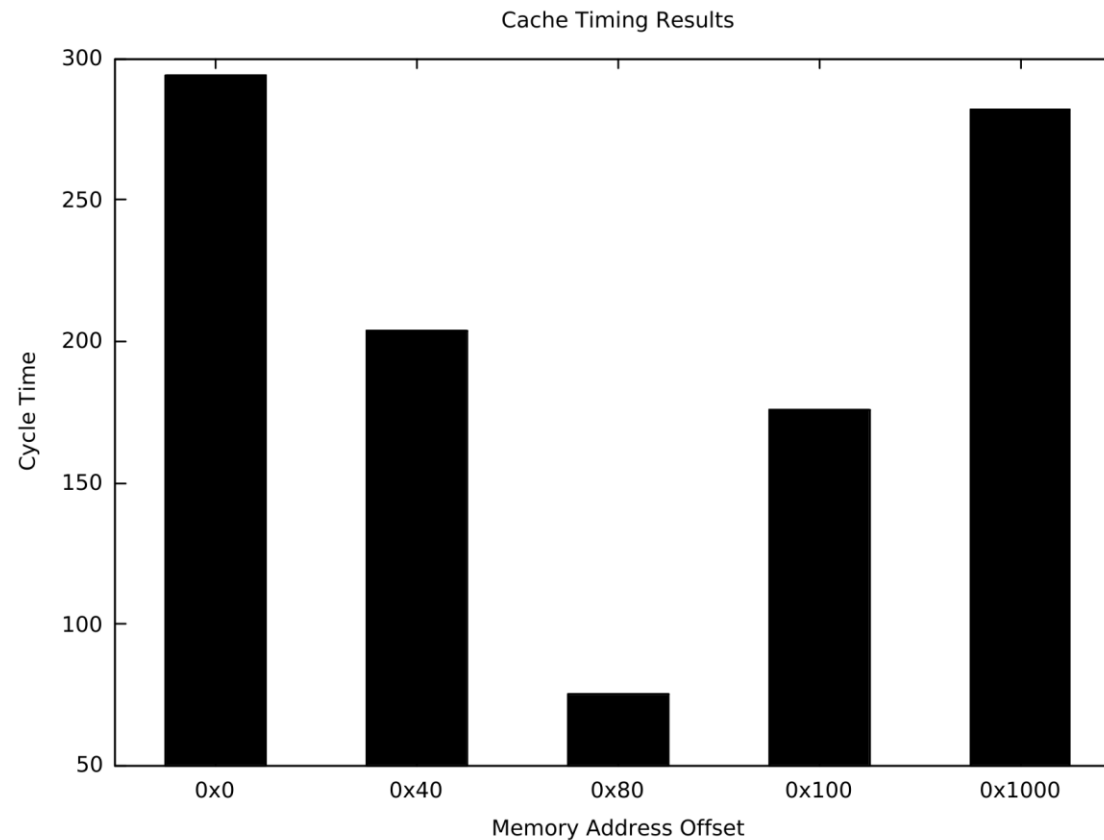
Results

Active Memory Intercession



Results

Passive Memory Intercession



Future Work

Next Steps

- This is a **first look** into hypervisor detection technology
- A **continuous detection environment** could be implemented
- Full **binary classification** could better categorize introspection
- A **response strategy** when introspection is discovered

Future Work

Technology on the Horizon

- New virtualization extensions limit detection via timing since they **reduce overhead**
 - *Virtualization Exceptions (#VE)*
 - *VMFUNC*
- Sub-page permissions permit **memory protections at a 128-byte granularity**

Conclusion

- Detection of instruction intercession is possible
- Active and passive memory monitoring is possible
 - Isolating a memory region to a specific process may be necessary
 - Active and passive monitoring used in conjunction could obfuscate results of timing techniques
- This work establishes the **limitations of hypervisor introspection detection**

Questions?

- We intend to open-source this work [whenever we get around to it]

This research is based upon work supported in part by the Office of the Director of National Intelligence (ODNI), Intelligence Advanced Research Projects Activity (IARPA). The views and conclusions contained herein are those of the authors and should not be interpreted as necessarily representing the official policies, either expressed or implied, of ODNI, IARPA, or the U.S. Government. The U.S. Government is authorized to reproduce and distribute reprints for governmental purposes notwithstanding any copyright annotation therein.